

NSQ

a realtime distributed messaging platform

<https://github.com/bitly/nsq>

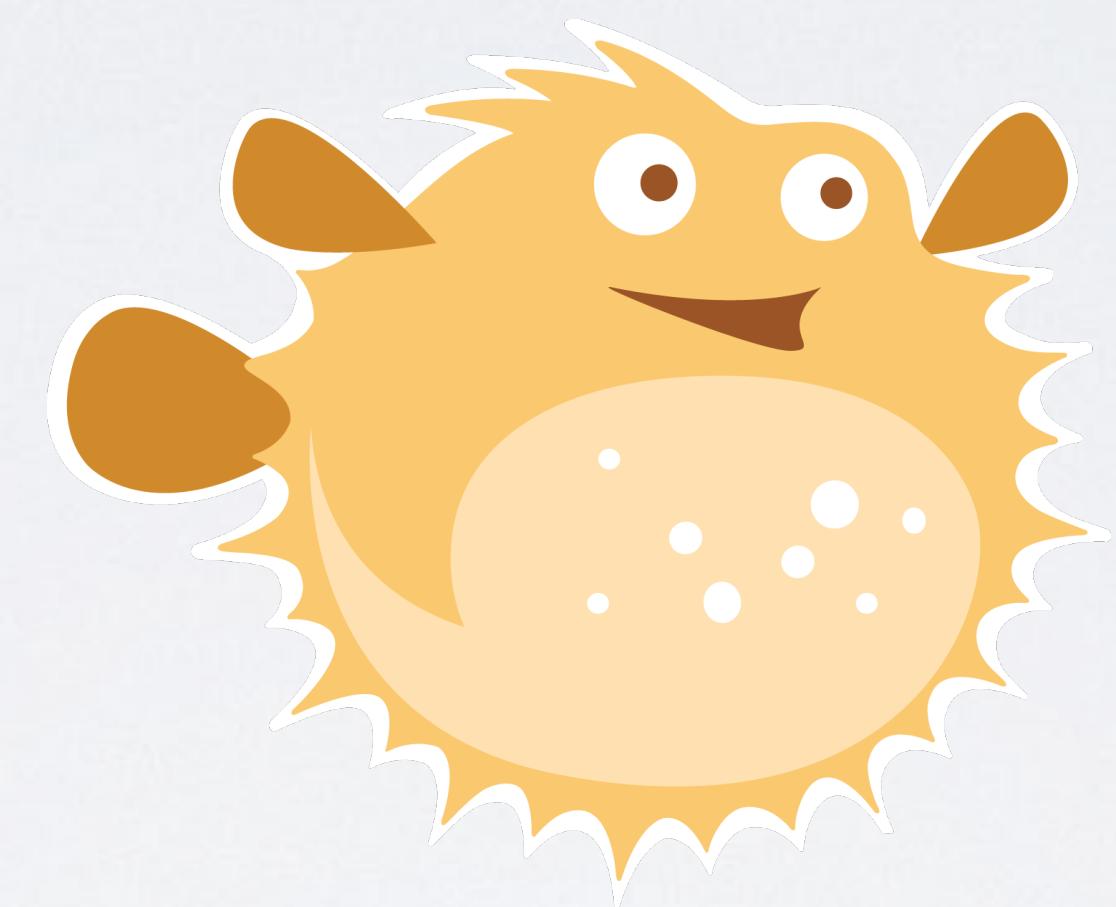
@imsnakes

April 24th 2014 - GopherCon

Matt Reiferson - CTO at Torando Labs



HI, I'M MATT



HI, I'M MATT

NAMING SUCKS QUEUE

WHAT EVEN IS NSQ?

NOT STABLE QUEUE

NEW SIMPLE QUEUE

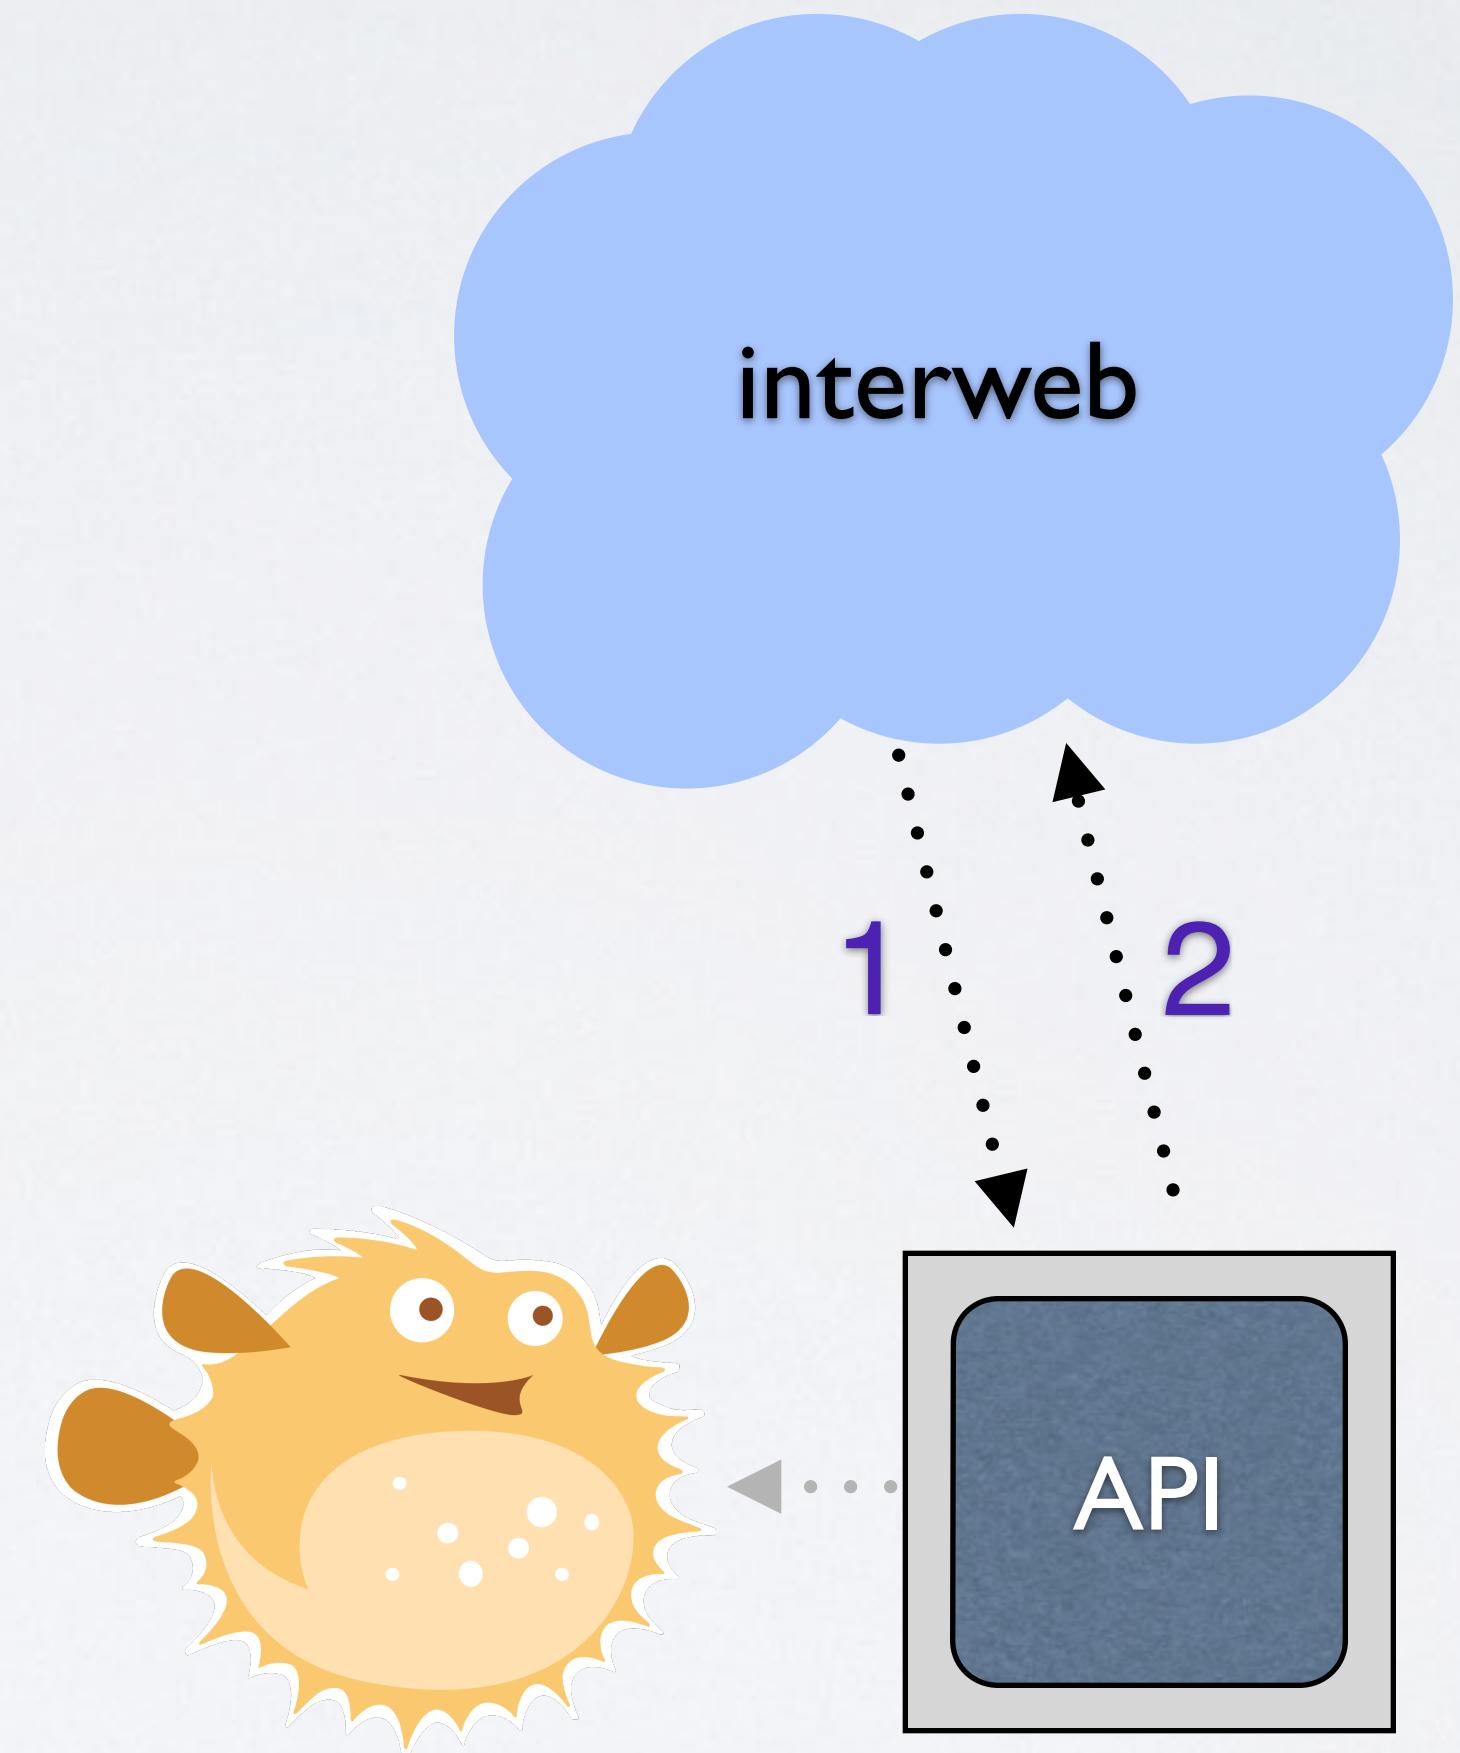


SPRAY SOME NSQ ON IT



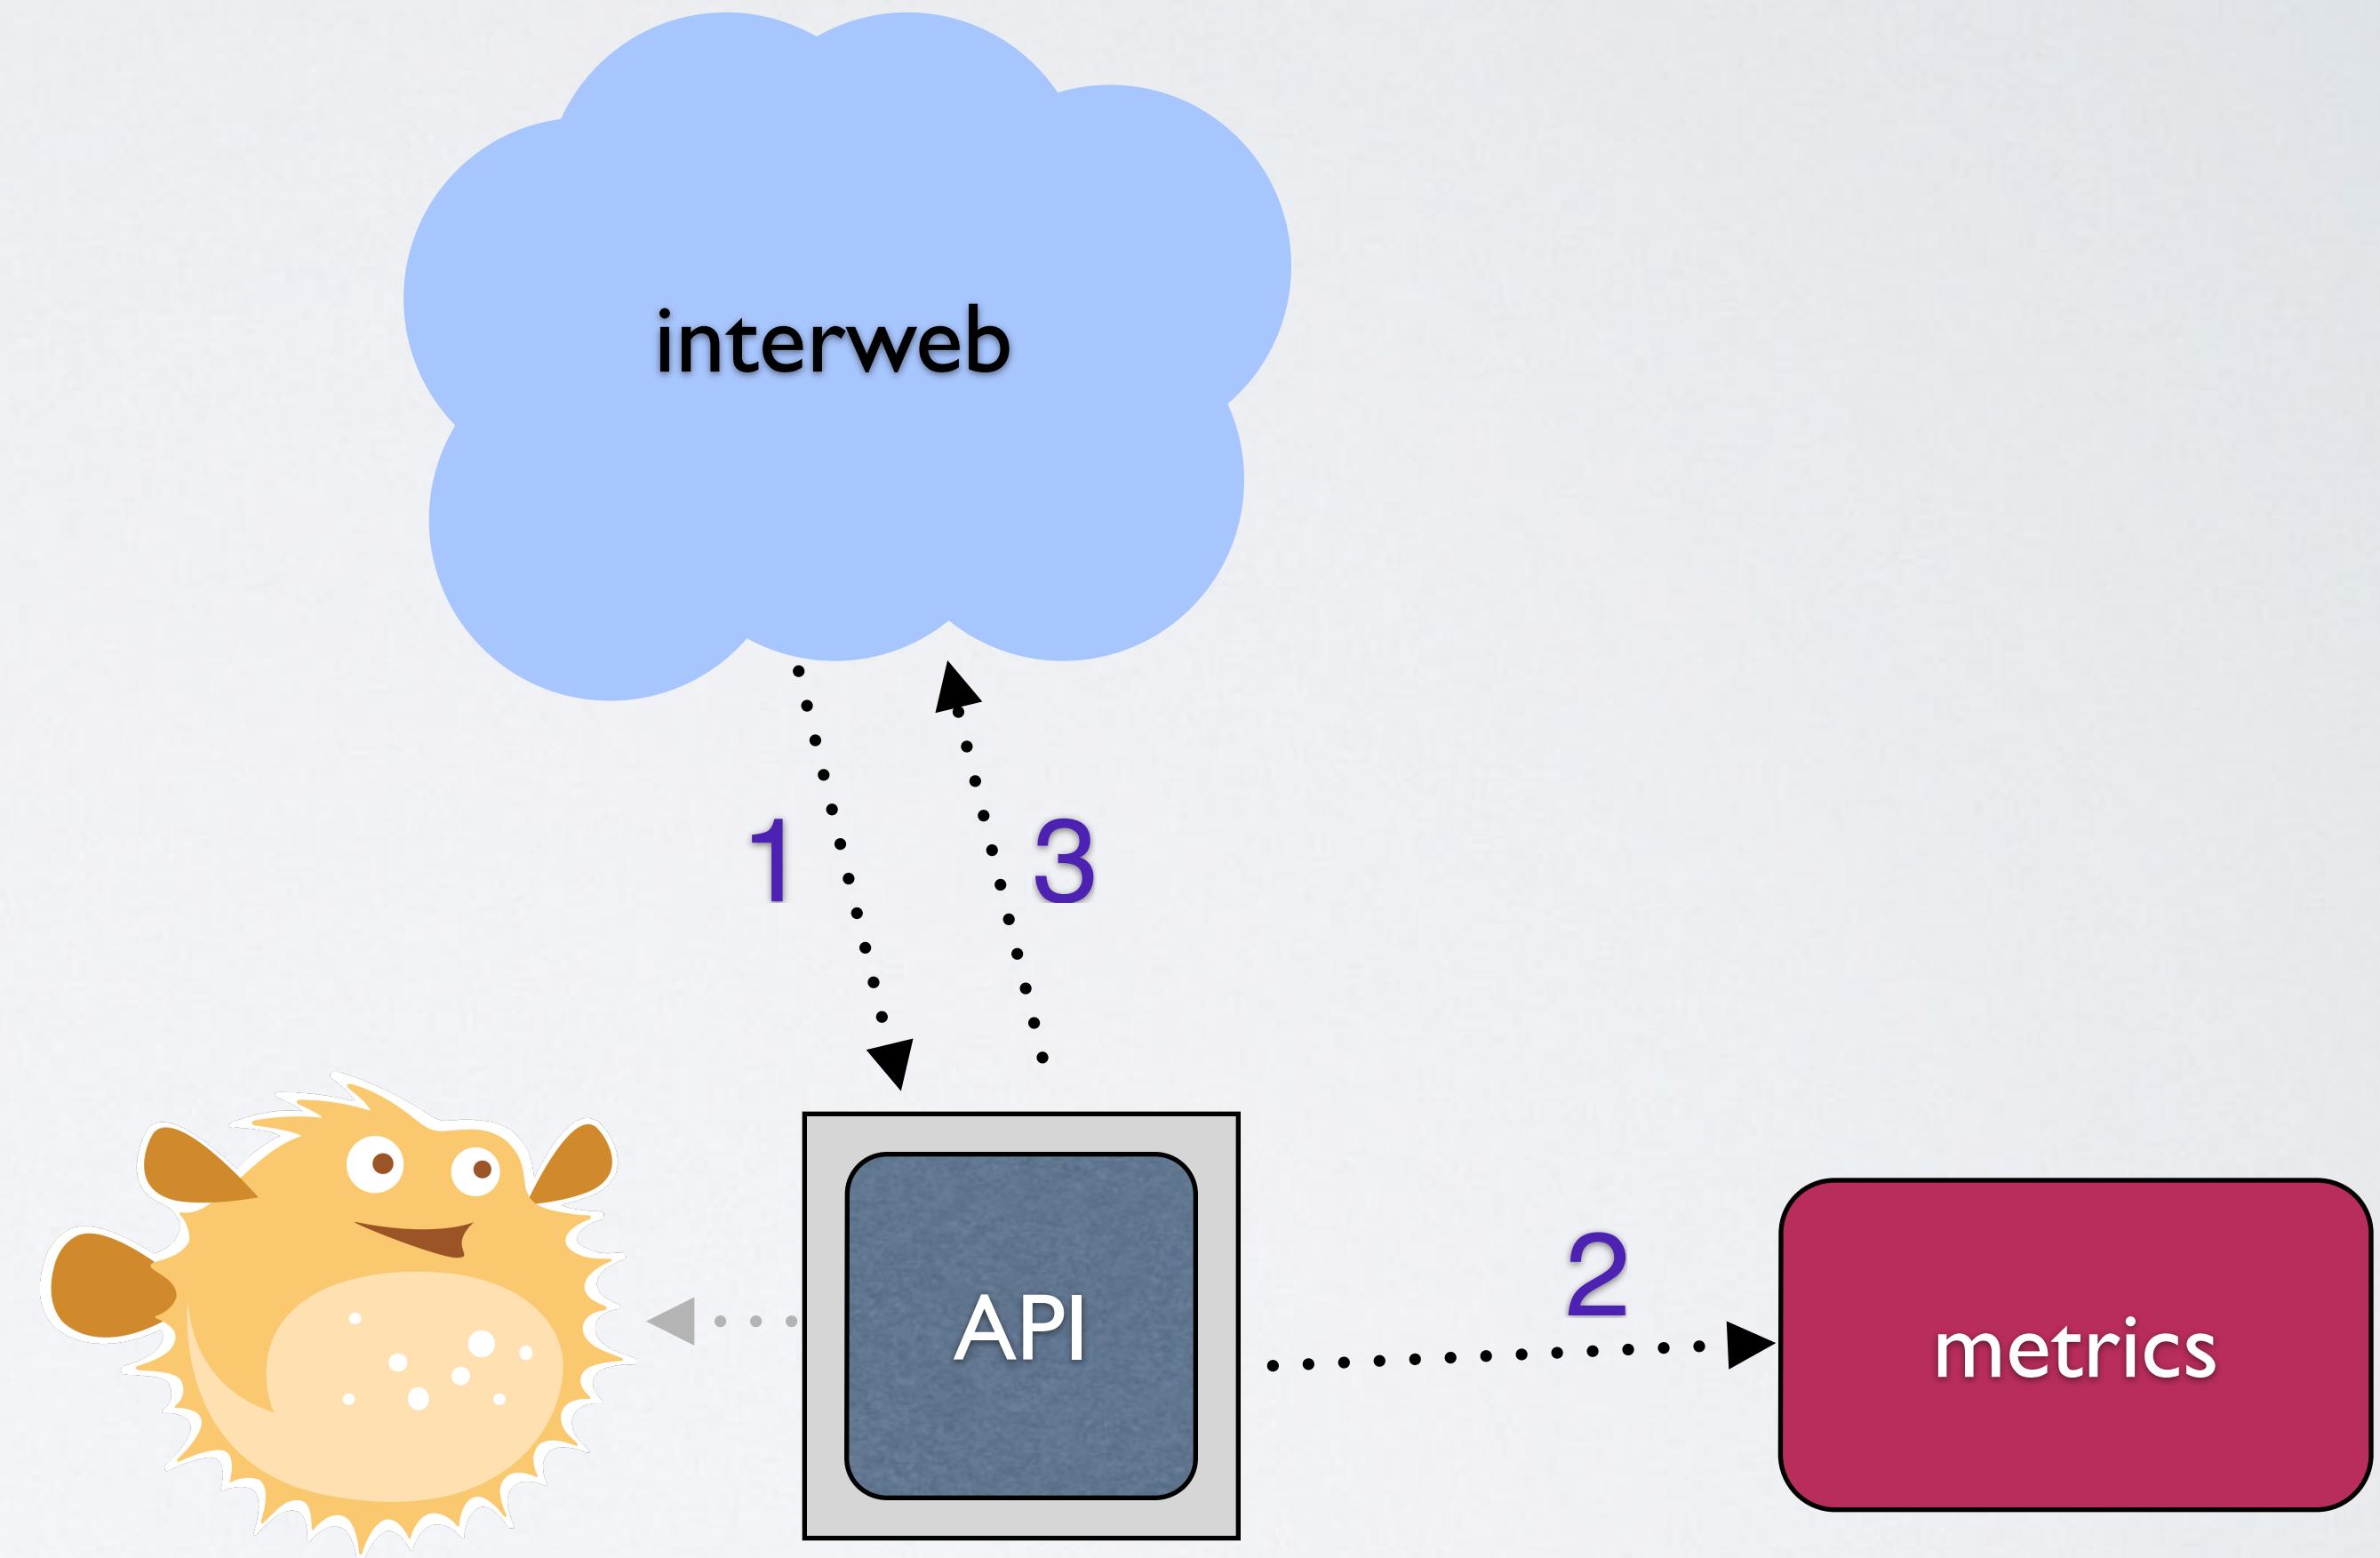
BIG THINGS HAVE SMALL BEGINNINGS

- boss says: “I want metrics!”
- single host
- synchronous writes

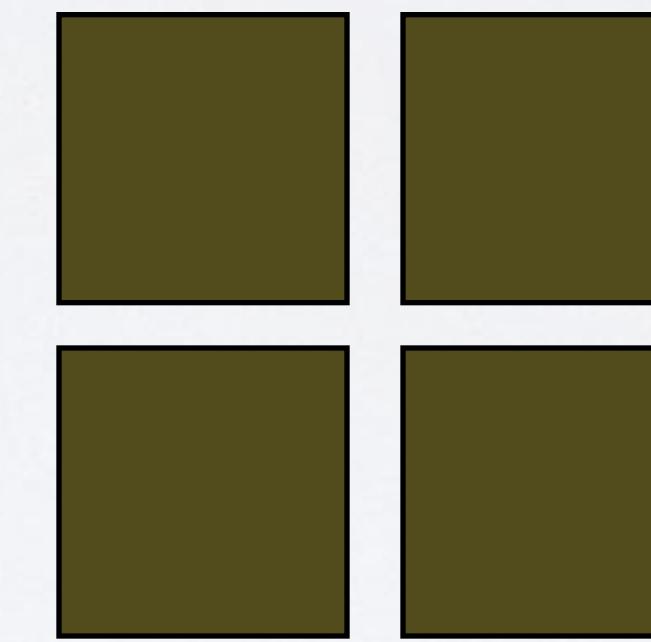
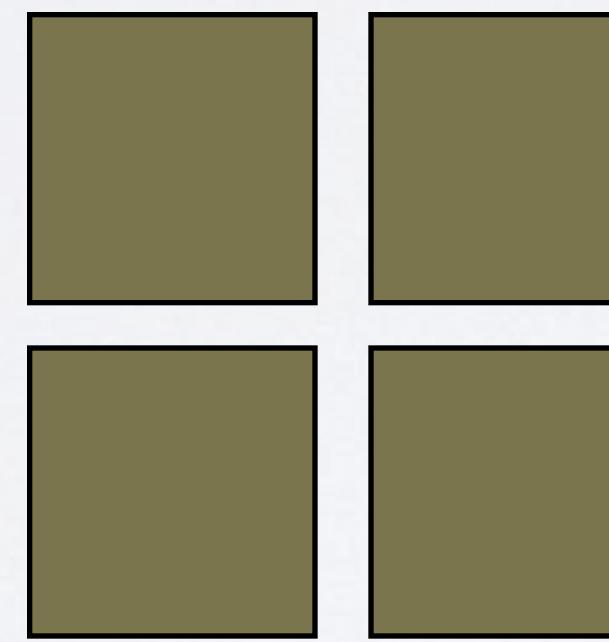
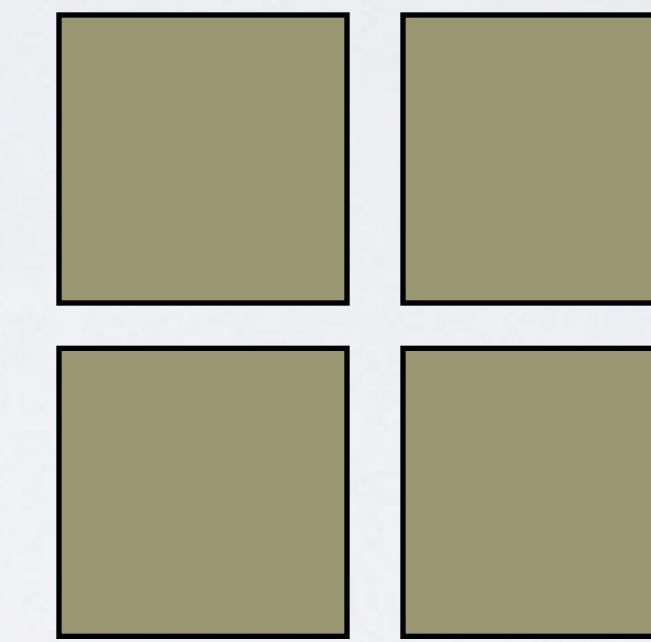
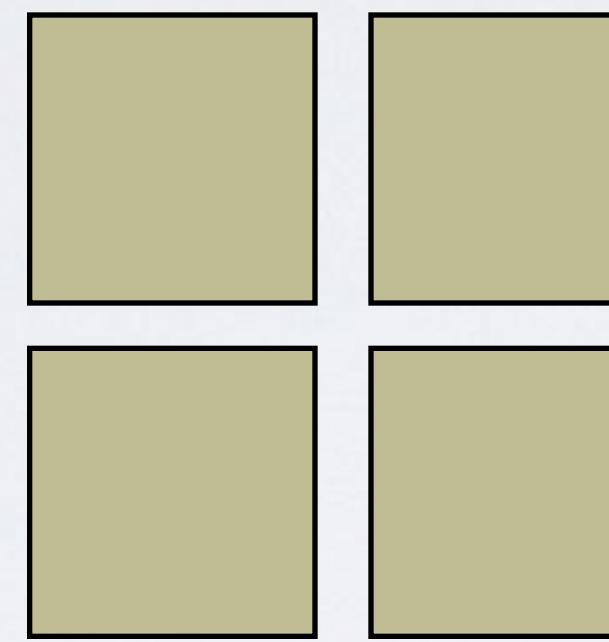


BIG THINGS HAVE SMALL BEGINNINGS

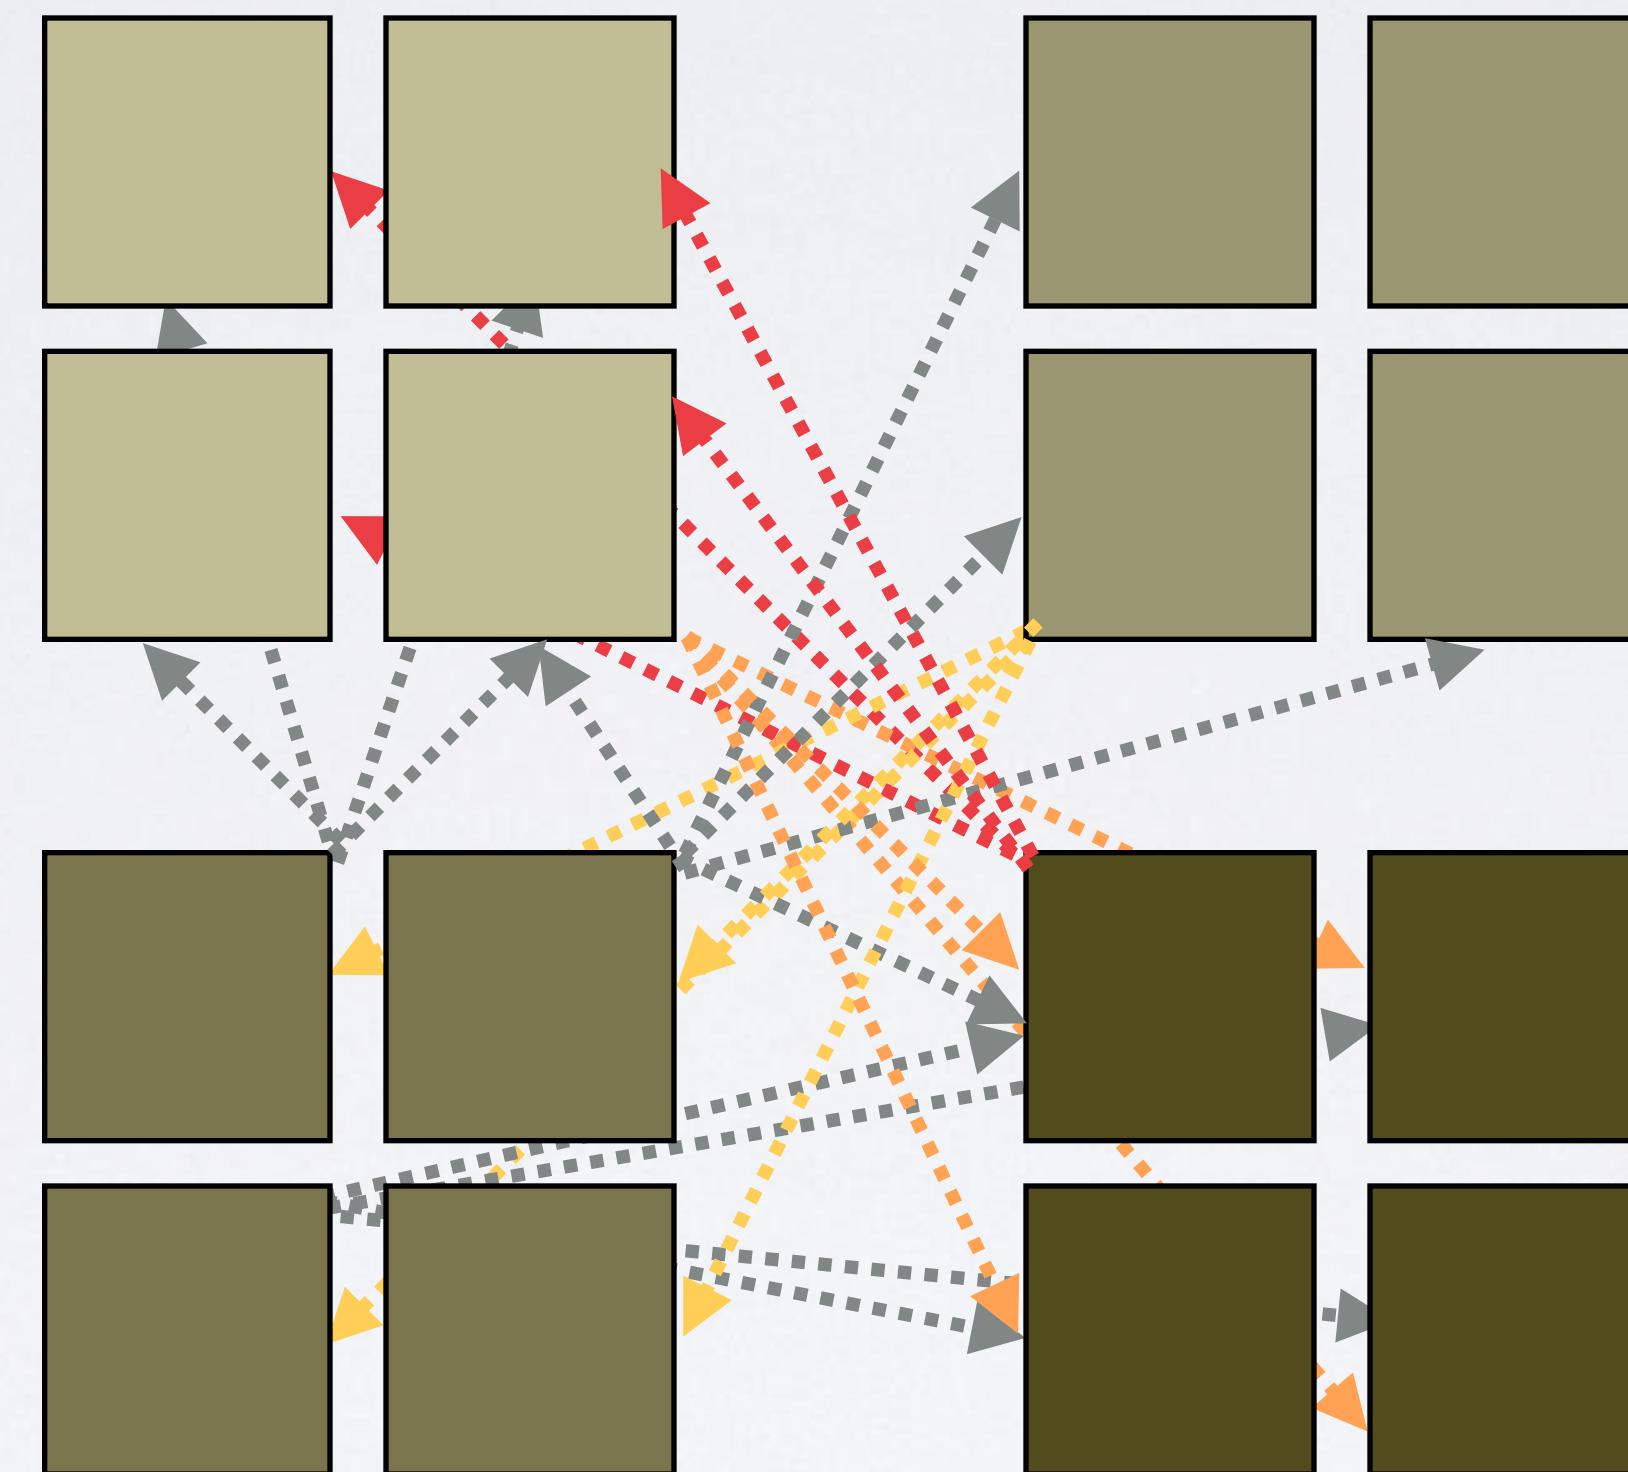
- boss says: “I want metrics!”
- single host
- synchronous writes



...AND THEN THEY EAT YOU



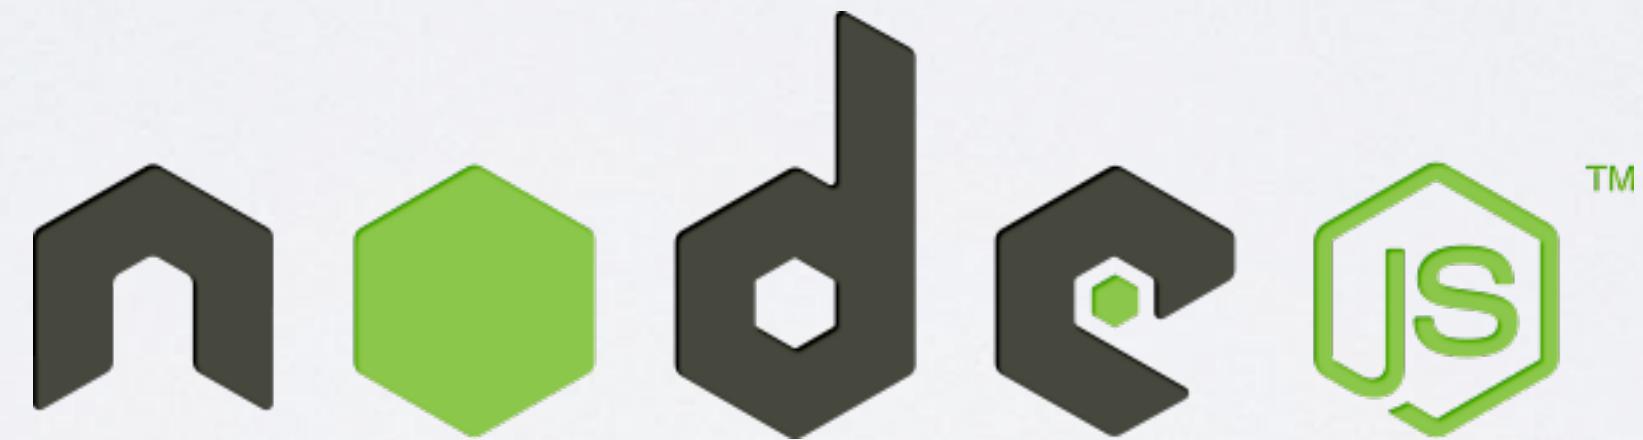
...AND THEN THEY EAT YOU



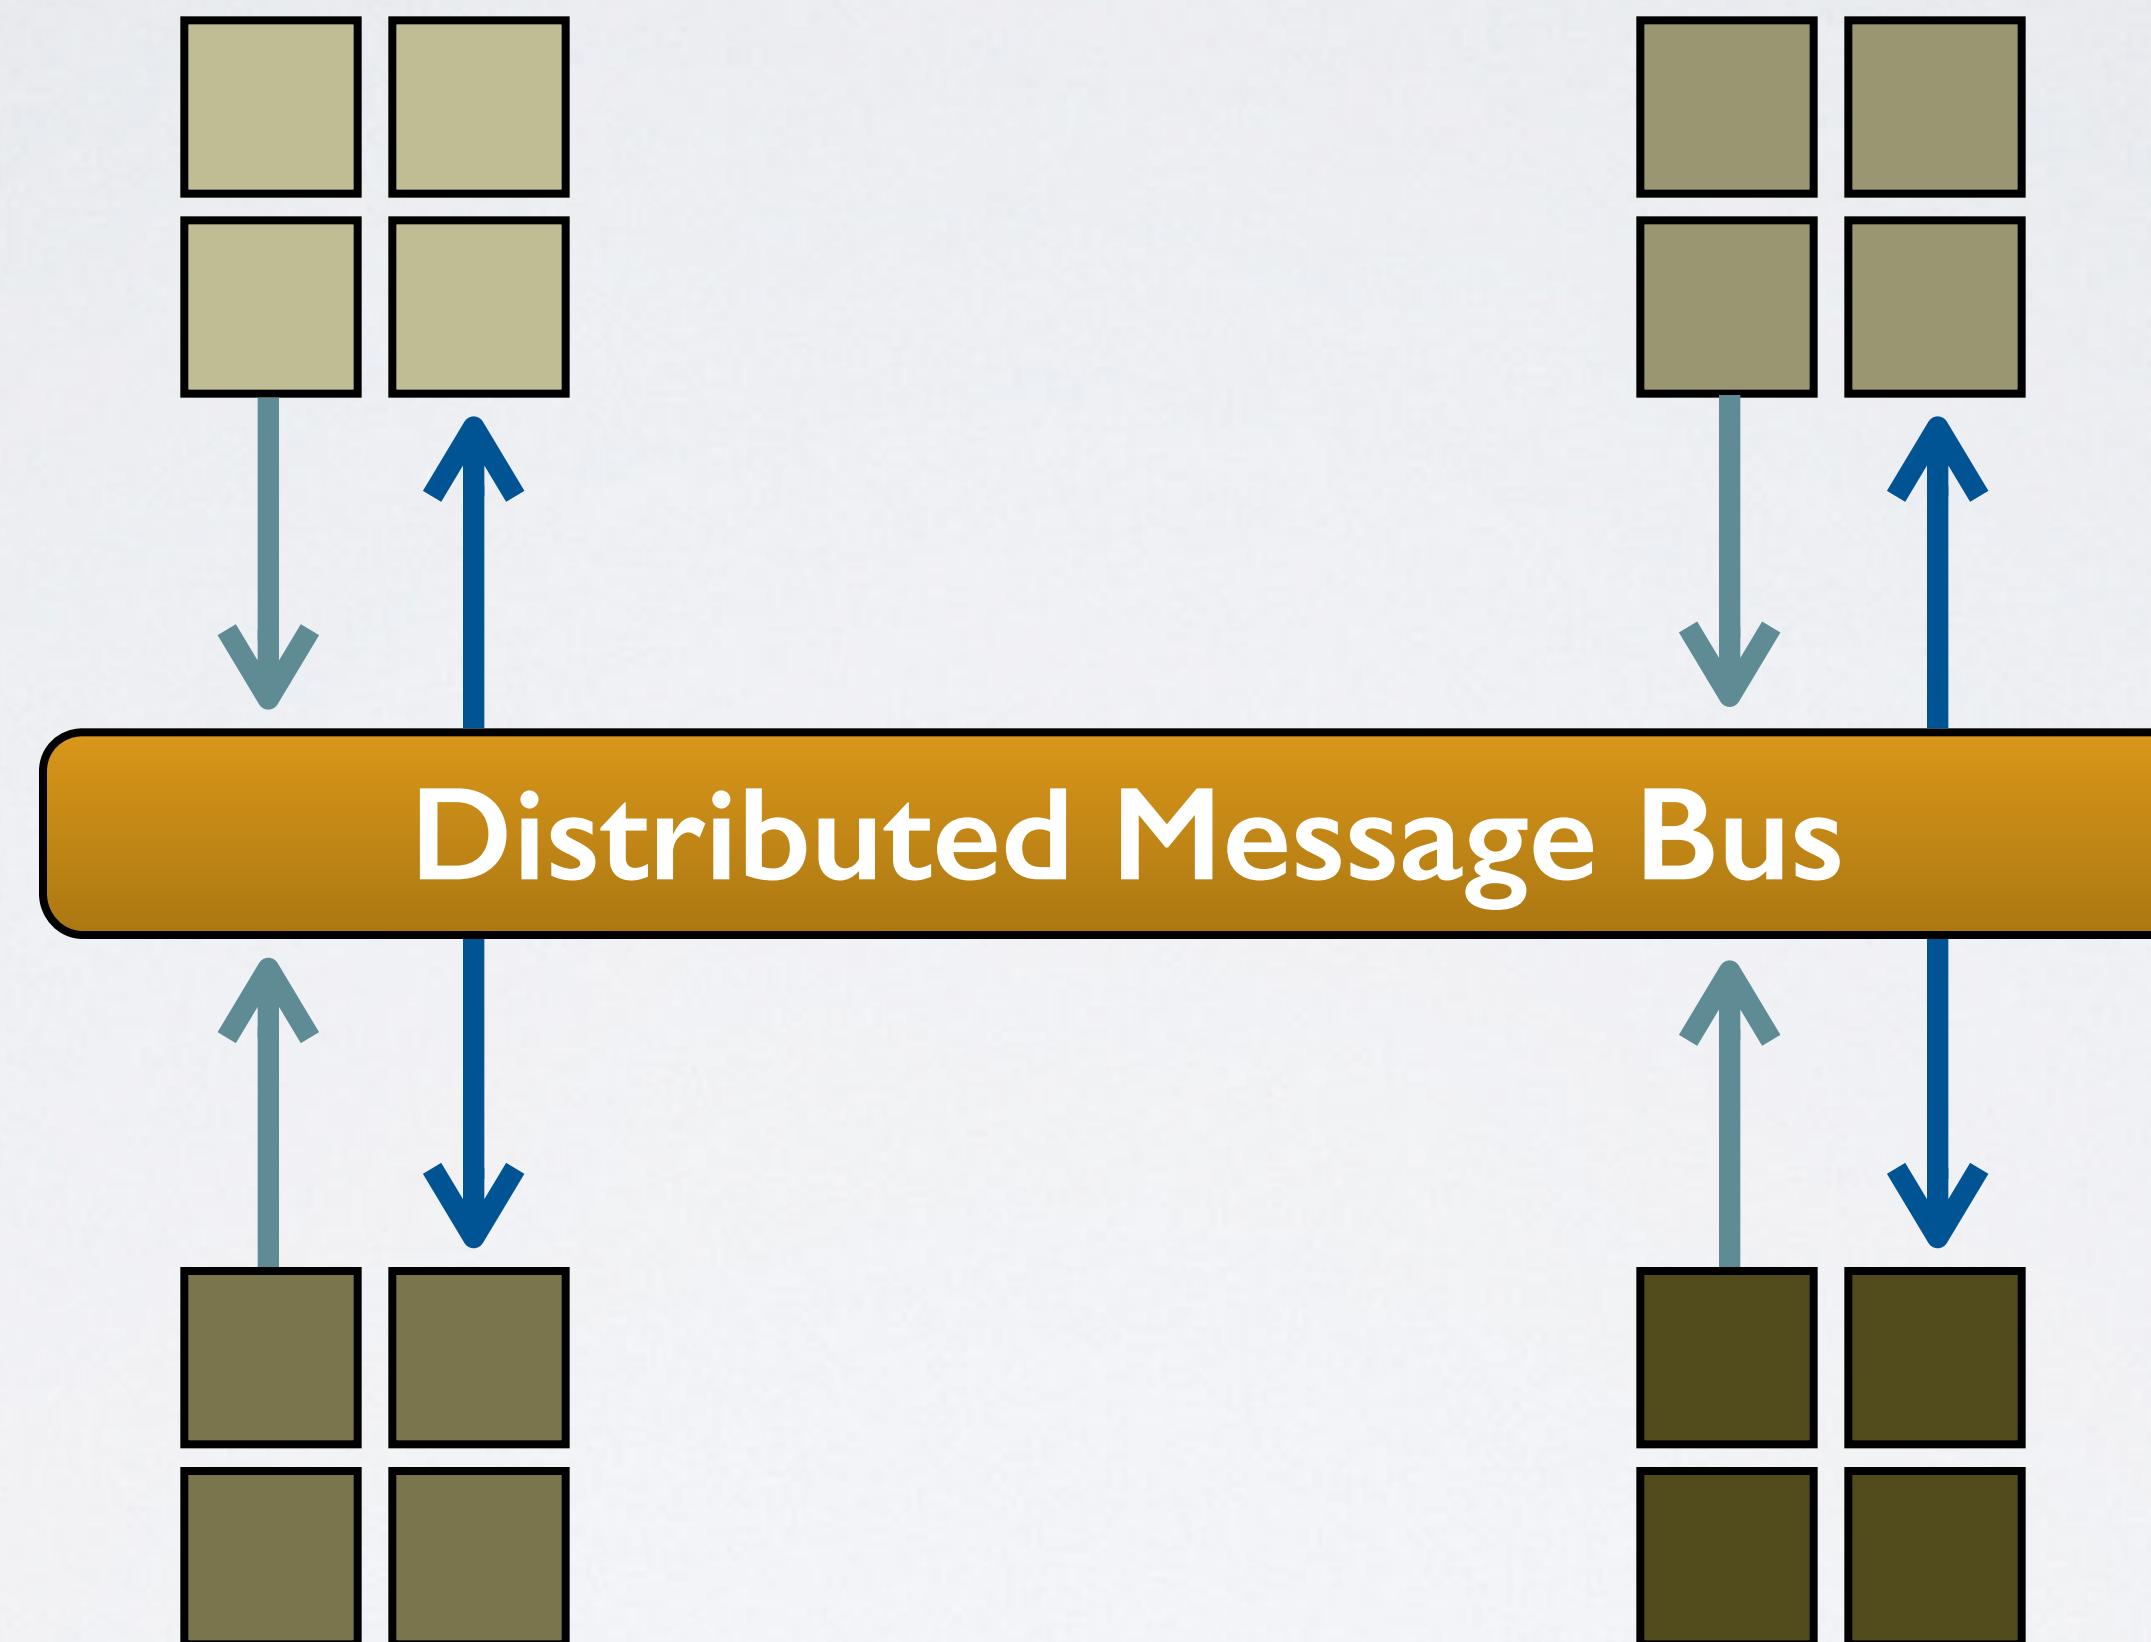
nightmare

BUT THERE'S HOPE

JUST USE



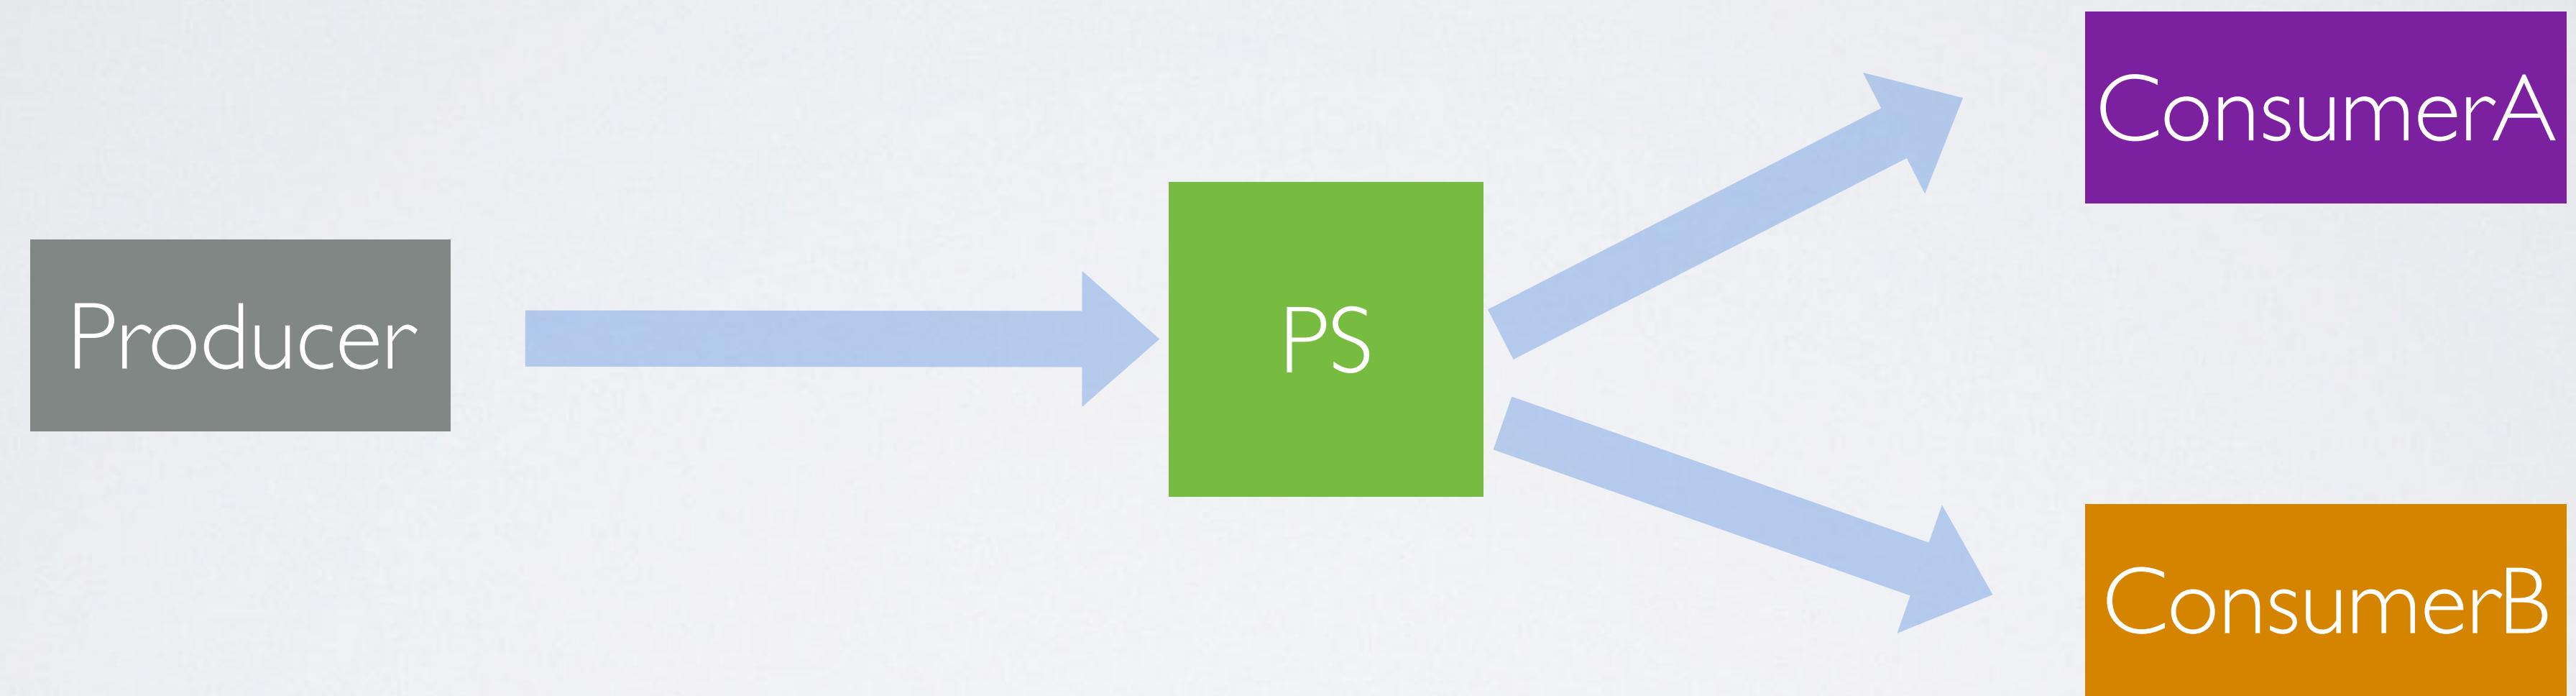
THE GOAL



provide a unifying distributed system to receive
and disseminate event data

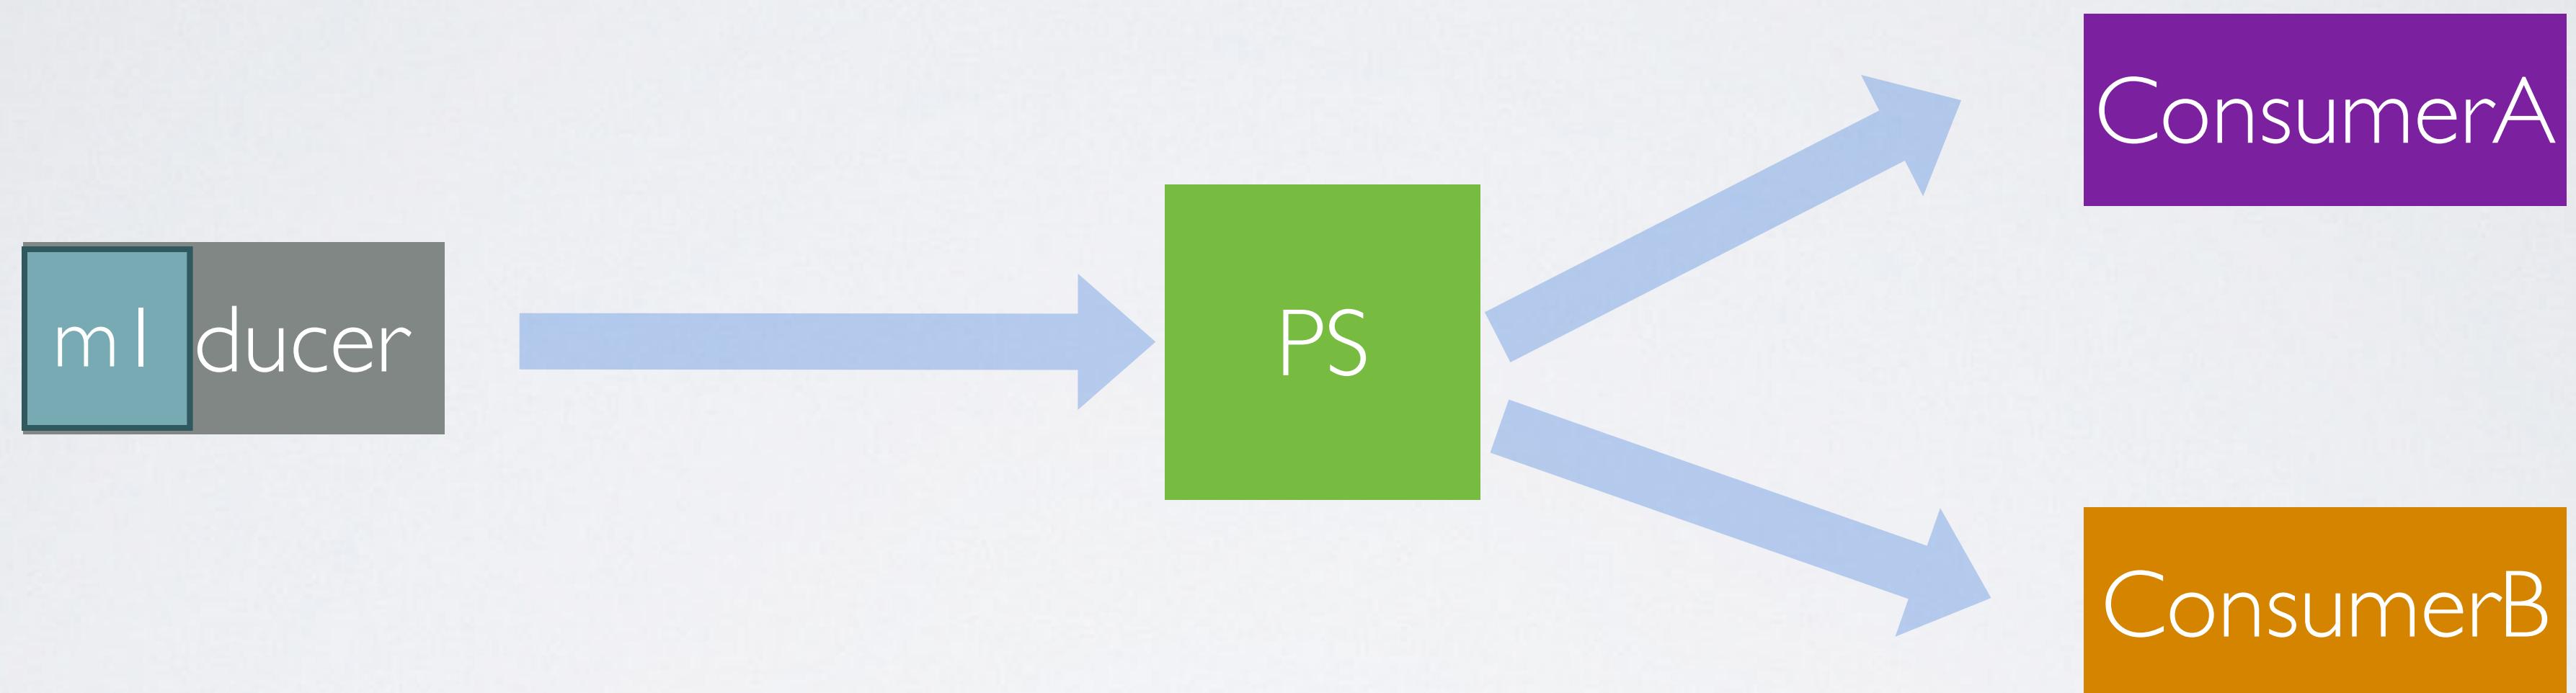
MESSAGING PATTERNS

BROADCAST



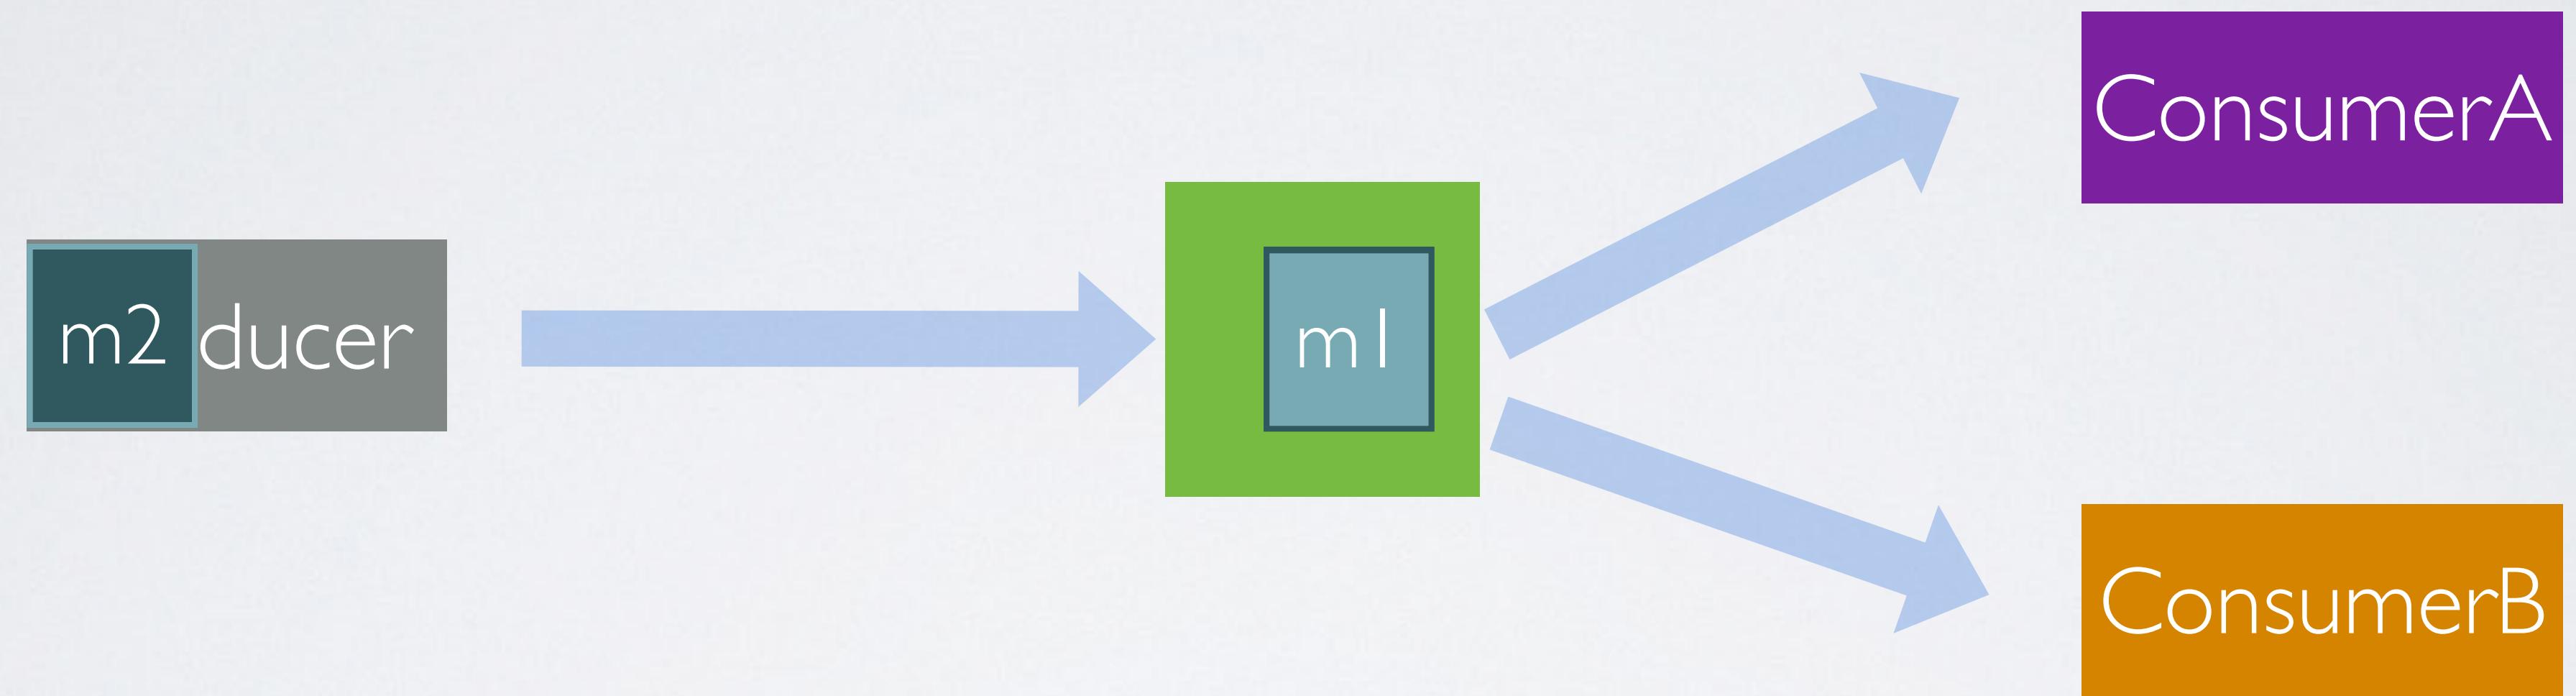
de-couple

BROADCAST



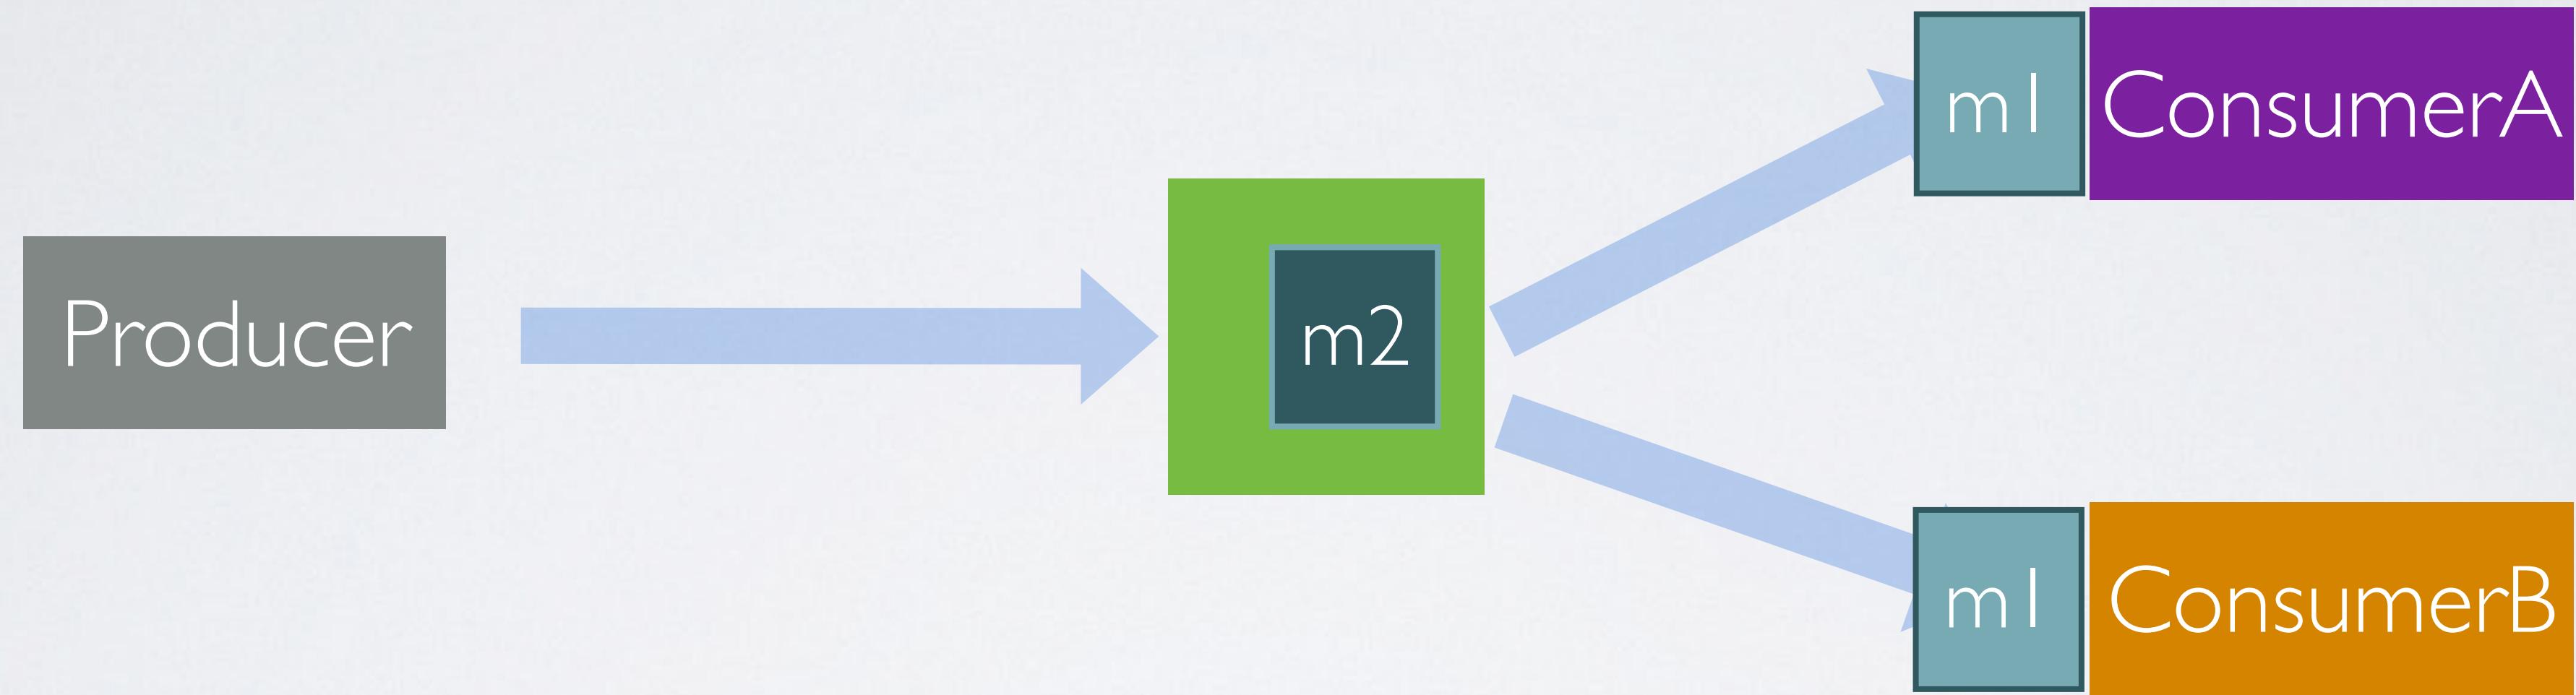
de-couple

BROADCAST



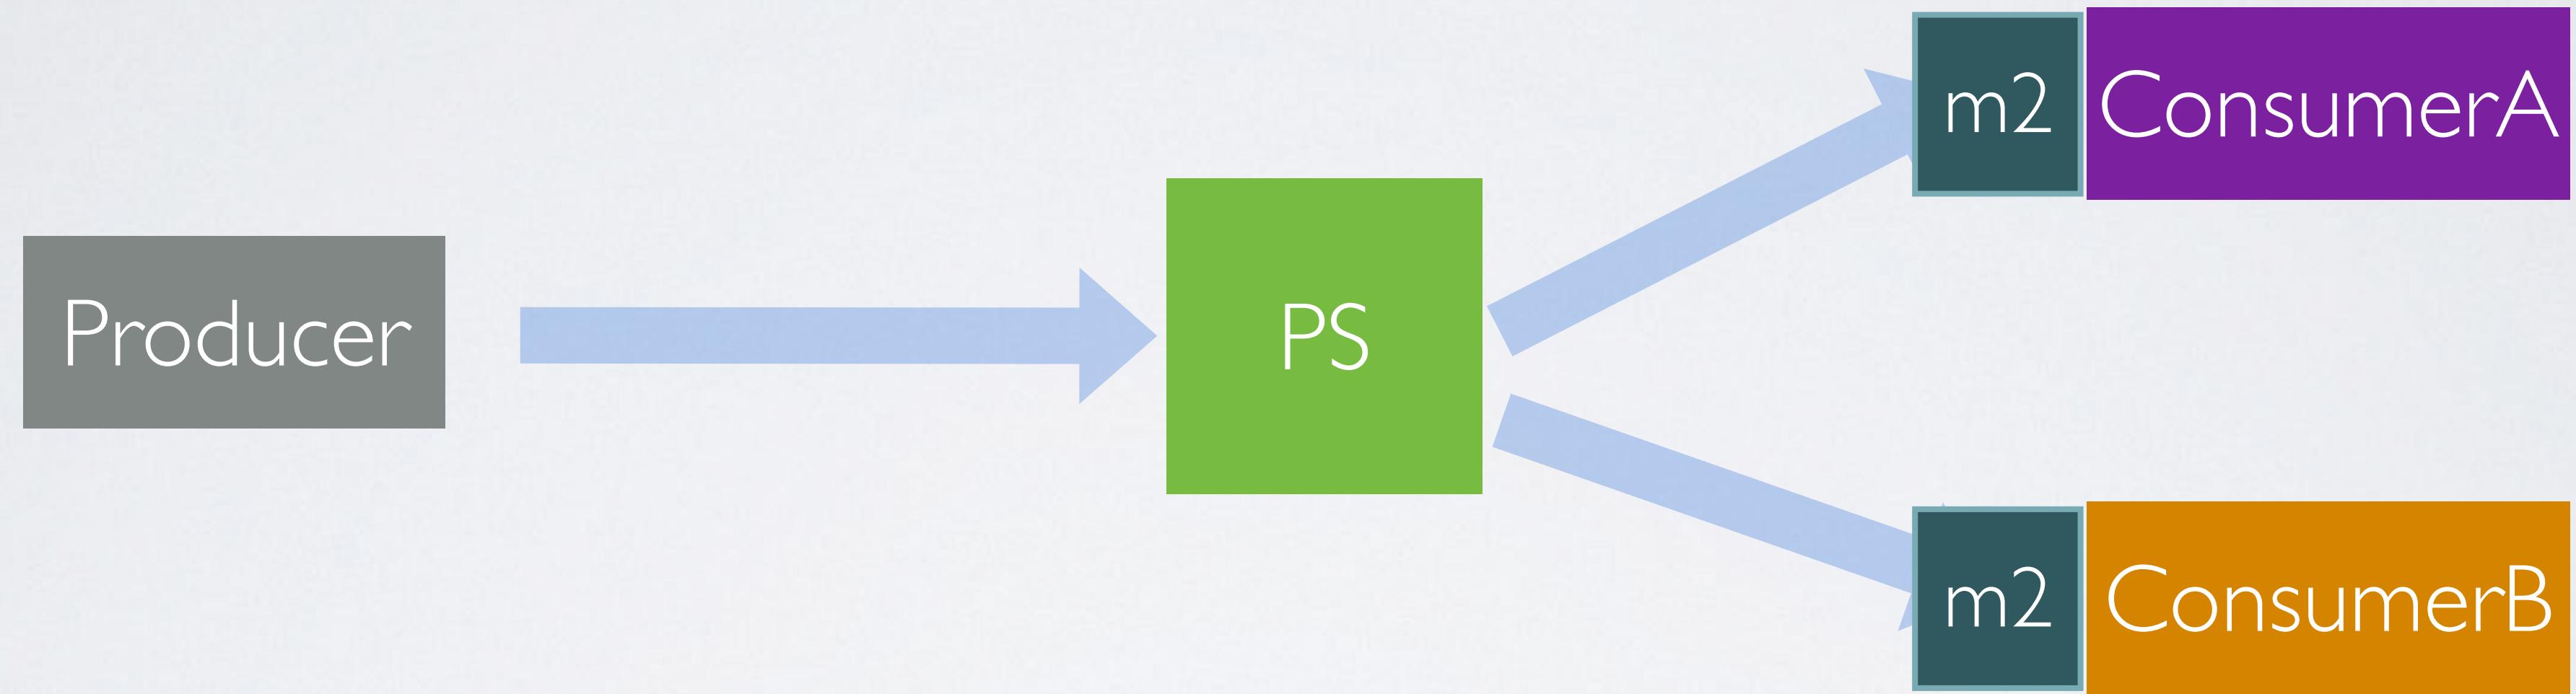
de-couple

BROADCAST



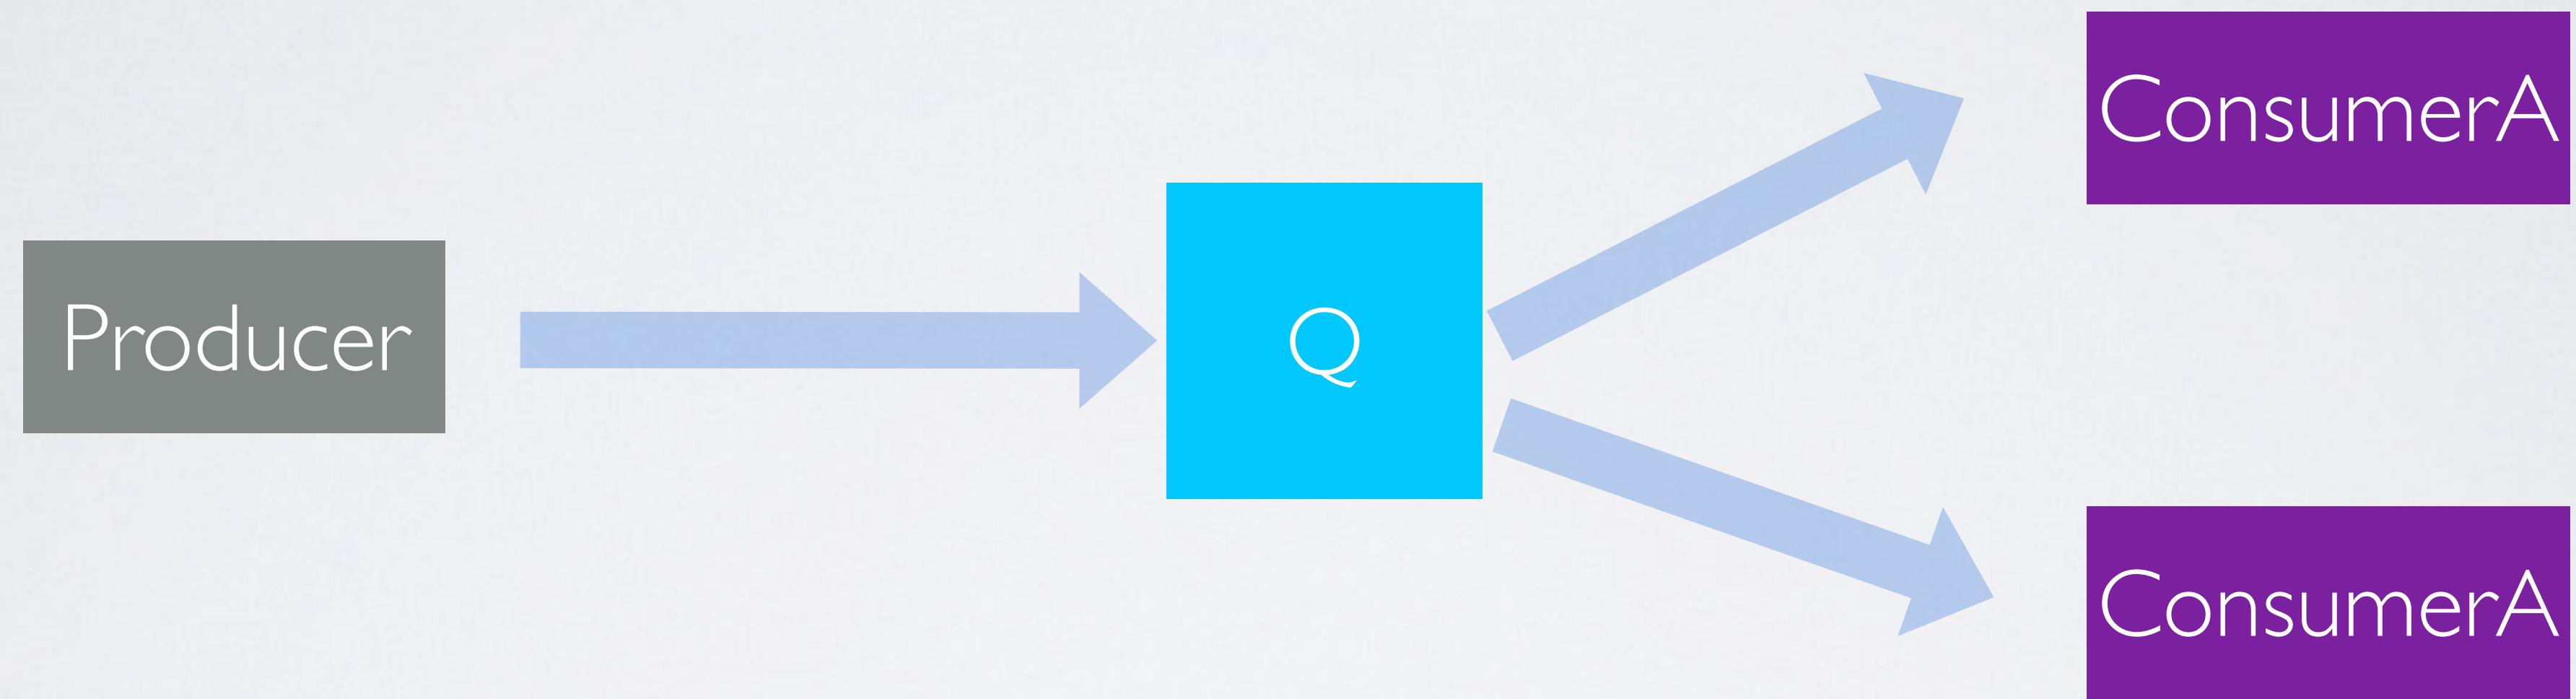
de-couple

BROADCAST



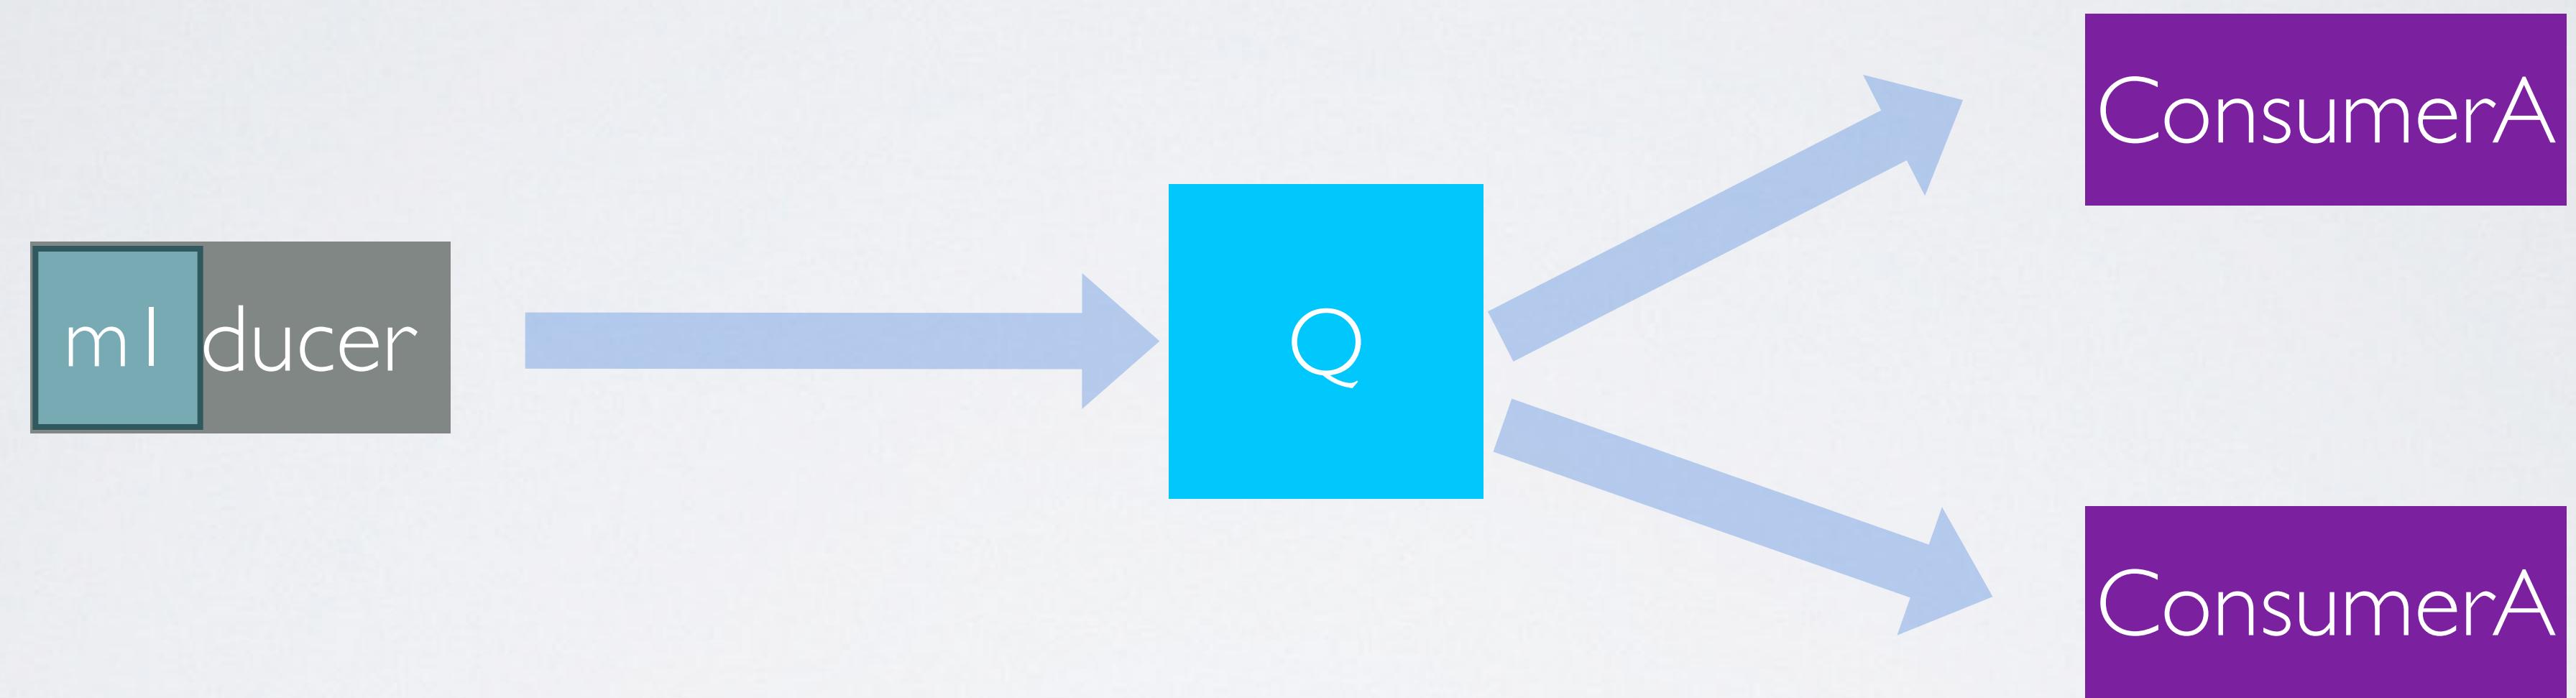
de-couple

DISTRIBUTION



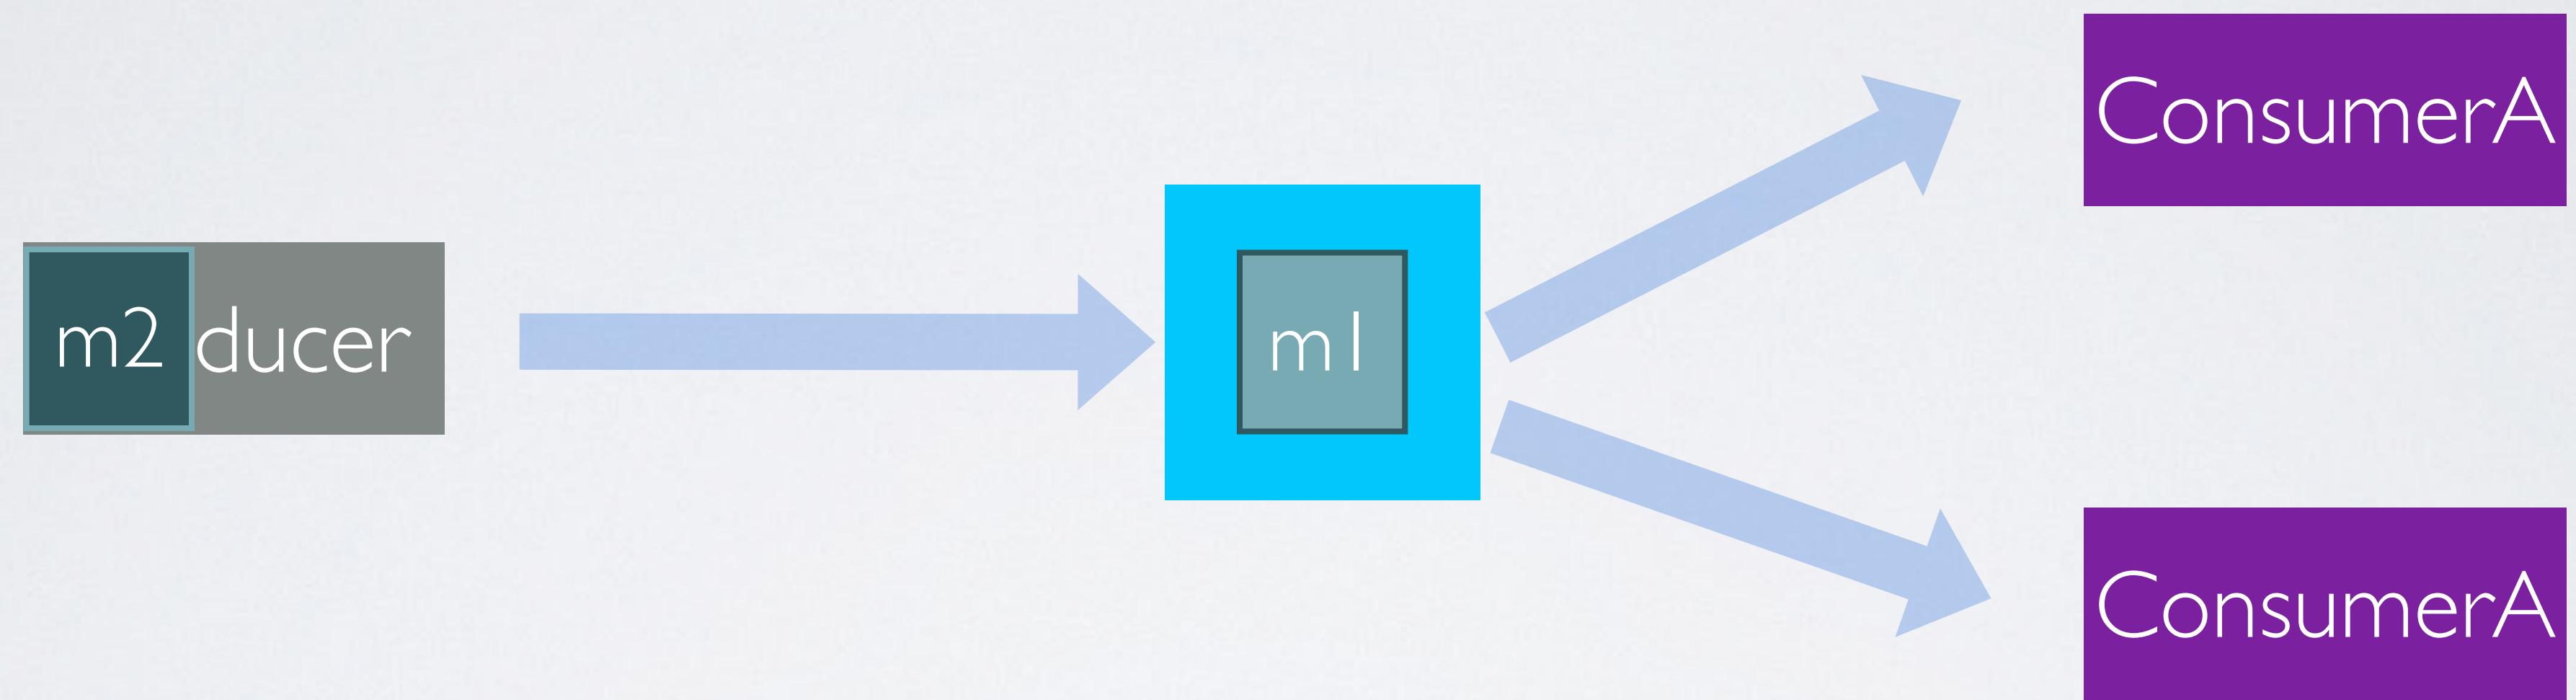
horizontal scalability

DISTRIBUTION



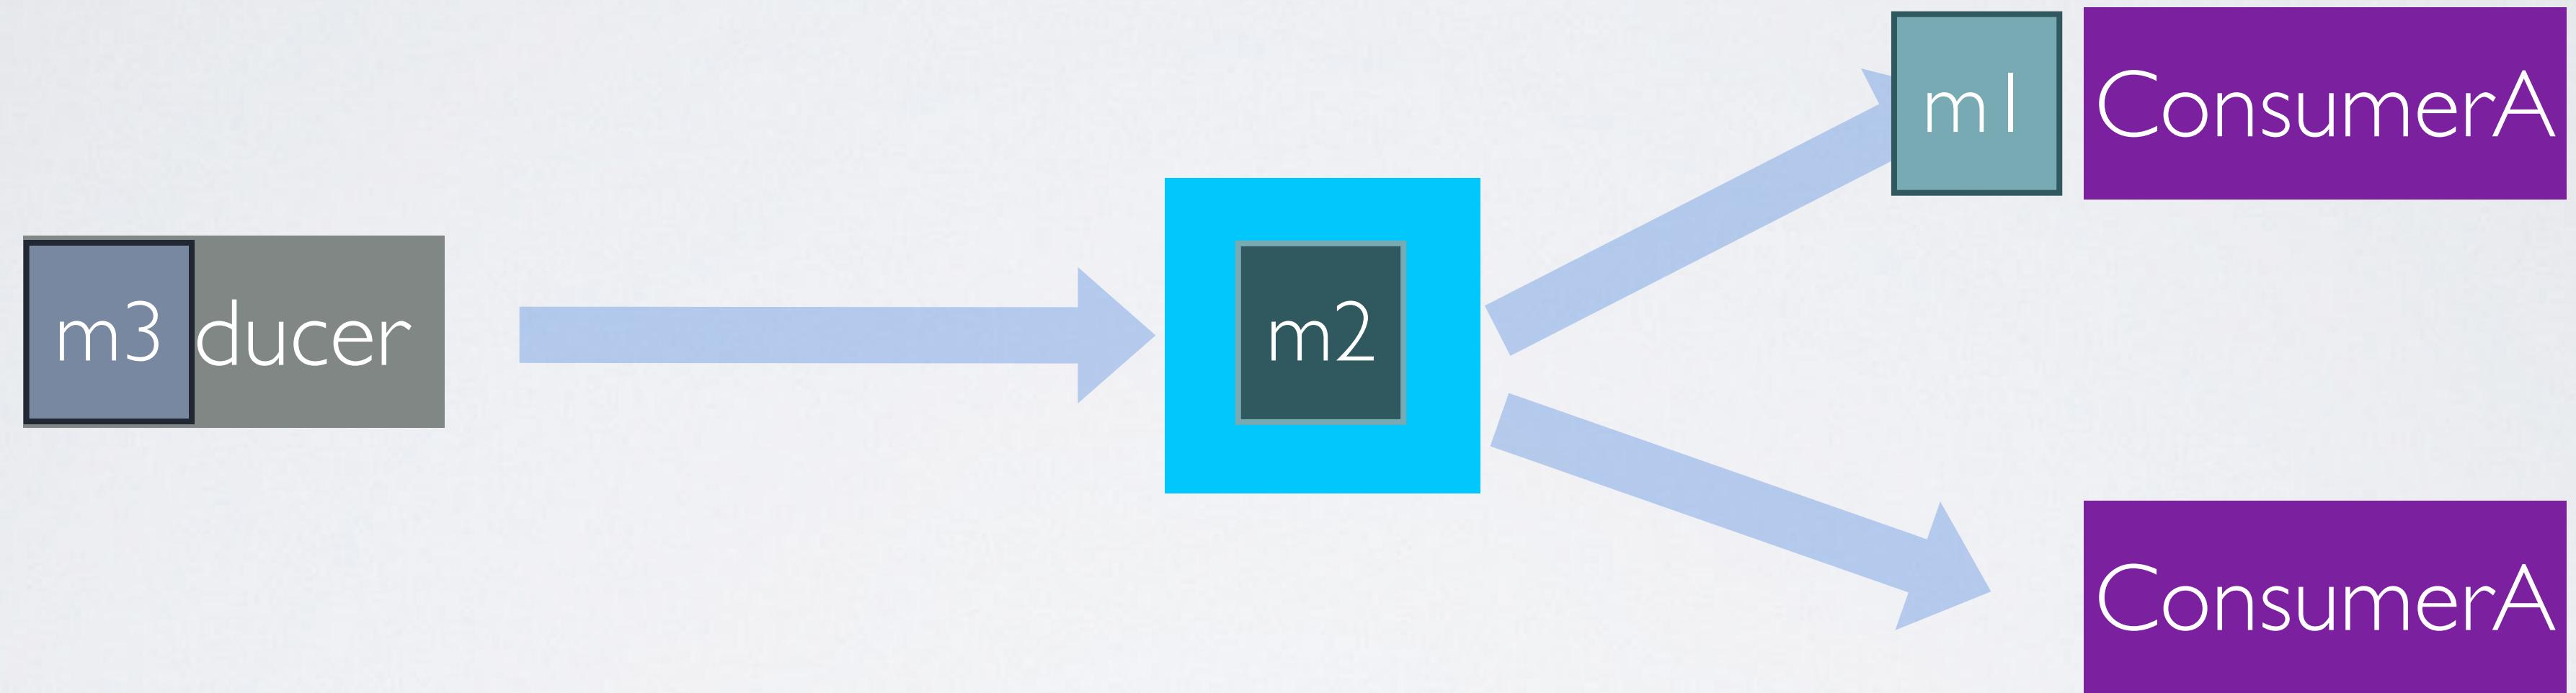
horizontal scalability

DISTRIBUTION



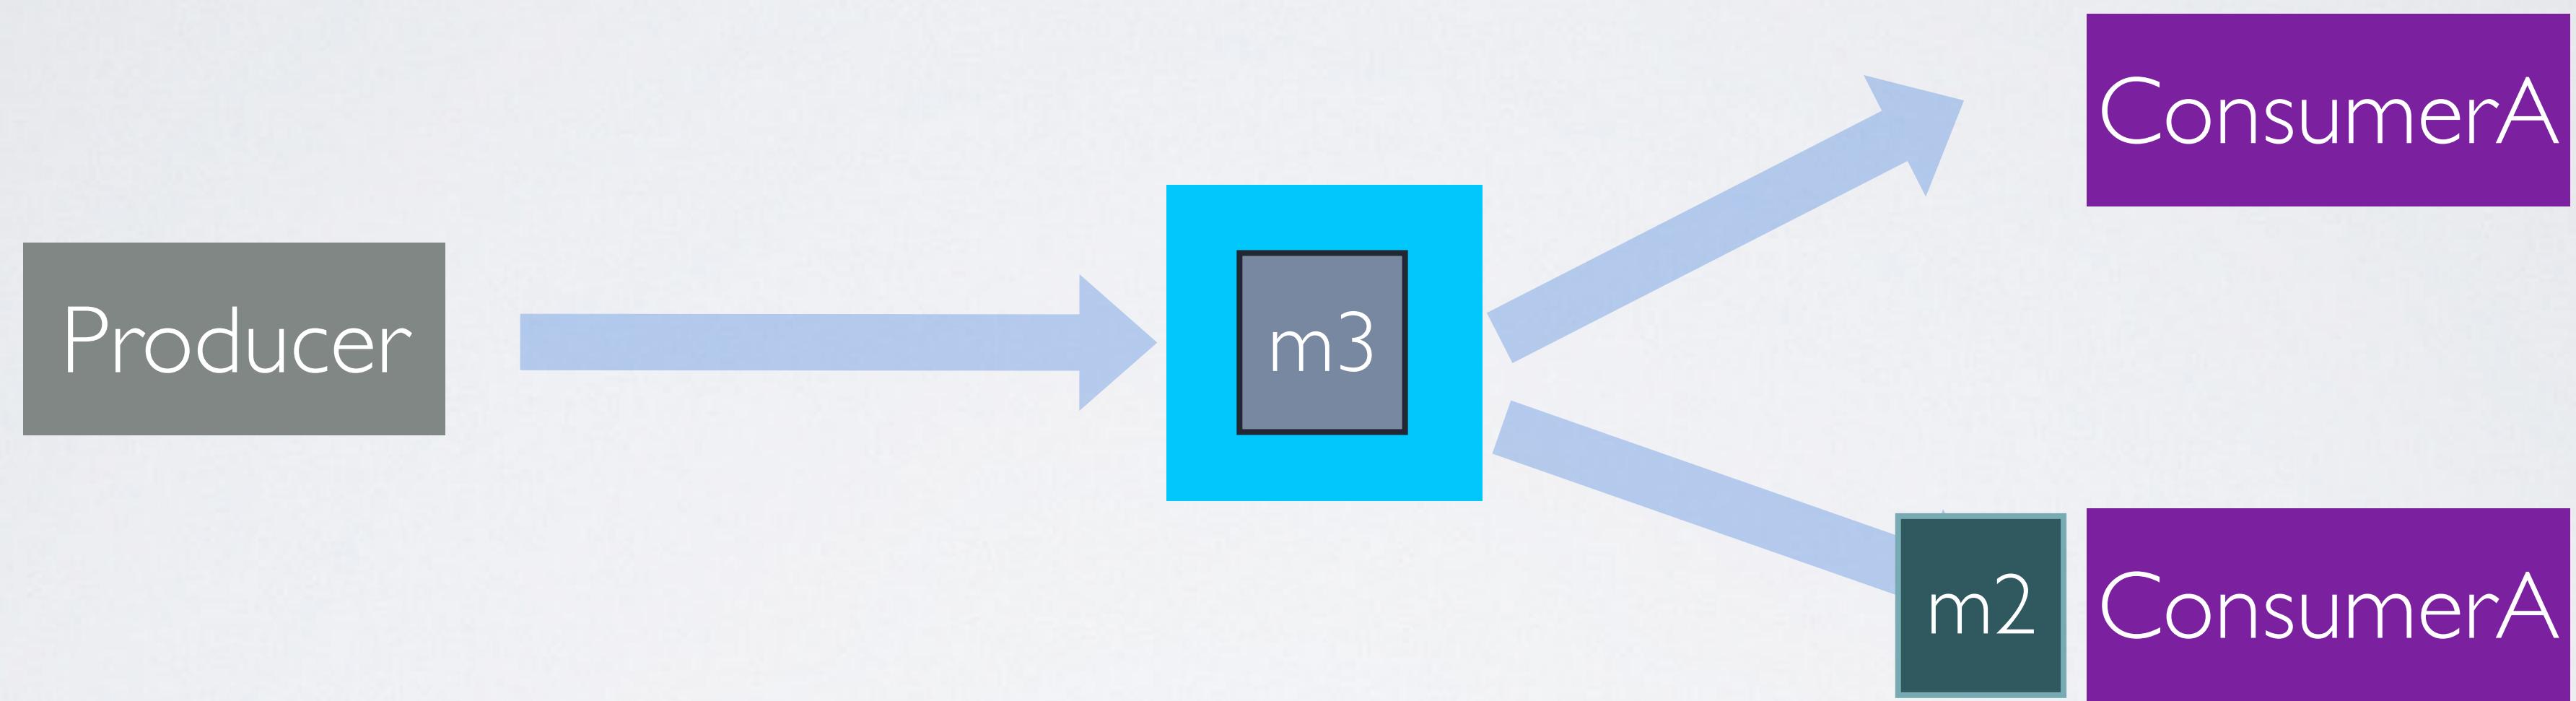
horizontal scalability

DISTRIBUTION



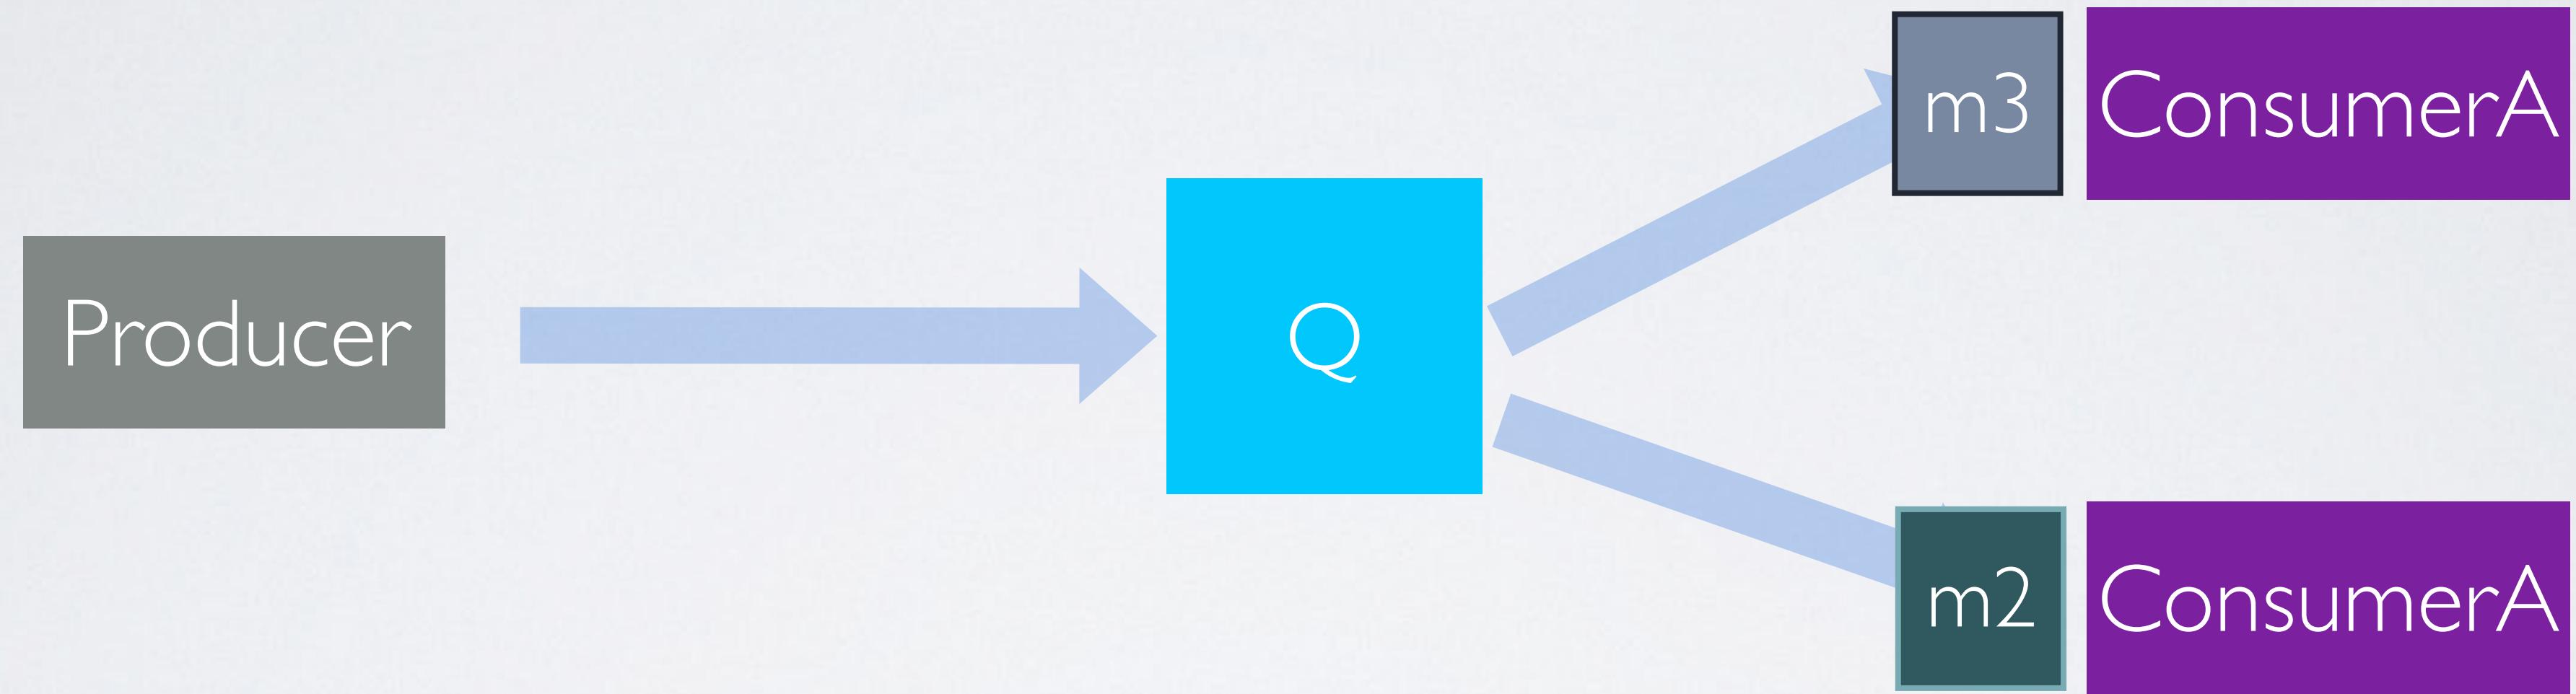
horizontal scalability

DISTRIBUTION



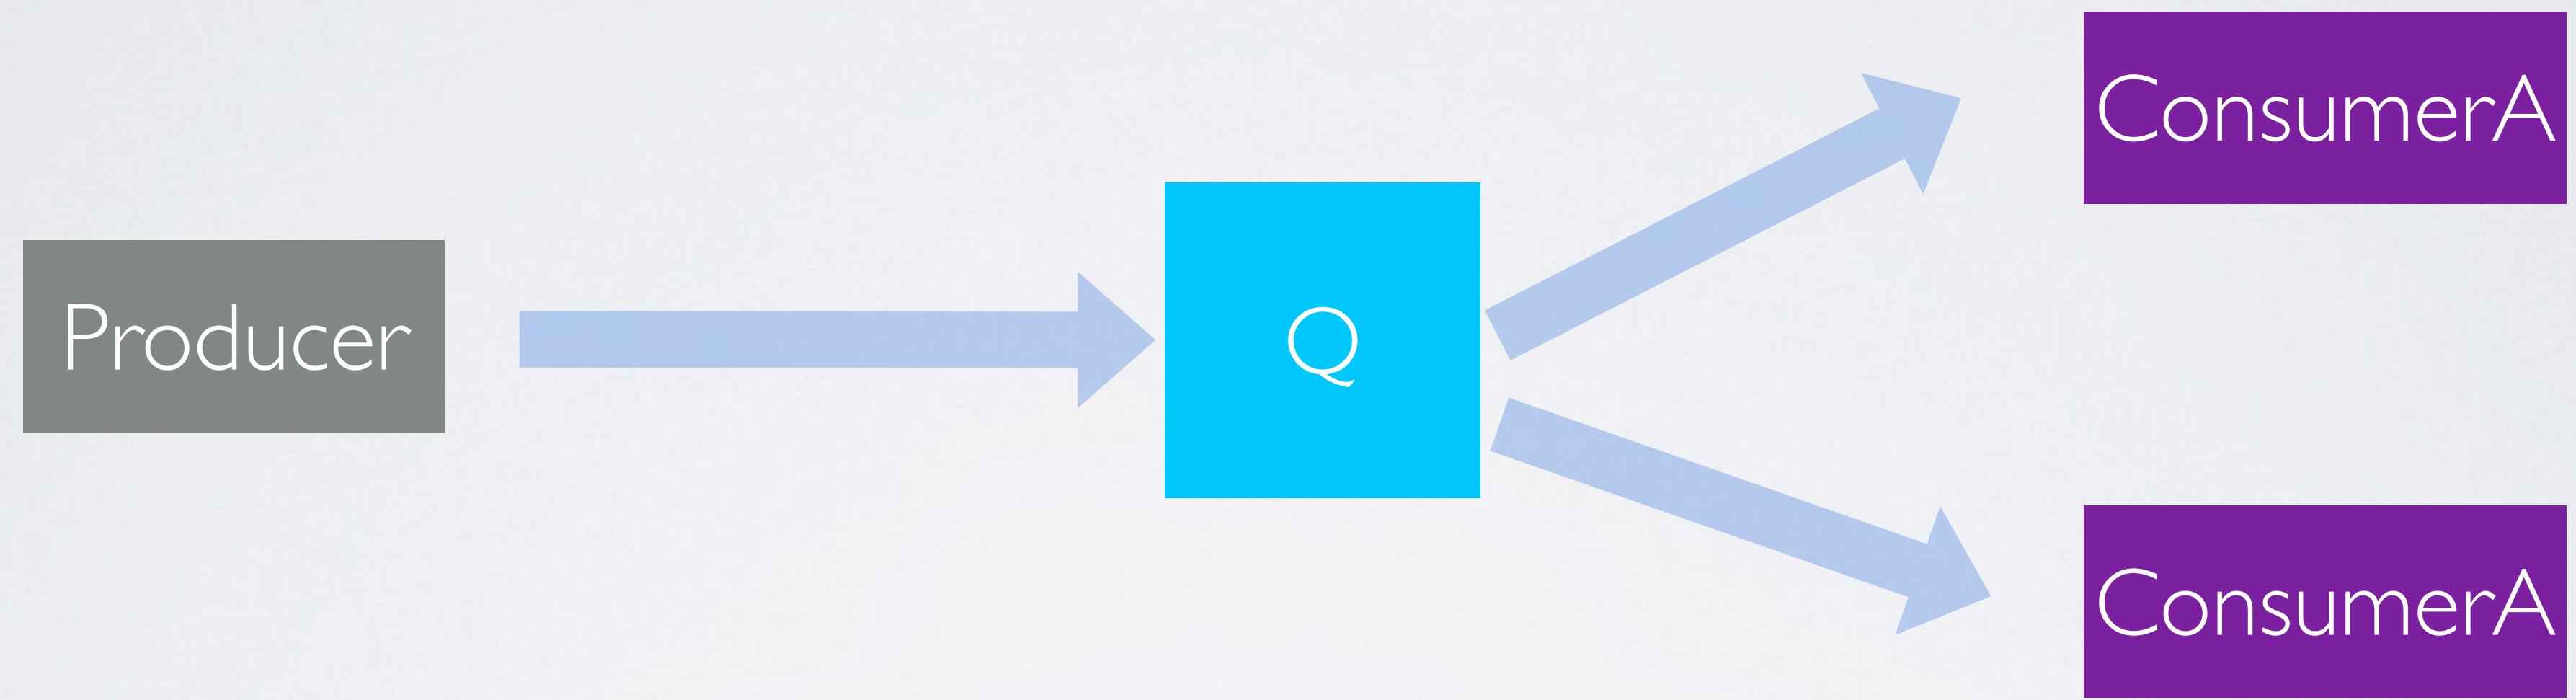
horizontal scalability

DISTRIBUTION



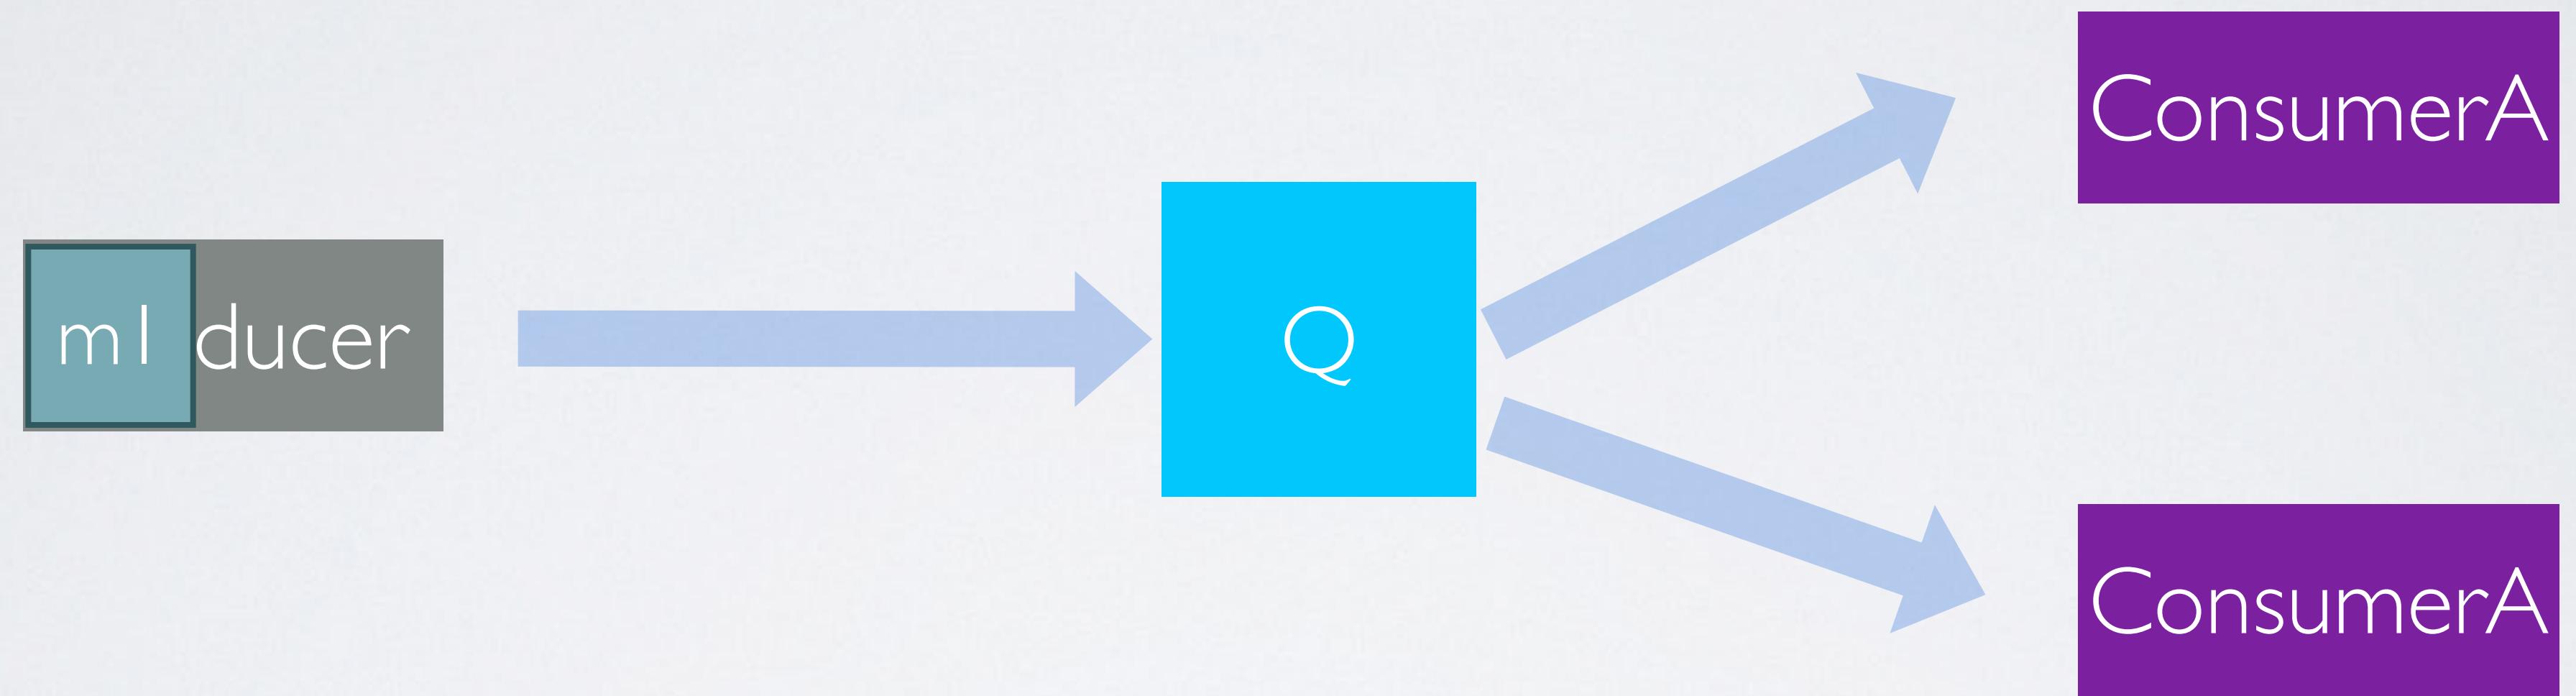
horizontal scalability

FAILURE



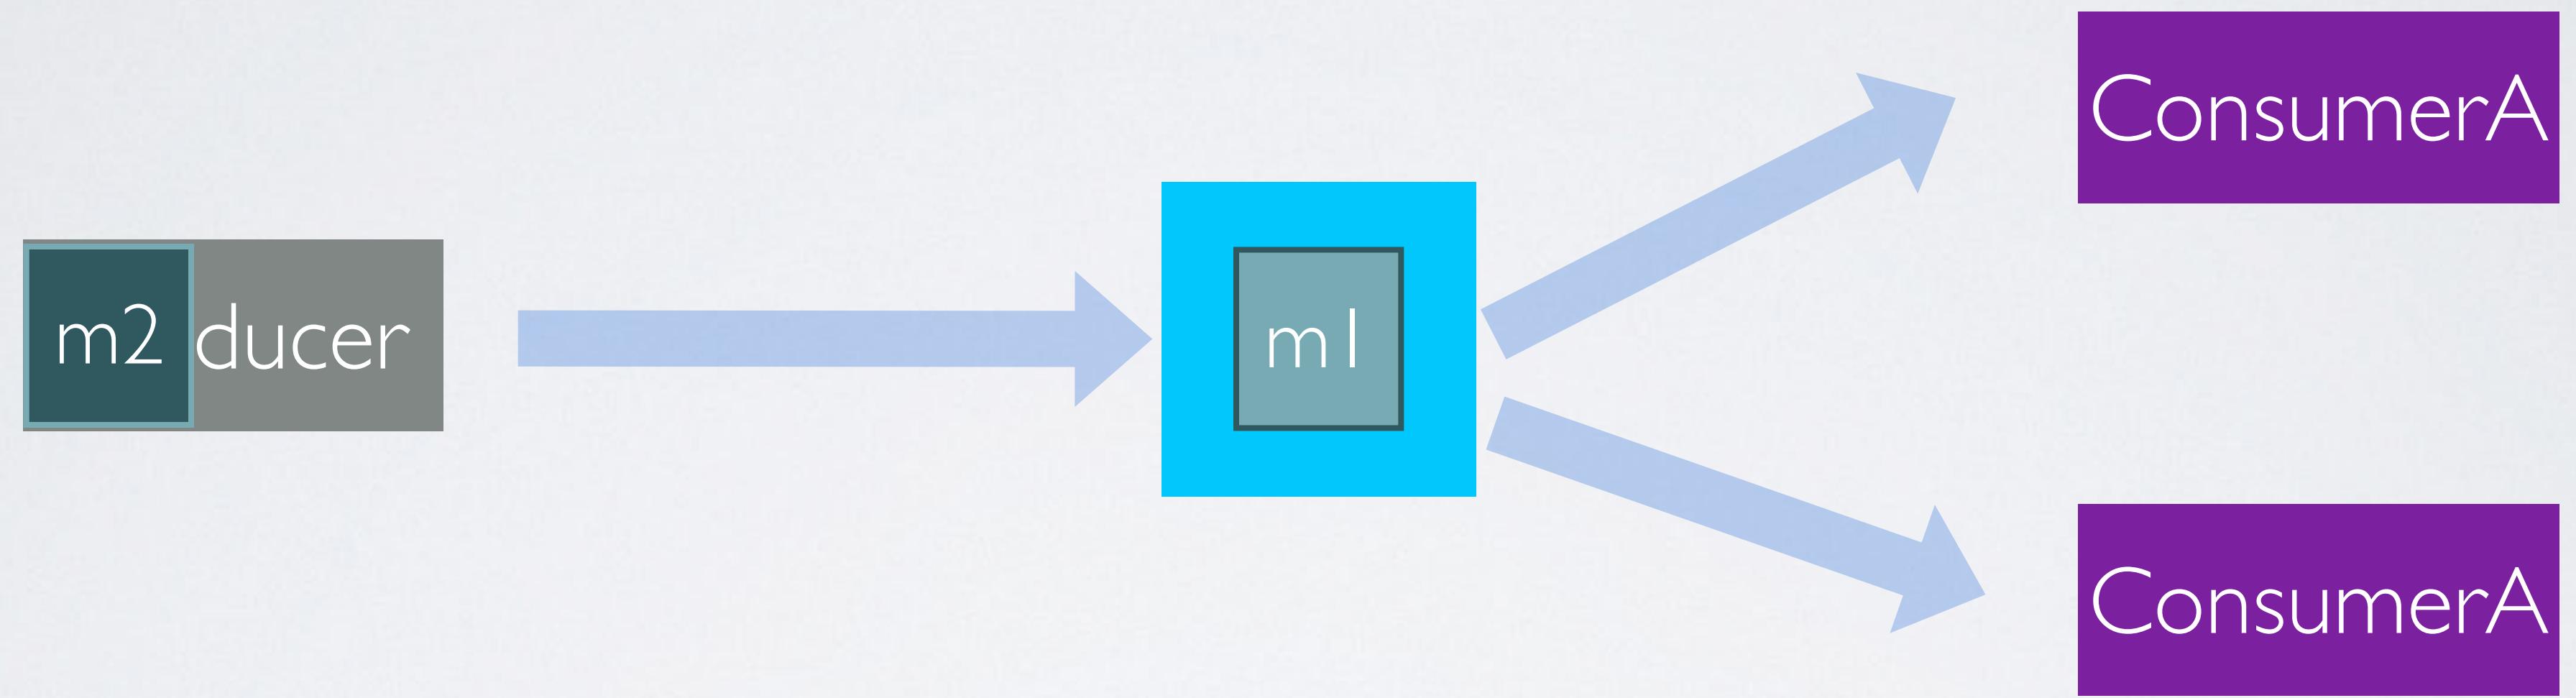
fault tolerance

FAILURE



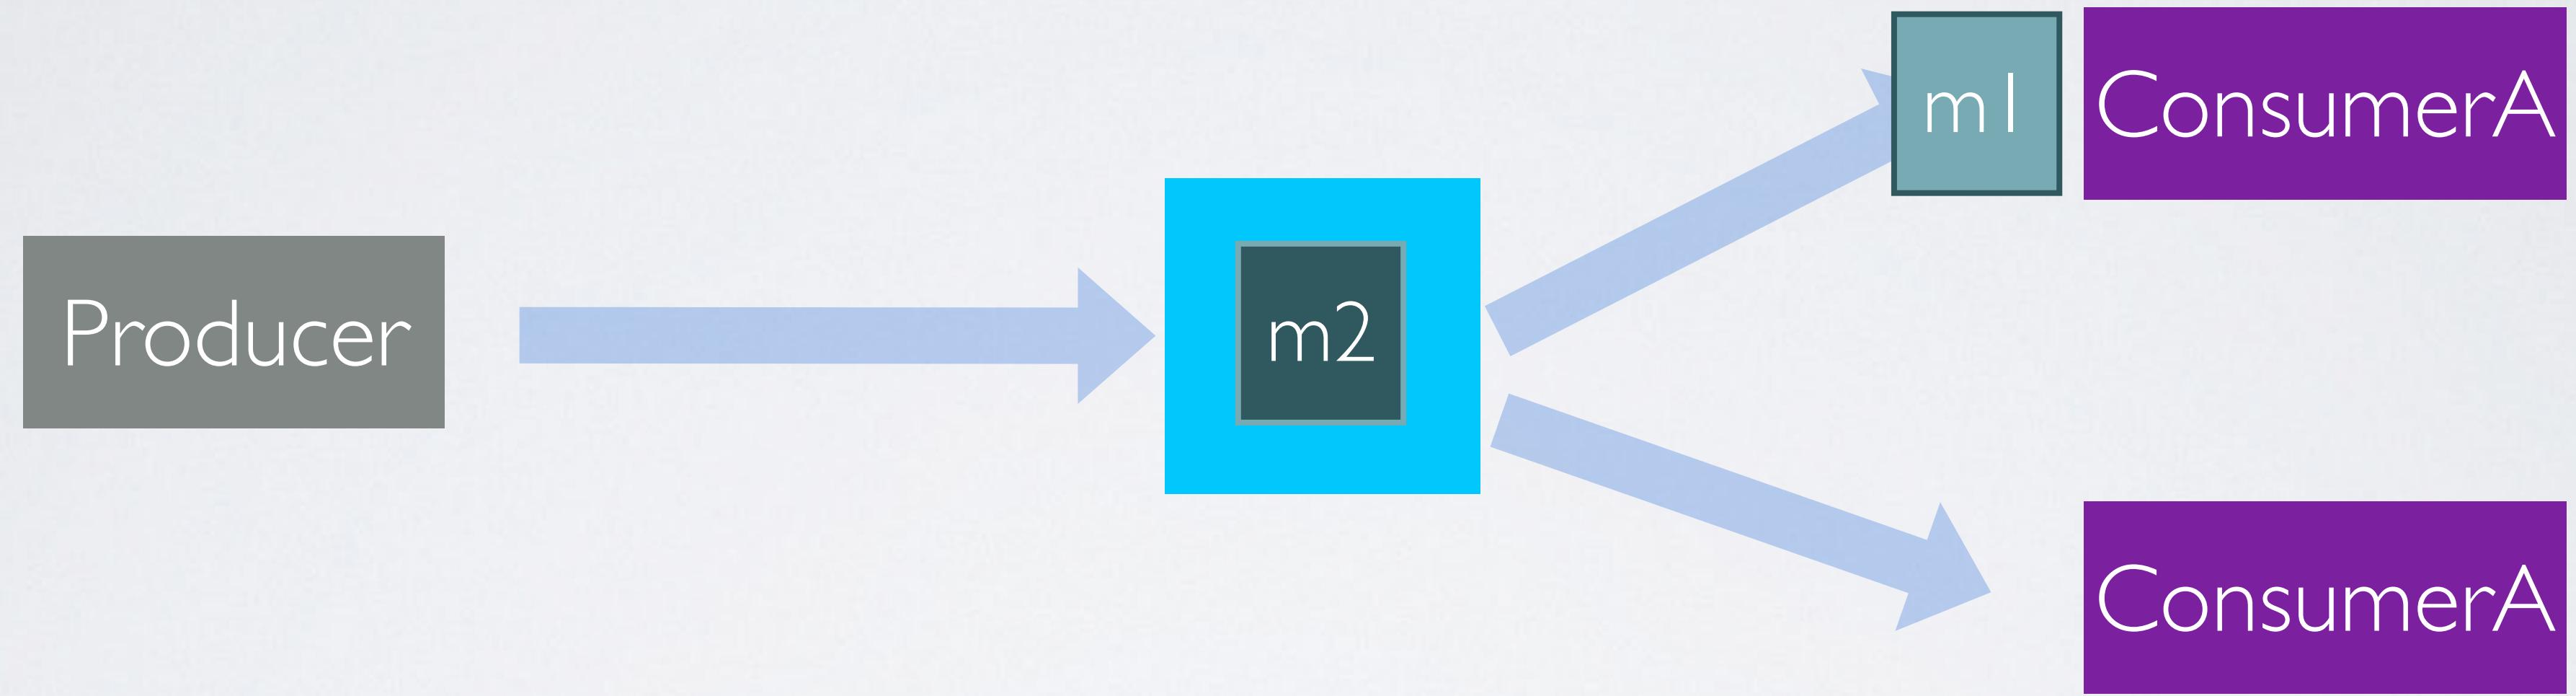
fault tolerance

FAILURE



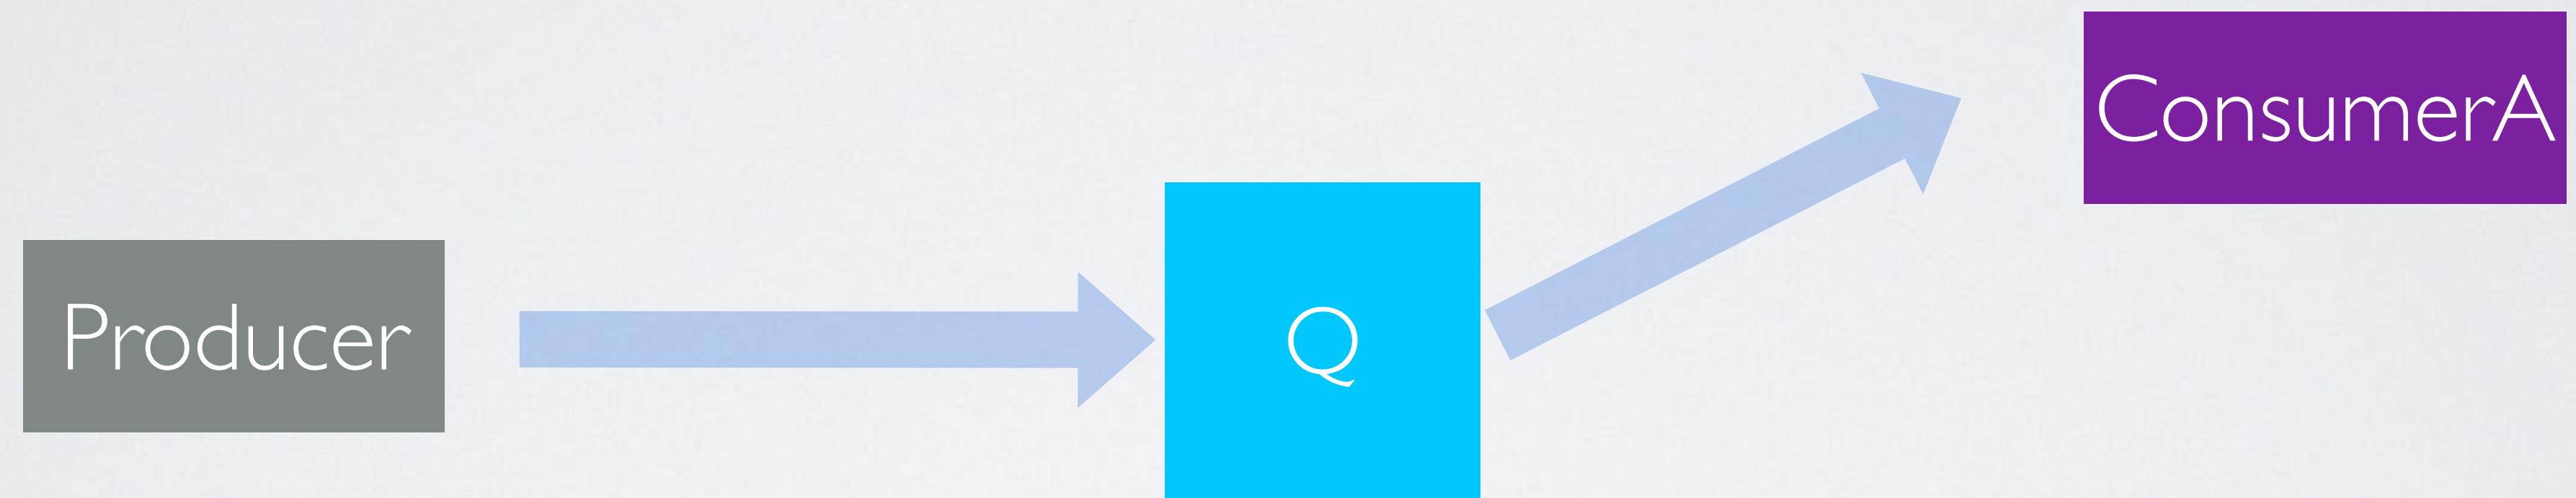
fault tolerance

FAILURE



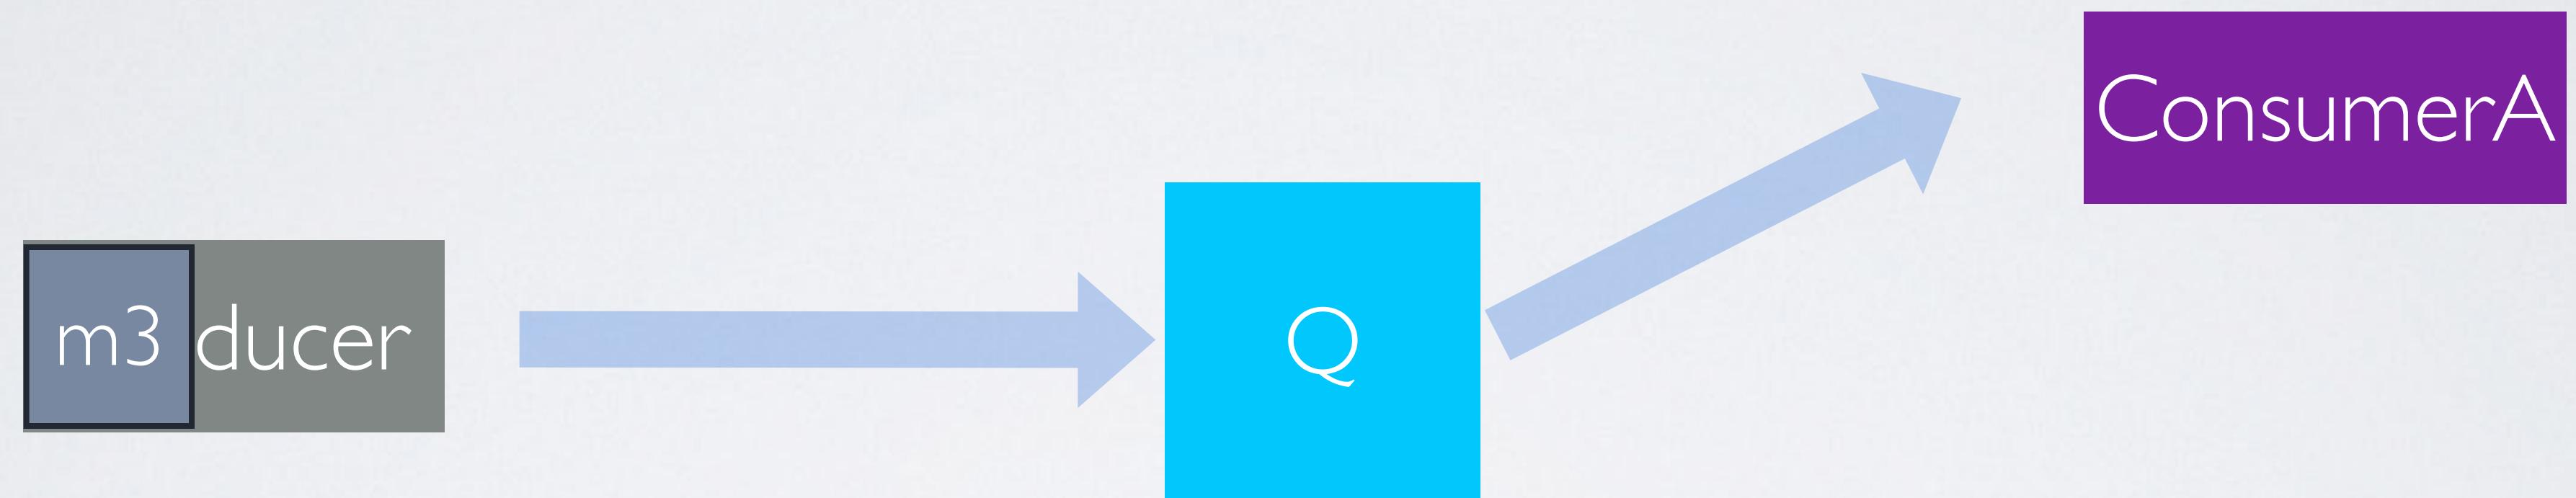
fault tolerance

FAILURE



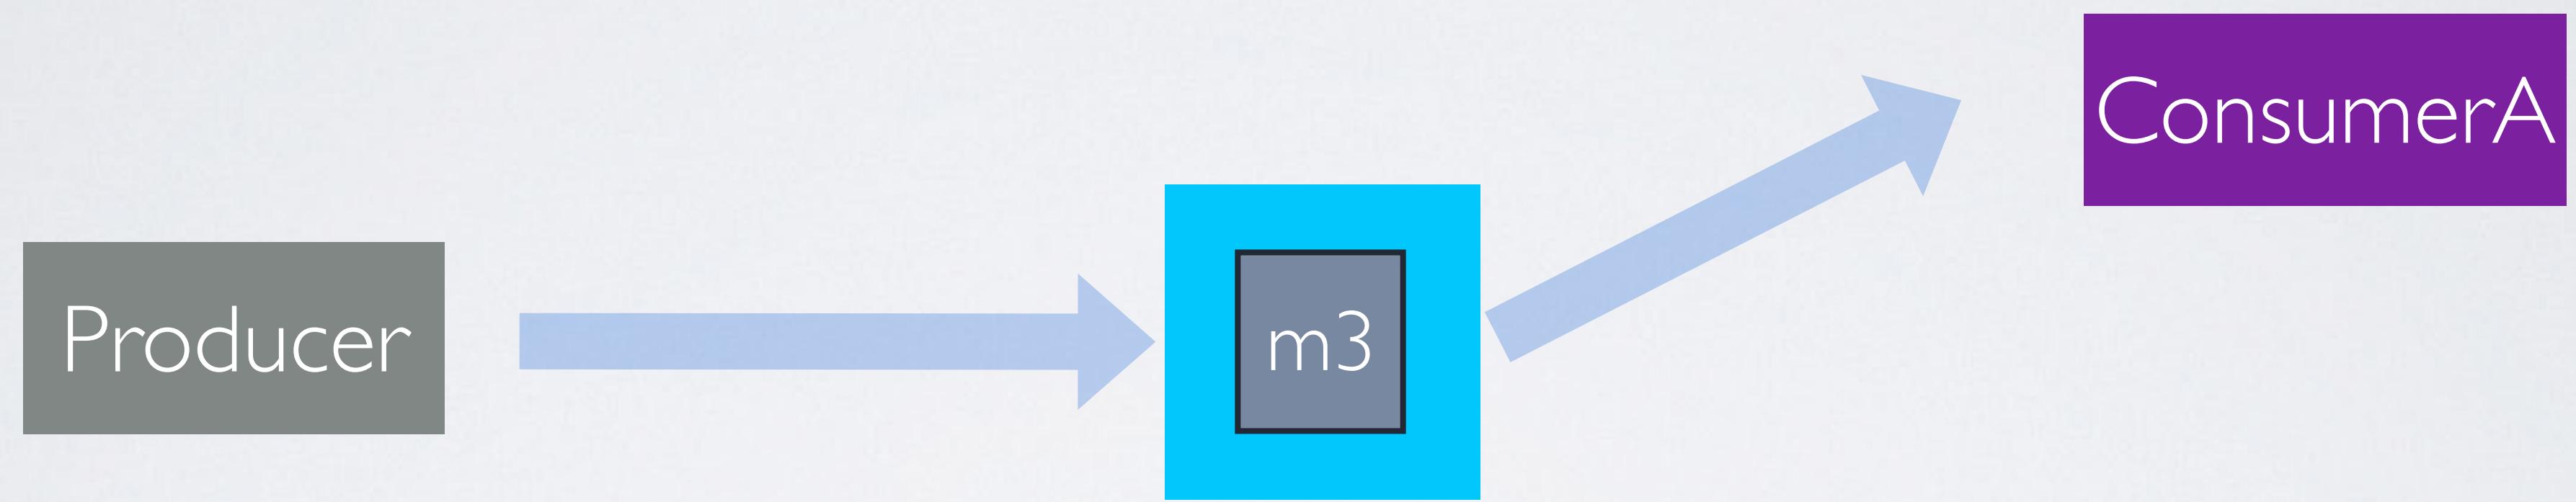
fault tolerance

FAILURE



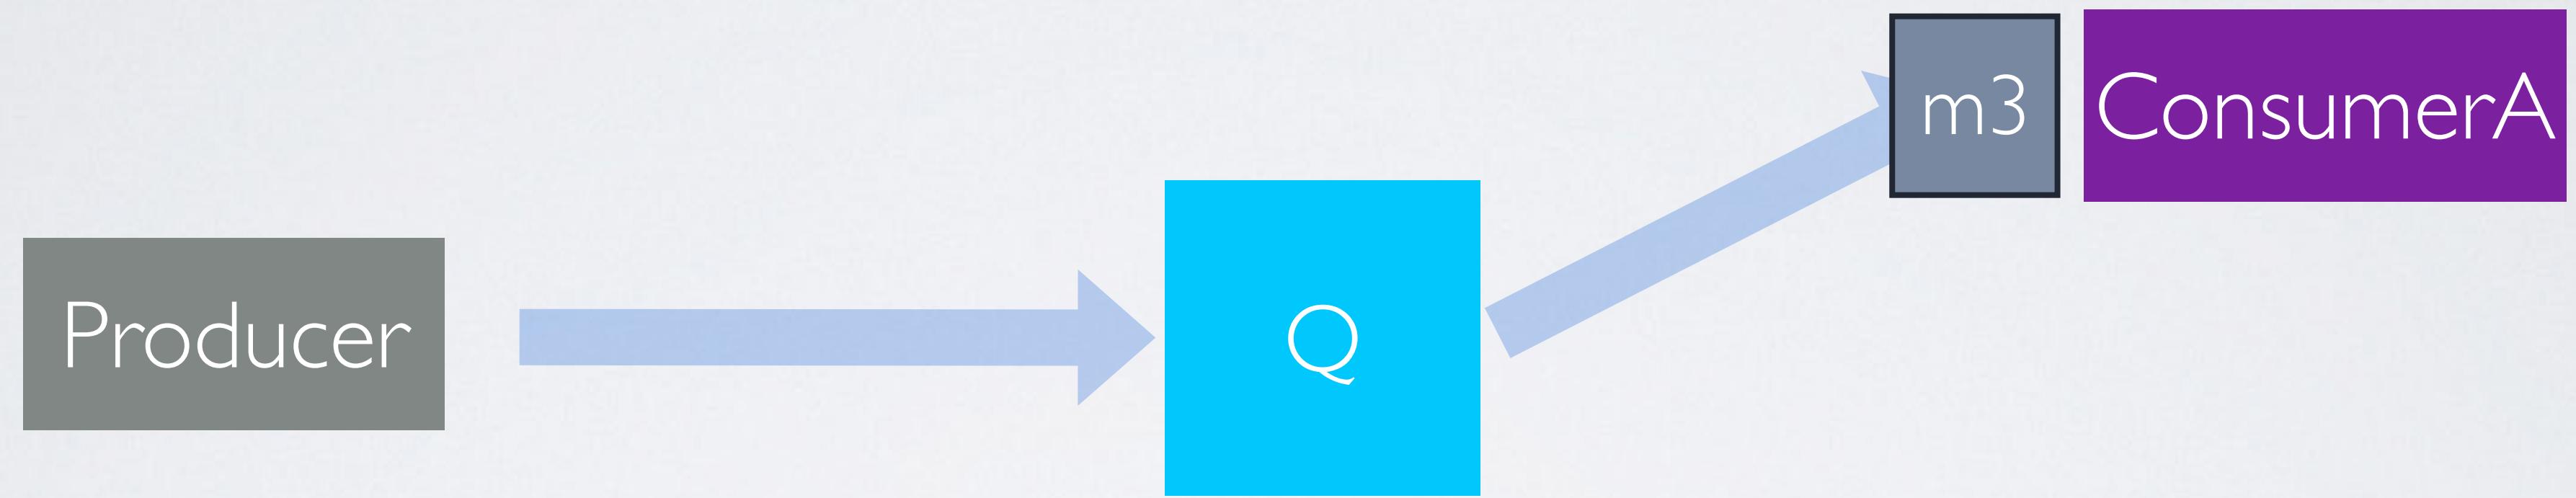
fault tolerance

FAILURE



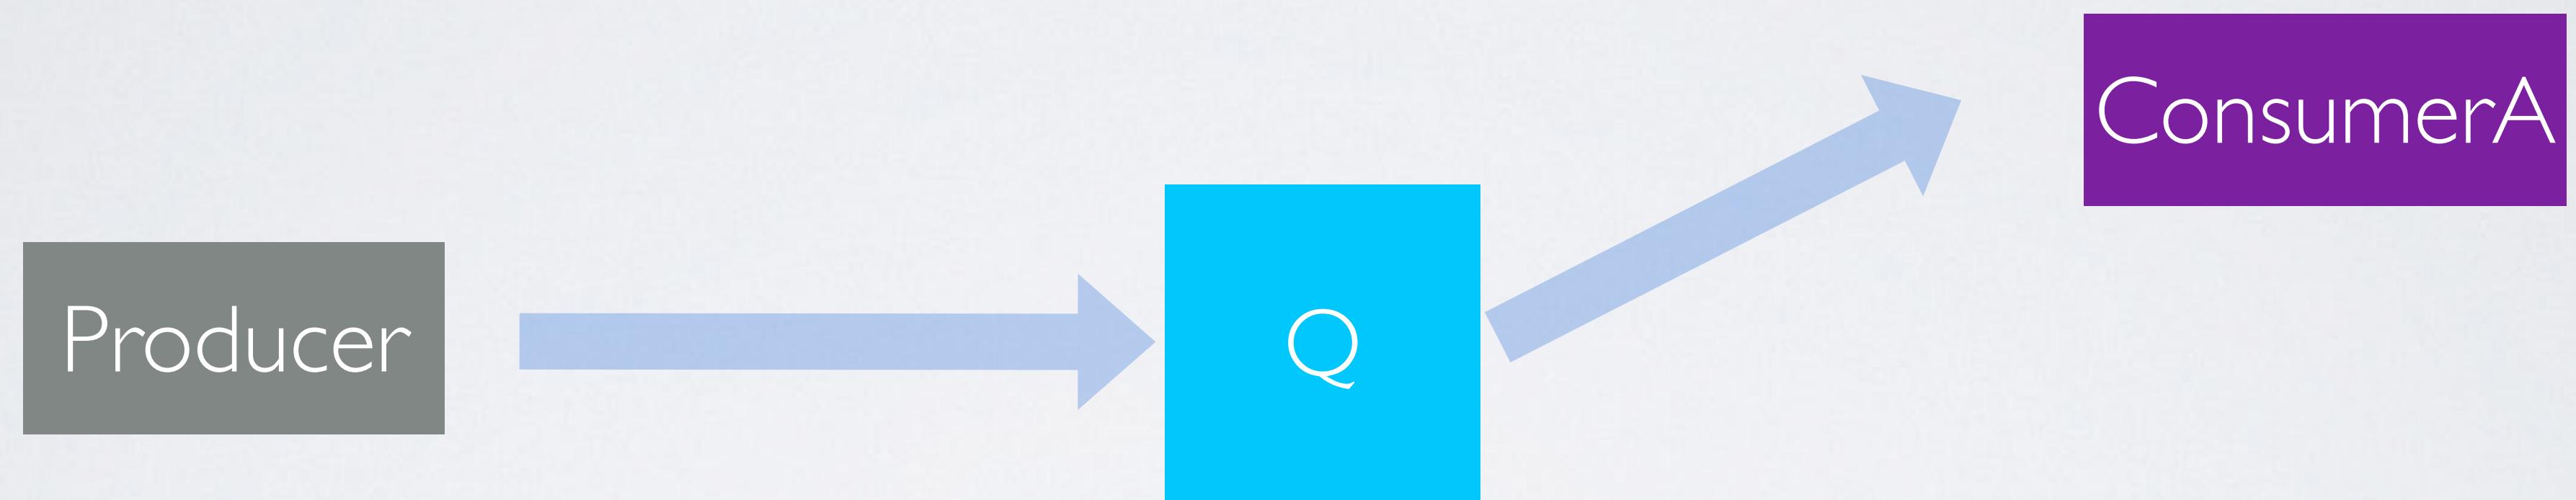
fault tolerance

FAILURE



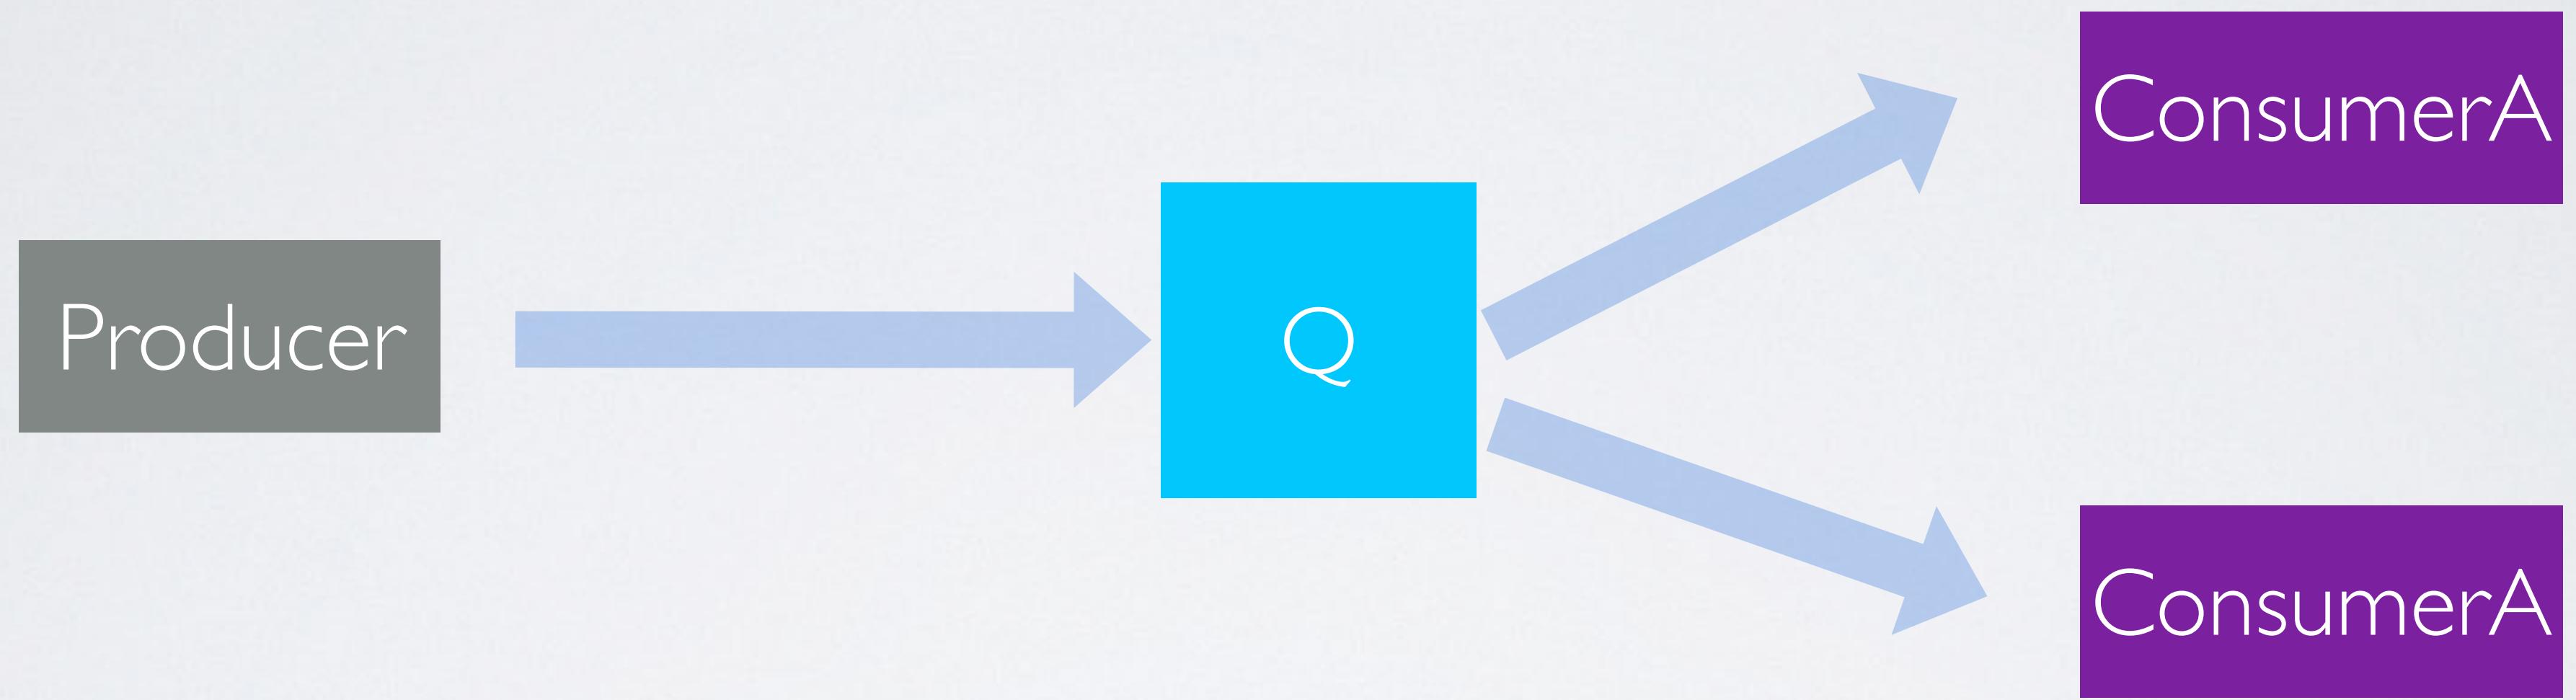
fault tolerance

FAILURE



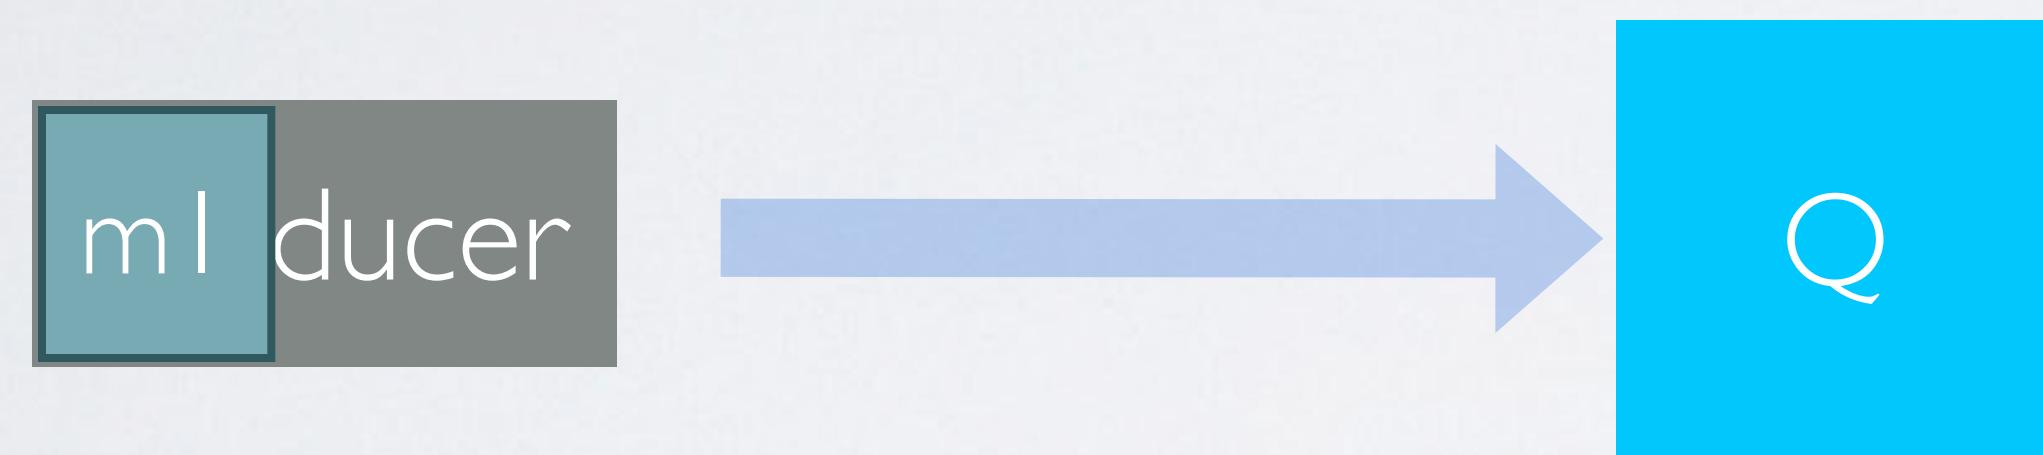
fault tolerance

EVEN MORE FAILURE



welp

EVEN MORE FAILURE



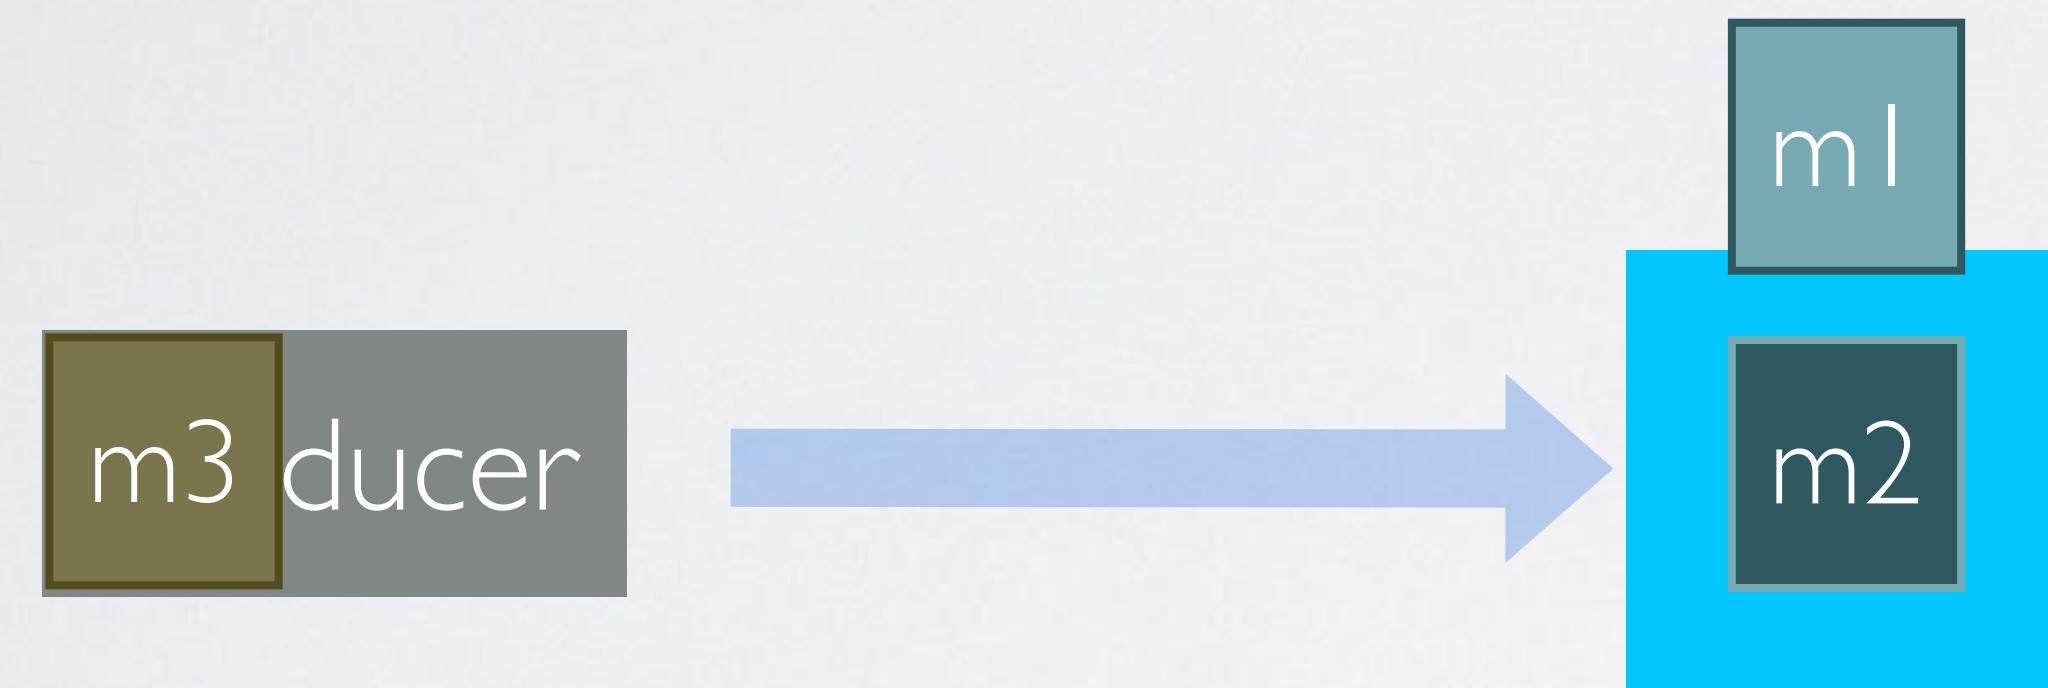
welp

EVEN MORE FAILURE



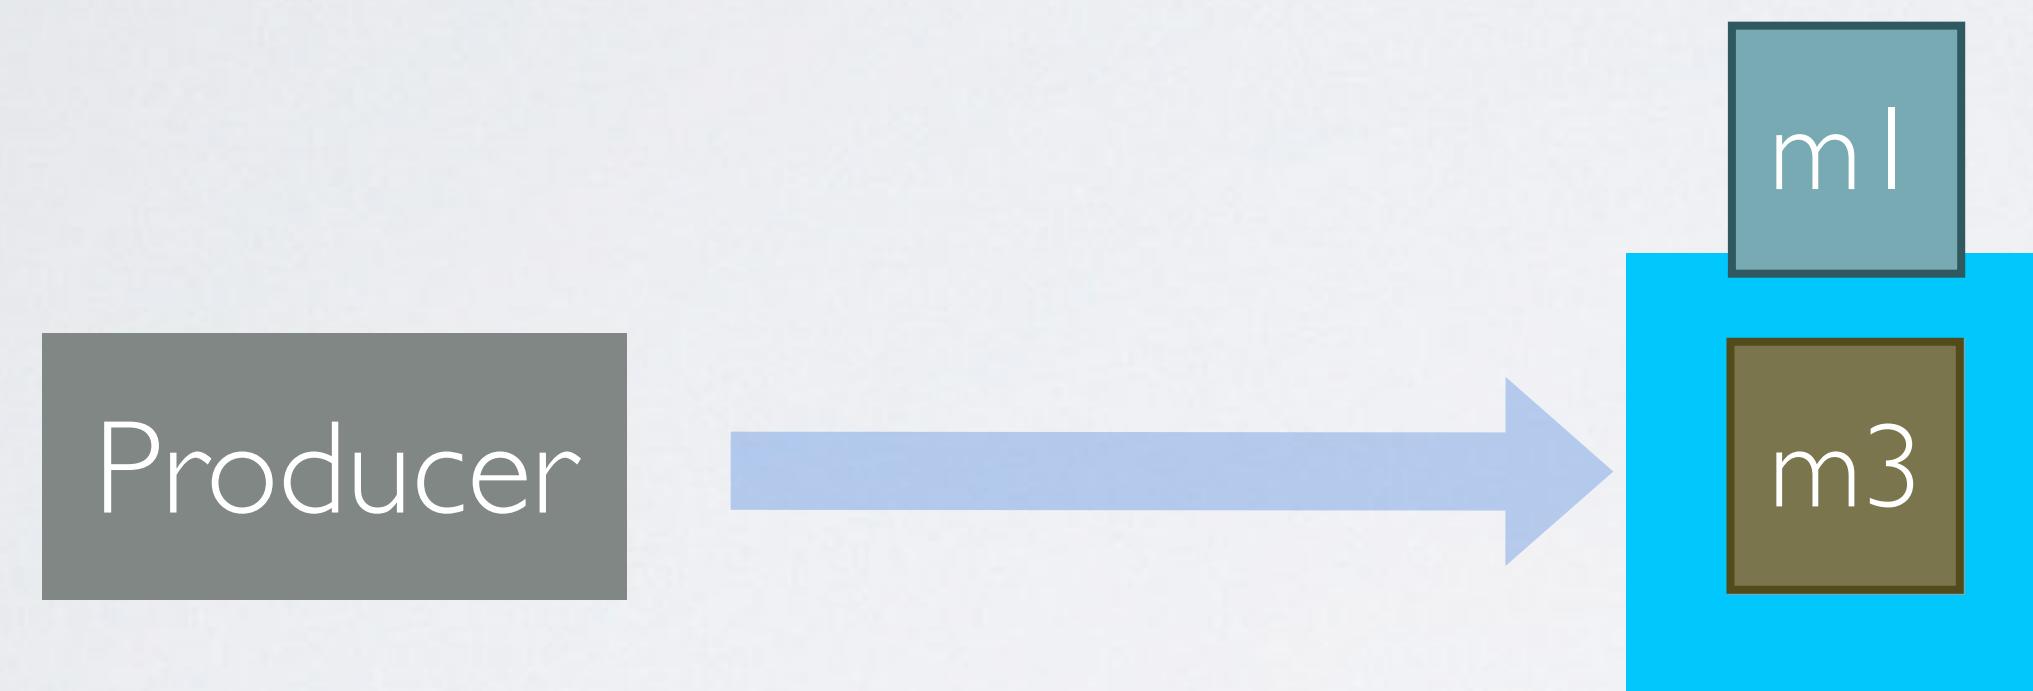
welp

EVEN MORE FAILURE



welp

EVEN MORE FAILURE



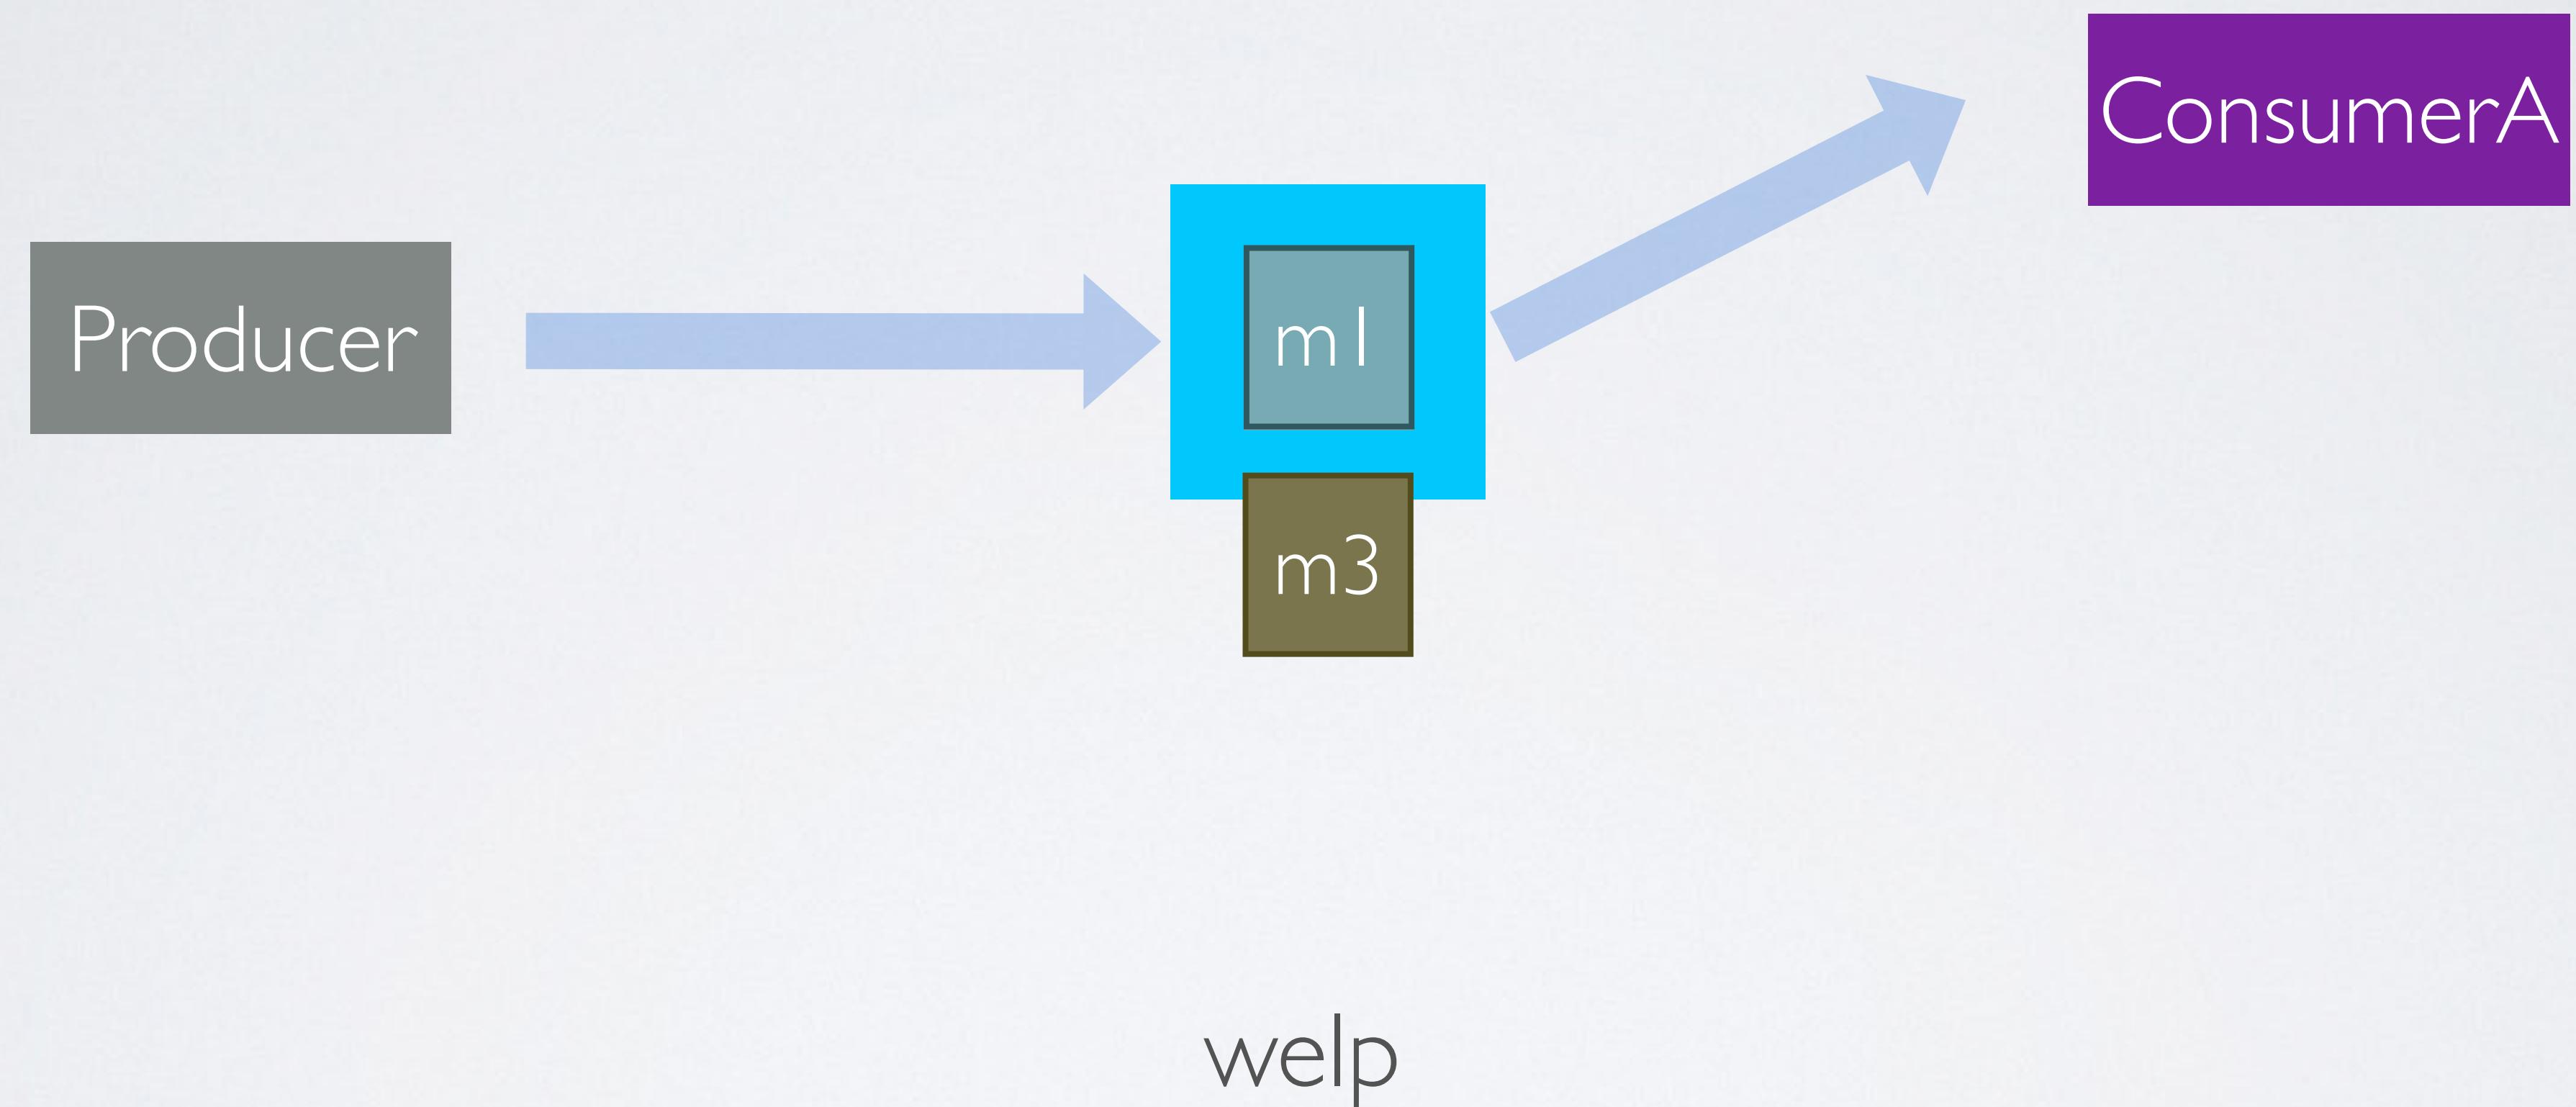
welp

EVEN MORE FAILURE

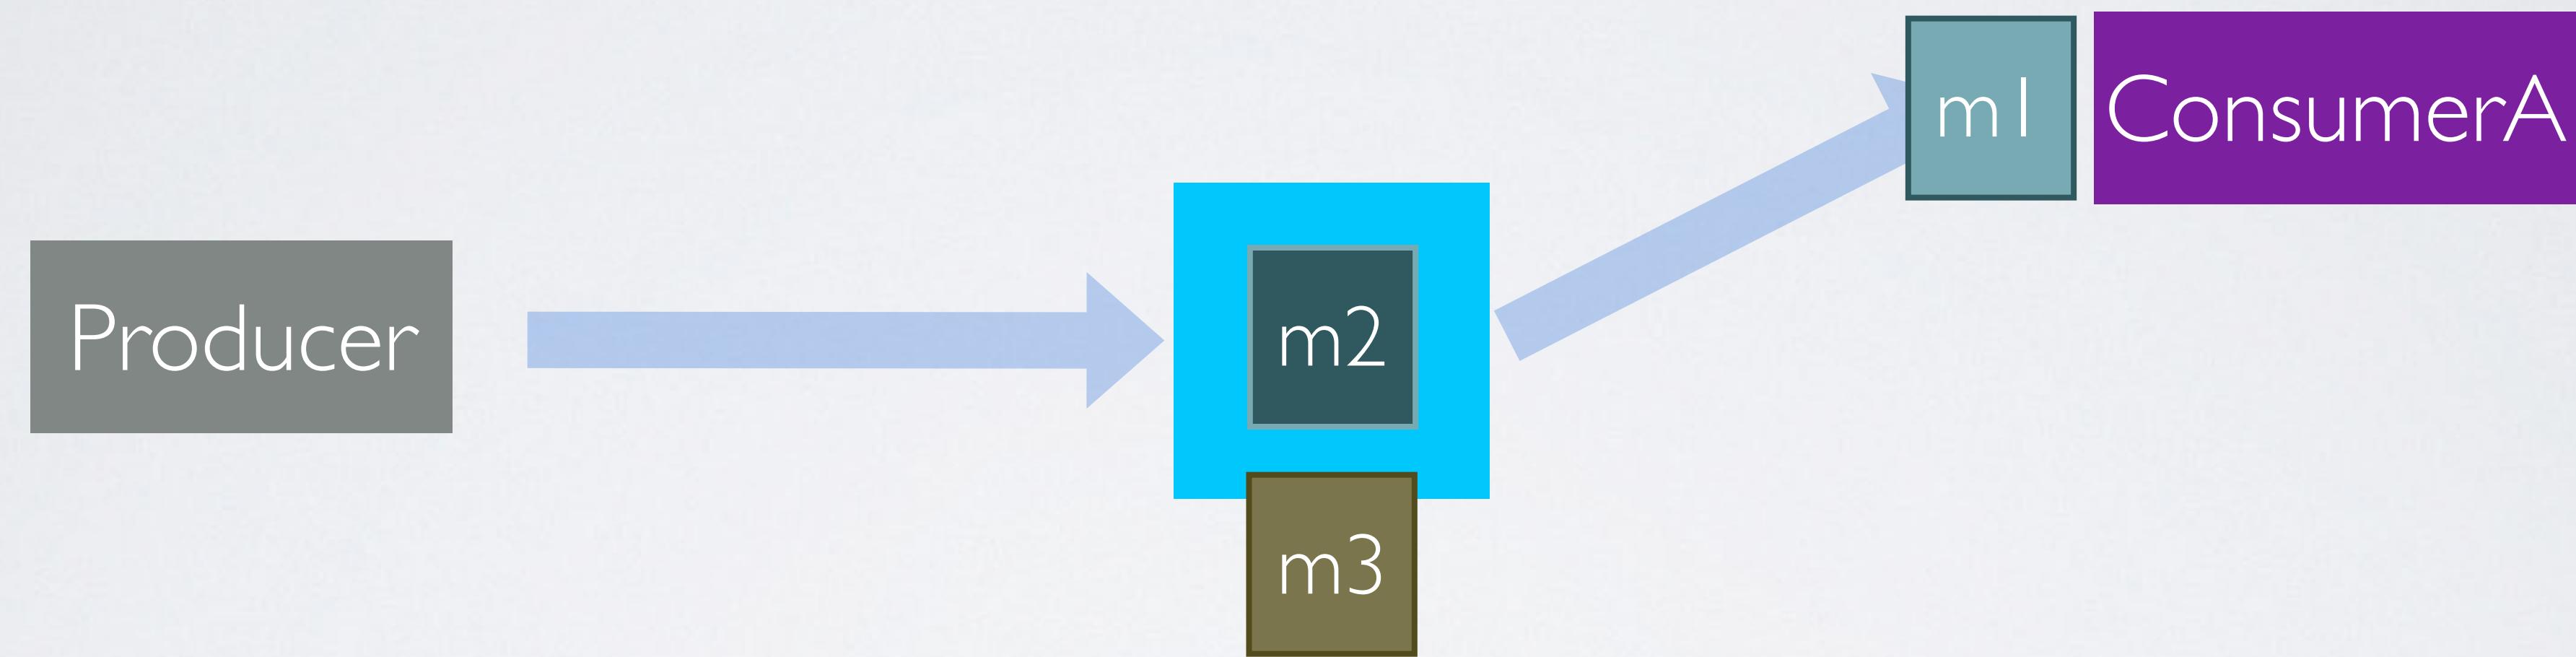


welp

EVEN MORE FAILURE

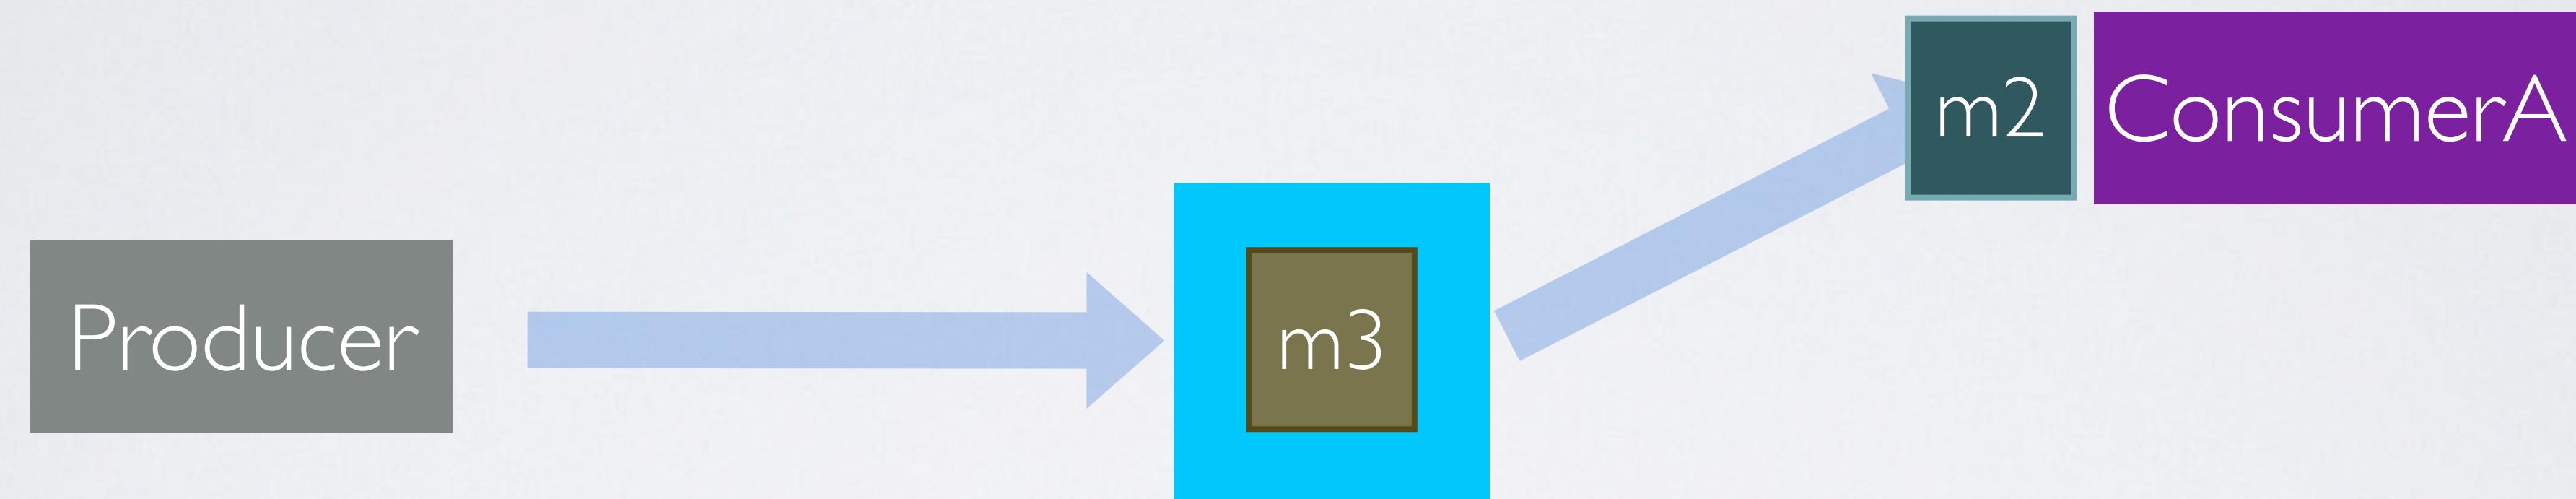


EVEN MORE FAILURE



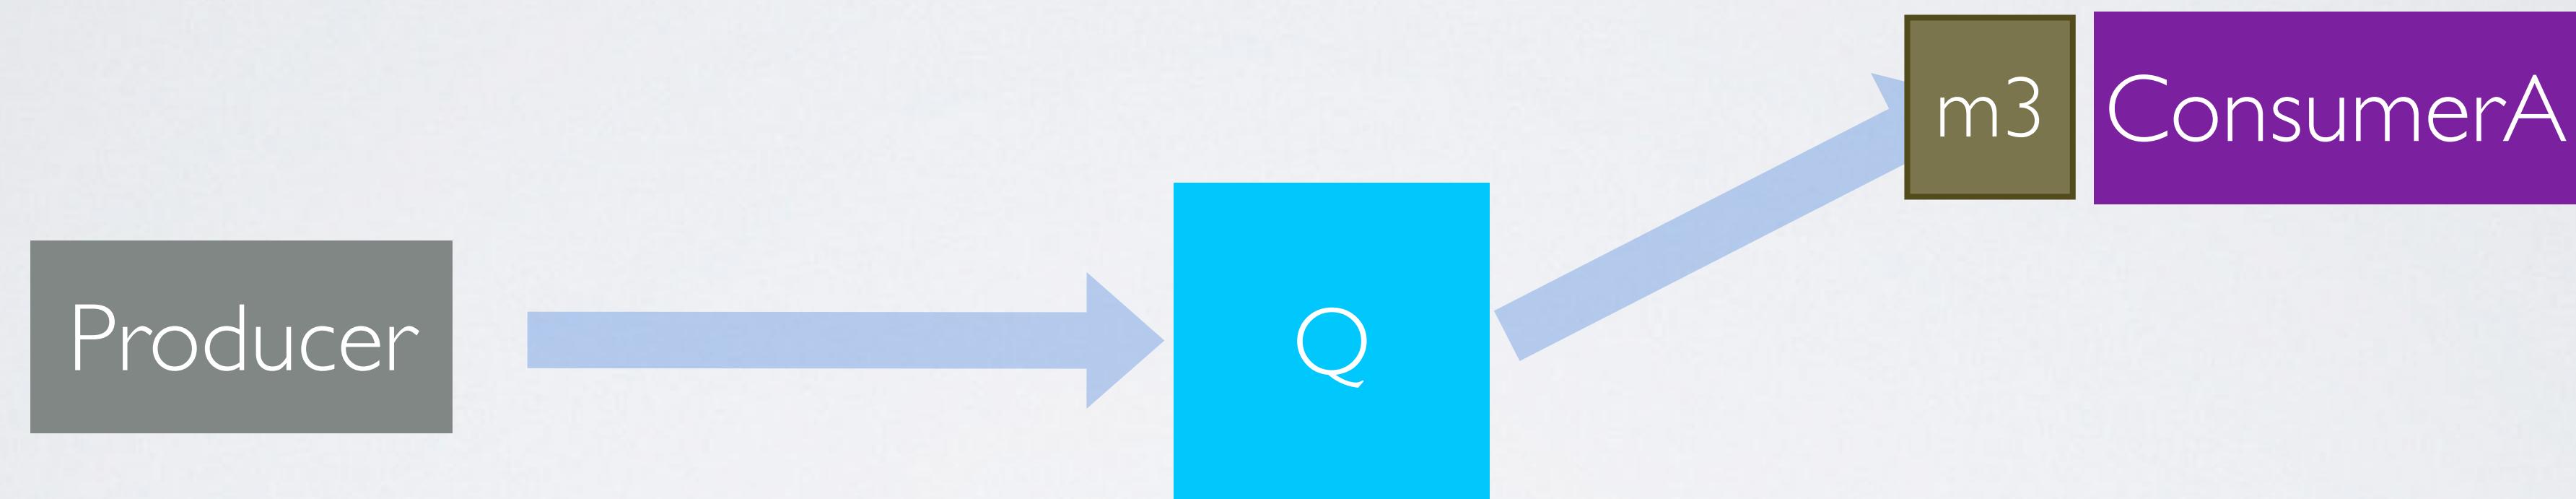
welp

EVEN MORE FAILURE



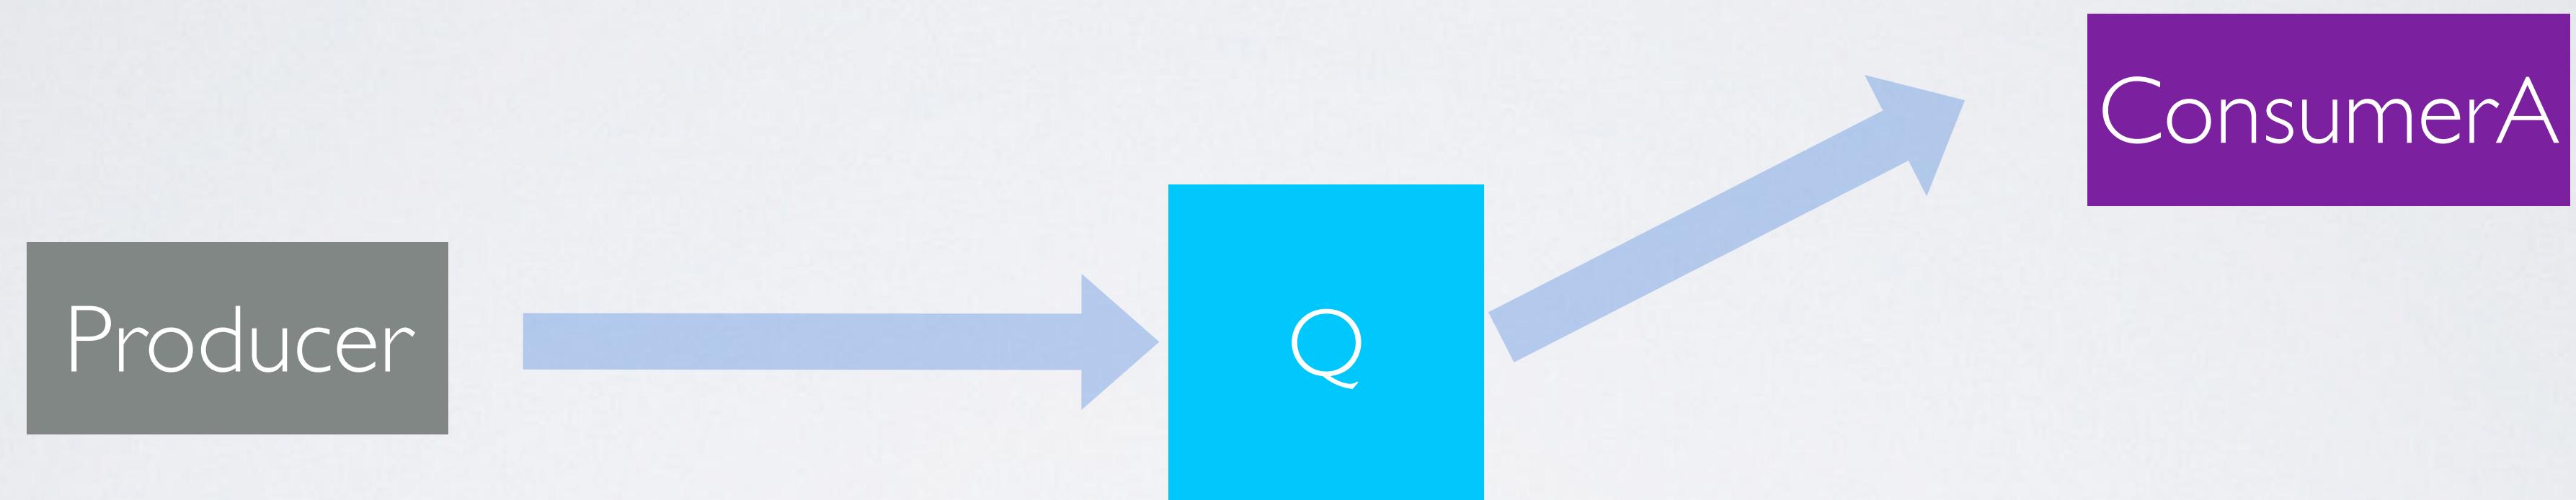
welp

EVEN MORE FAILURE



welp

EVEN MORE FAILURE



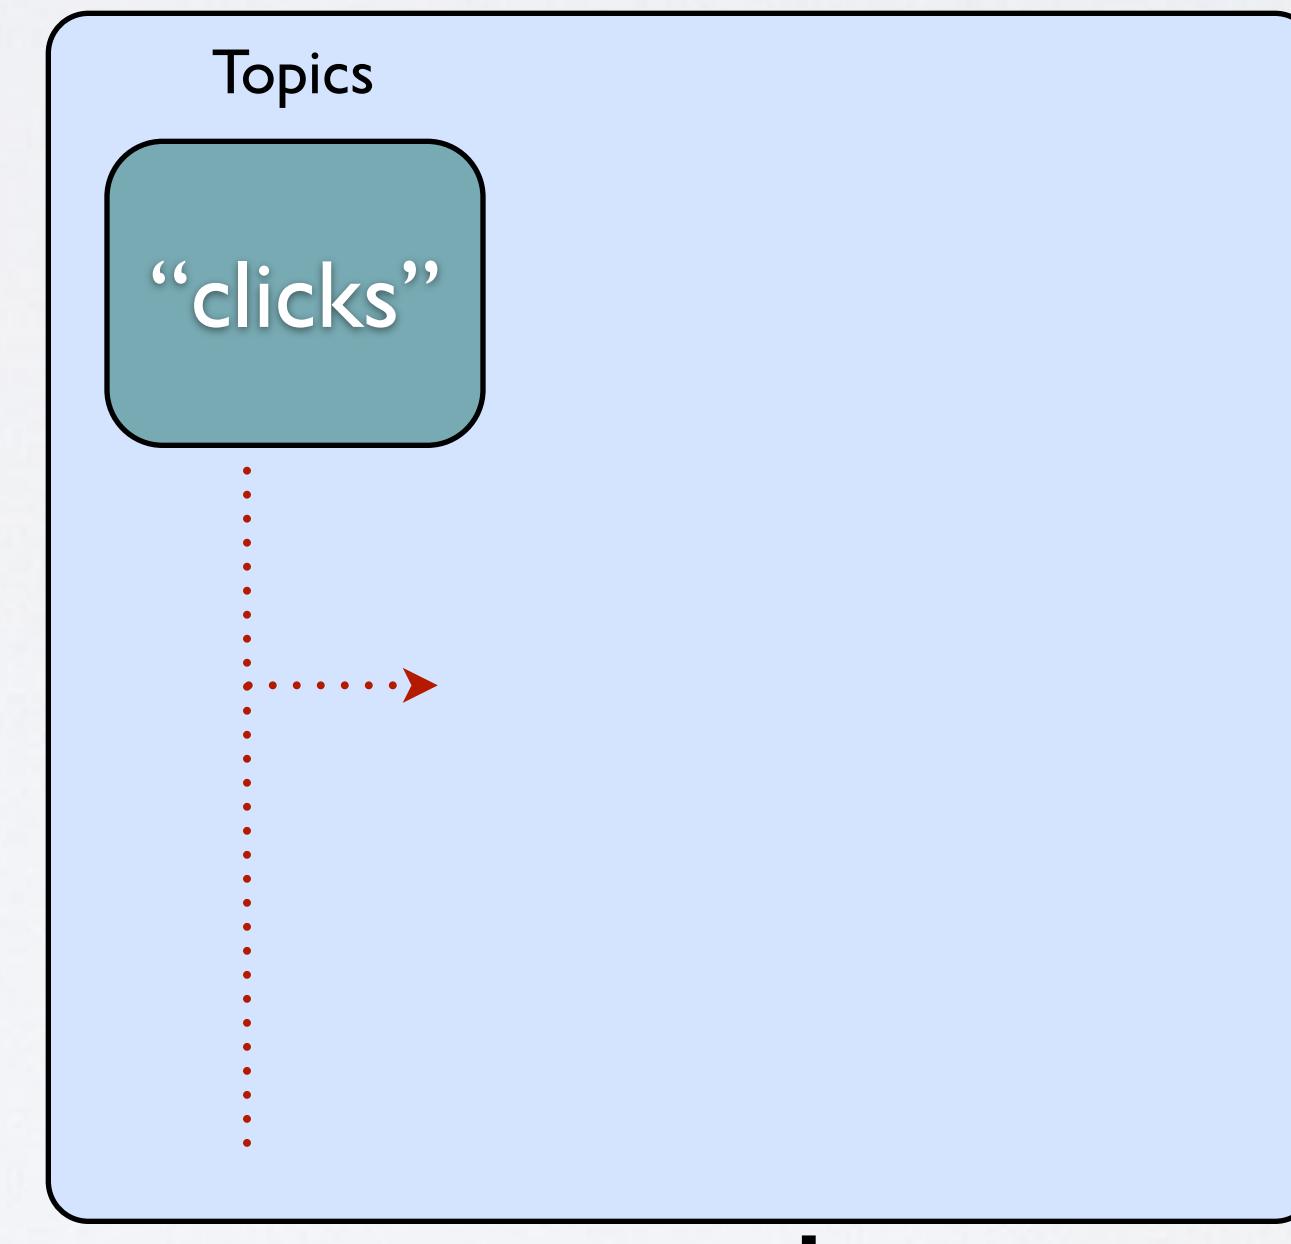
welp

NSQD

TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

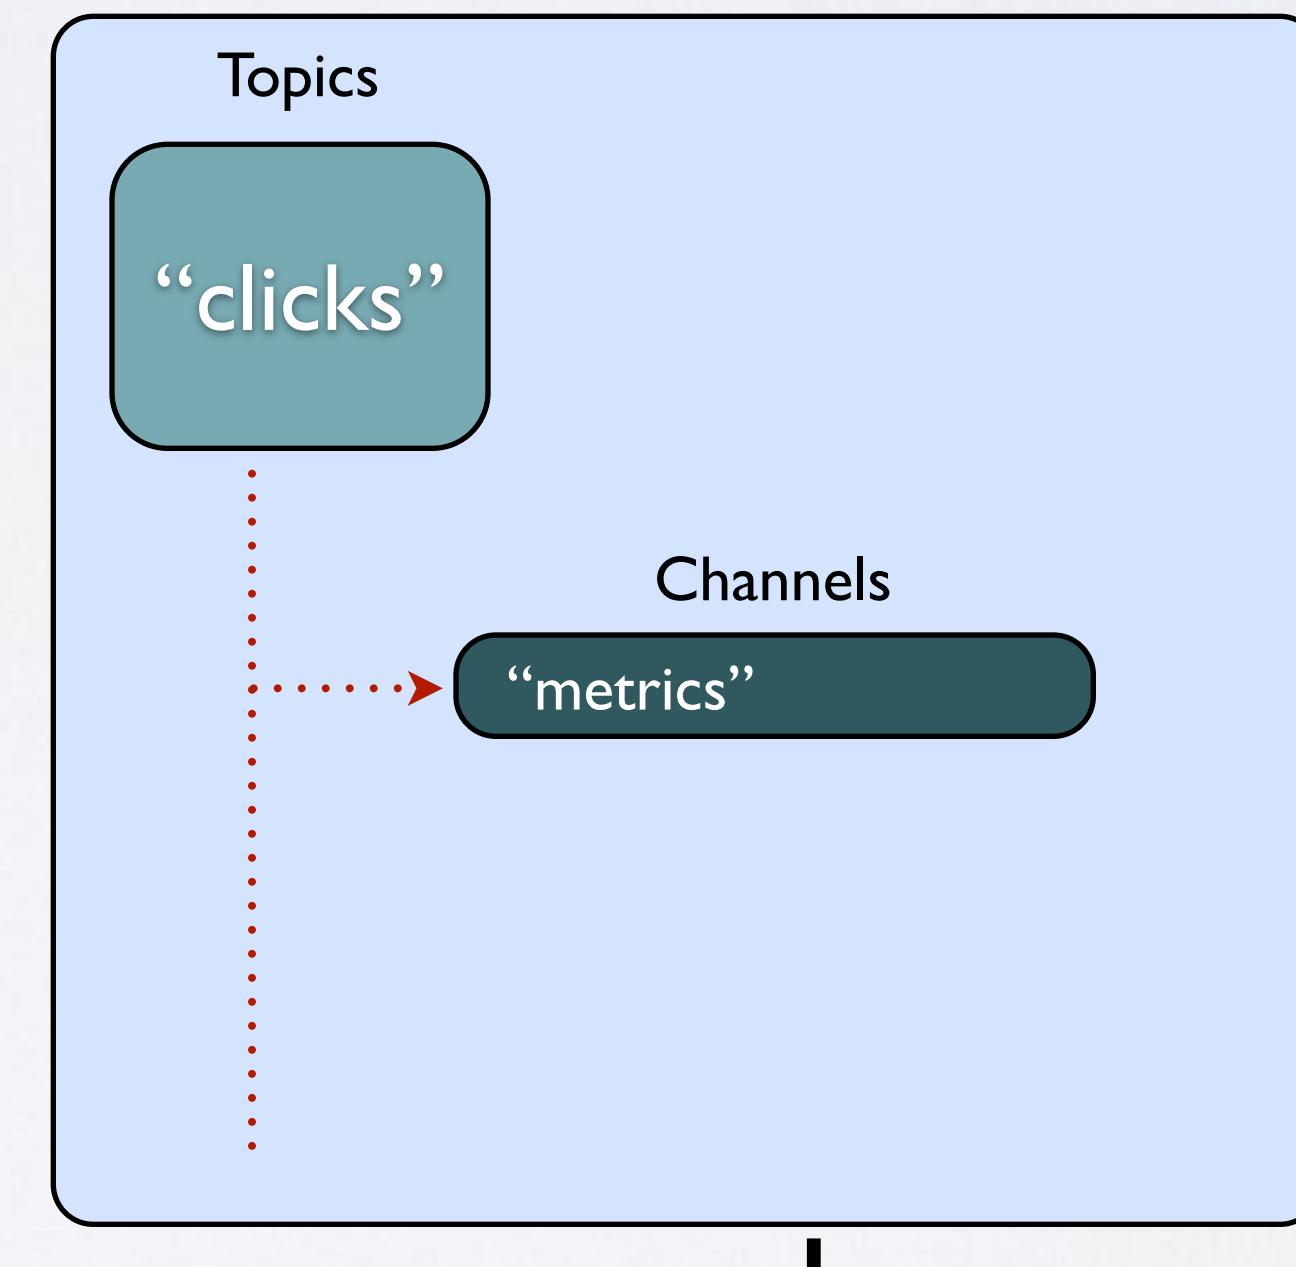
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

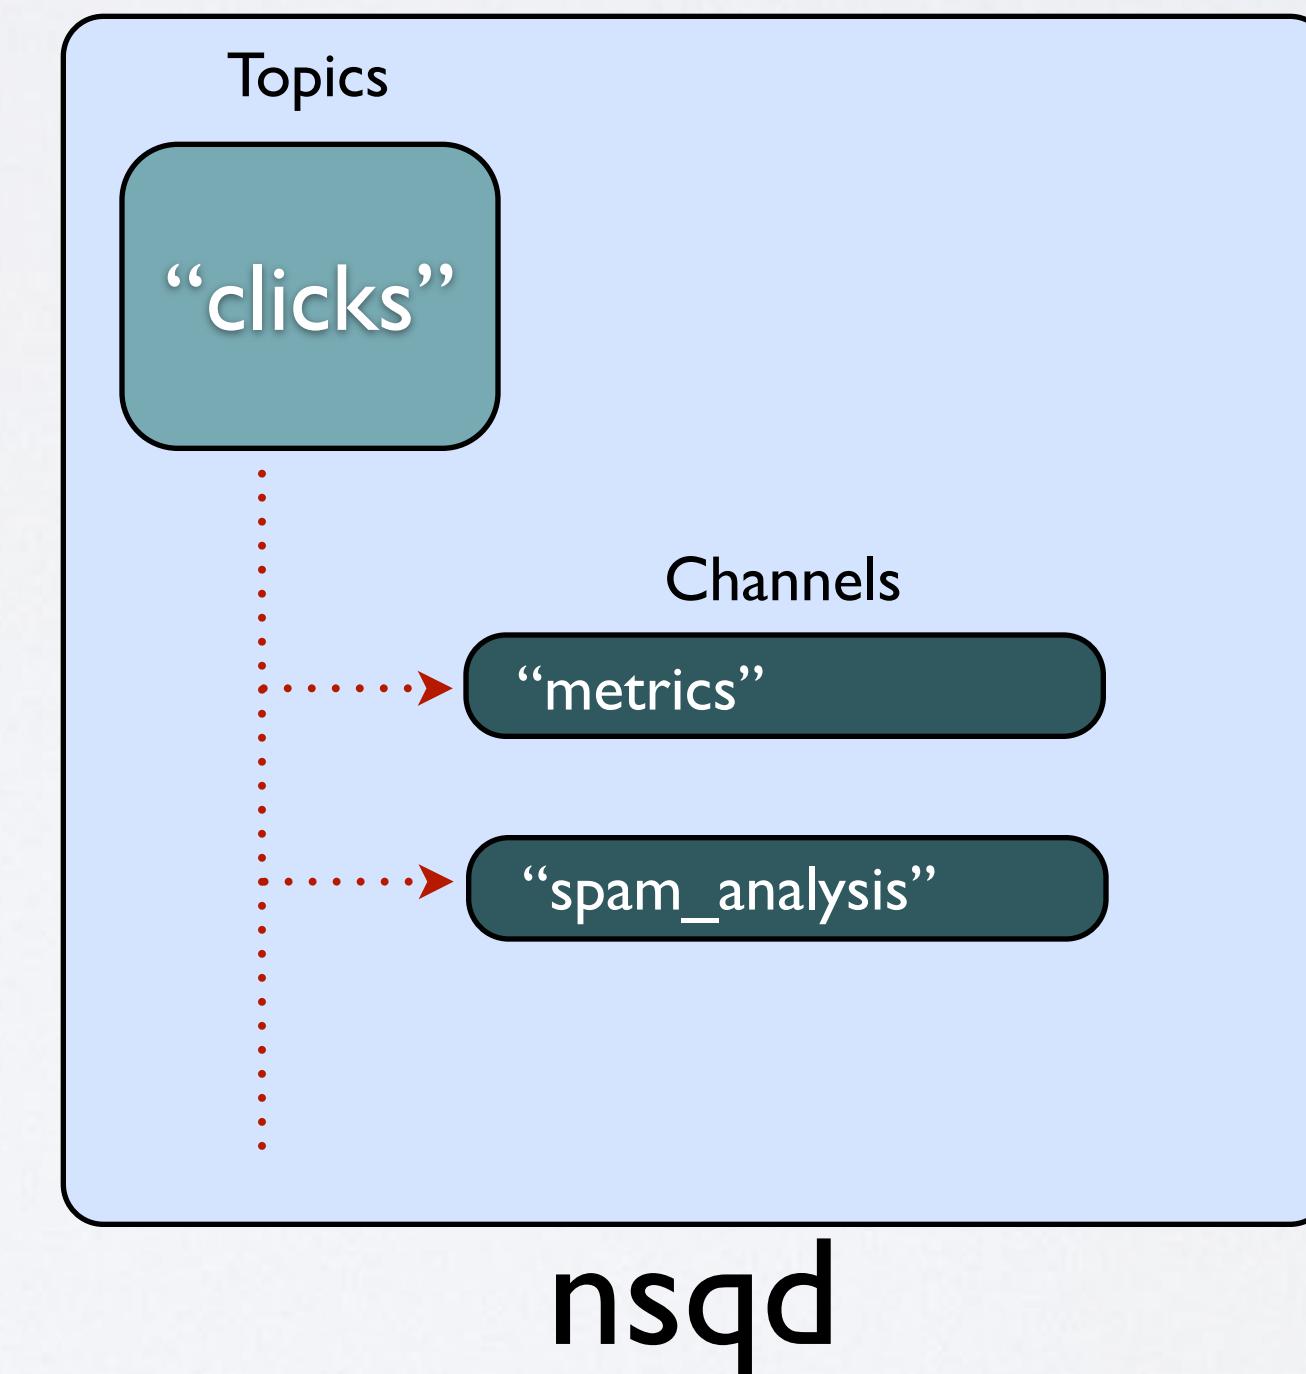
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

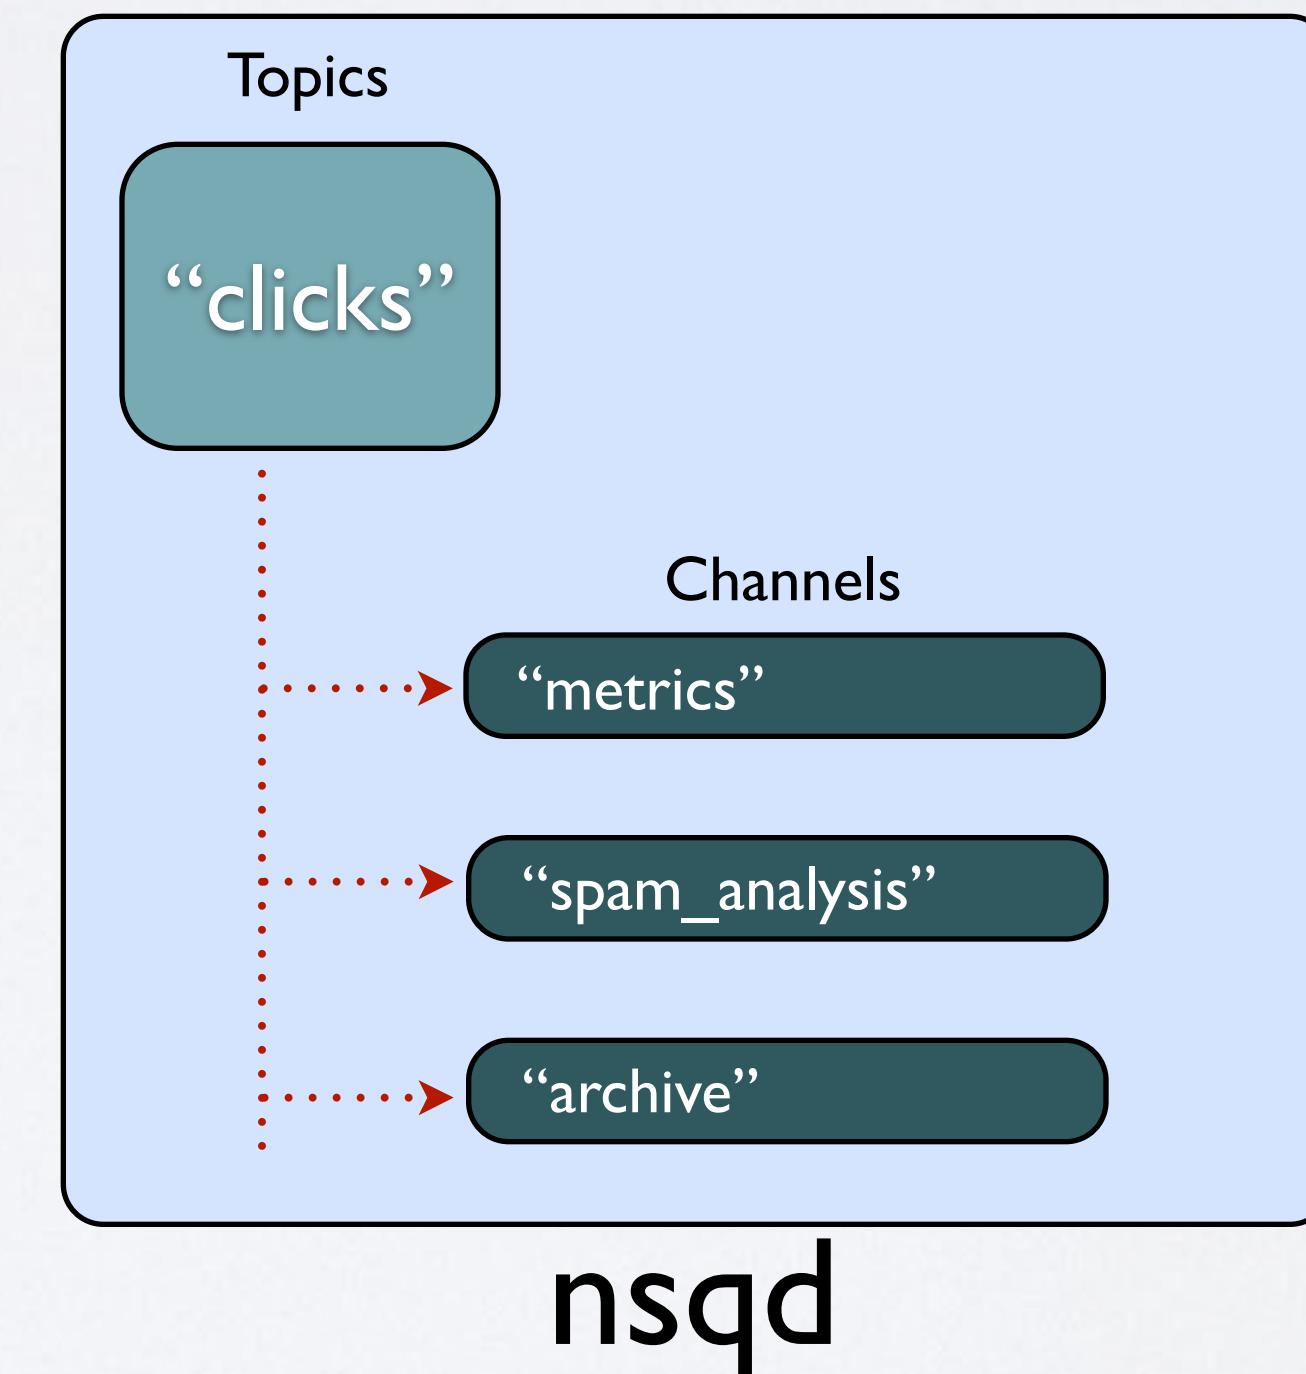
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

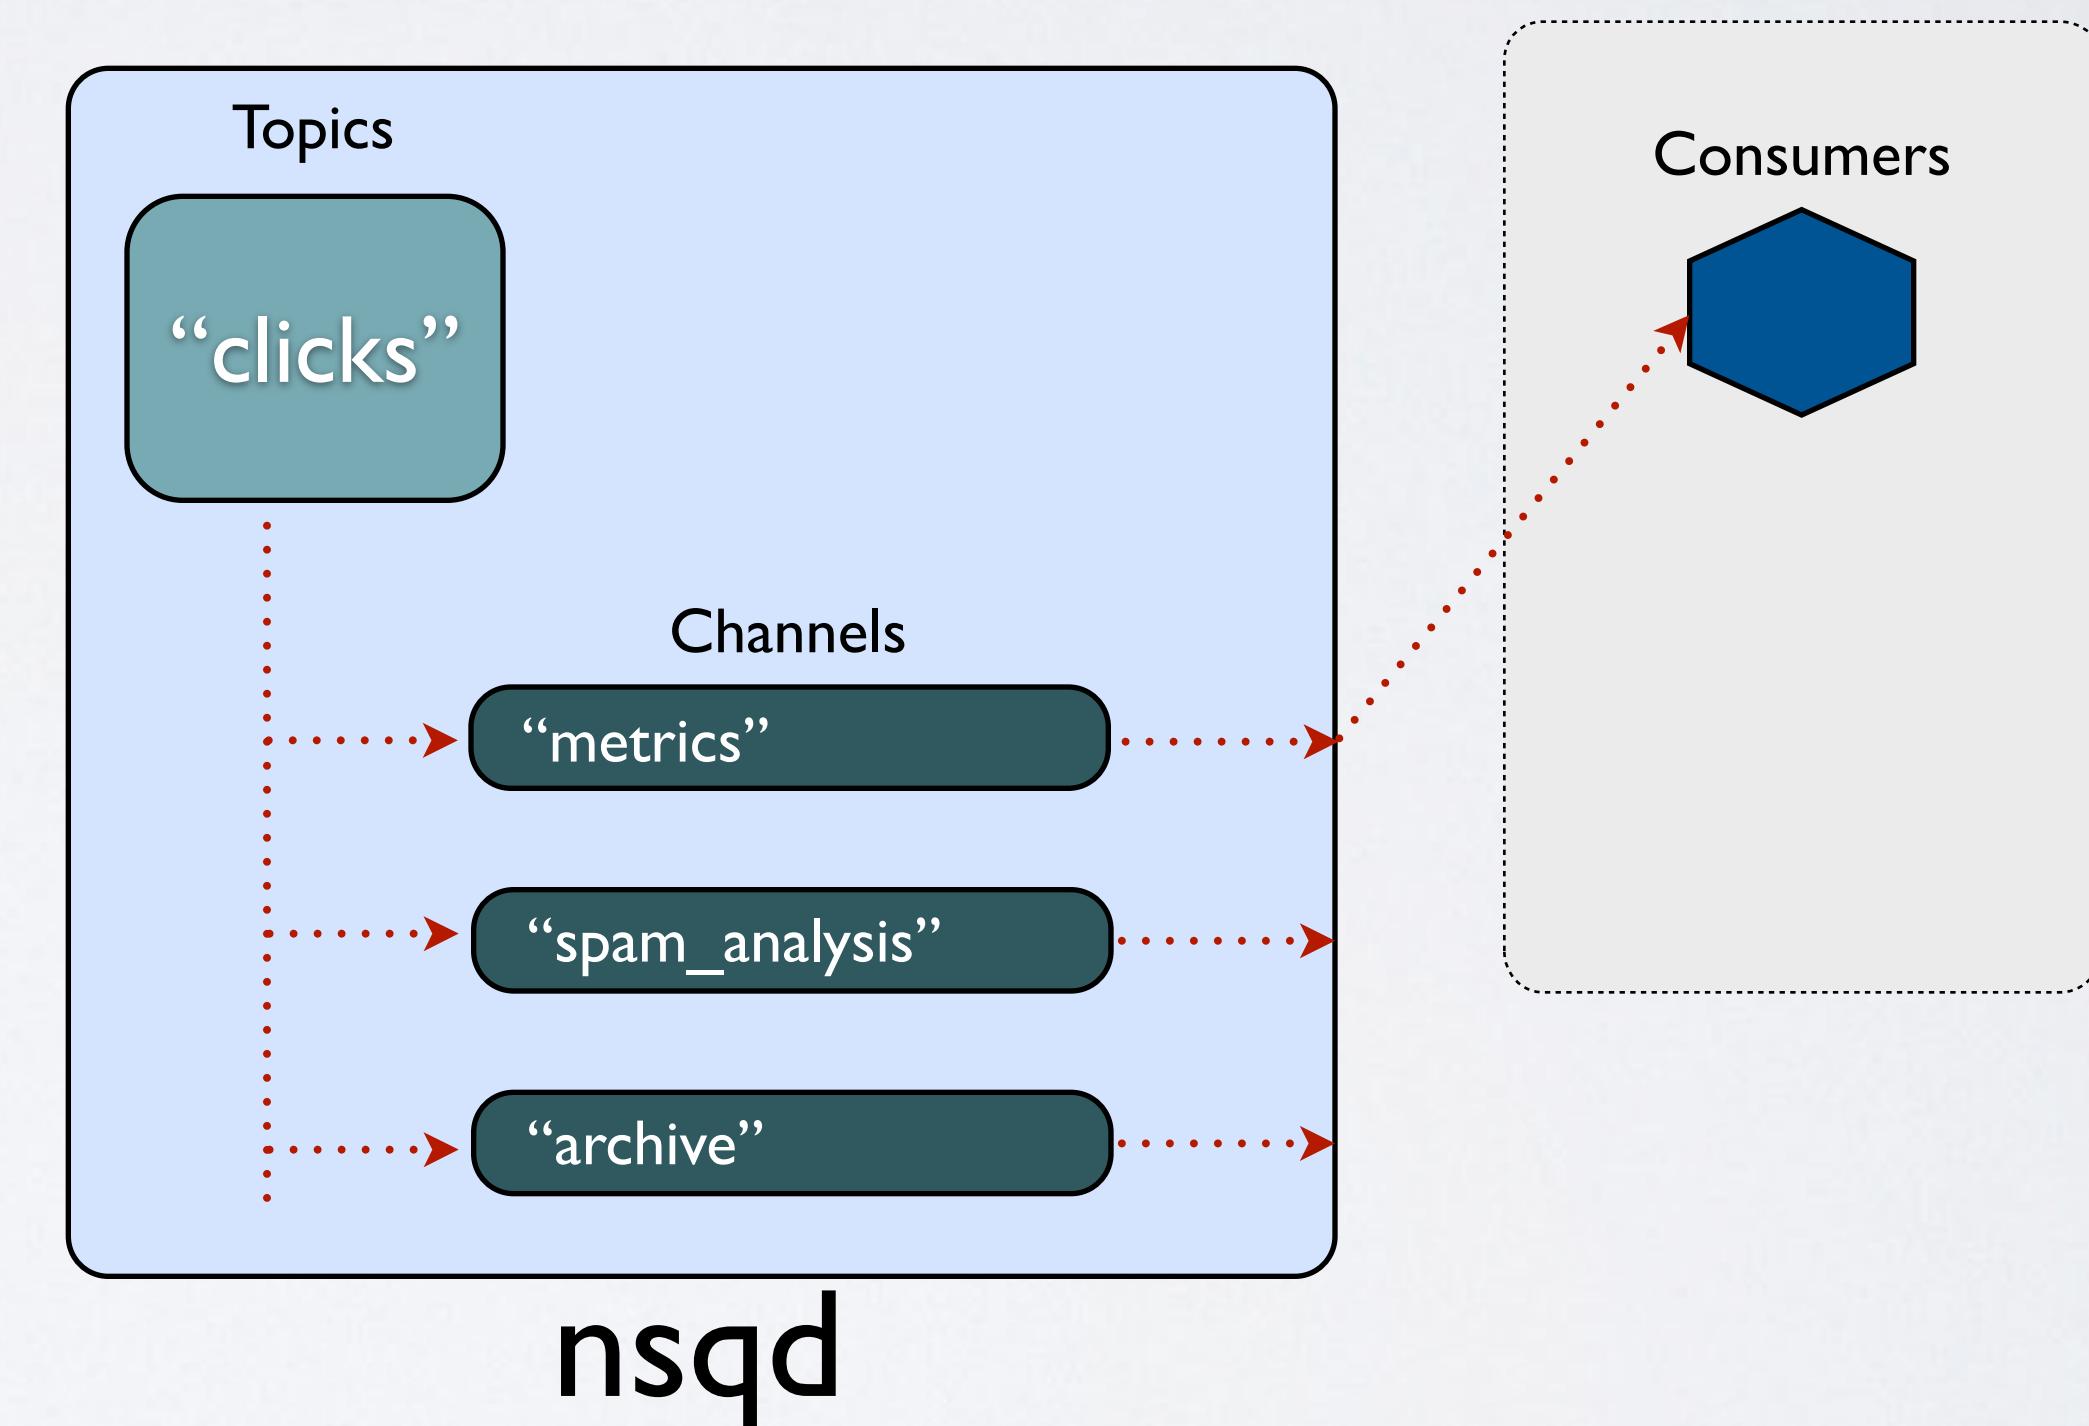
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

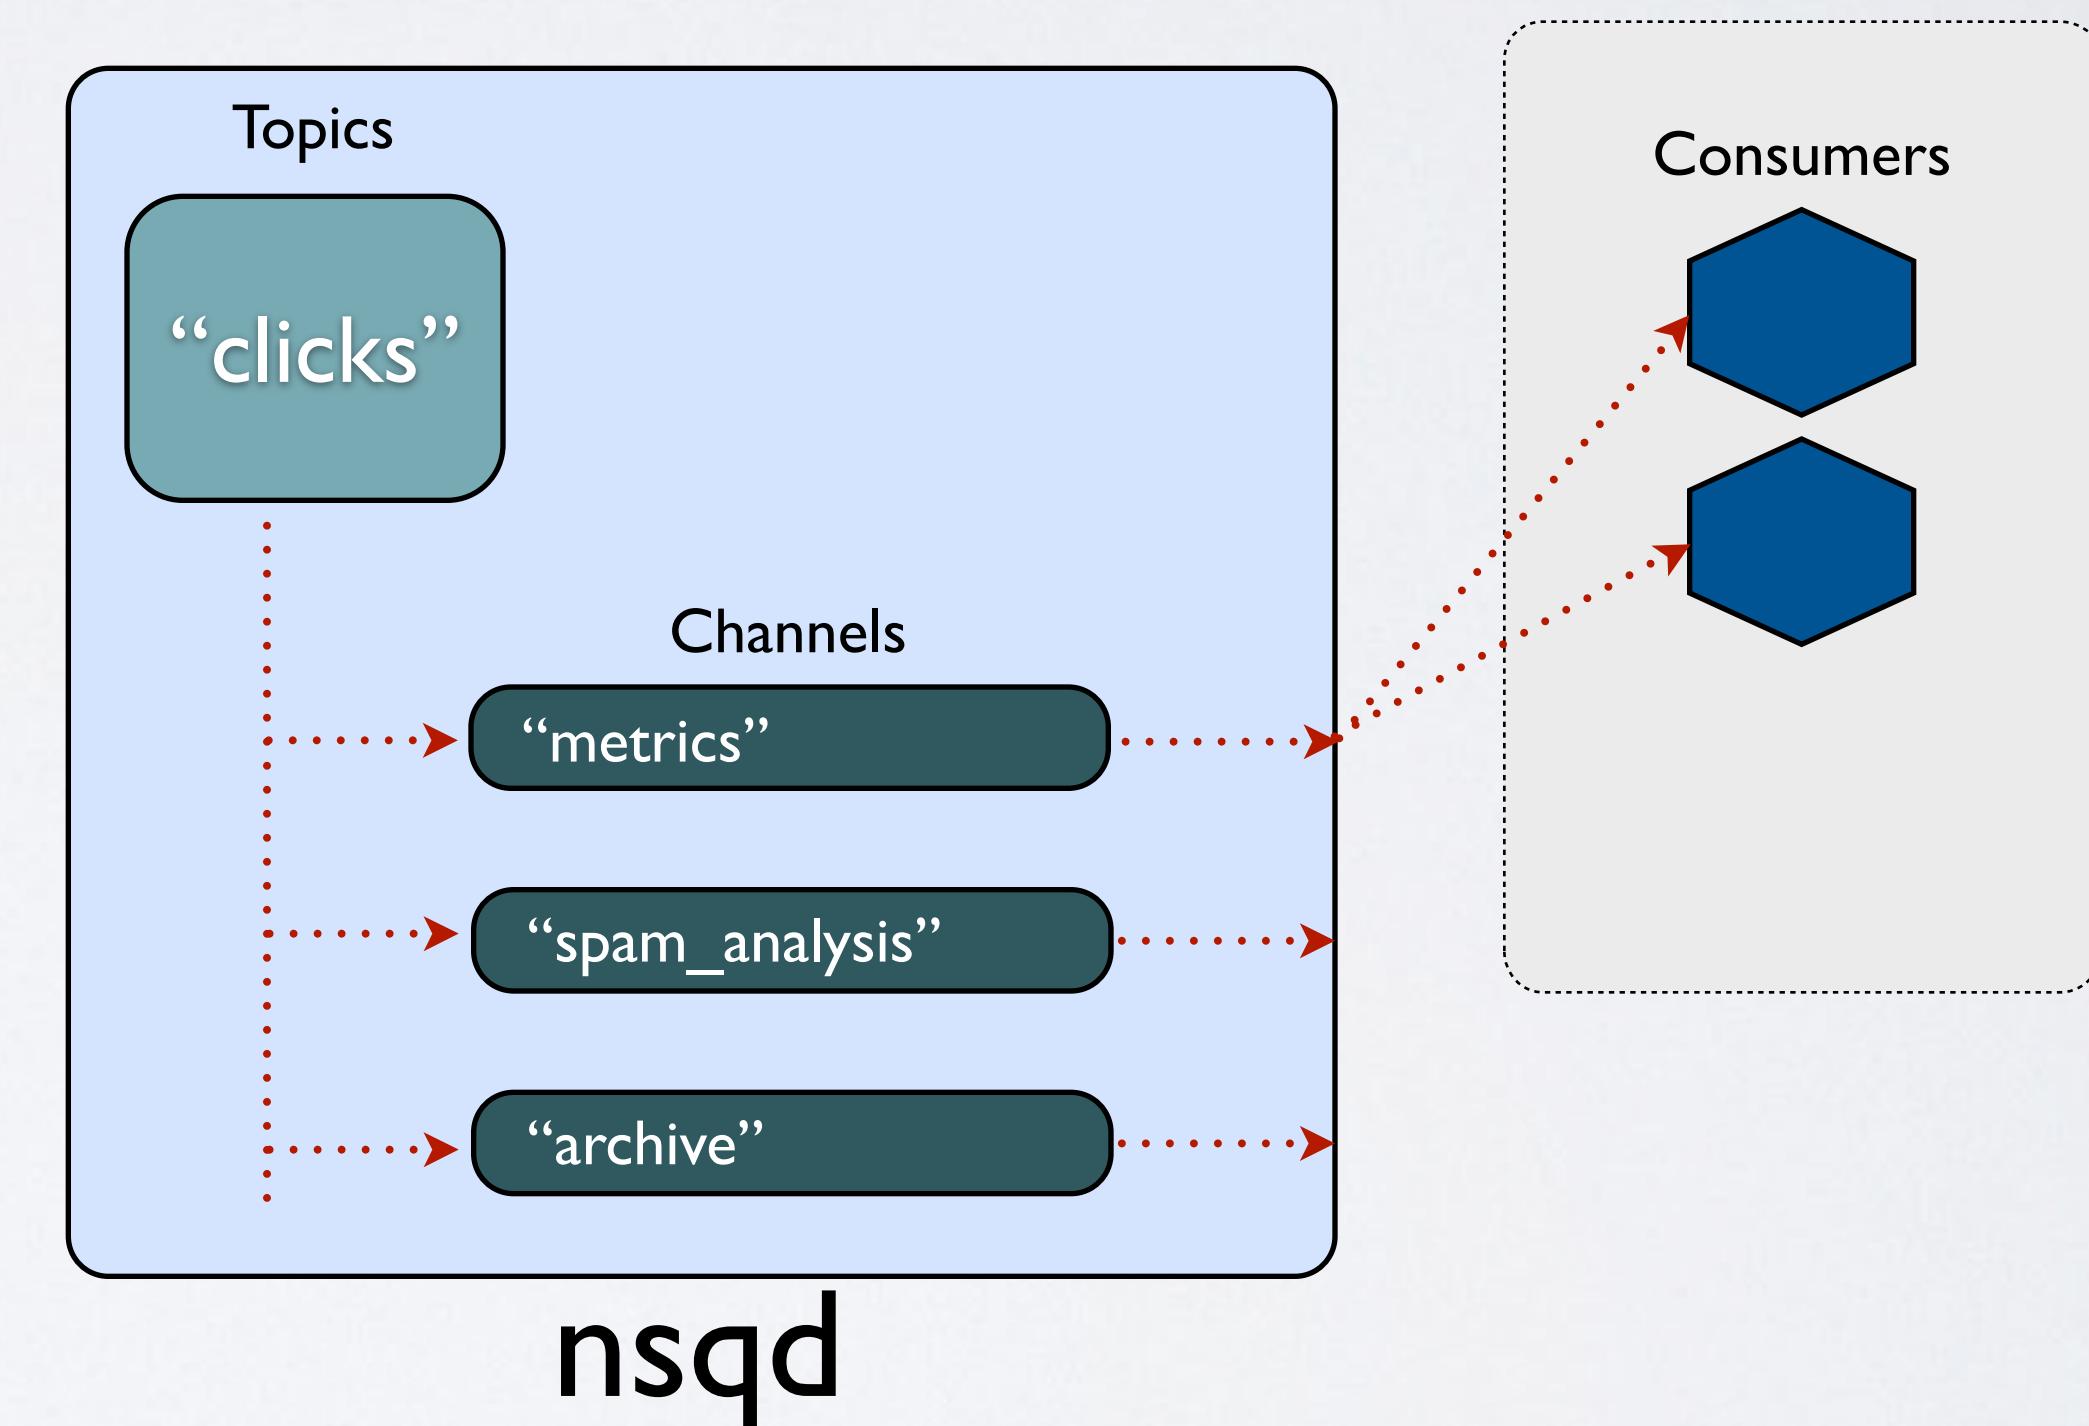
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

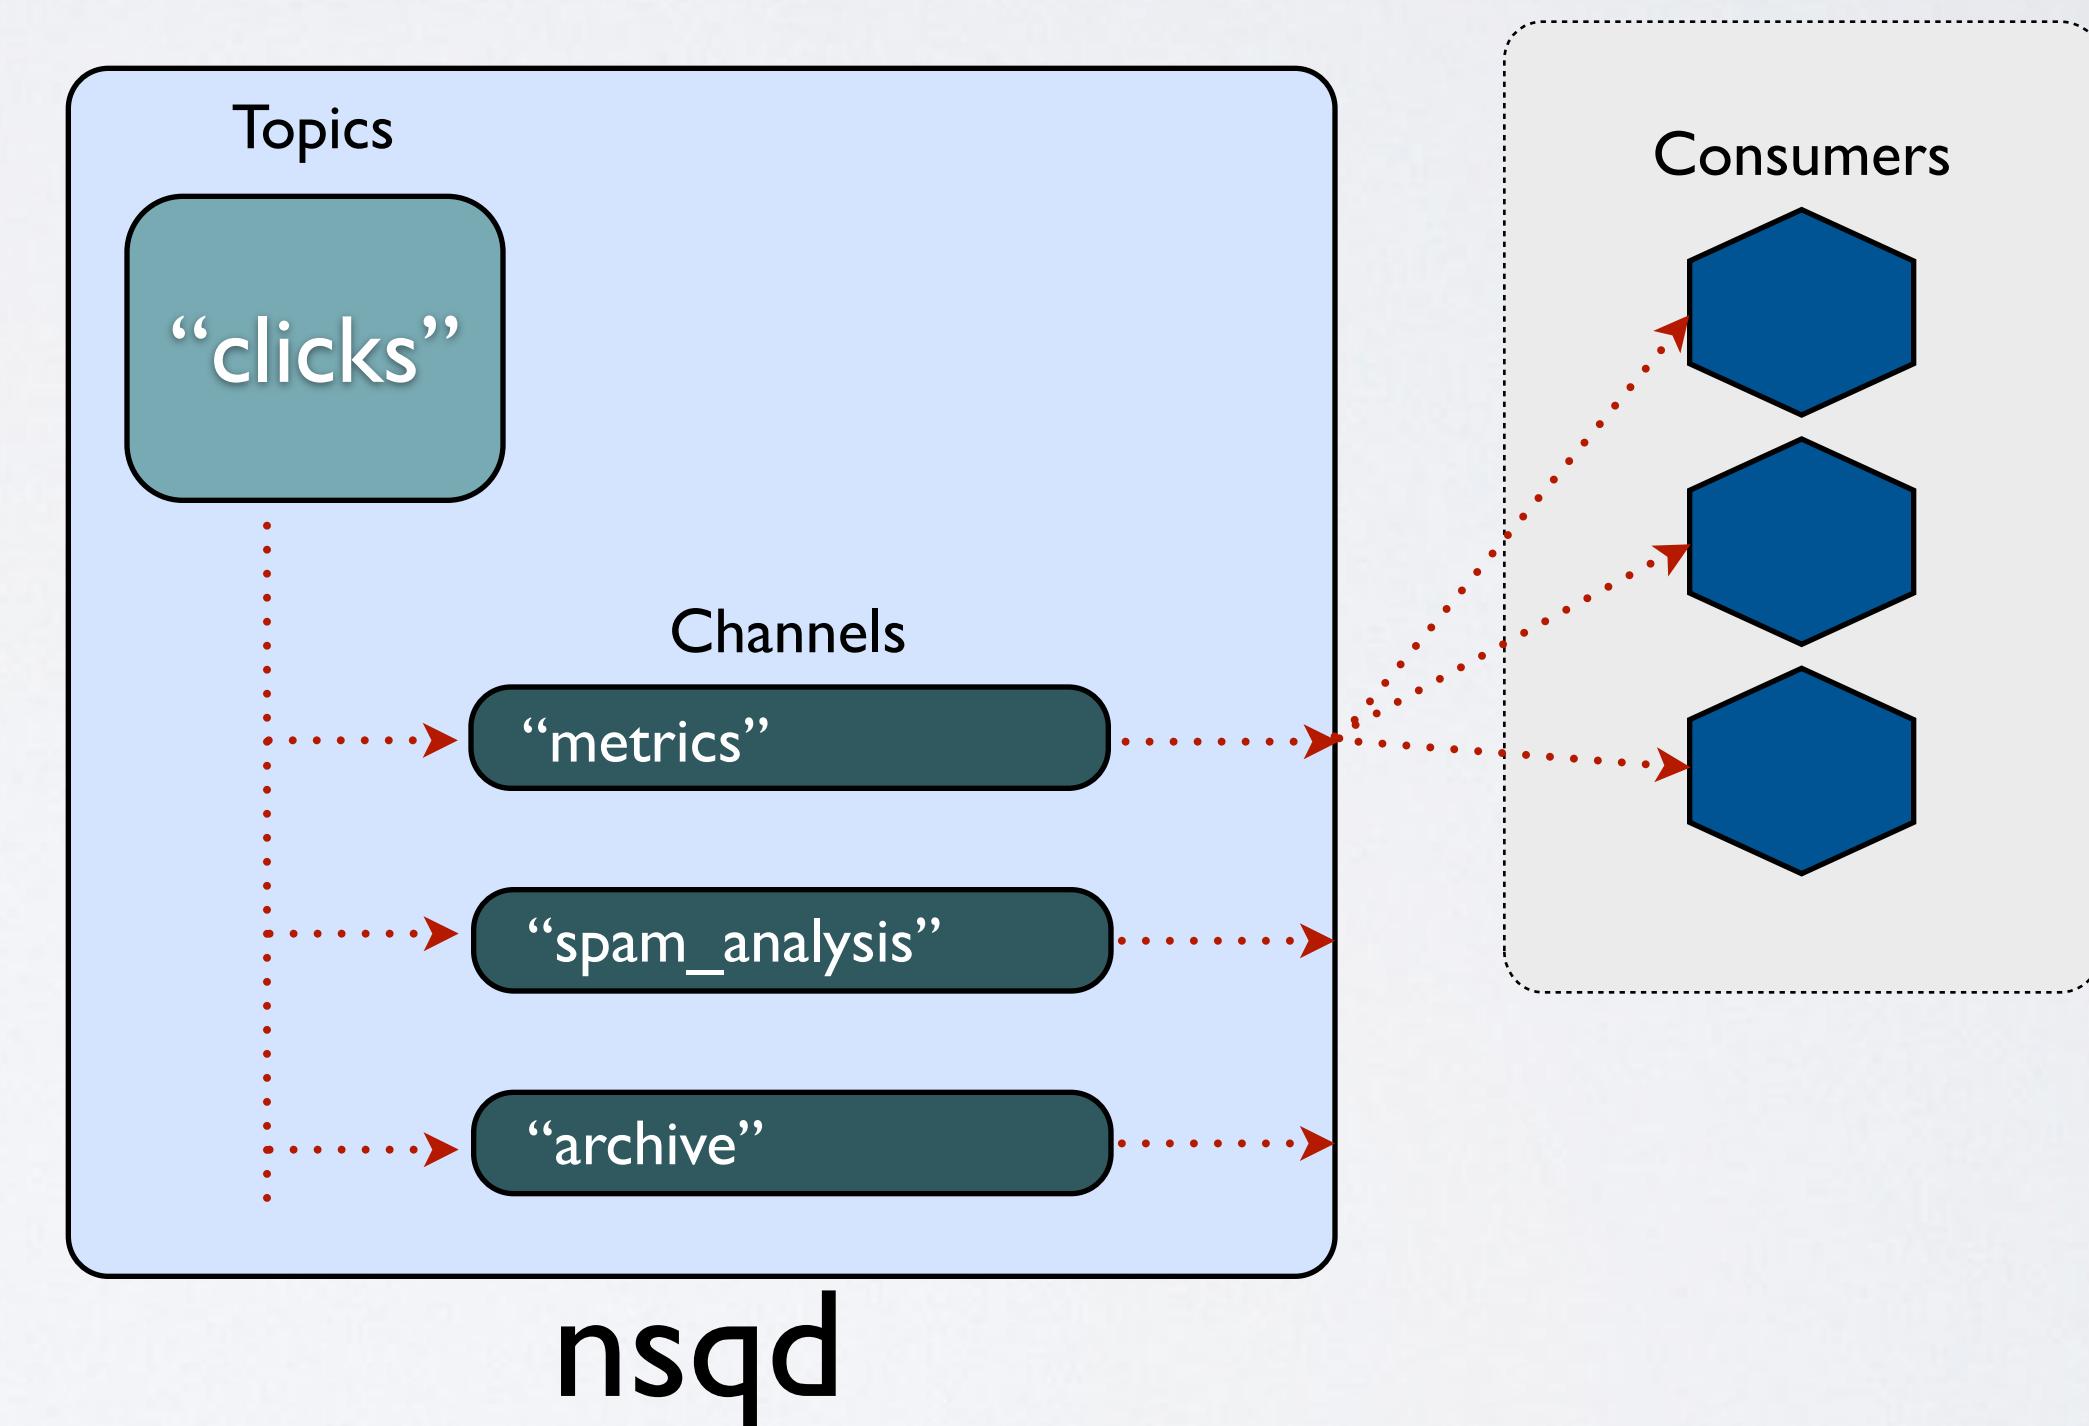
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

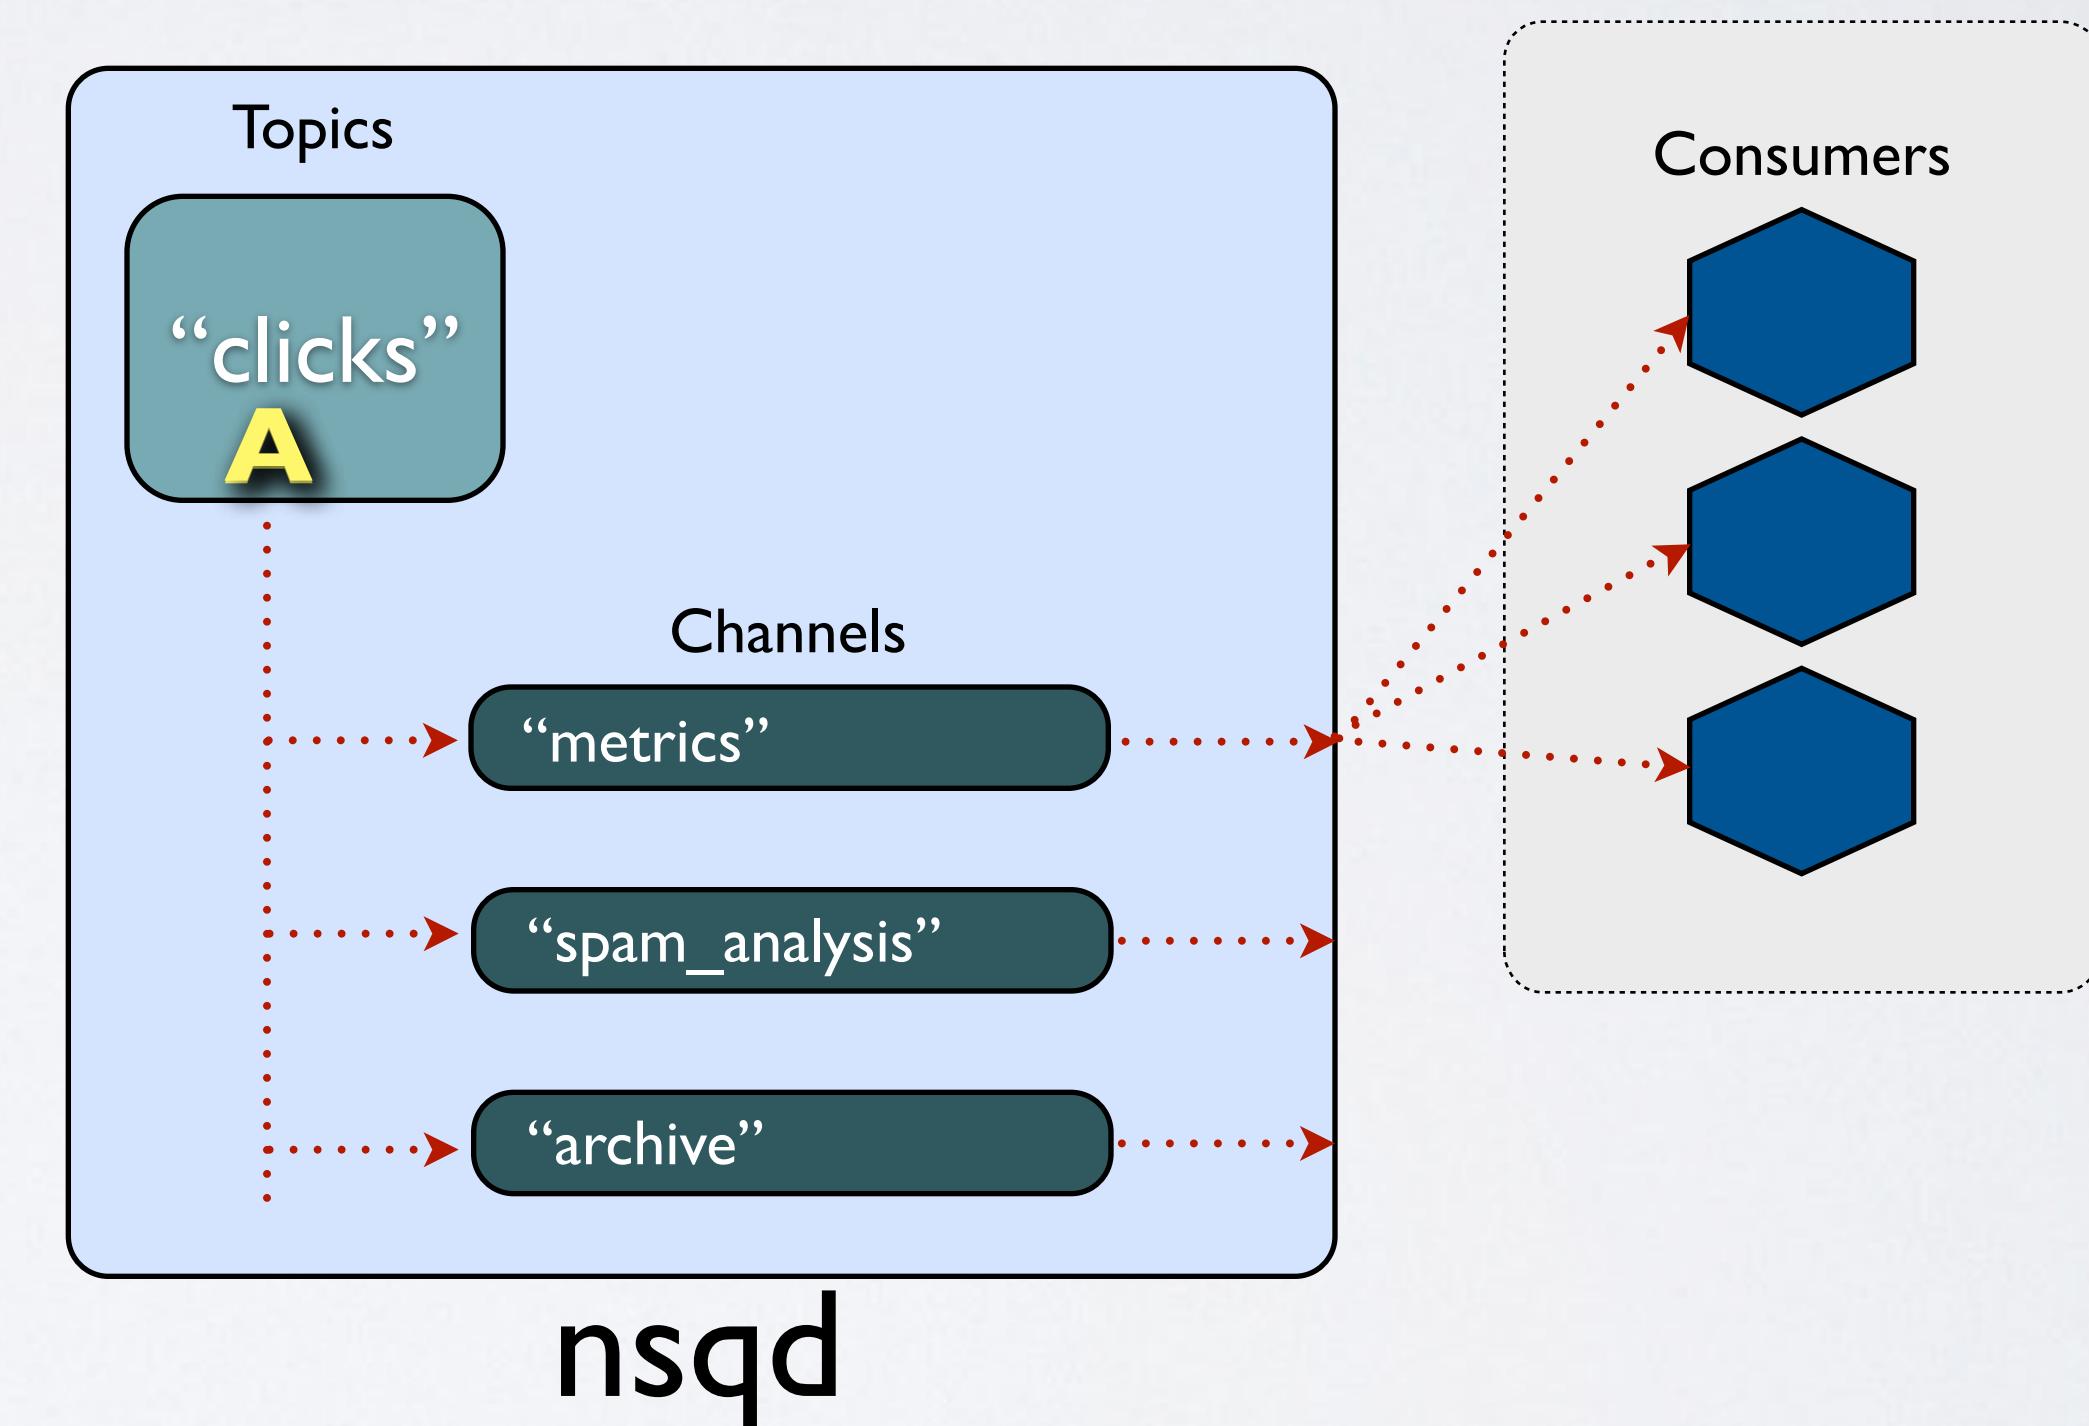
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

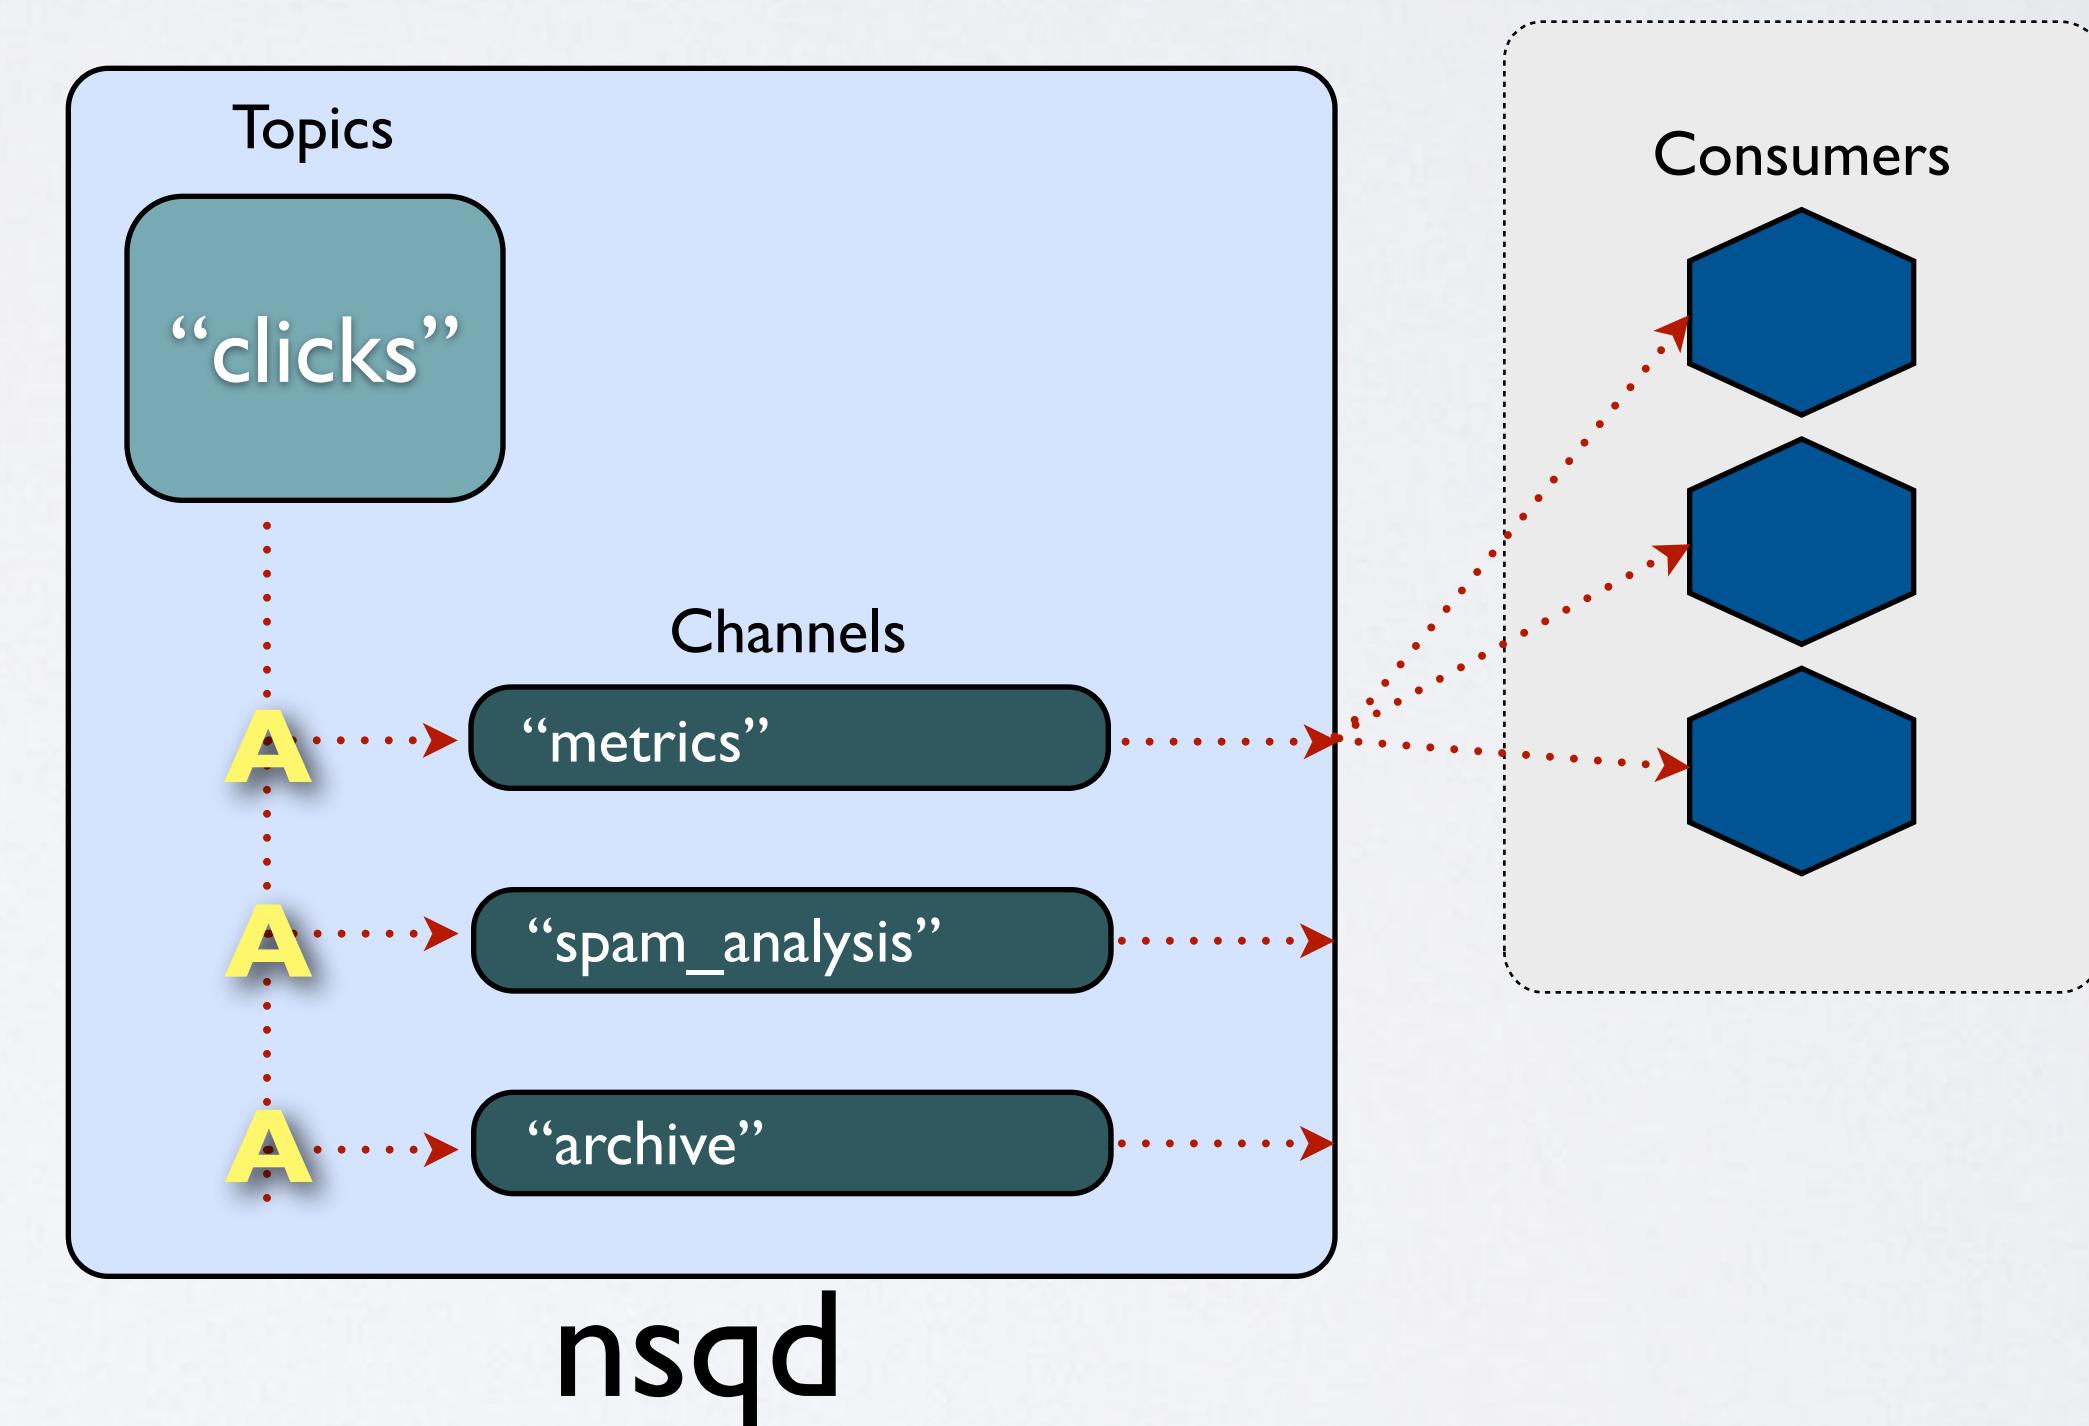
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

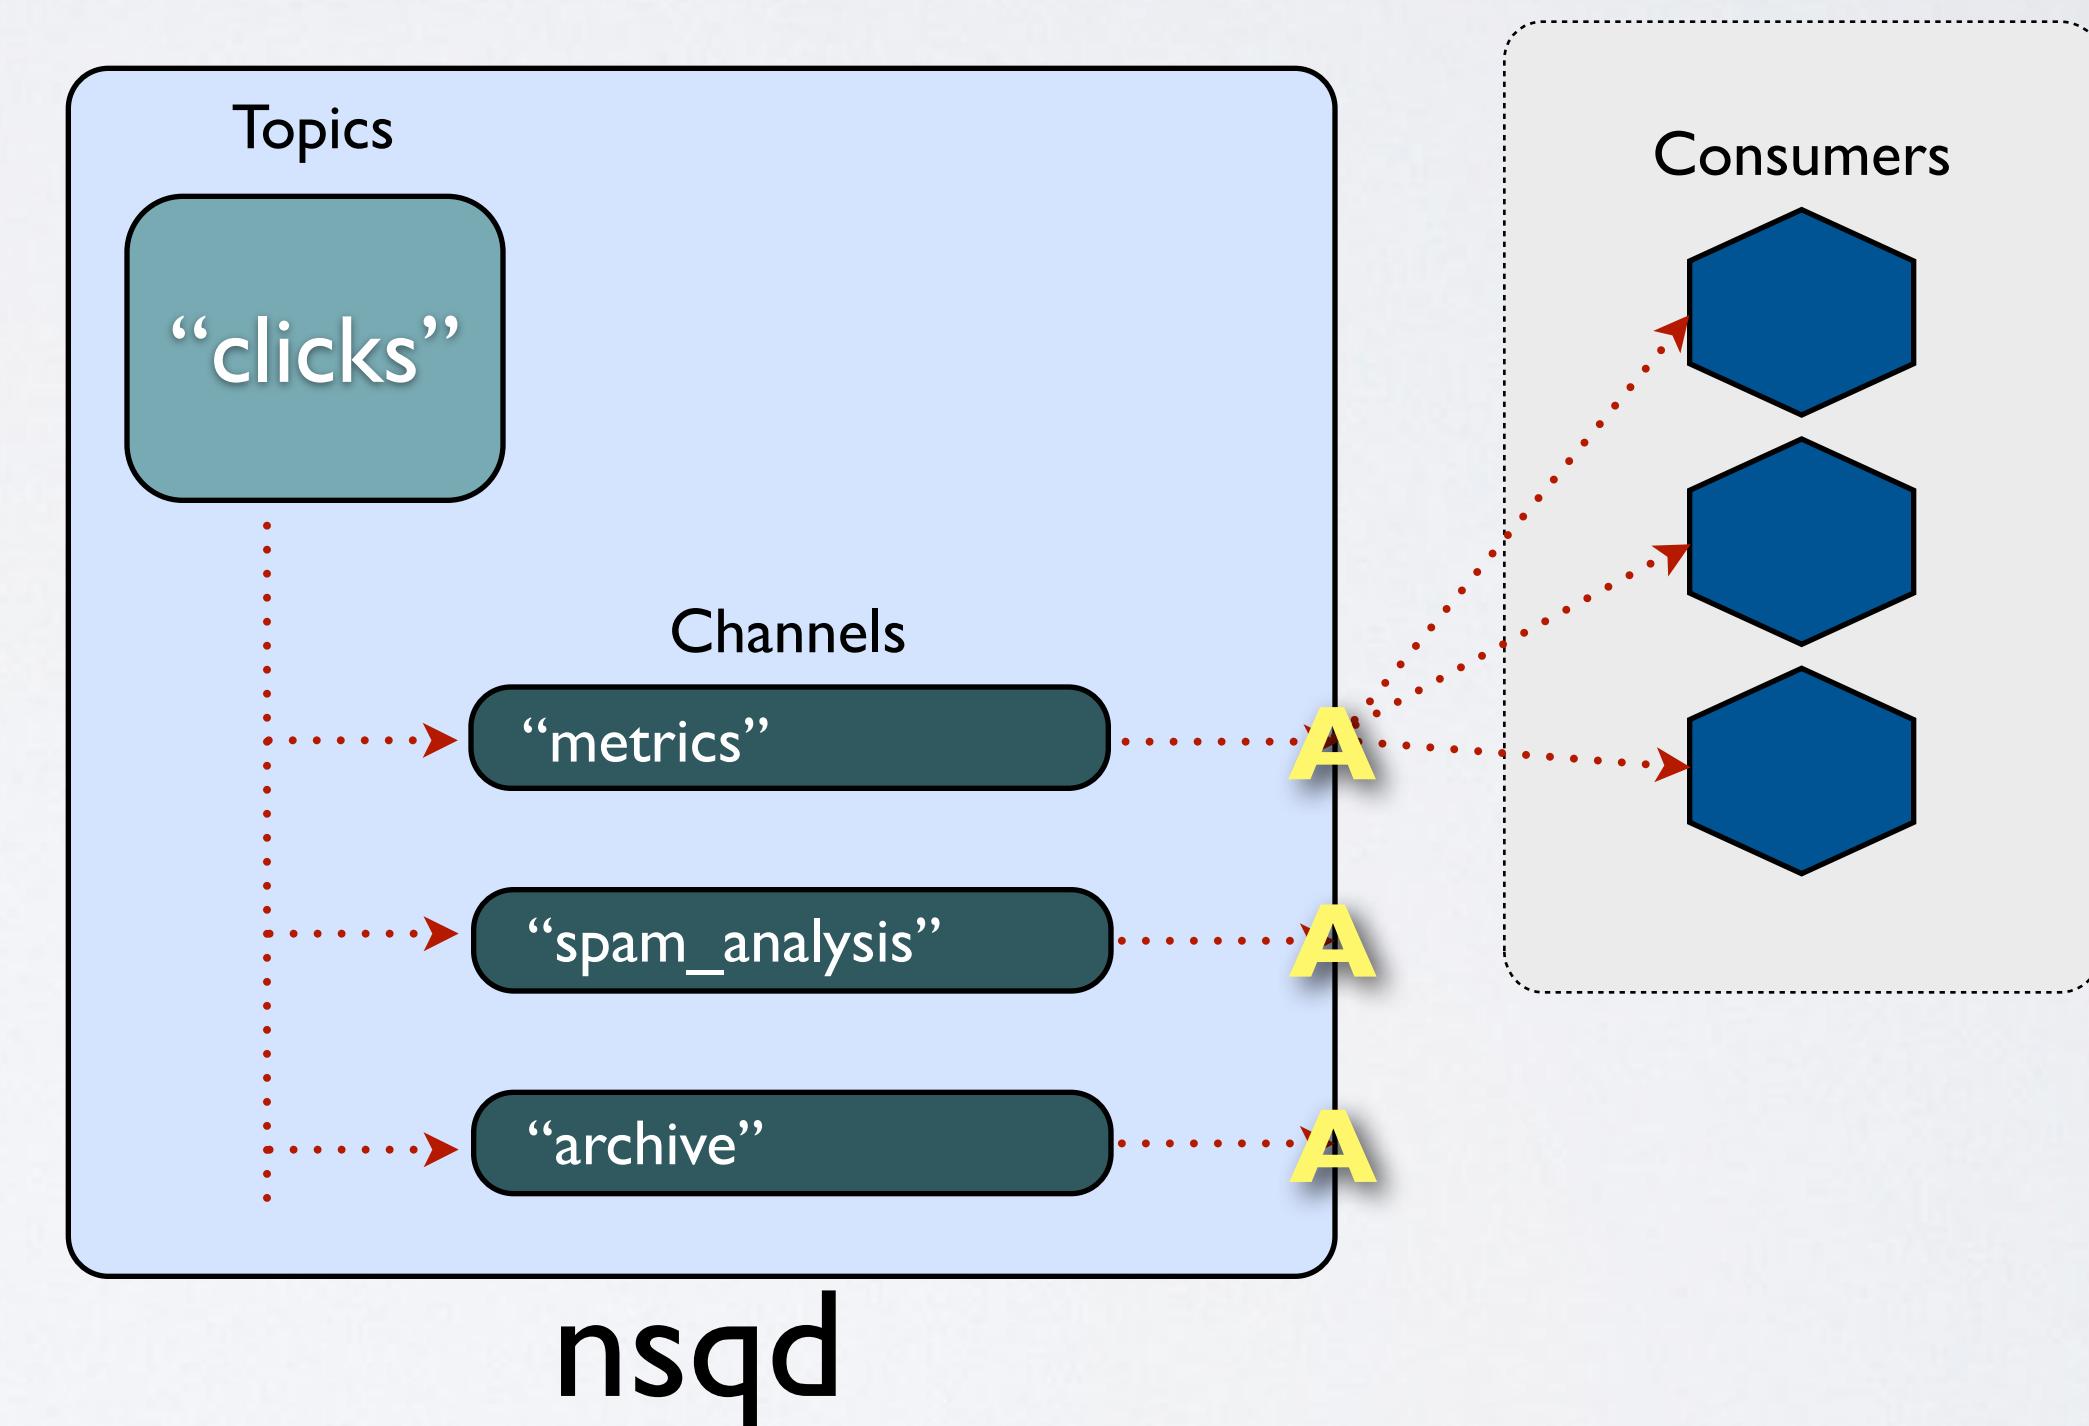
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

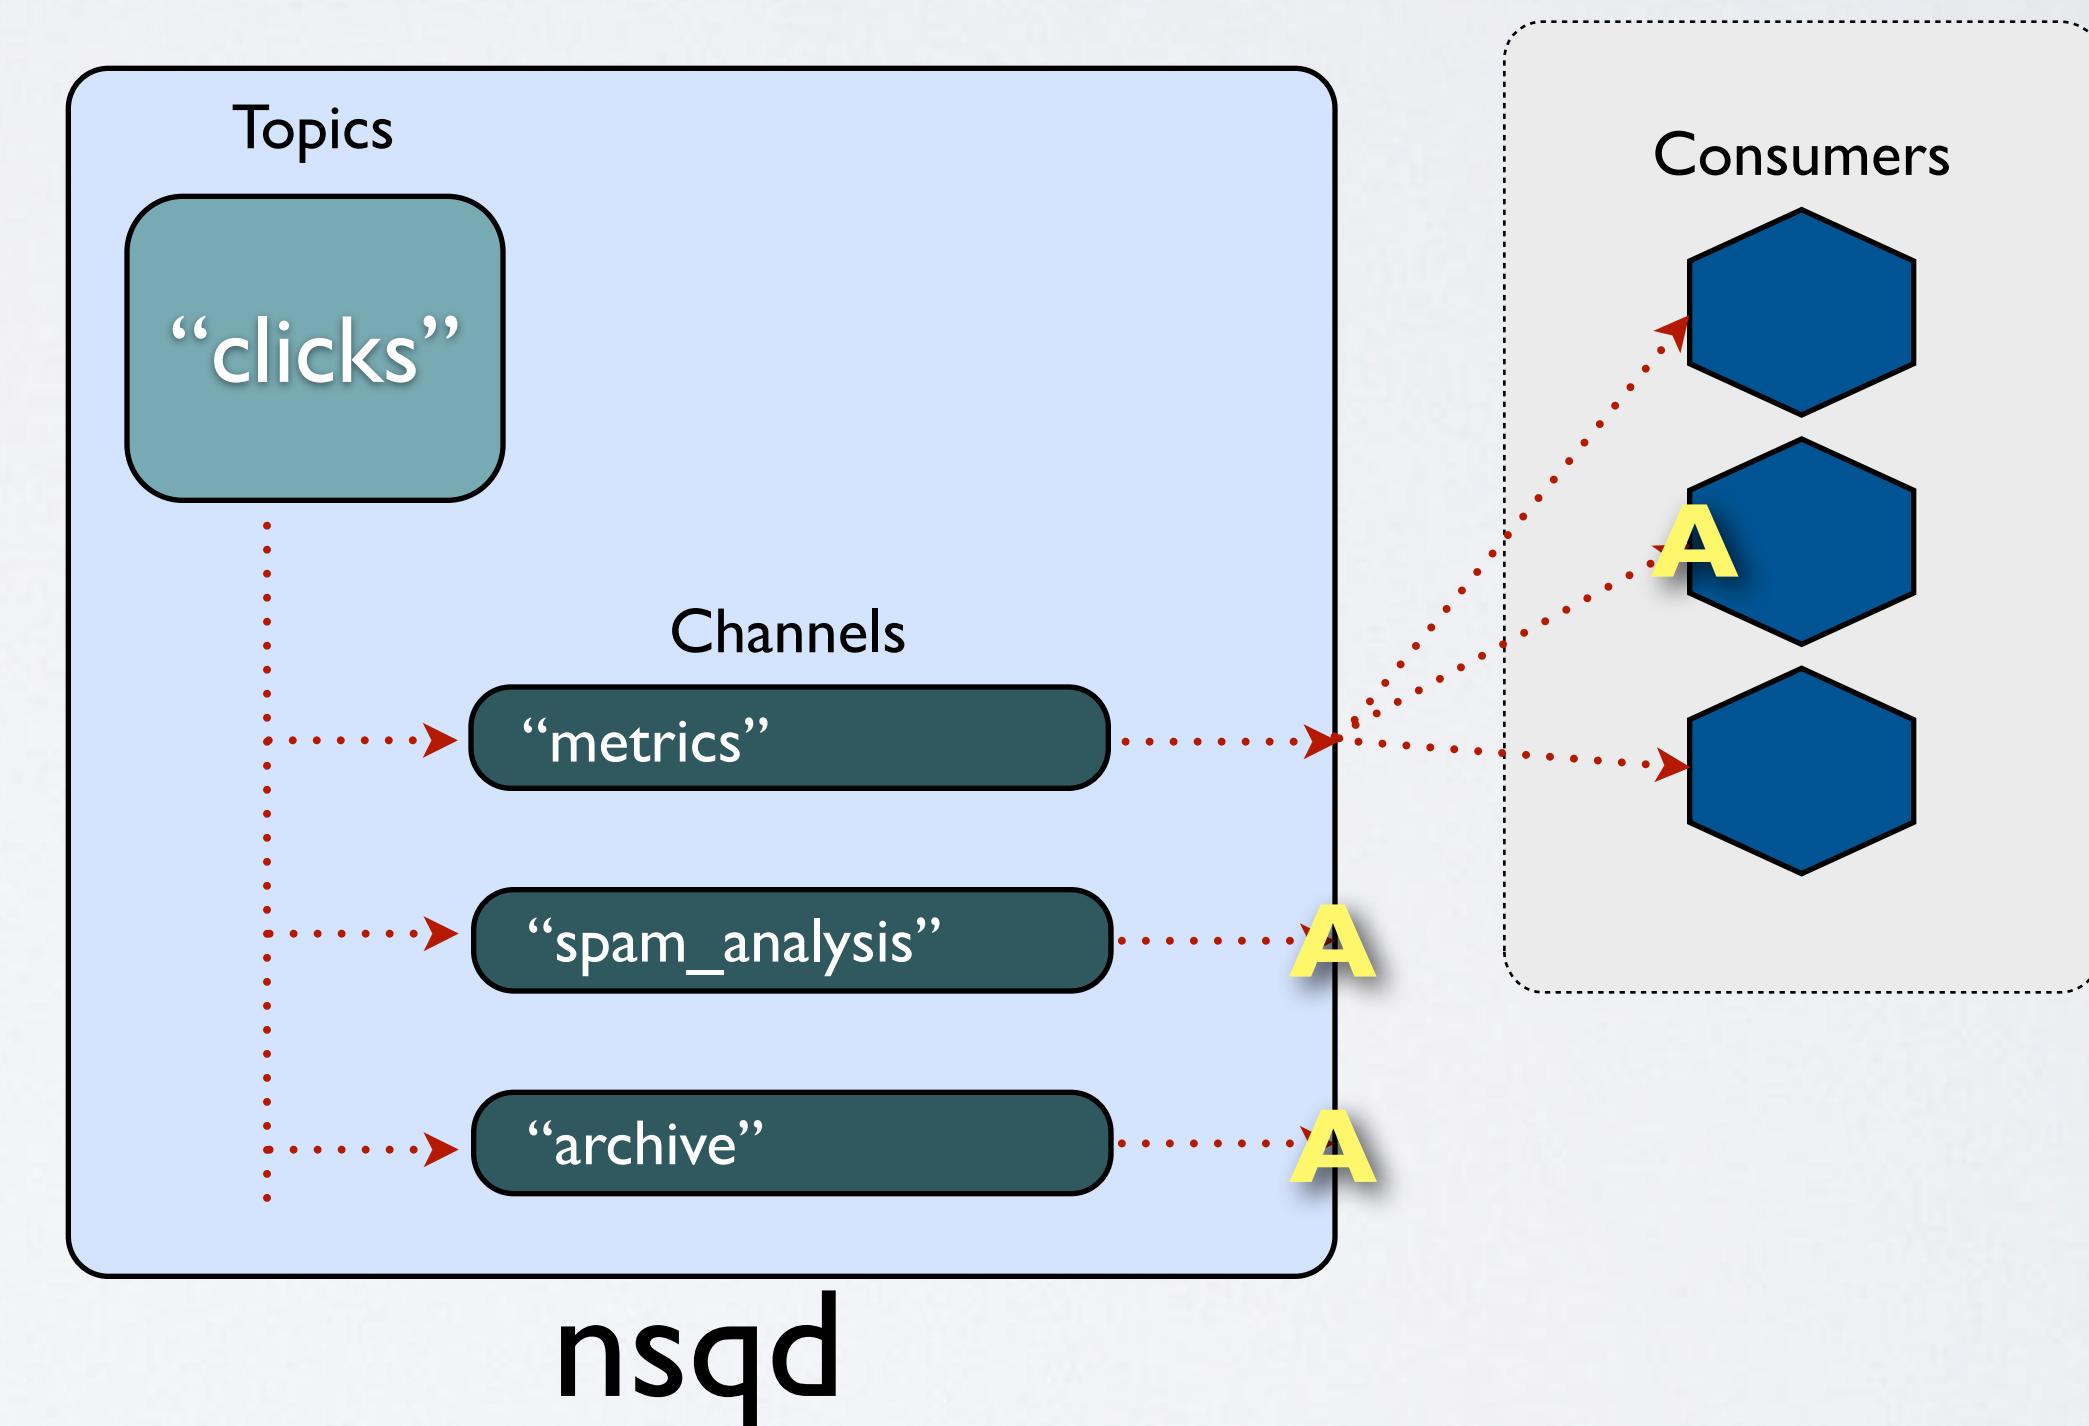
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

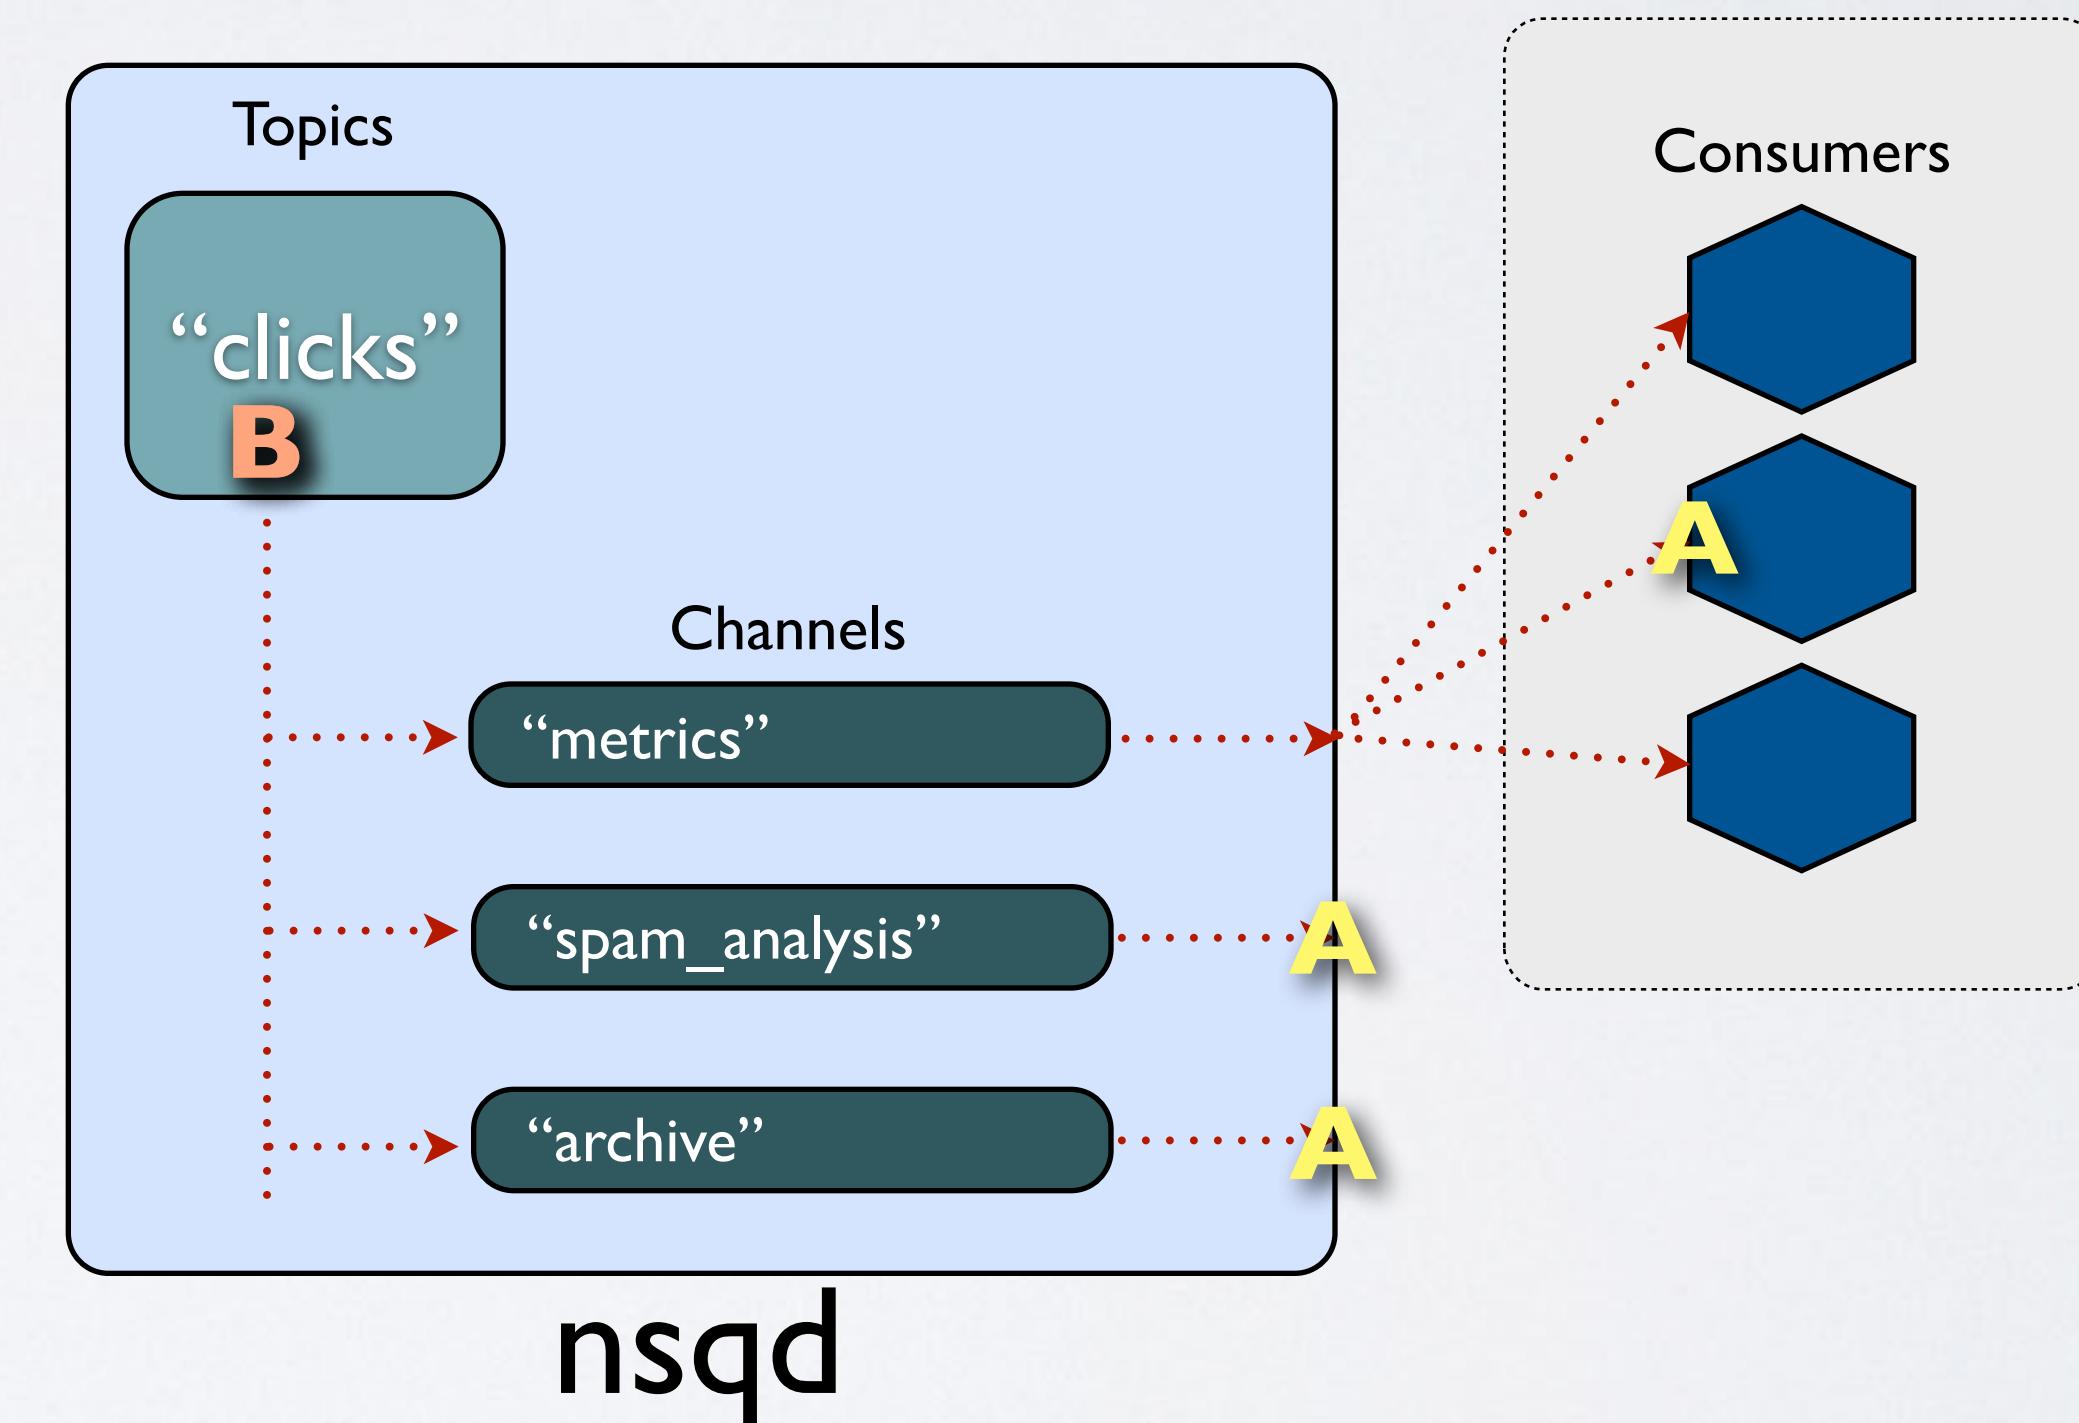
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

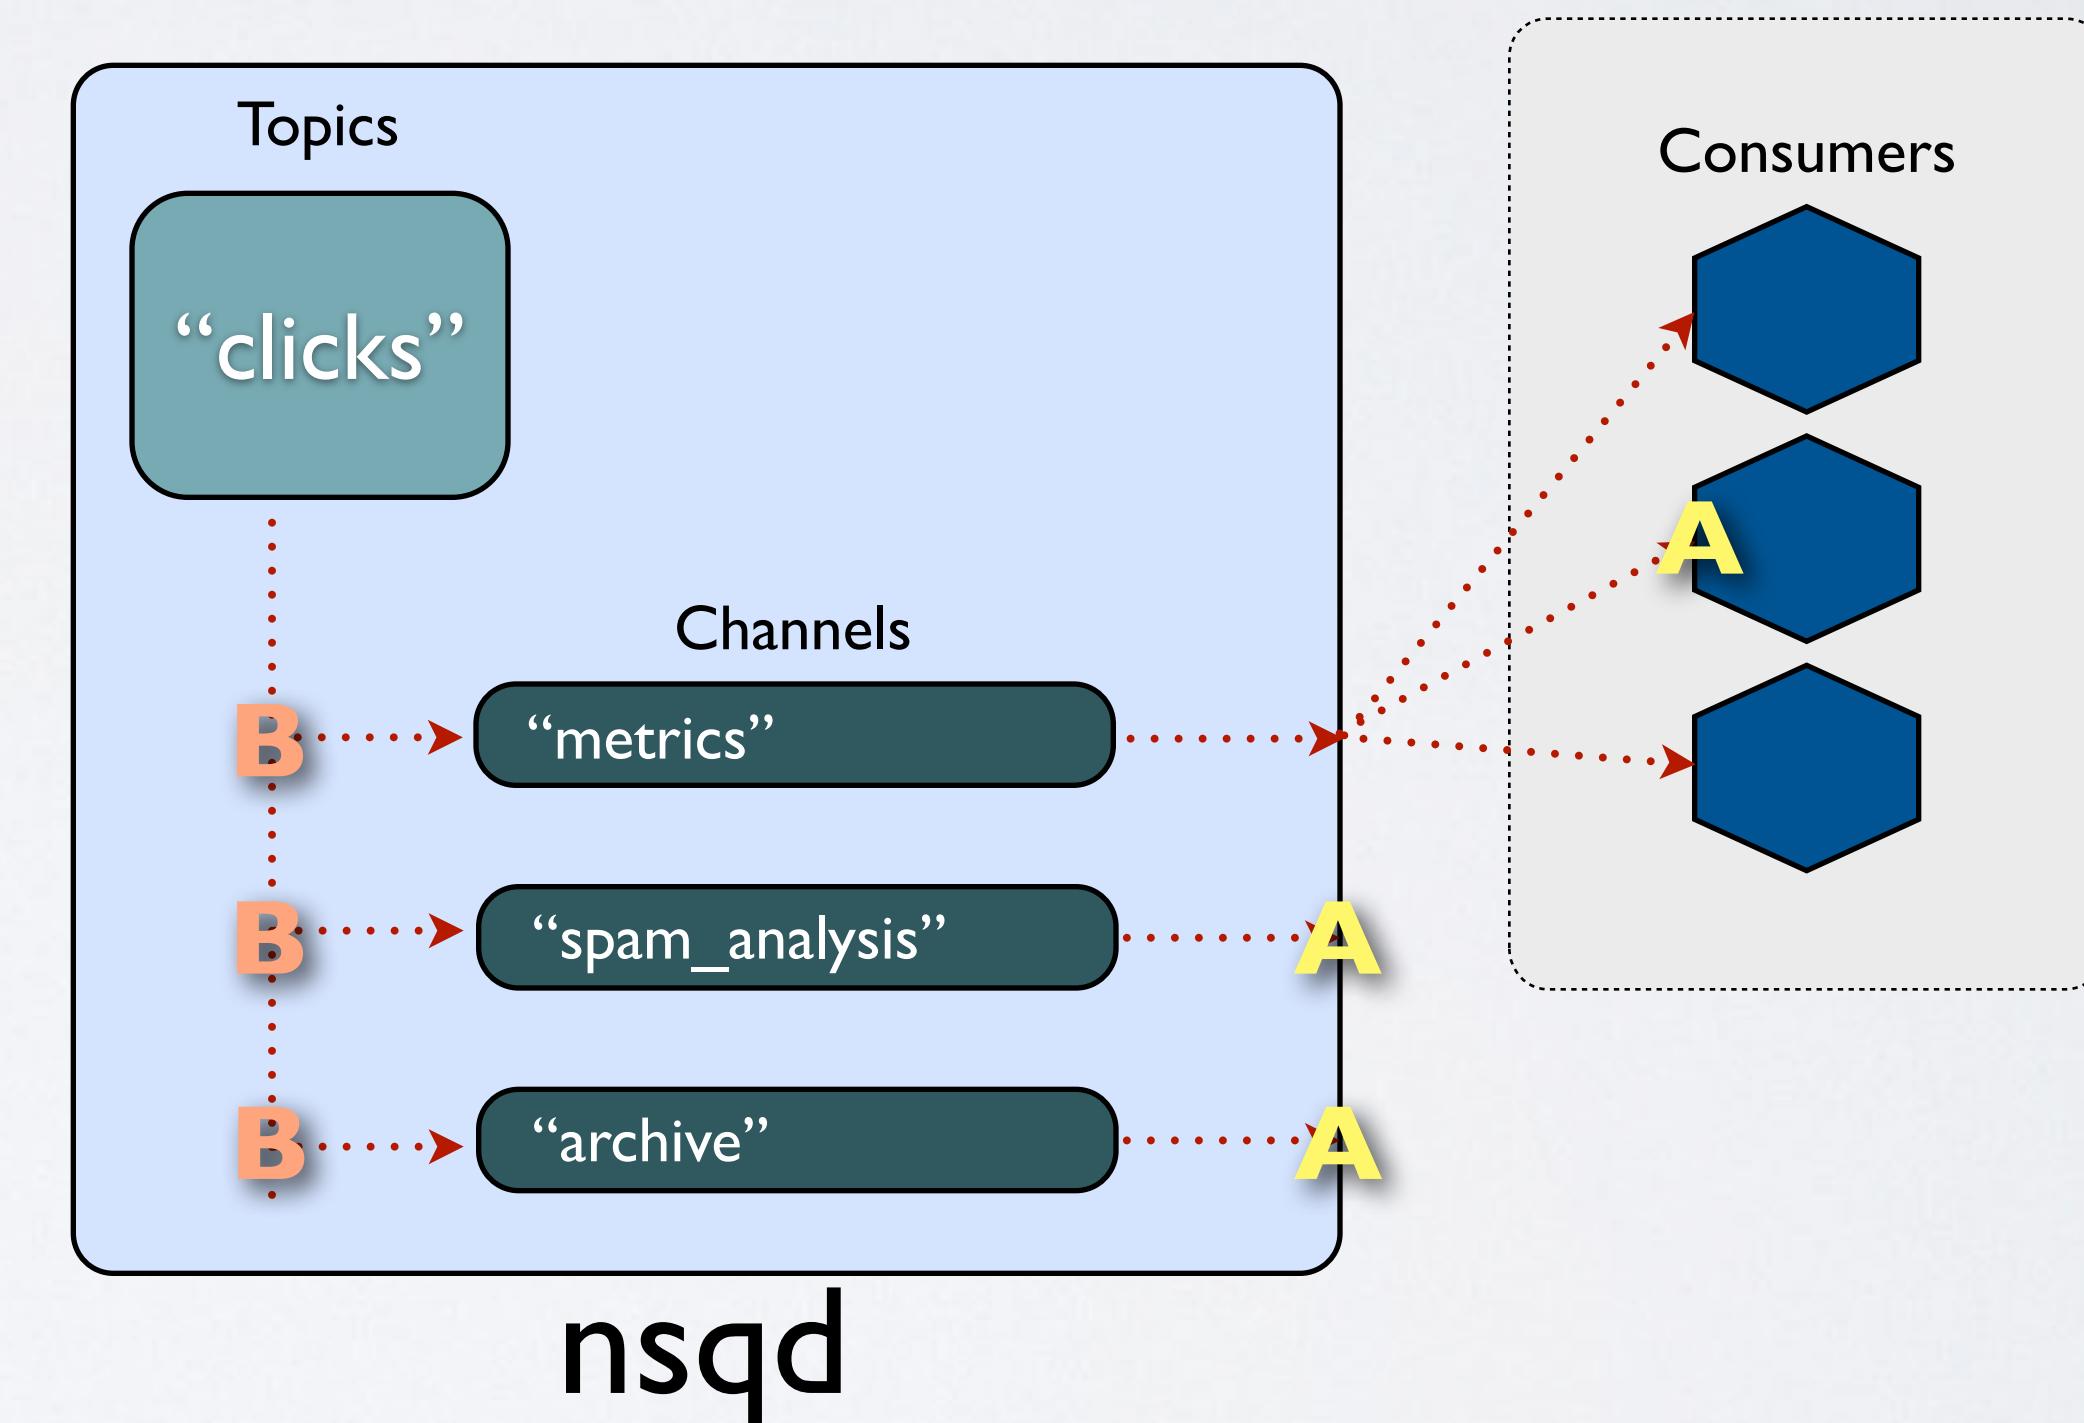
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

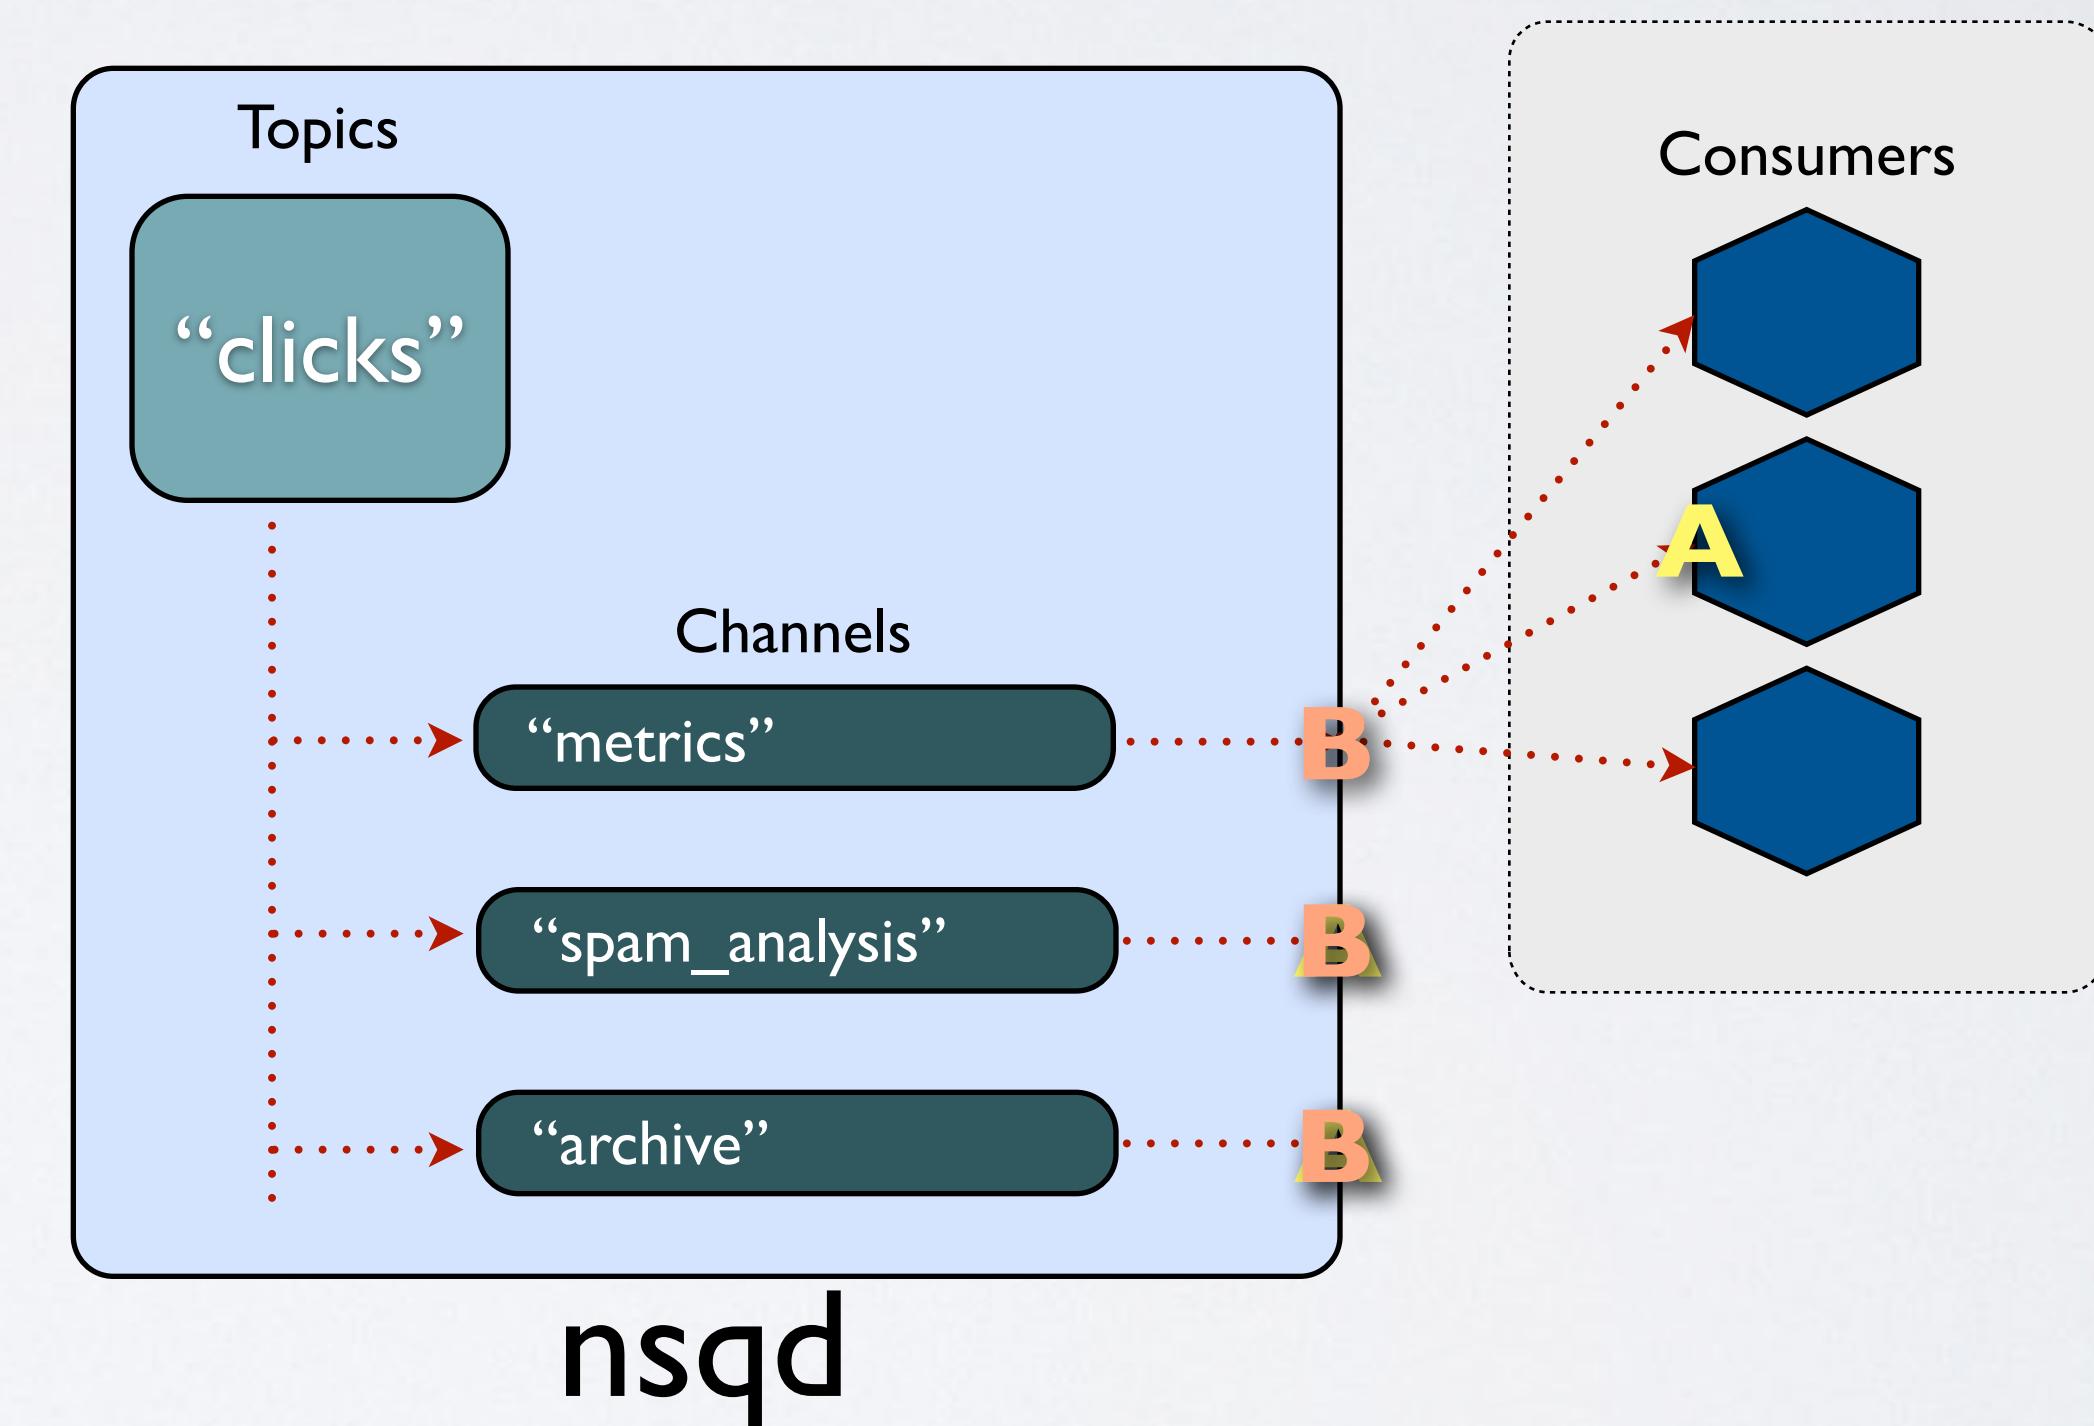
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

combine pubsub, distribution, and queueing

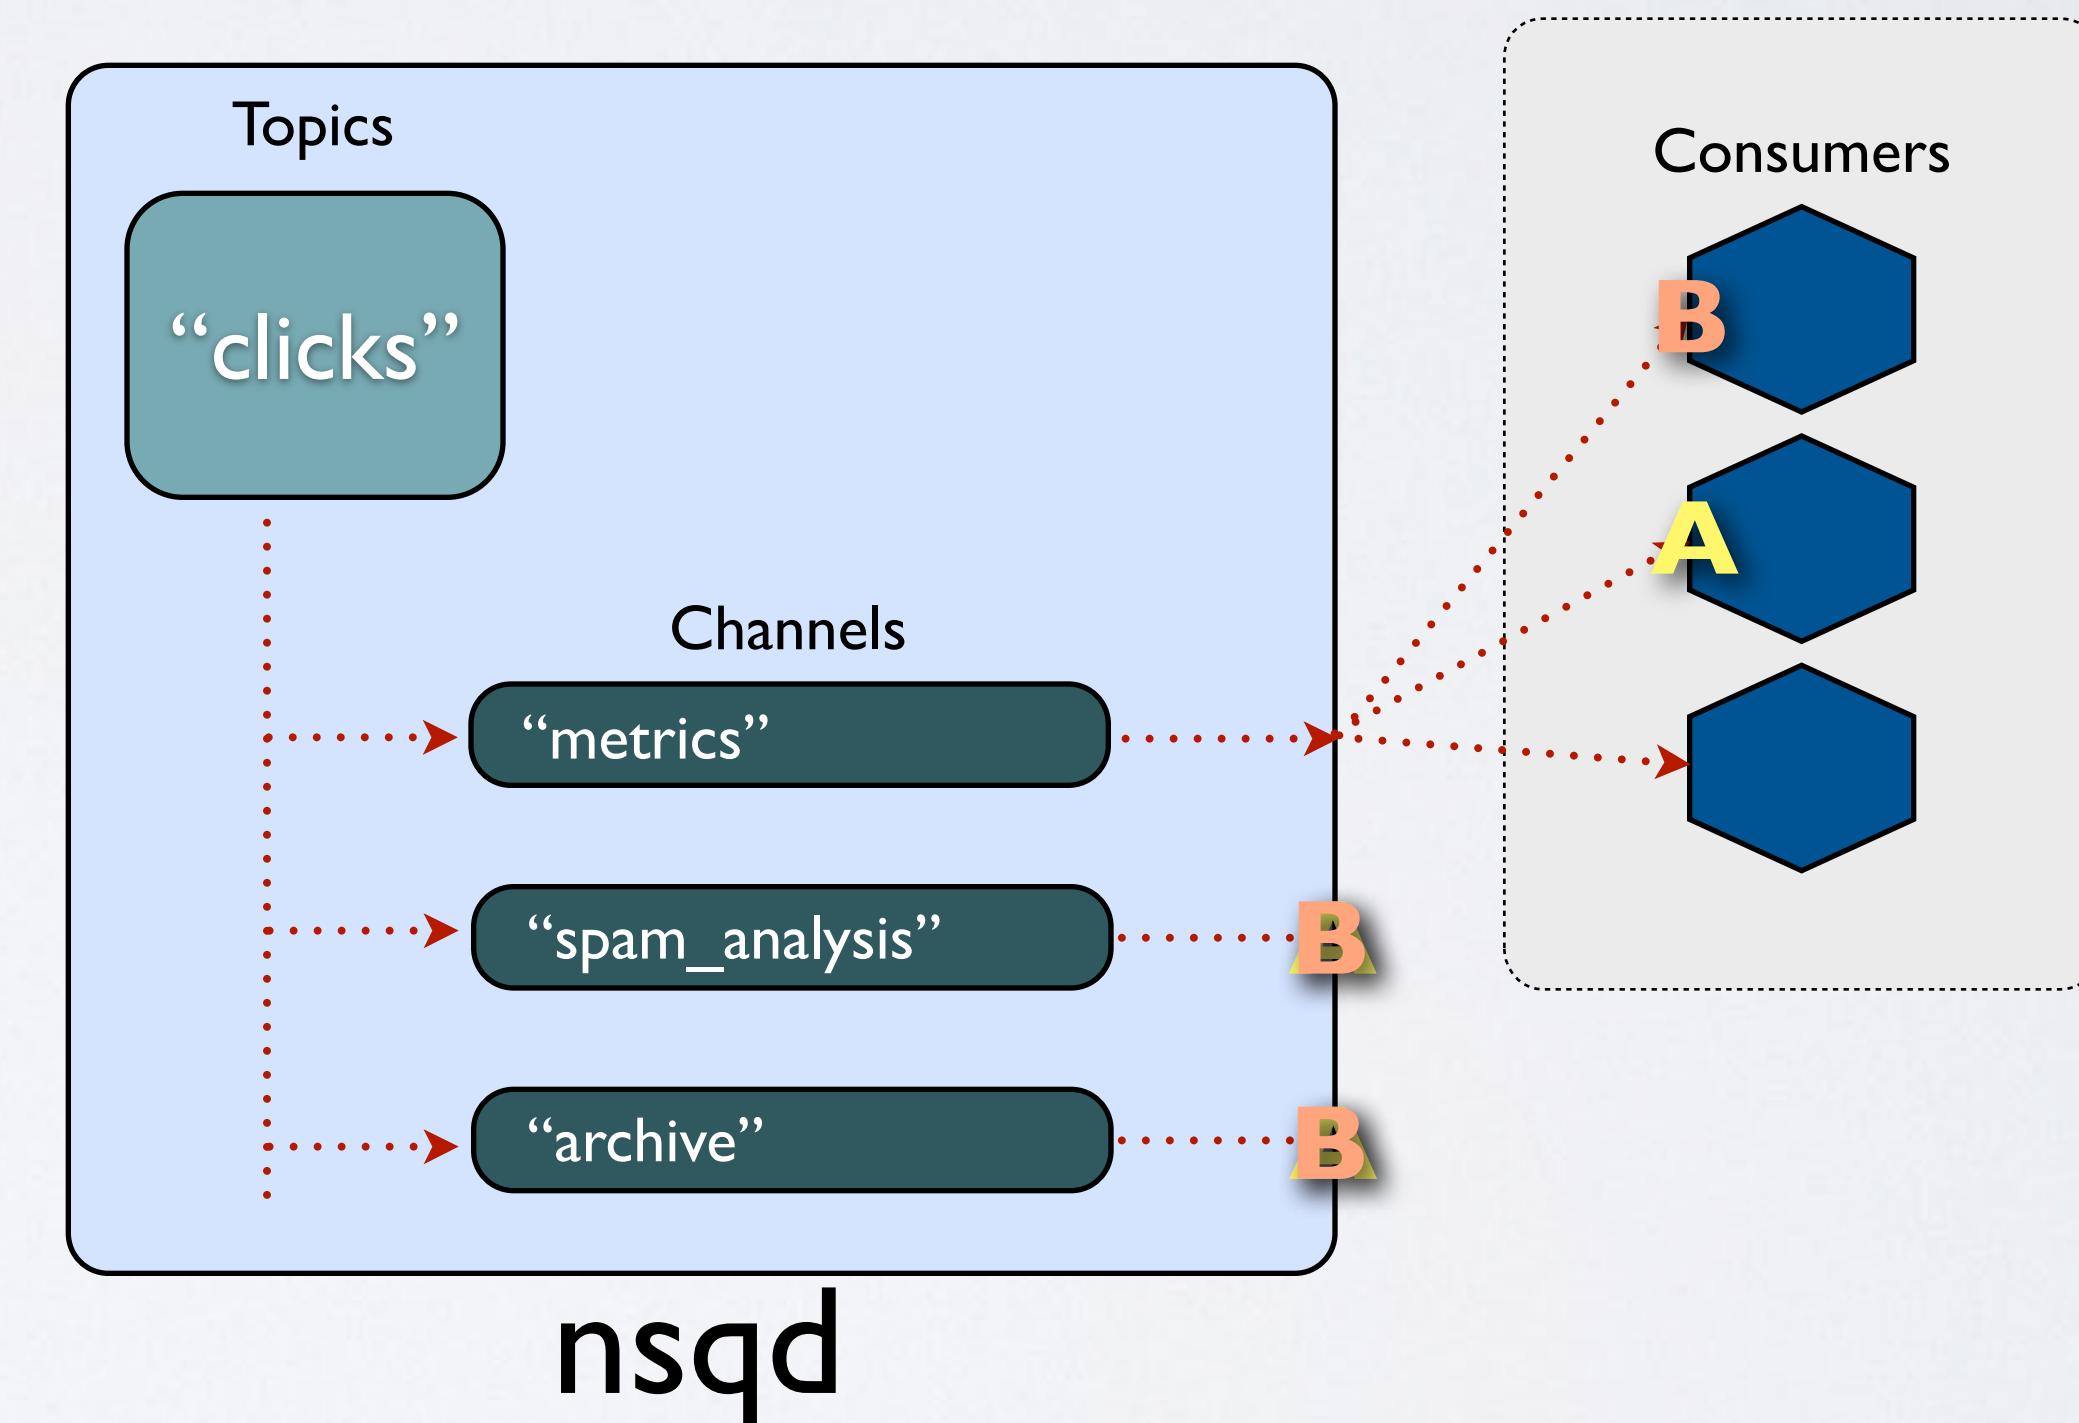
- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



TOPICS AND CHANNELS

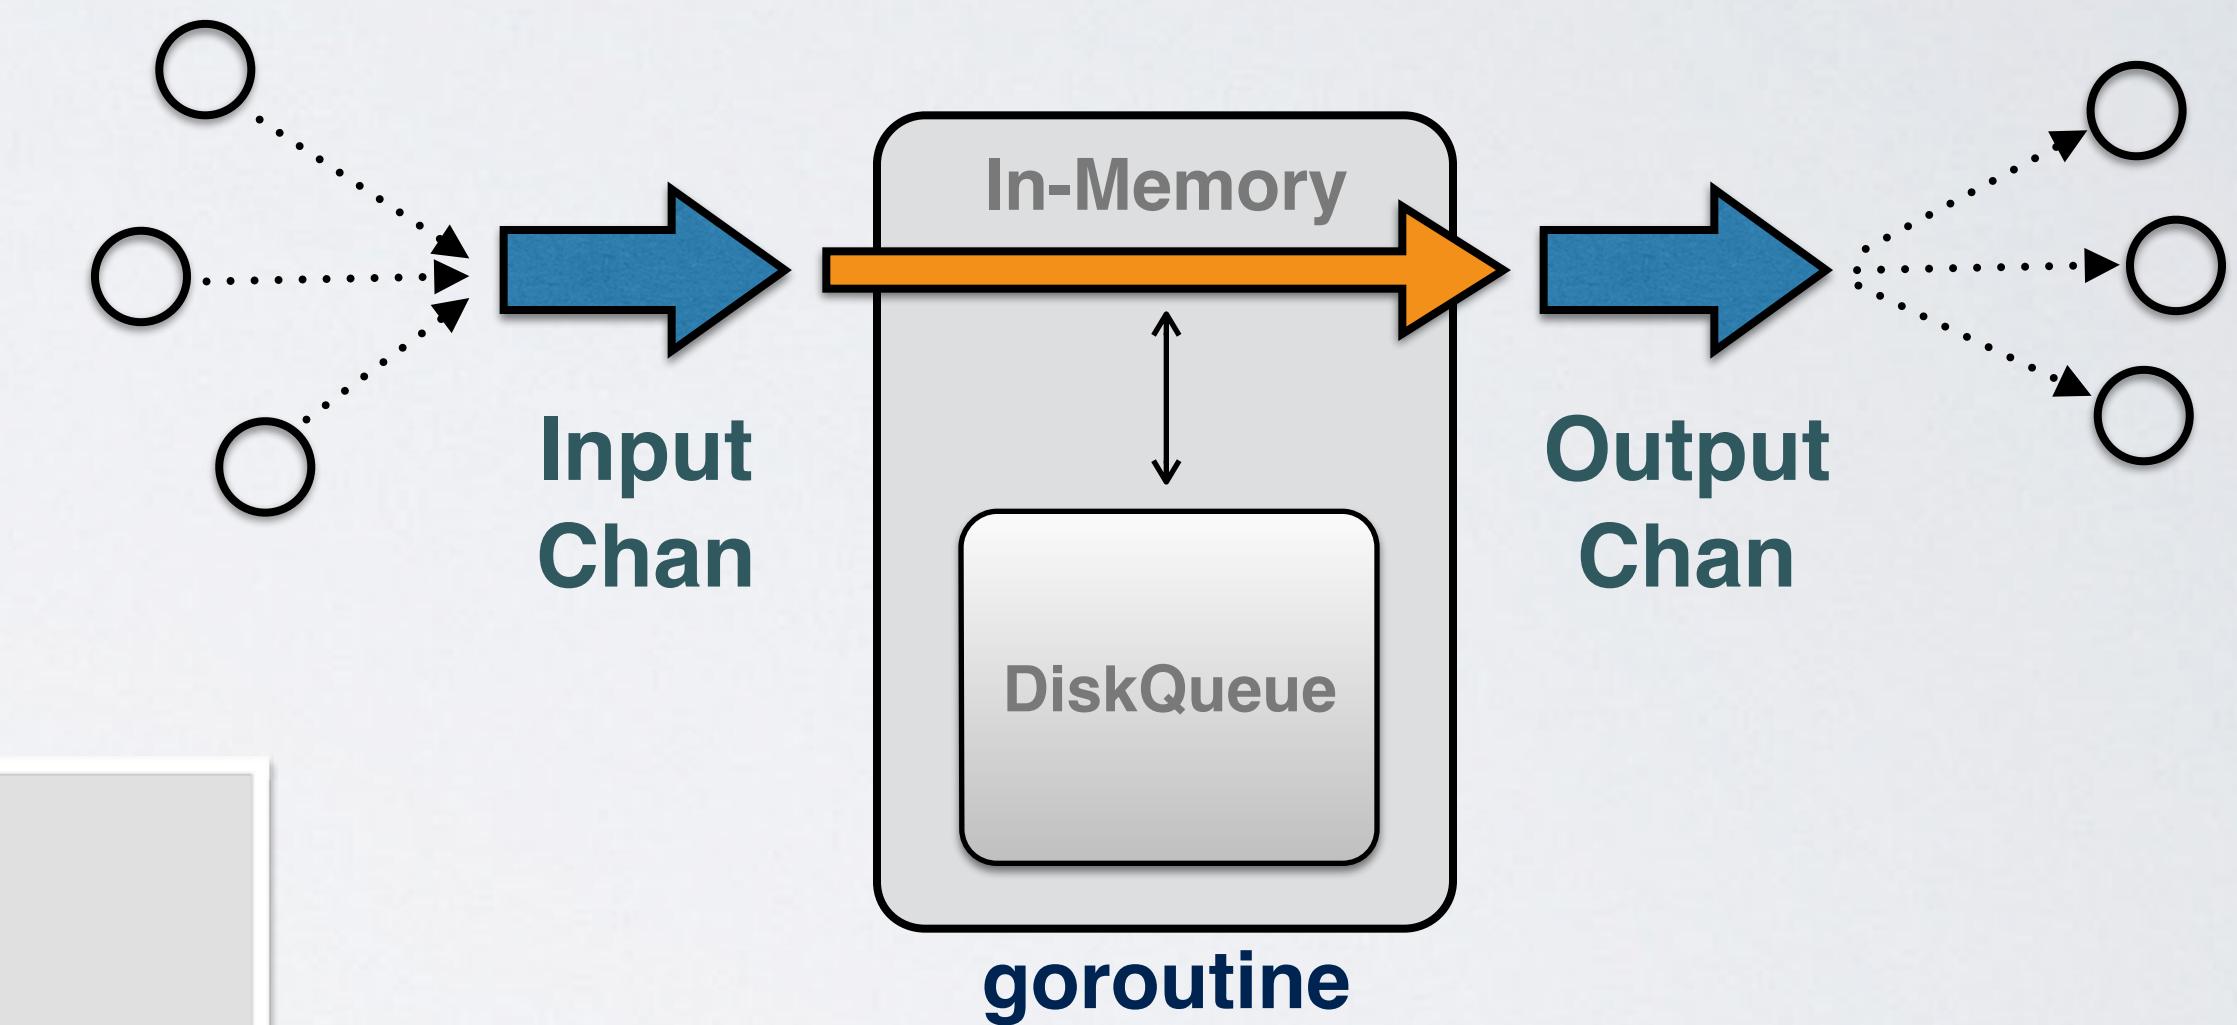
combine pubsub, distribution, and queueing

- a **topic** is a distinct stream of messages
- a **topic** has one or more **channels**
- topics and channels are created at **runtime**
- messages are **pushed** to consumers



UNDER THE HOOD

- topics and channels are *independent*
- configurable high water mark (disk persistence)
- “bounded” memory footprint

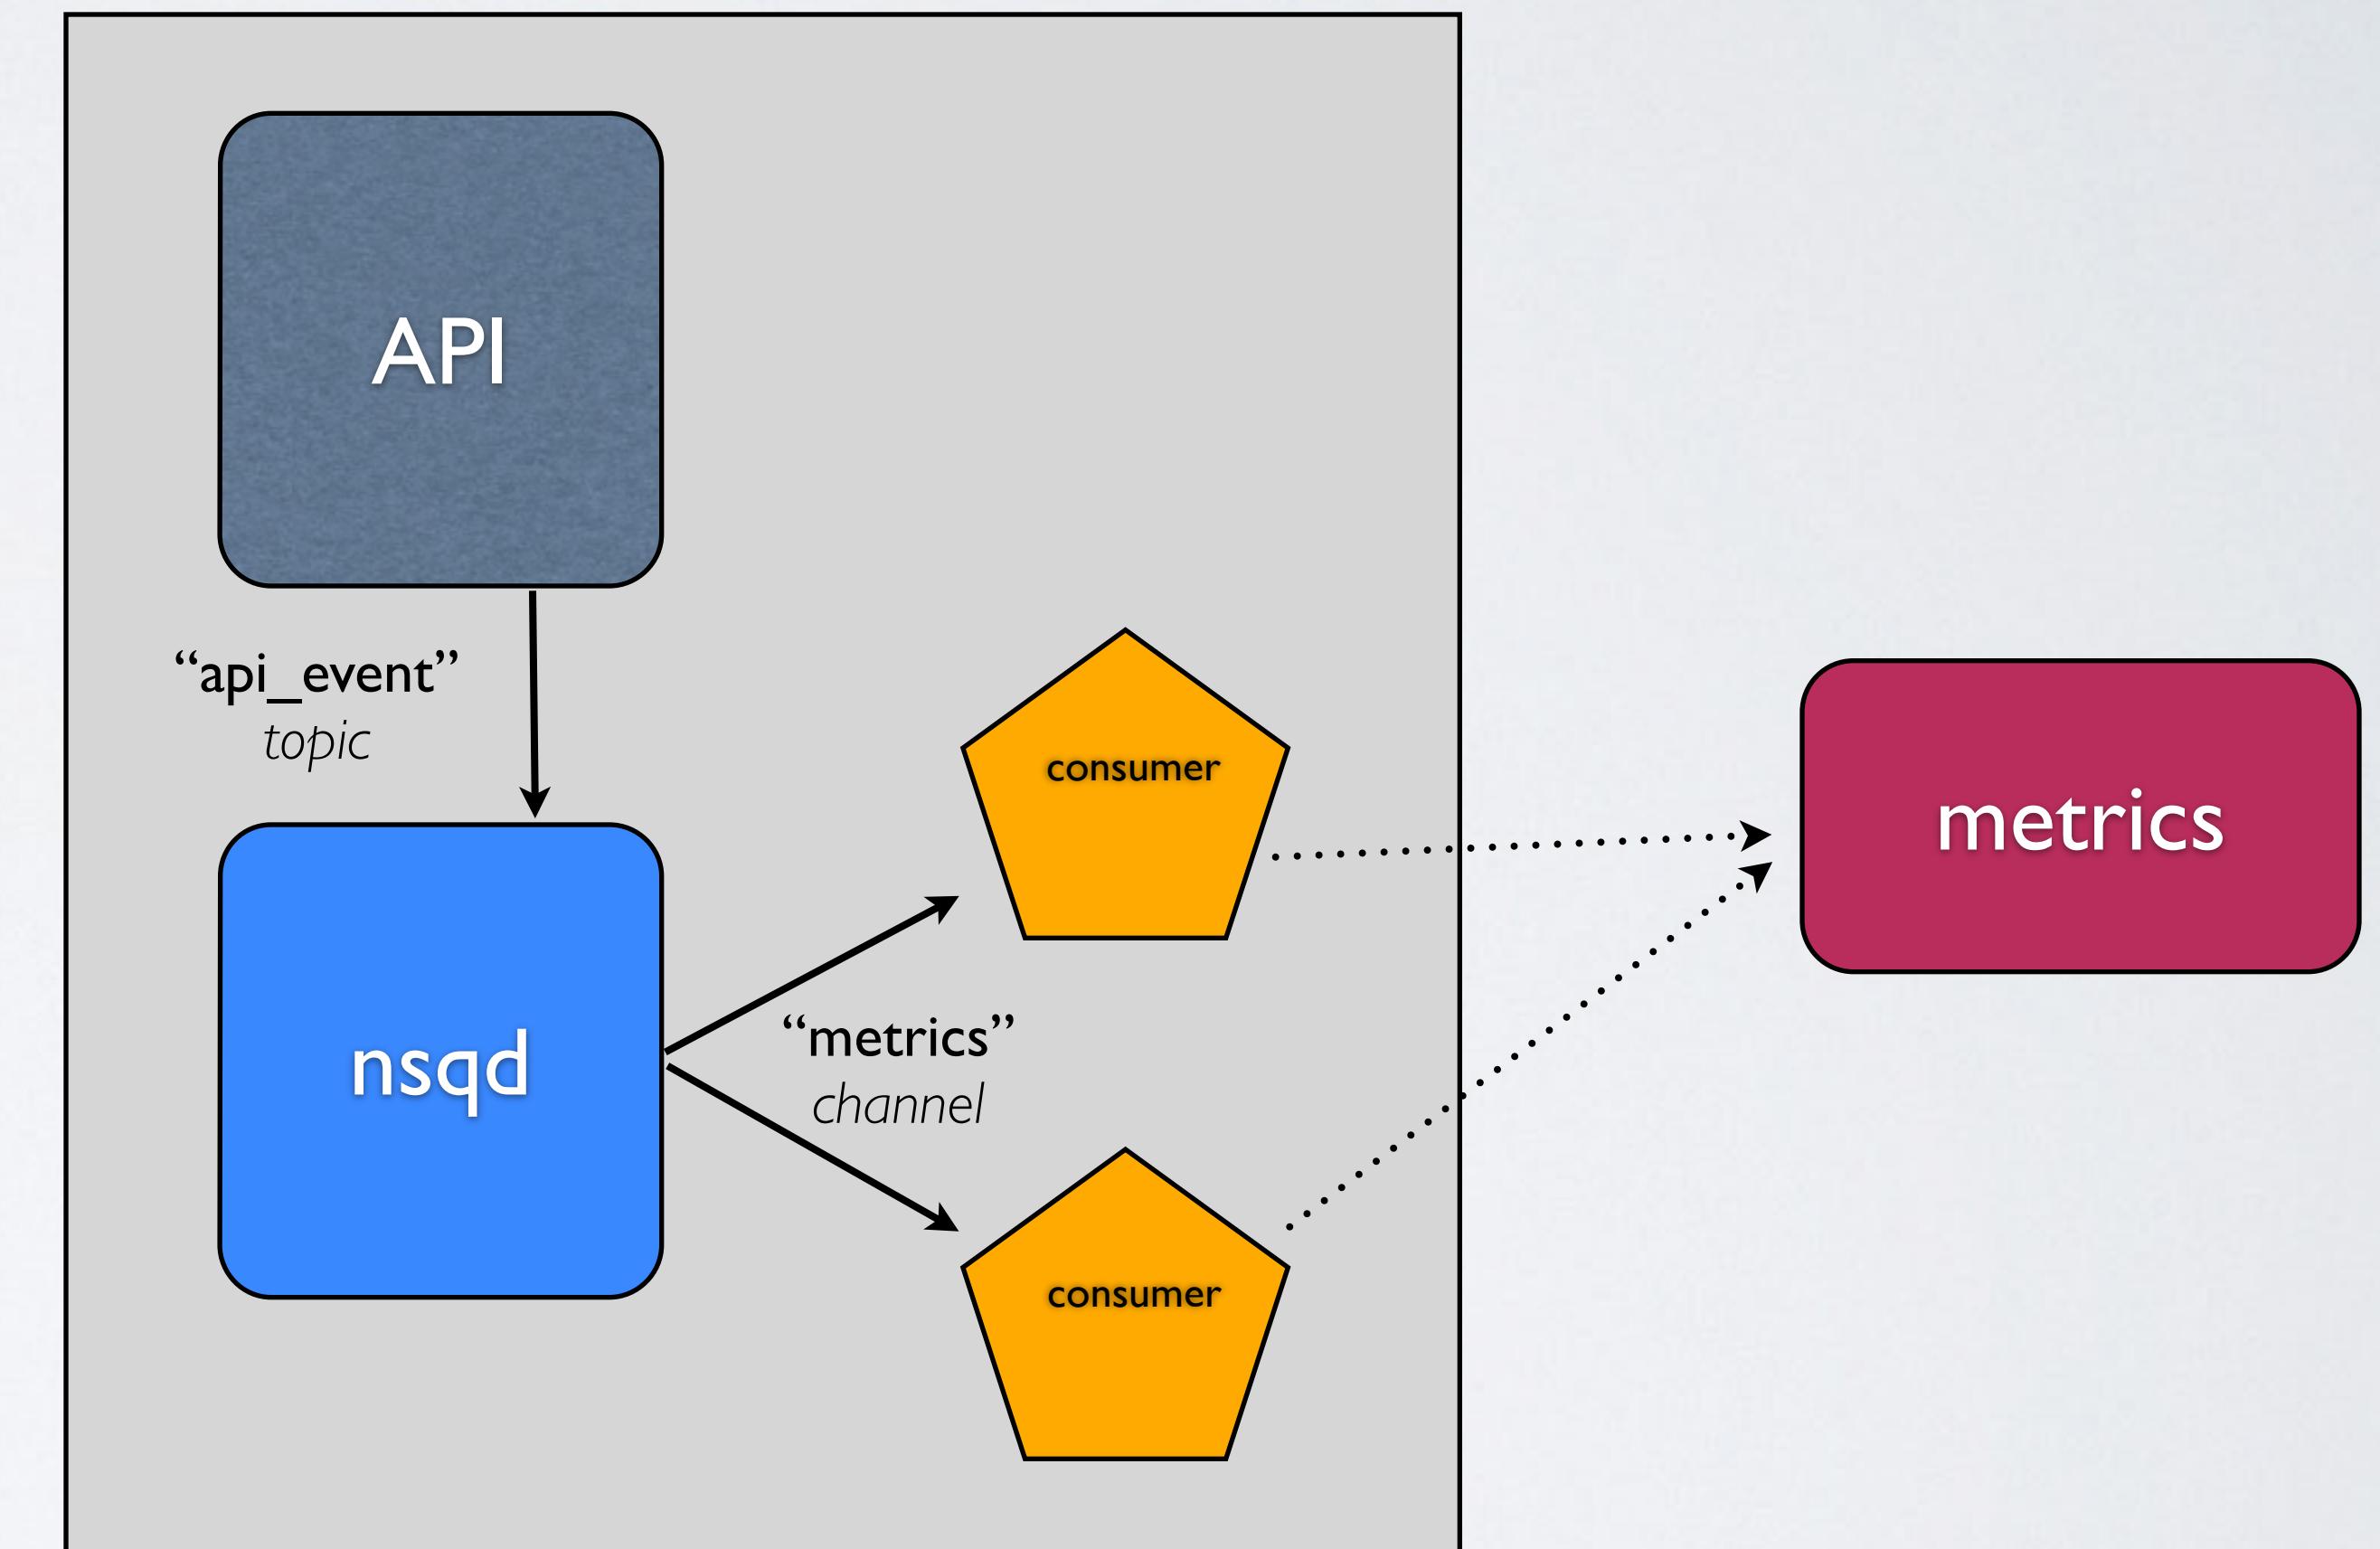


```
for msg := range c.incomingMsgChan {  
    select {  
        case c.memoryMsgChan <- msg:  
        default:  
            err := WriteMessageToBackend(&msgBuf, msg, c)  
            if err != nil {  
                // log whatever  
            }  
    }  
}
```

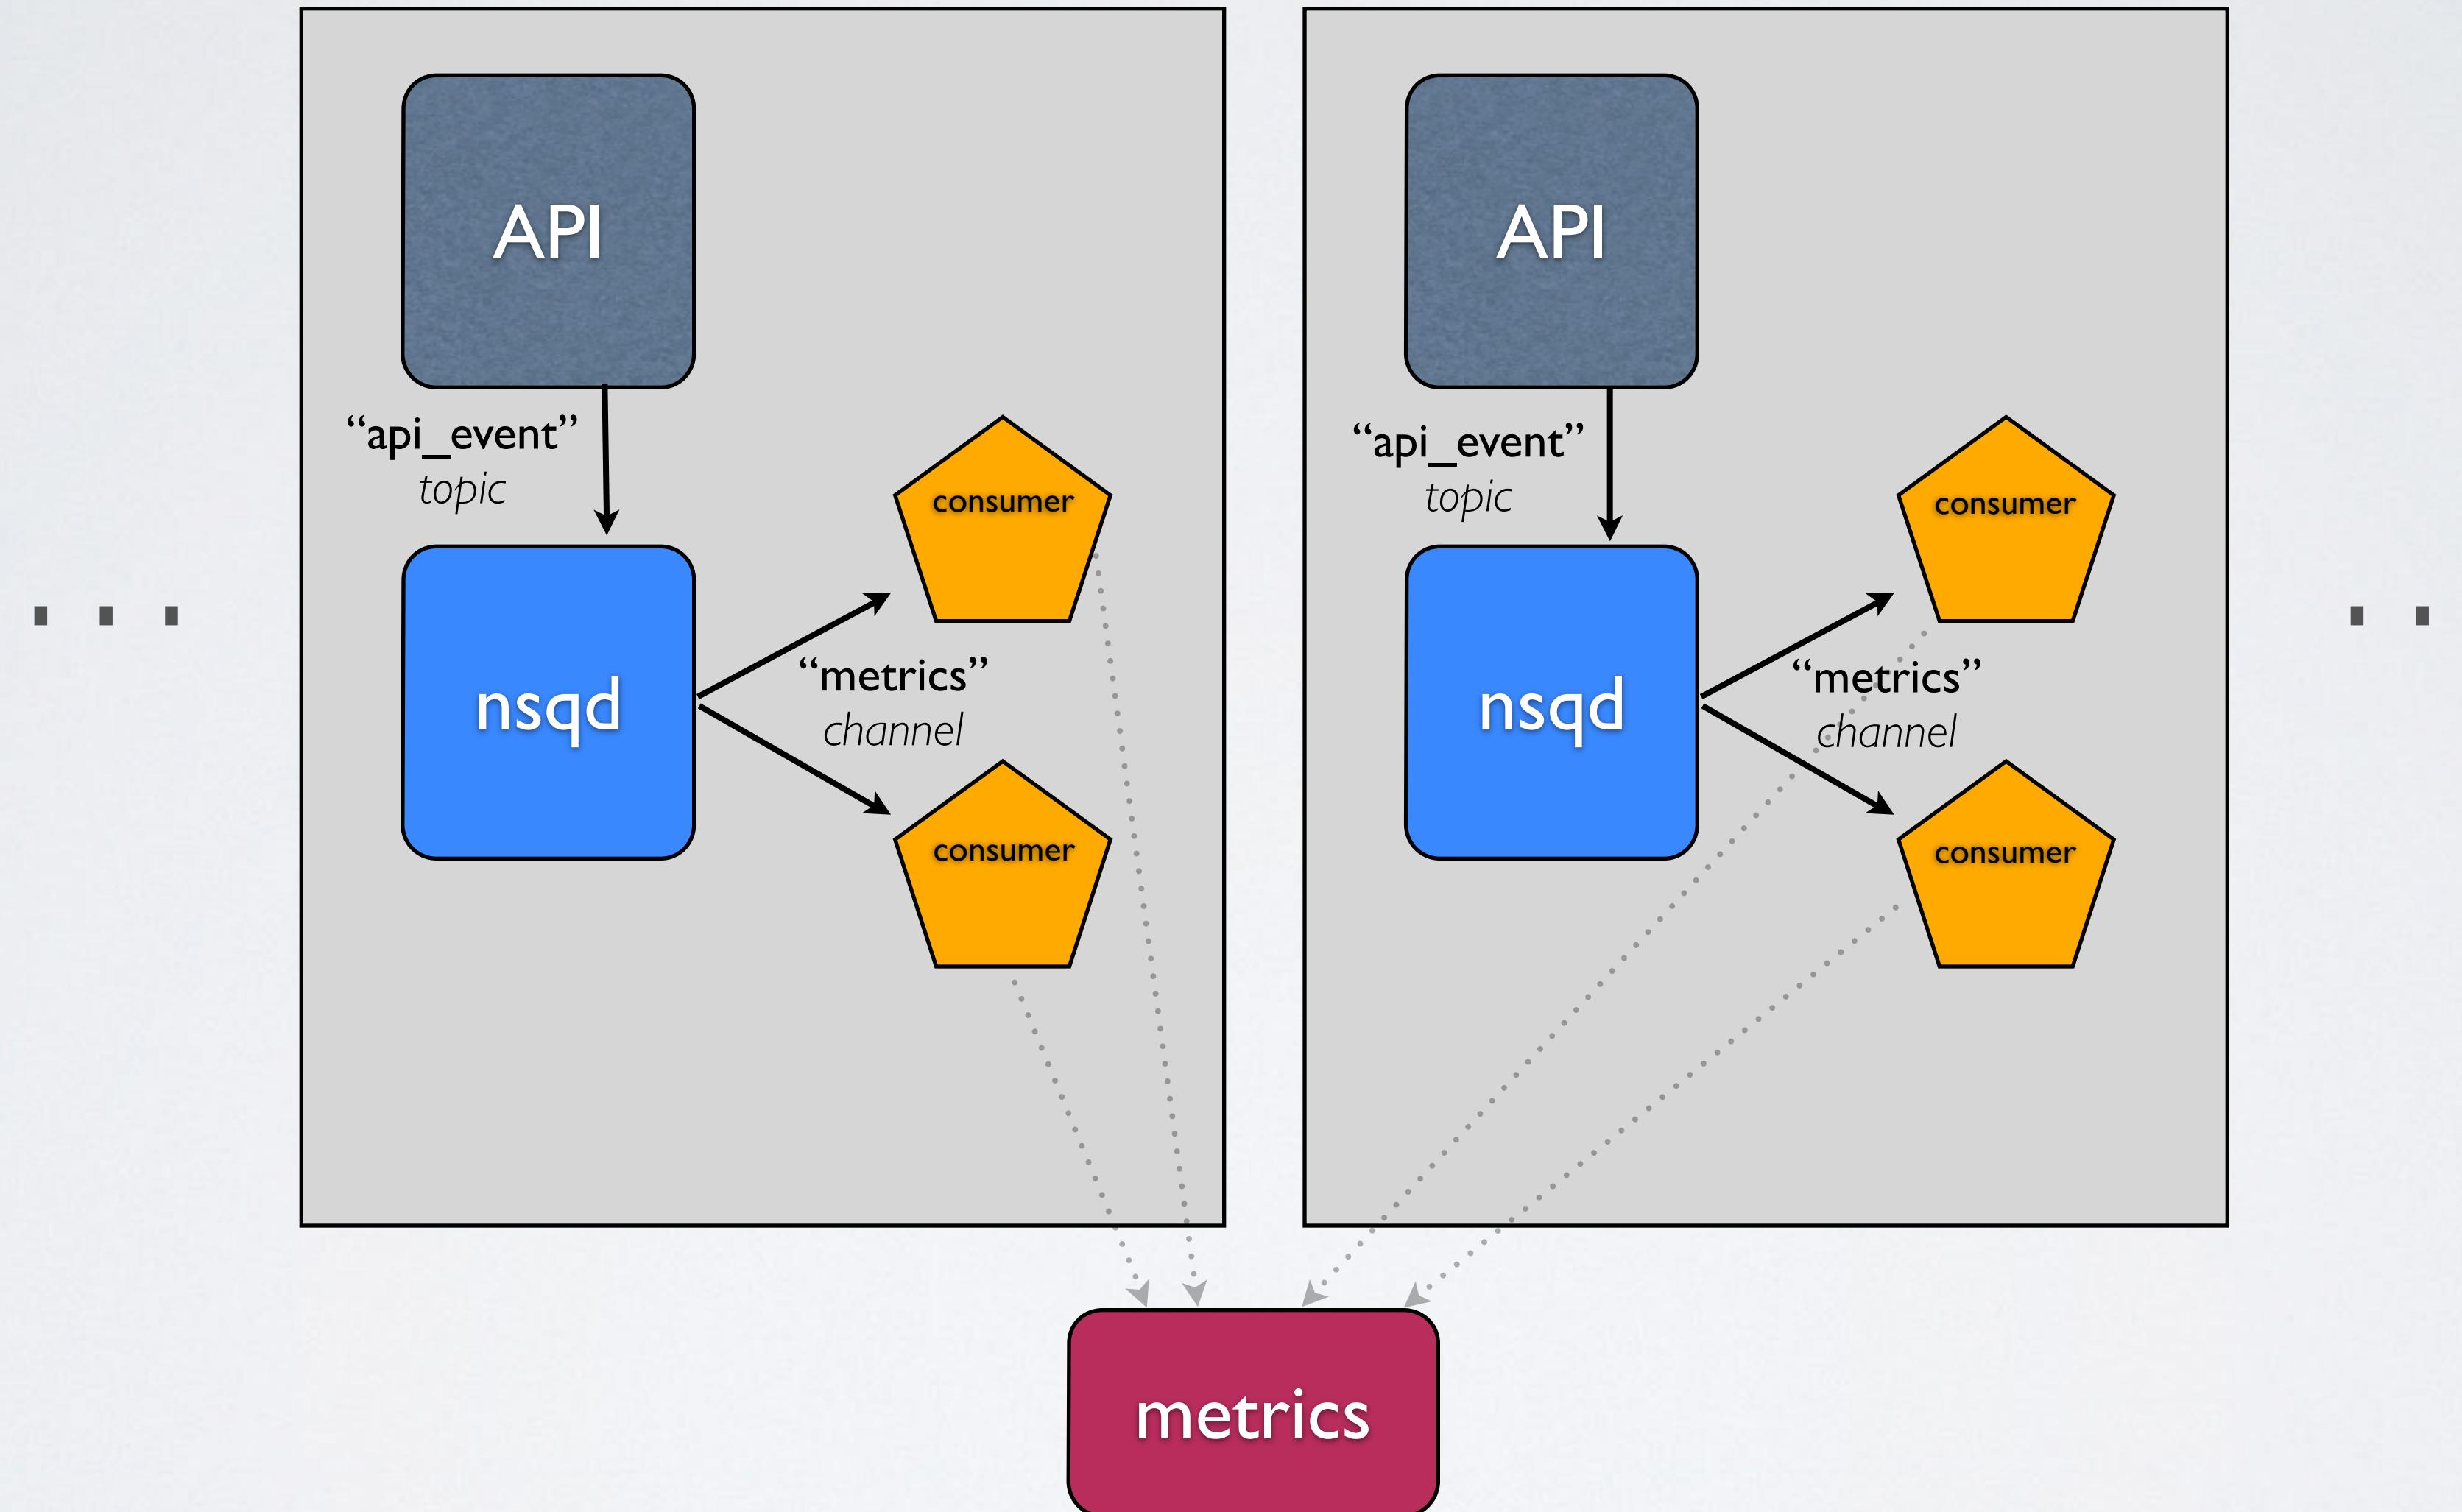
THERE AND BACK AGAIN

DE-COUPLE

- PUB locally to nsqd via HTTP
- perform work **async**
- co-locate everything (**silo**)



SCALE HORIZONTALLY

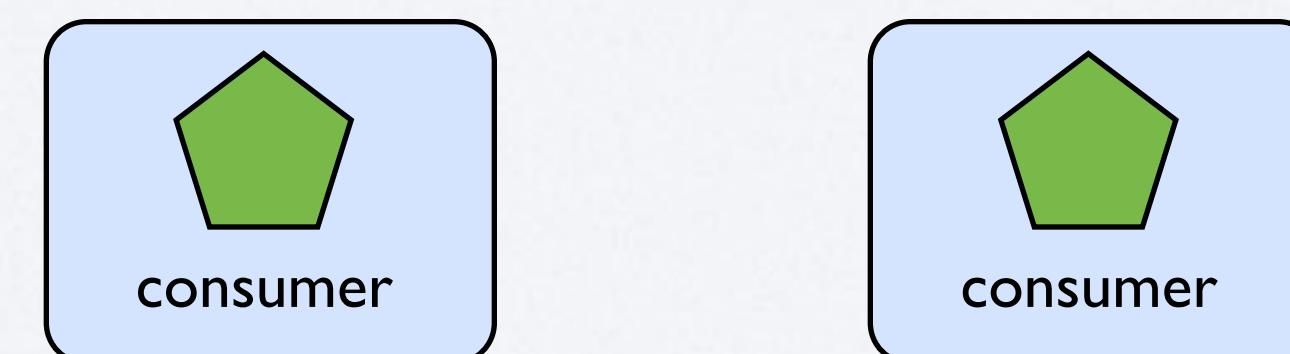
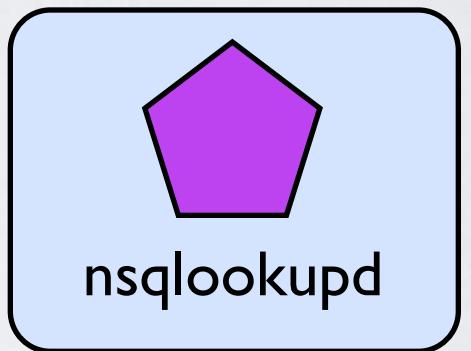
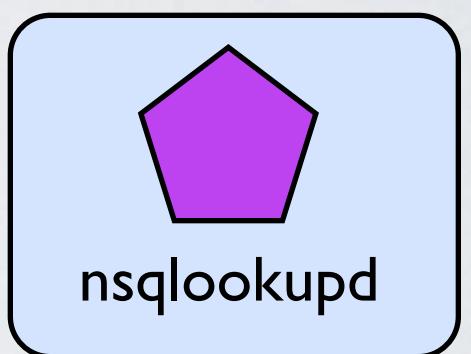
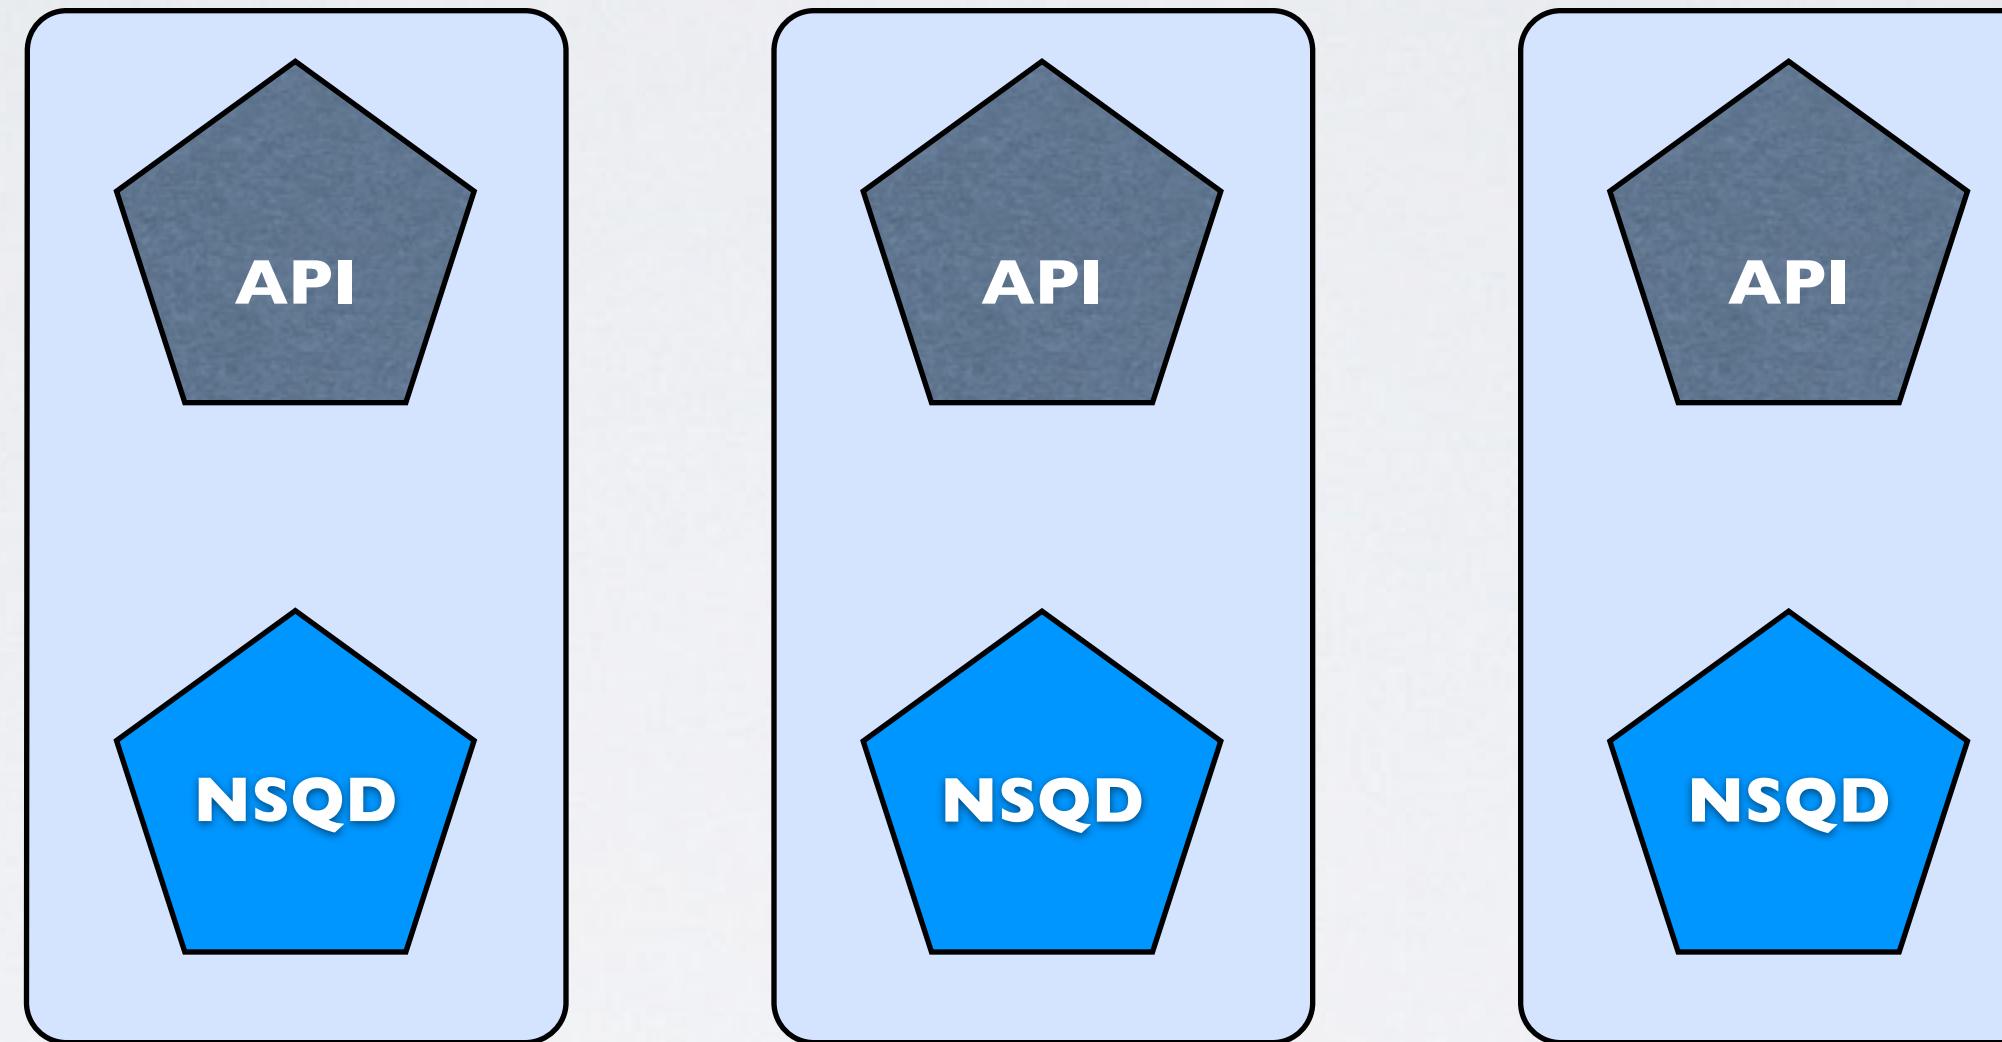


A MESSAGE QUEUE
IS BORING
(IN ISOLATION)

NSQLOOKUPD

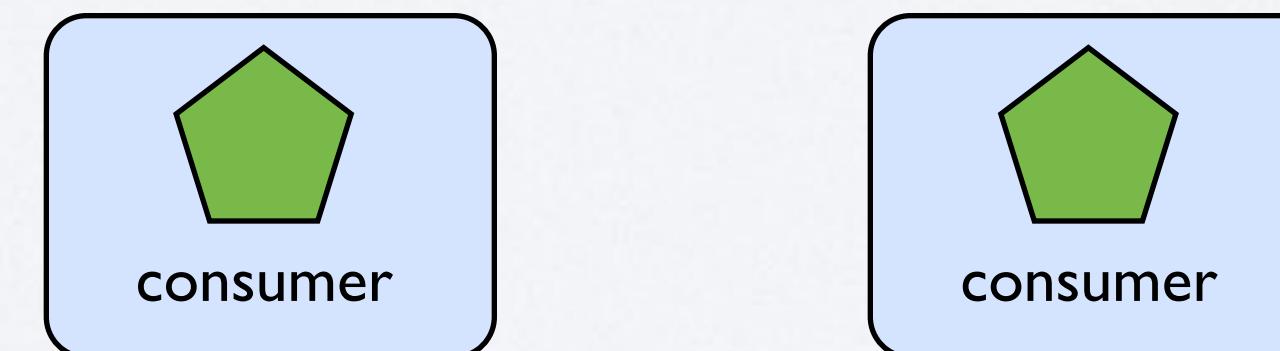
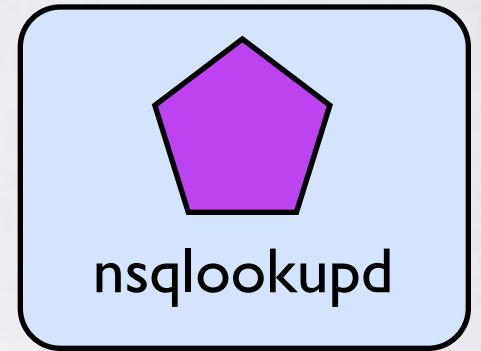
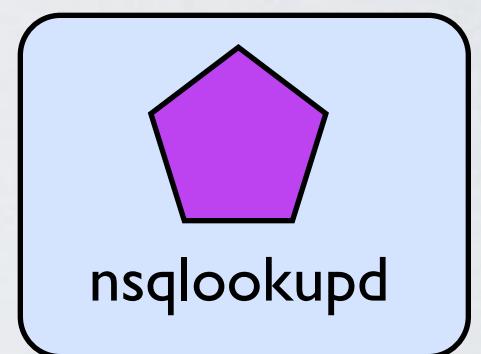
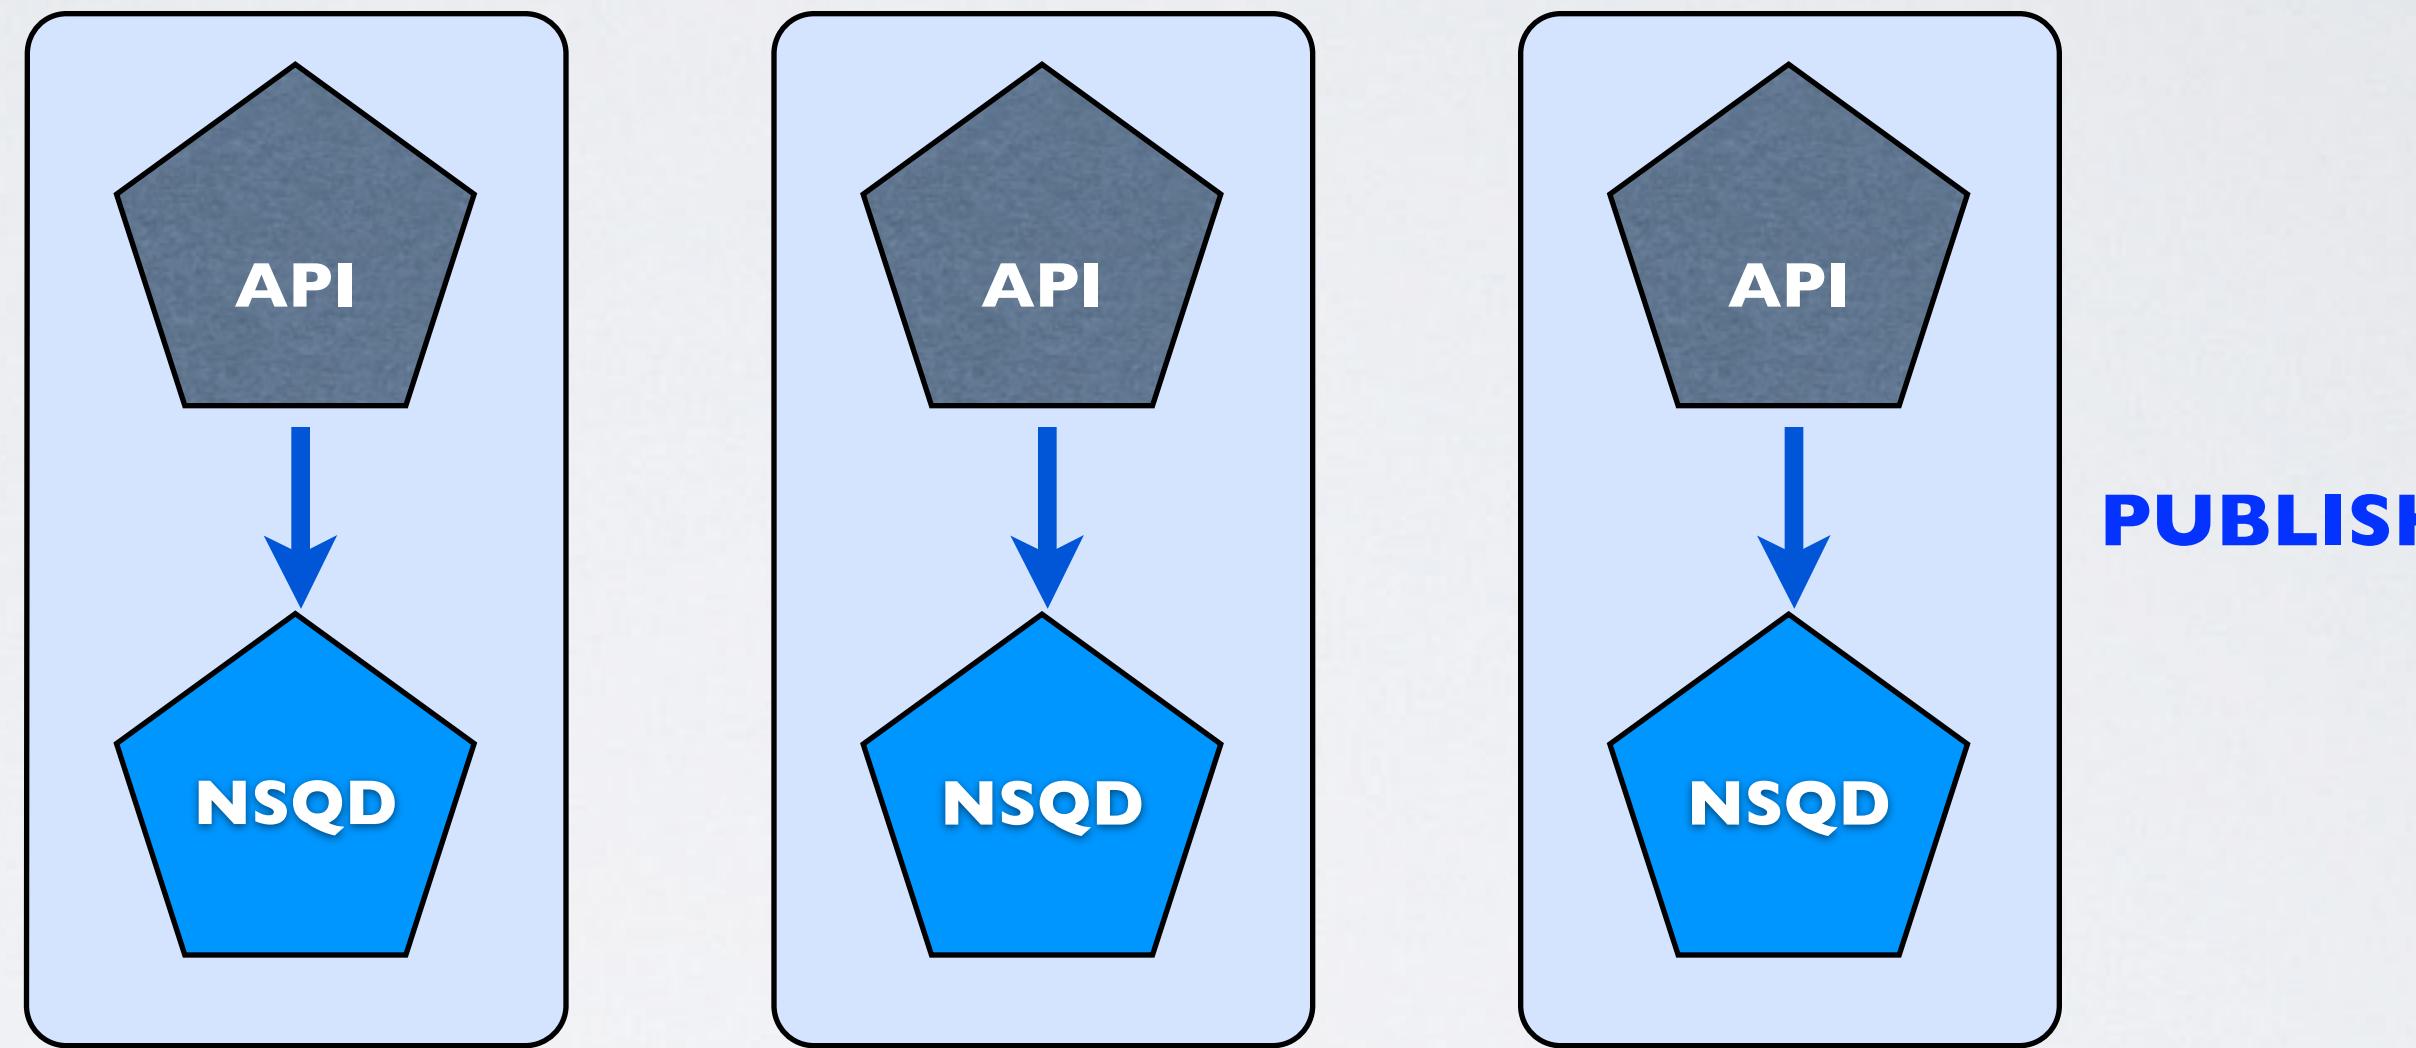
TYPICAL NSQ CLUSTER

- enable *distributed* and *decentralized* topologies
- no centralized broker
- **nsqlookupd** instances are *independent* (**no coordinatation**)



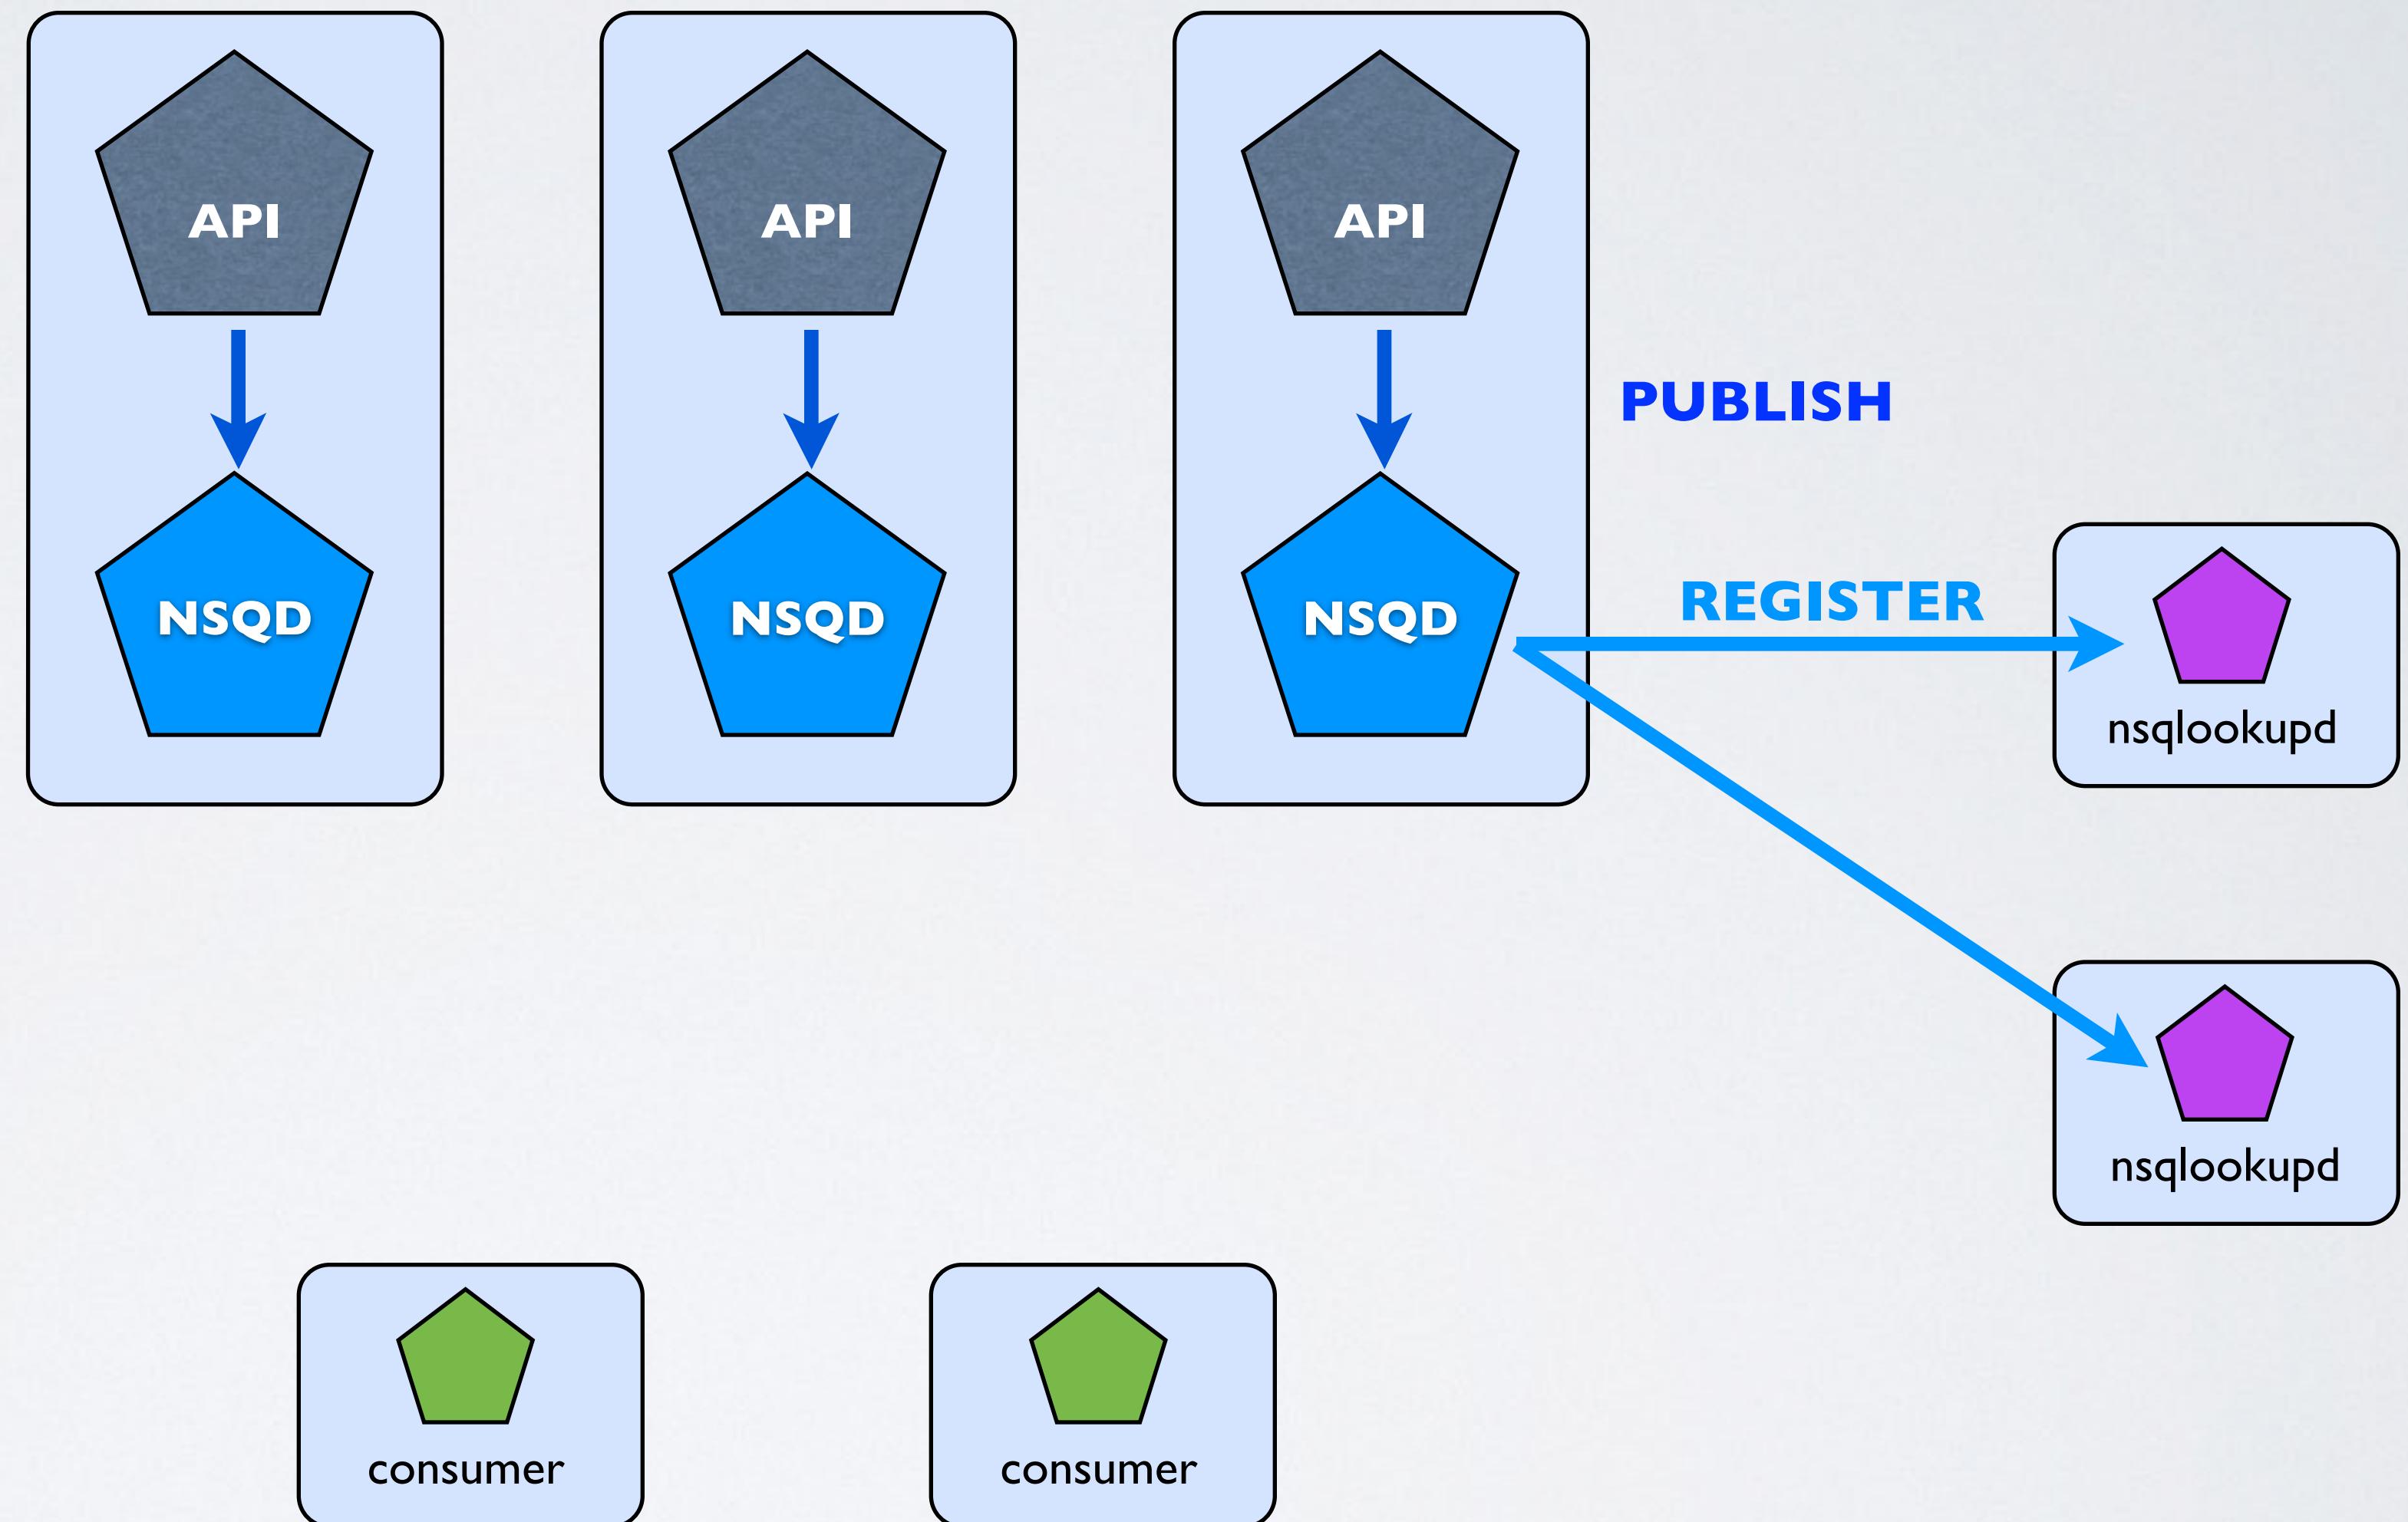
TYPICAL NSQ CLUSTER

- enable *distributed* and *decentralized* topologies
- no centralized broker
- **nsqlookupd** instances are *independent* (**no coordinatation**)



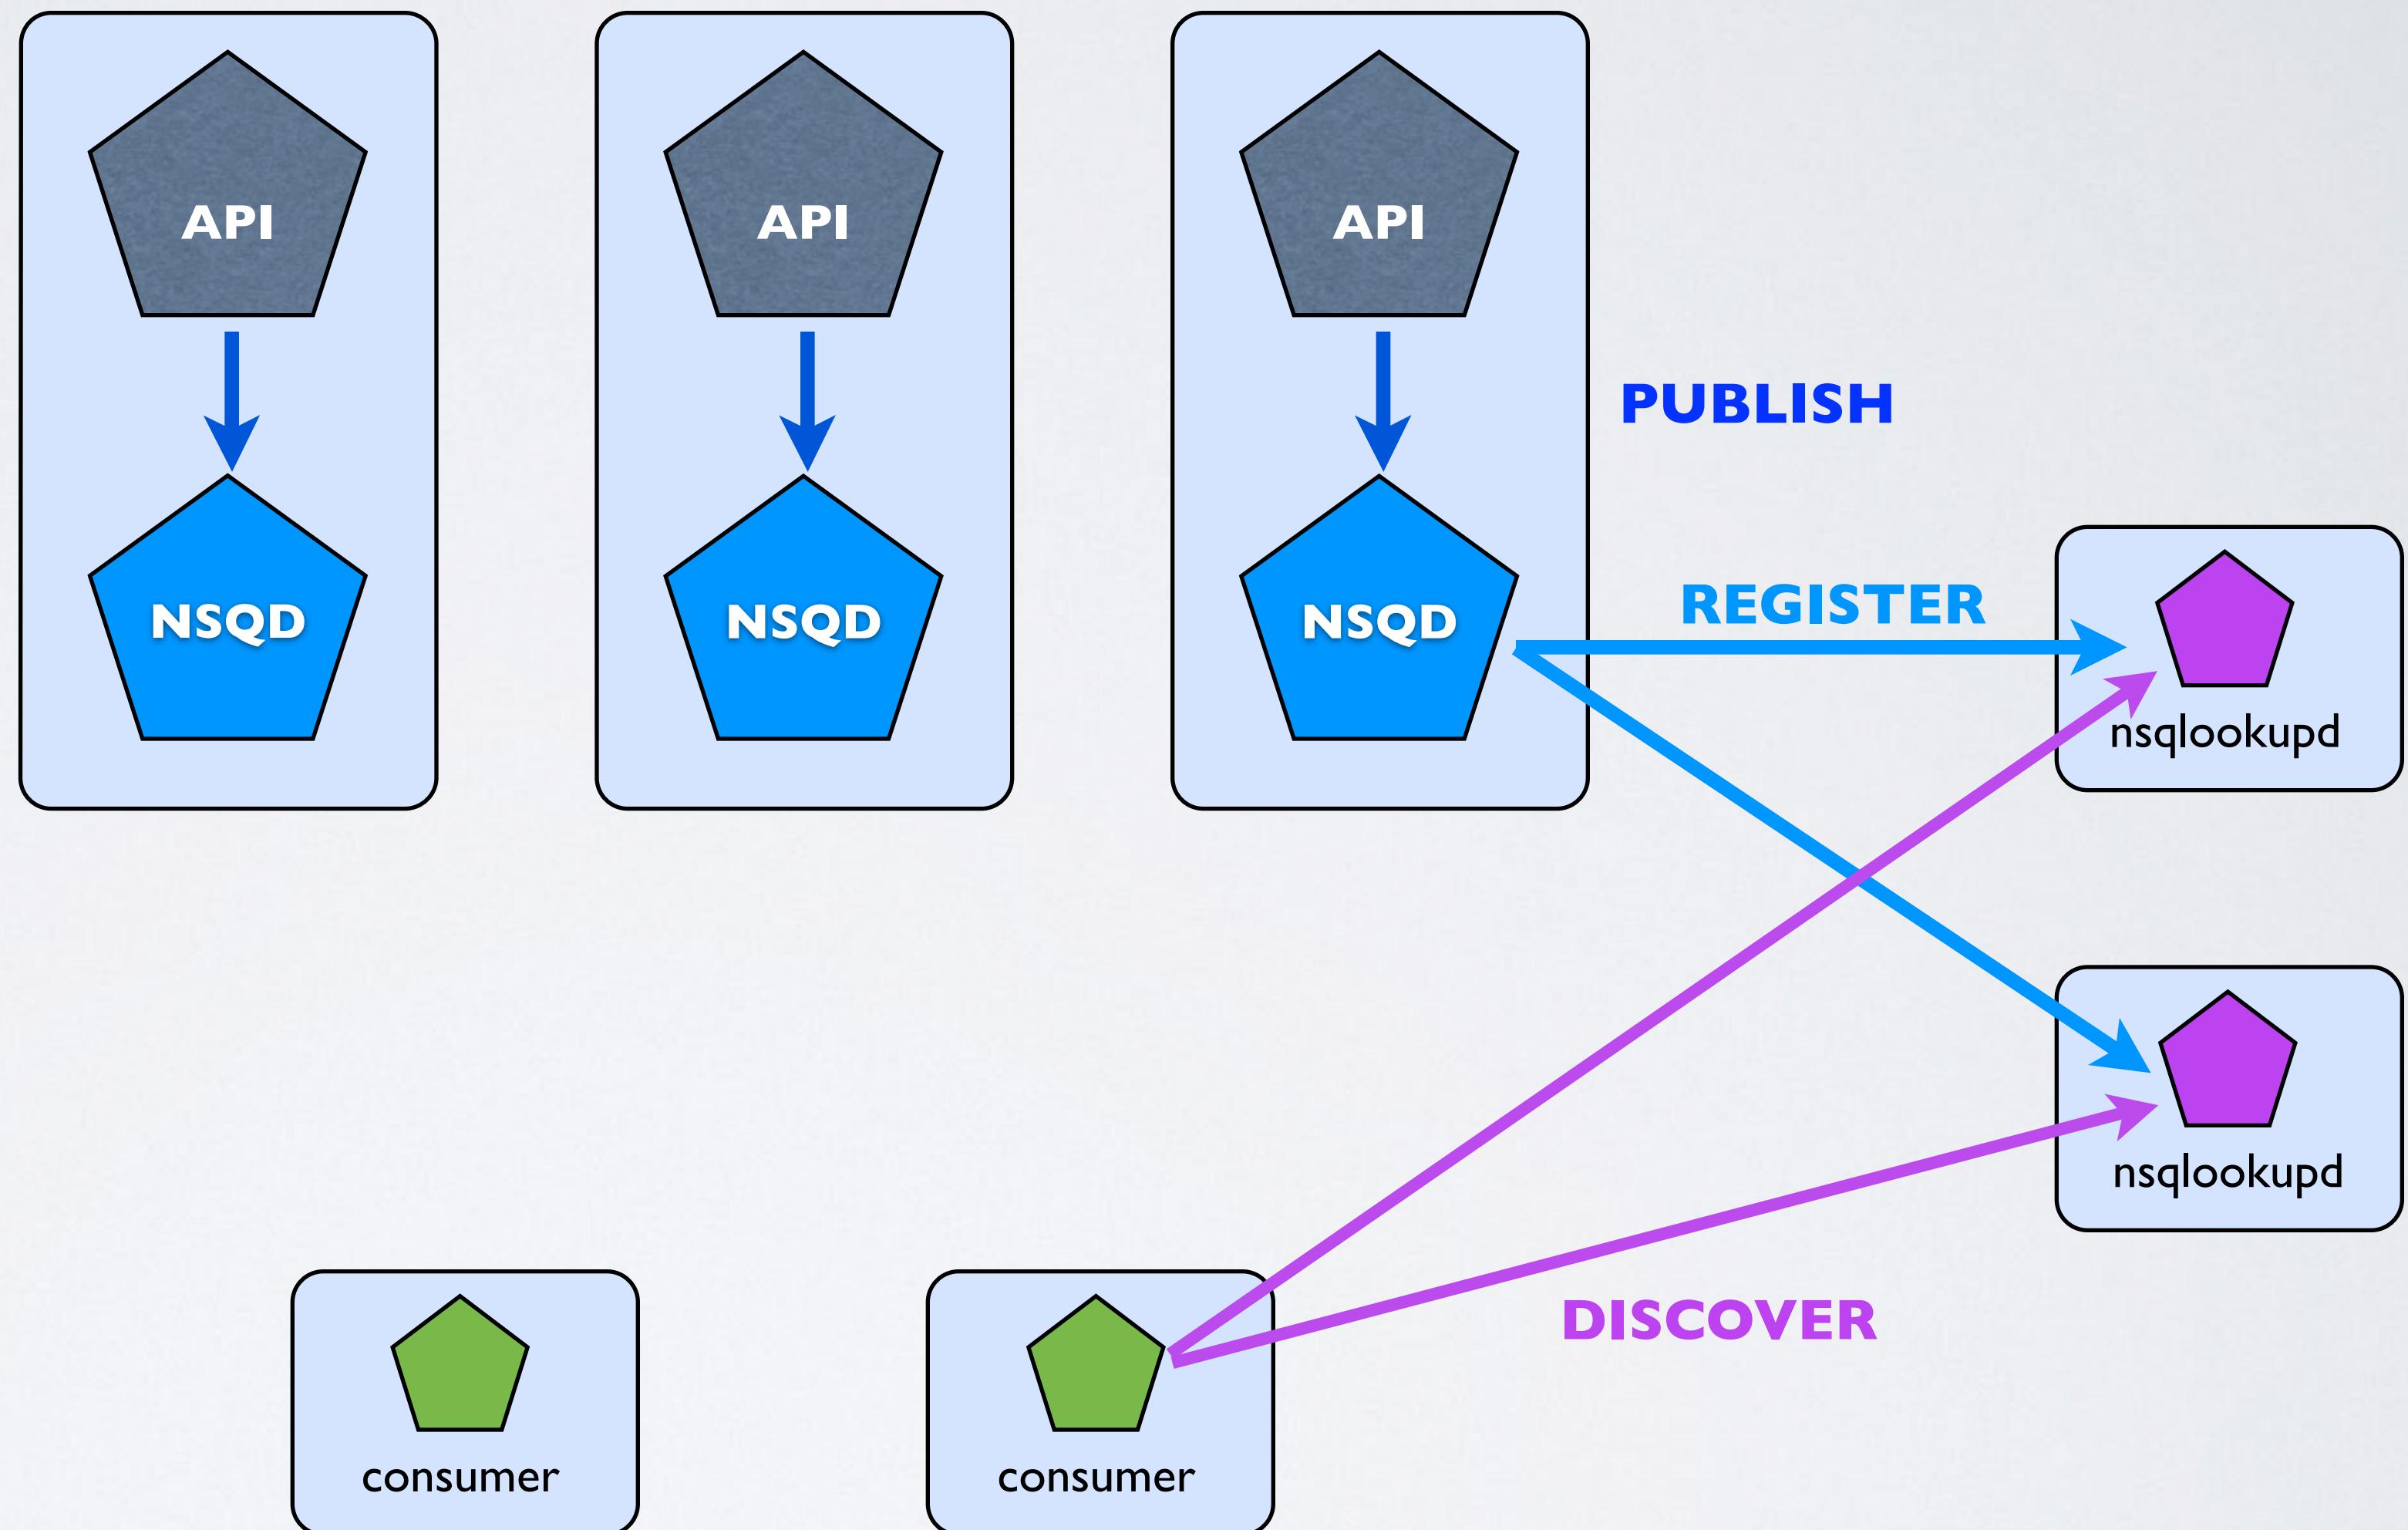
TYPICAL NSQ CLUSTER

- enable *distributed* and *decentralized* topologies
- no centralized broker
- nsqlookupd instances are *independent* (**no coordinatation**)



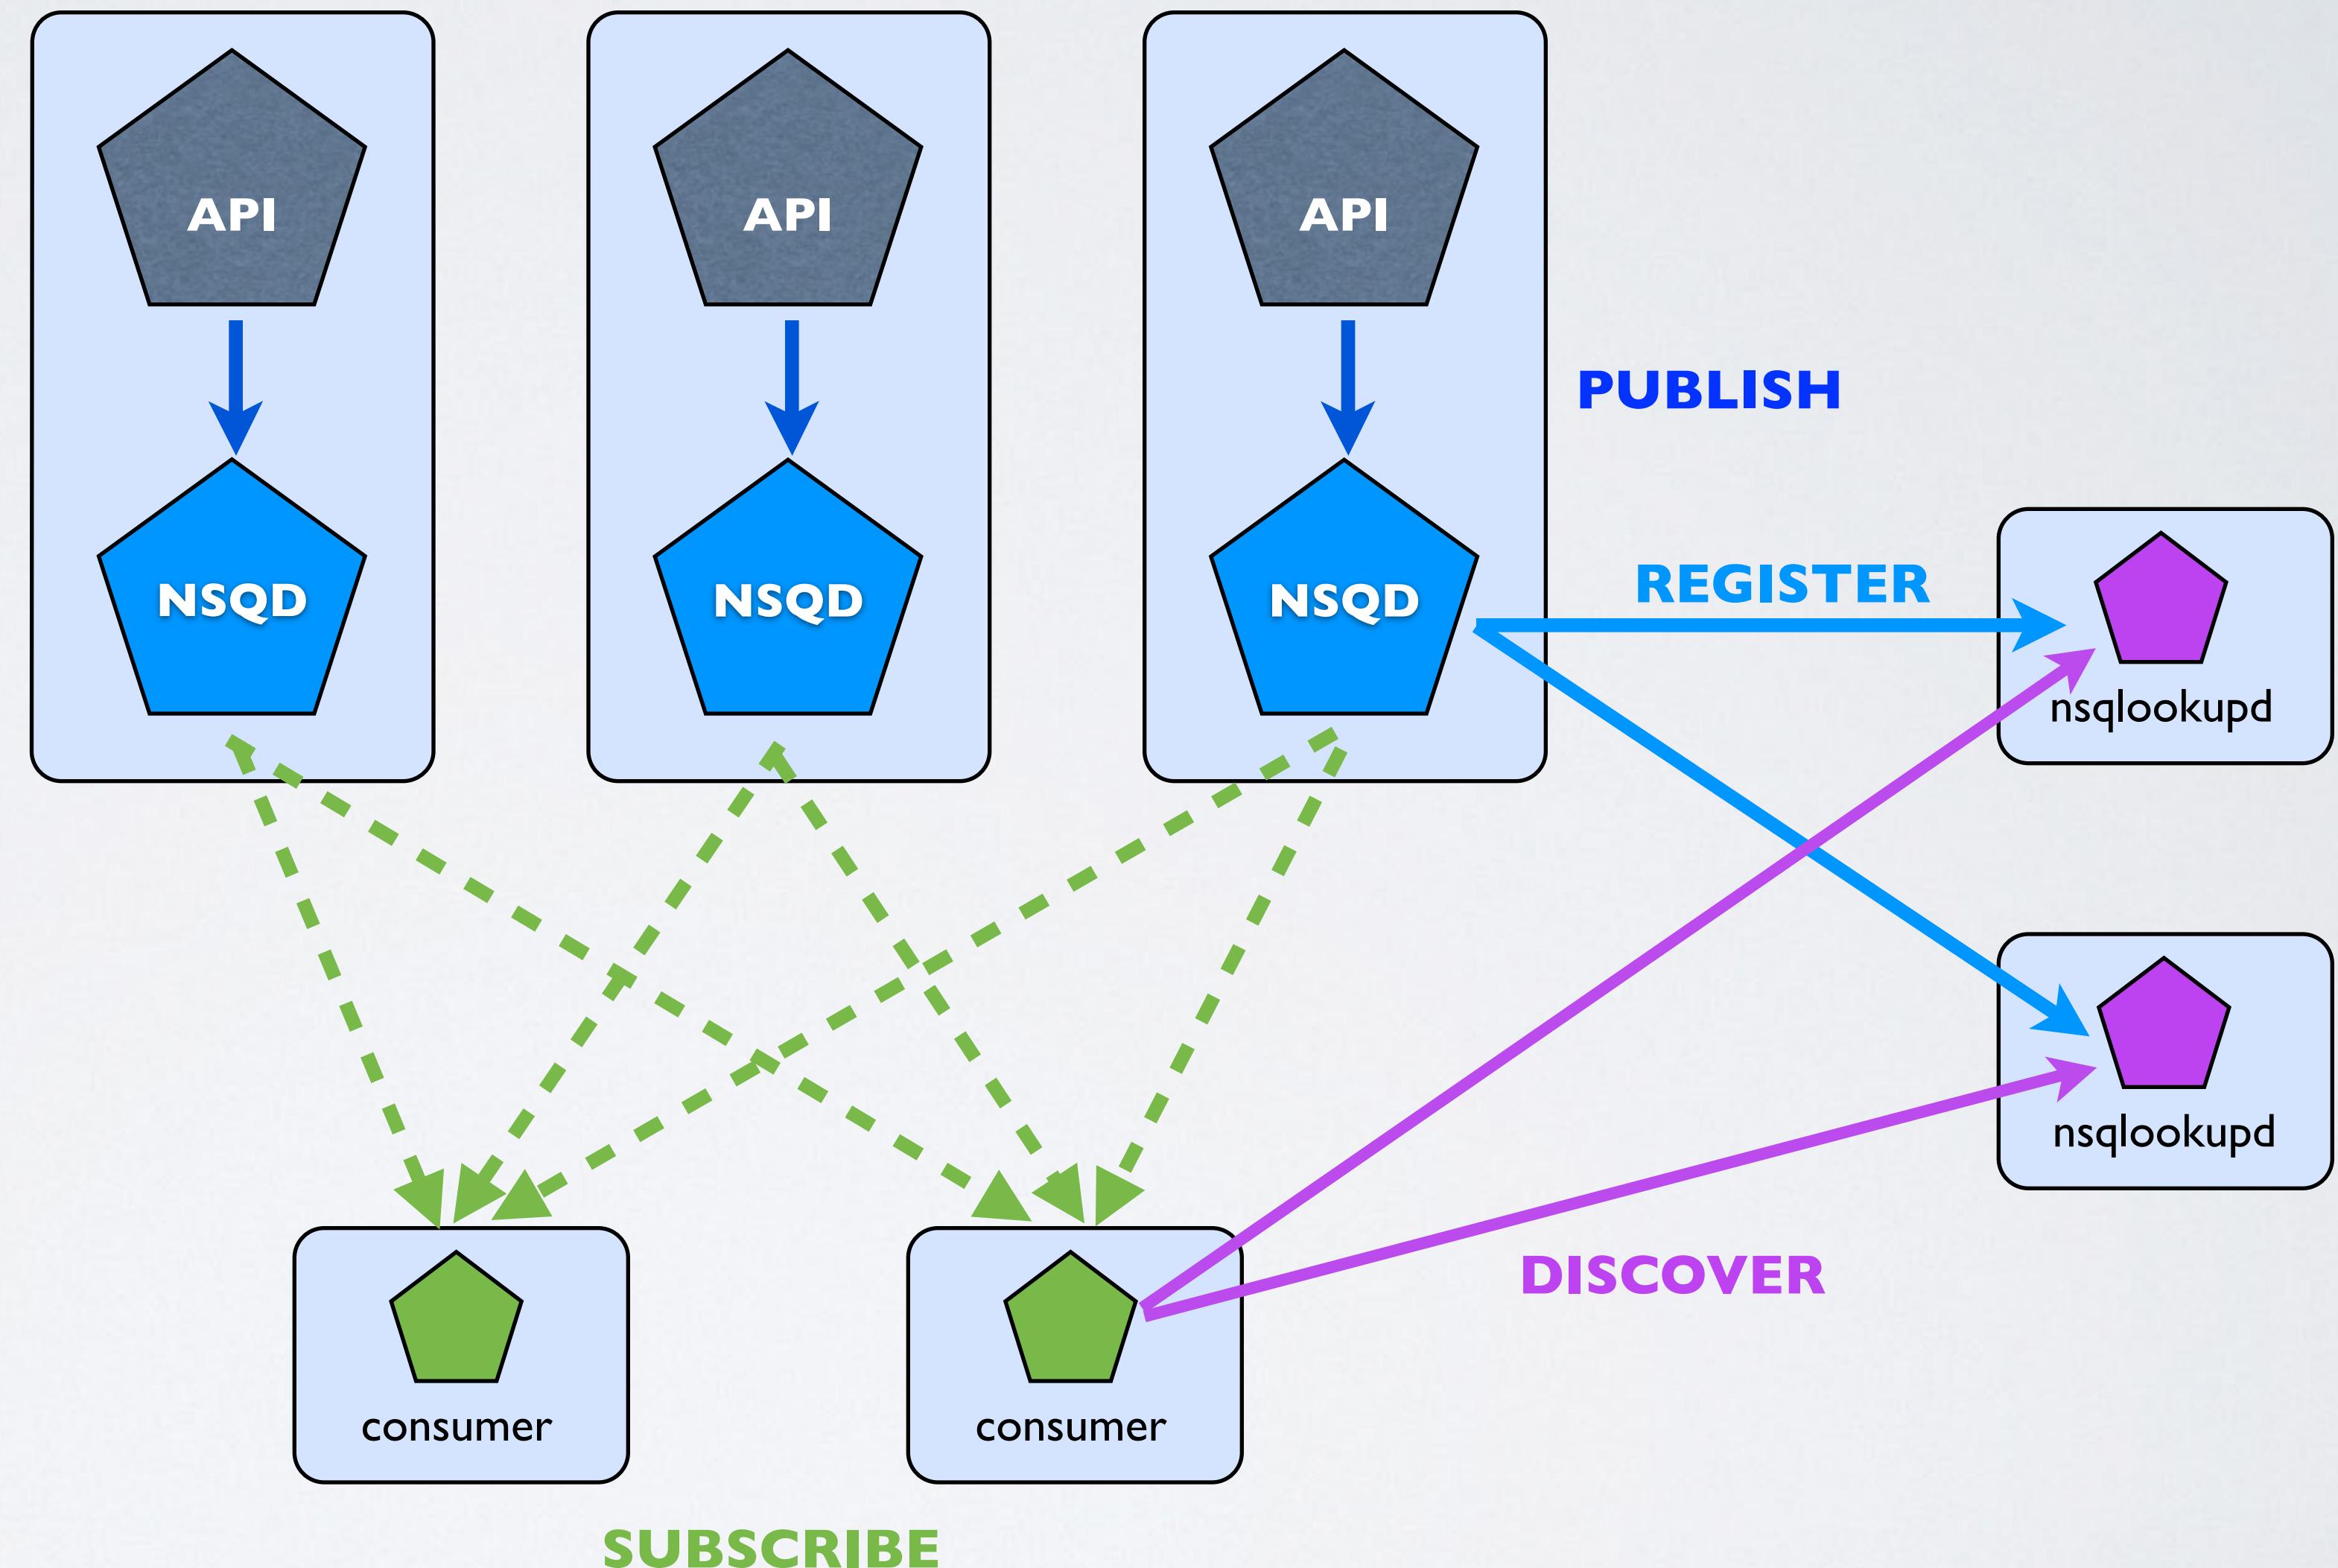
TYPICAL NSQ CLUSTER

- enable *distributed* and *decentralized* topologies
- no centralized broker
- nsqlookupd instances are *independent* (**no coordinatation**)



TYPICAL NSQ CLUSTER

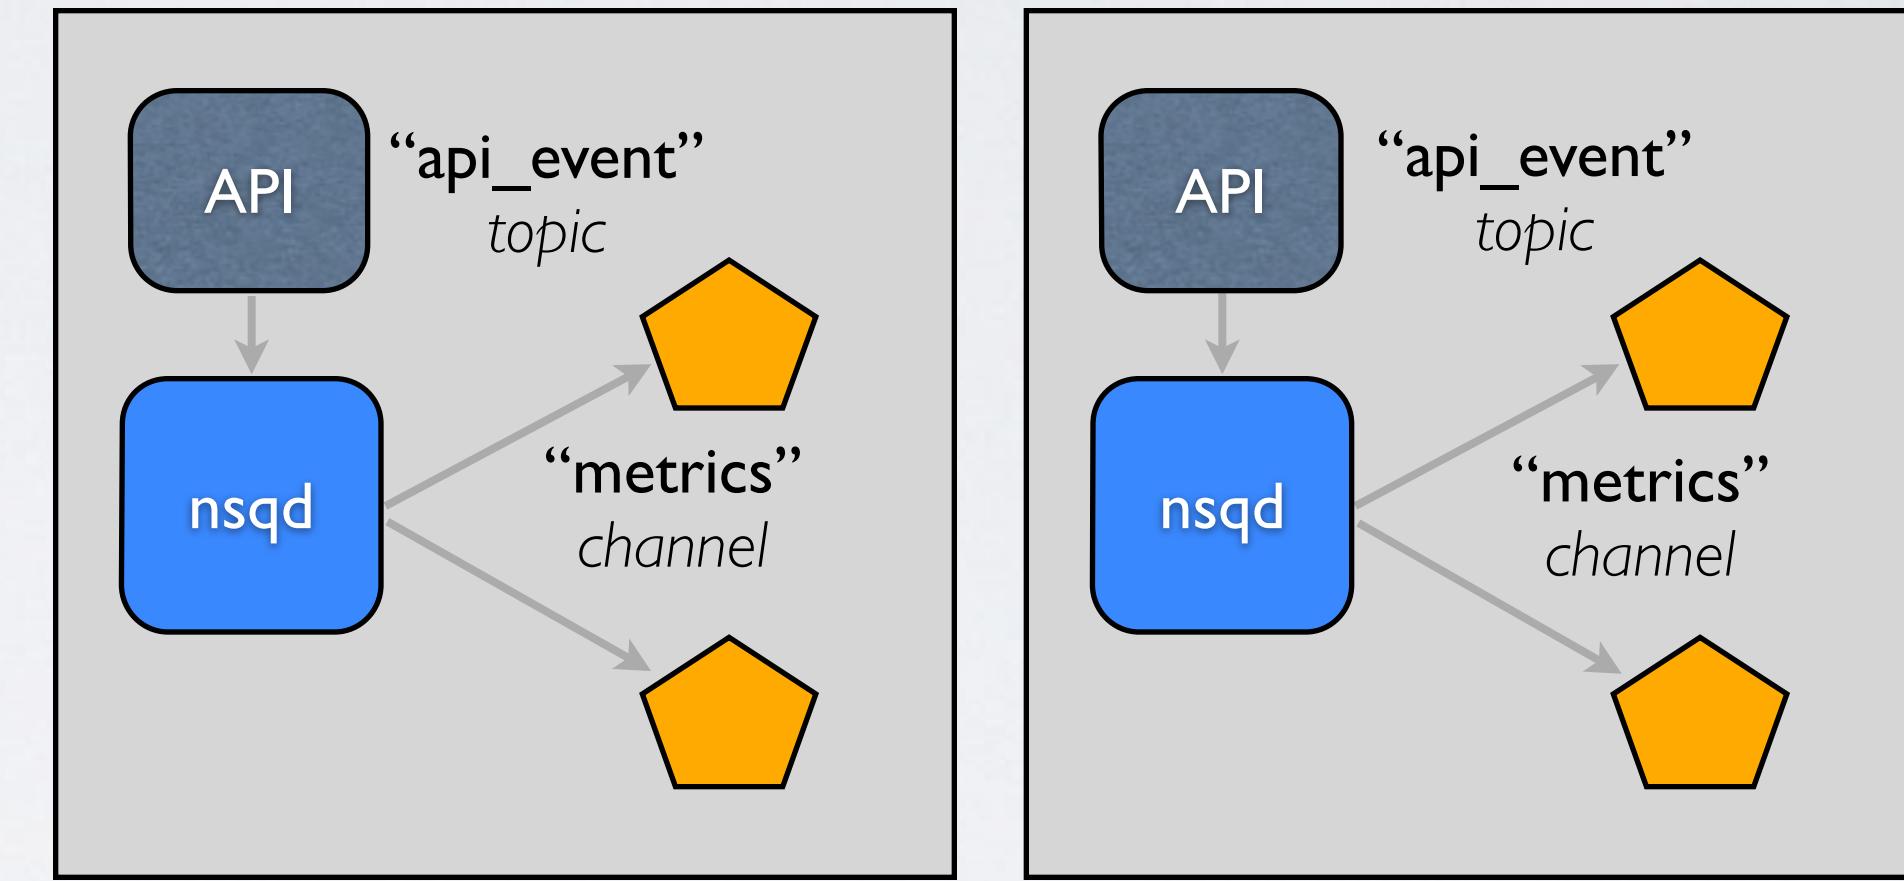
- enable *distributed* and *decentralized* topologies
- no centralized broker
- nsqlookupd instances are *independent* (**no coordination**)



... AND AGAIN

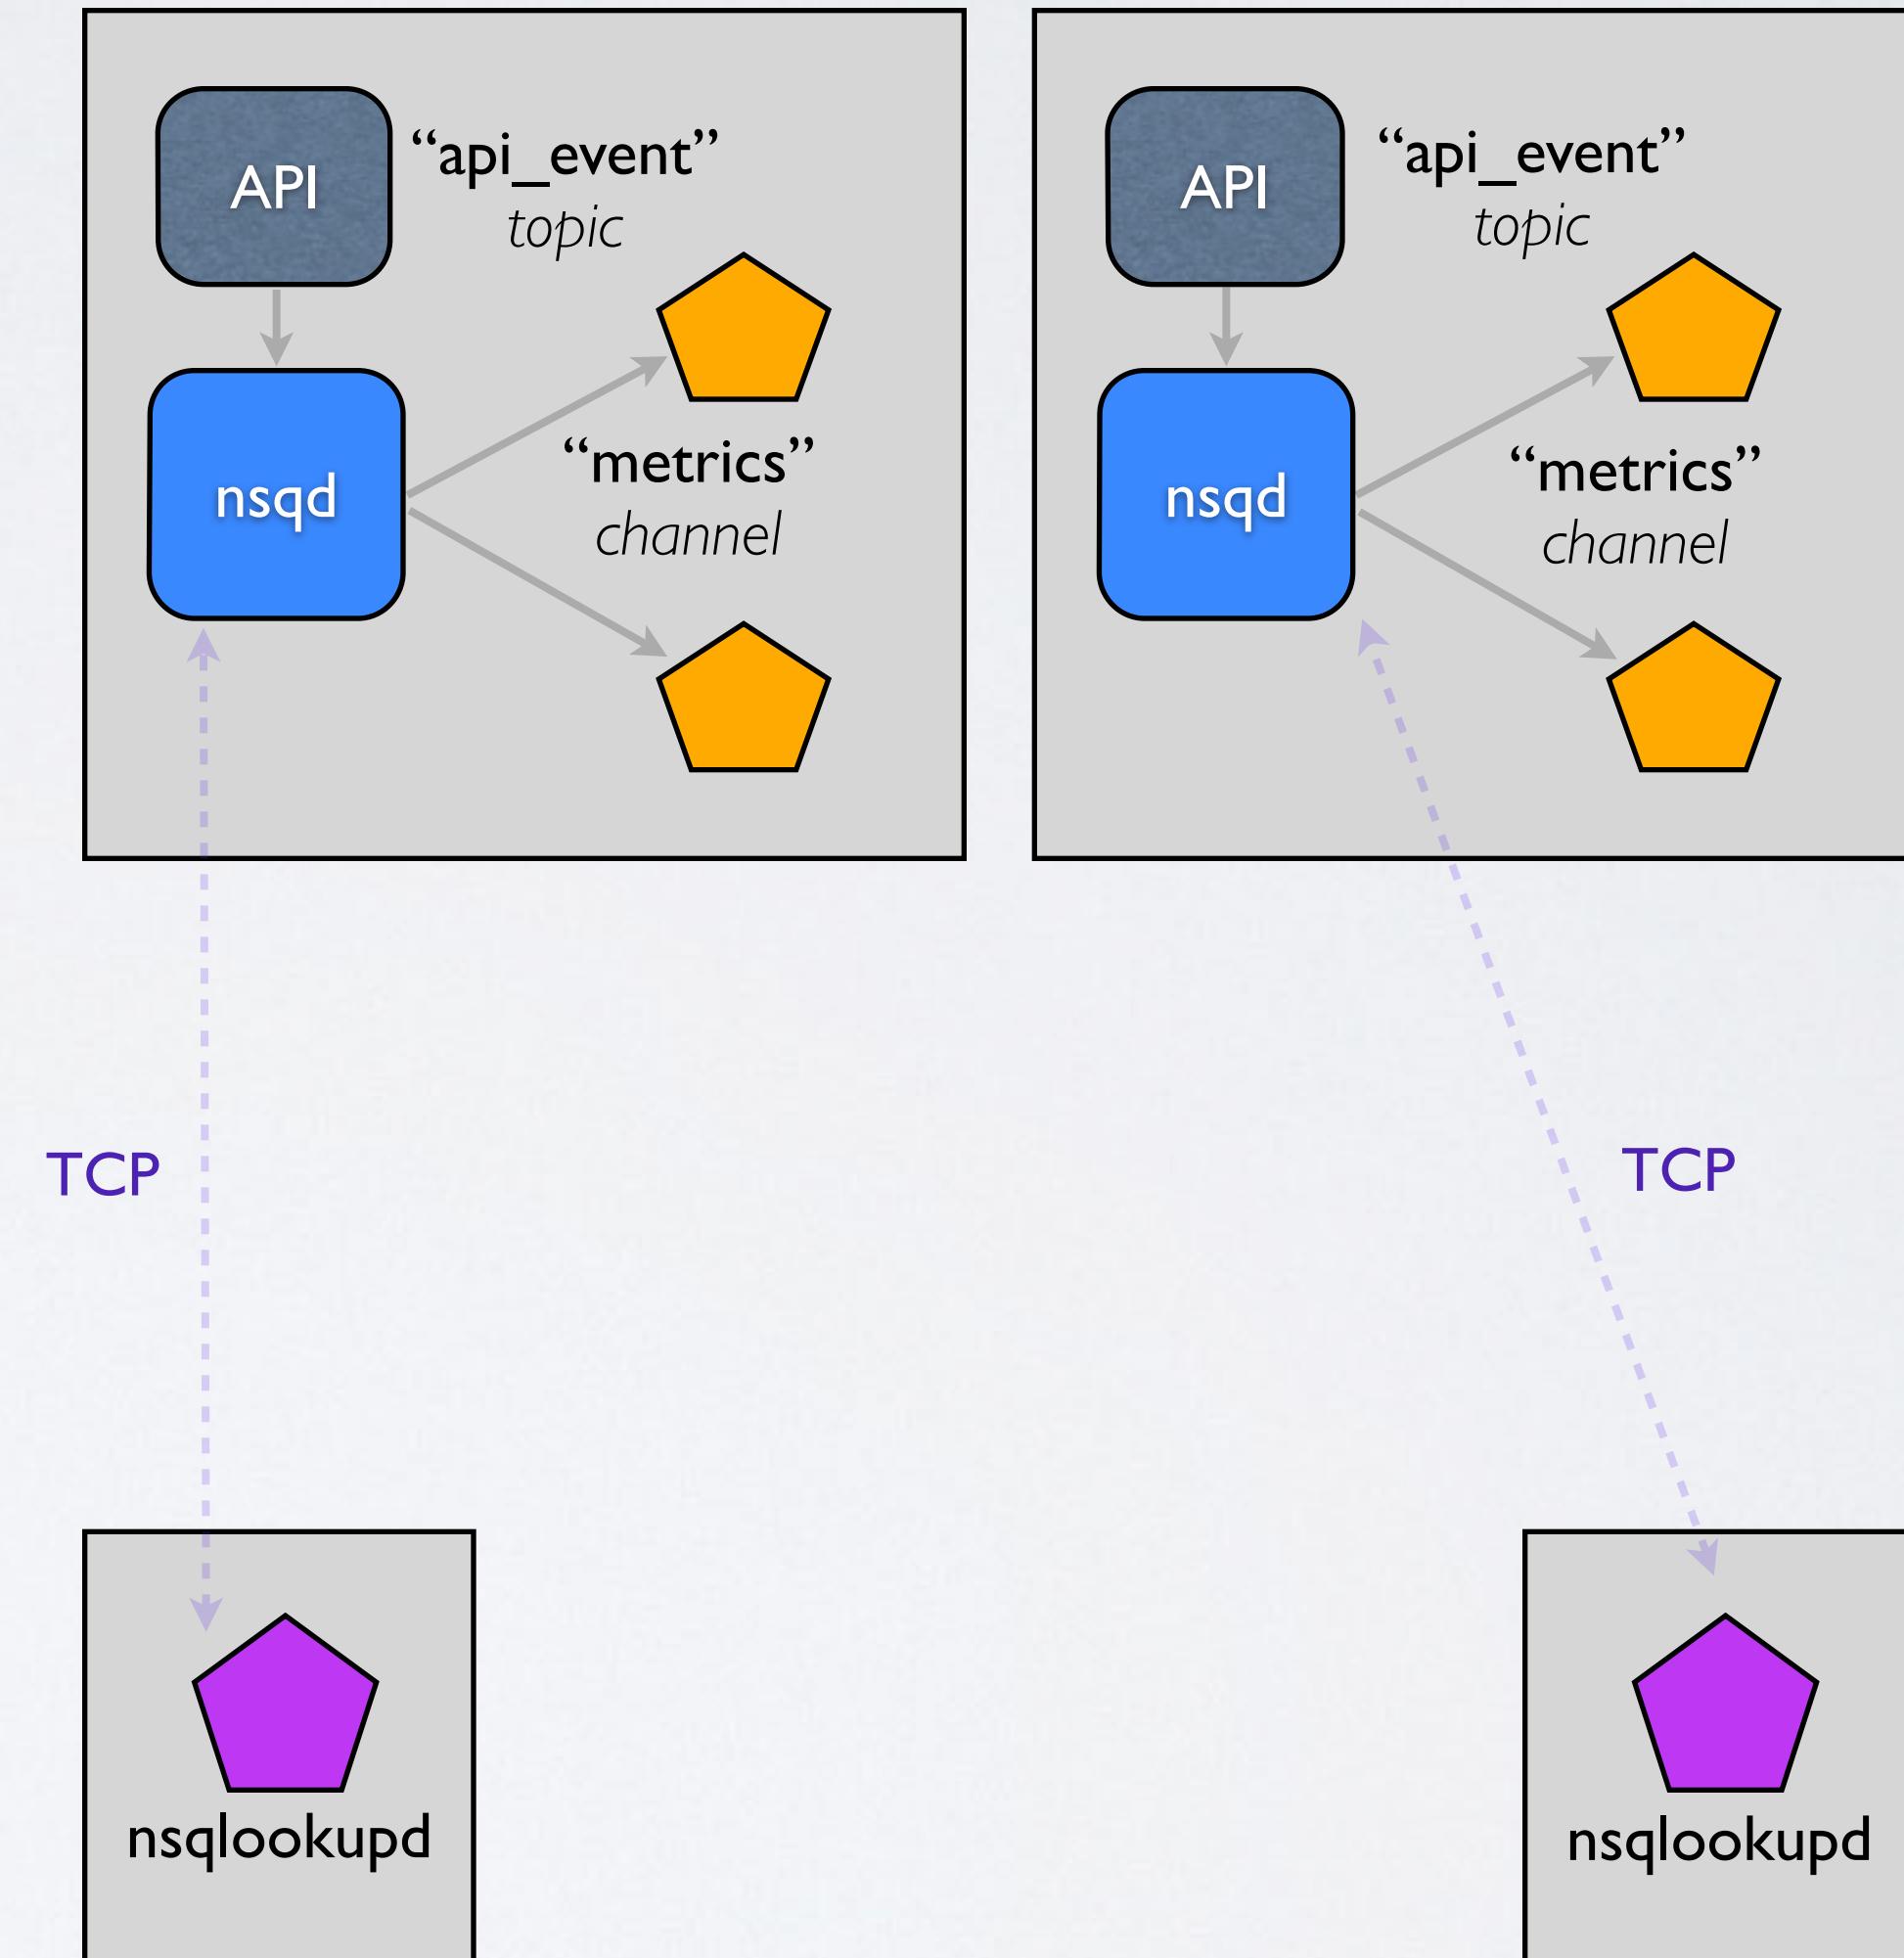
DISCOVERABILITY

- introduce `nsqlookupd`
- producers and consumers come and go
- other services can discover and subscribe to this topic



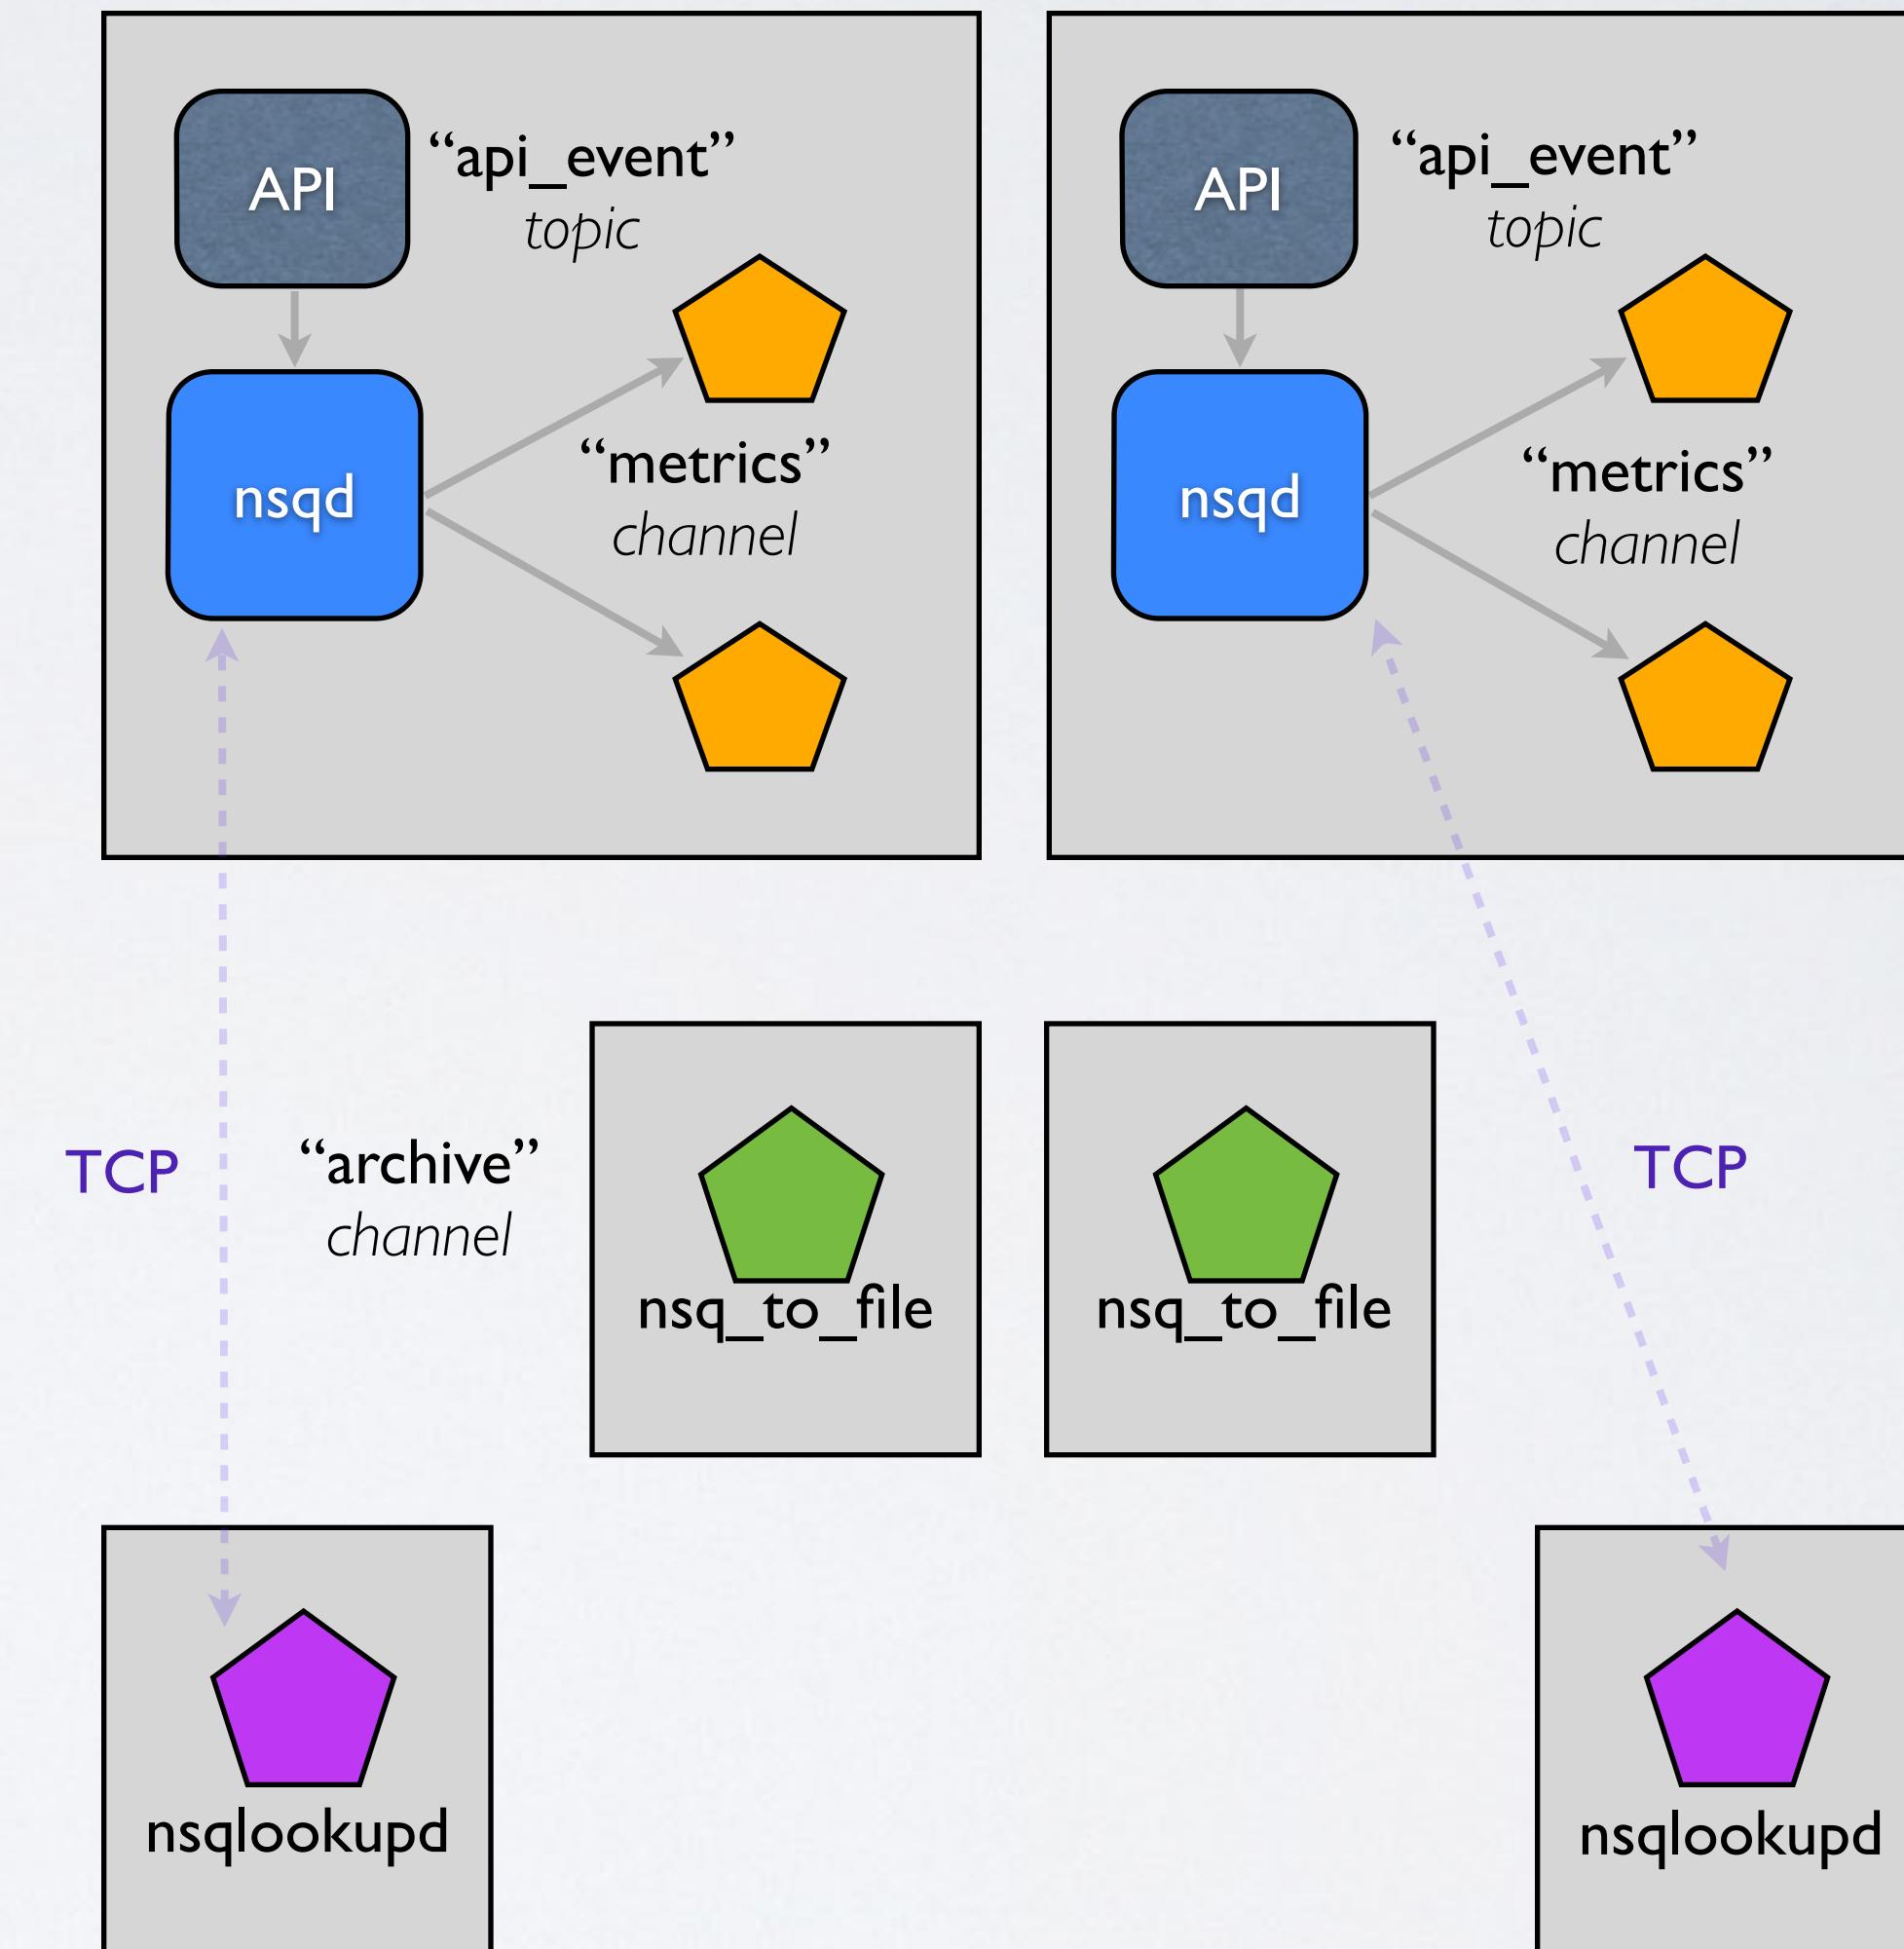
DISCOVERABILITY

- introduce nsqlookupd
- producers and consumers come and go
- other services can discover and subscribe to this topic



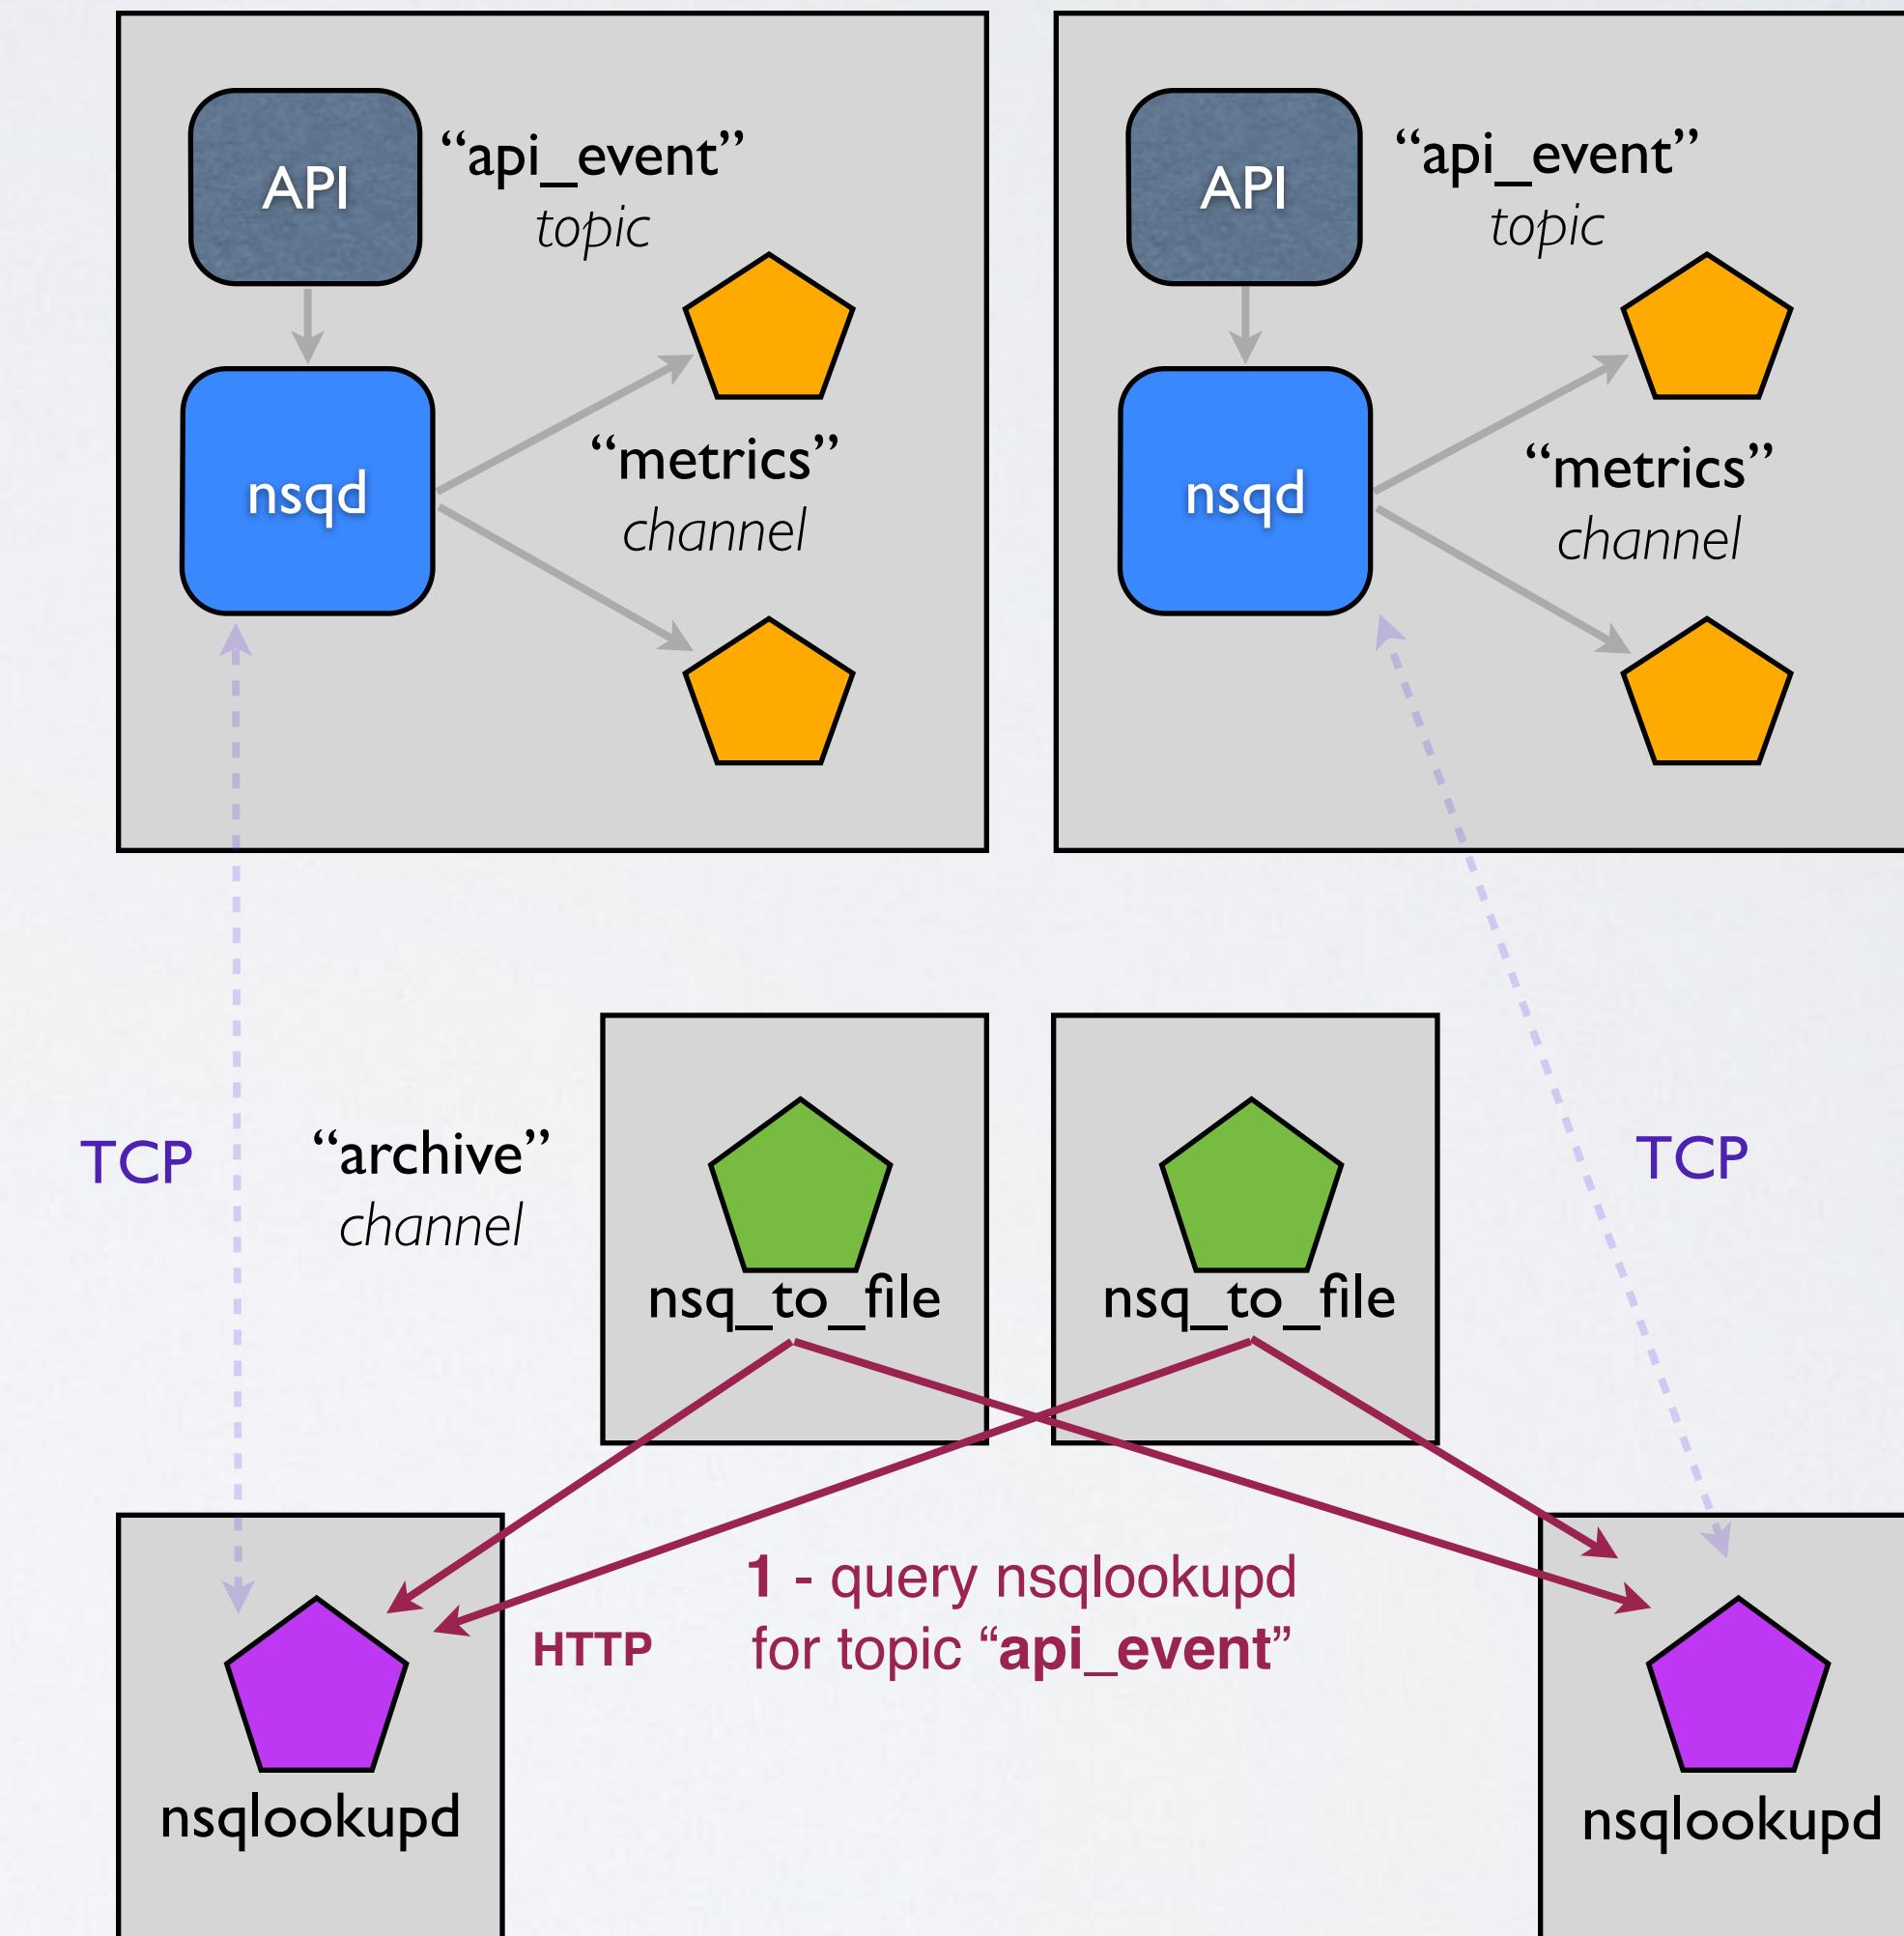
DISCOVERABILITY

- introduce nsqlookupd
- producers and consumers come and go
- other services can discover and subscribe to this topic



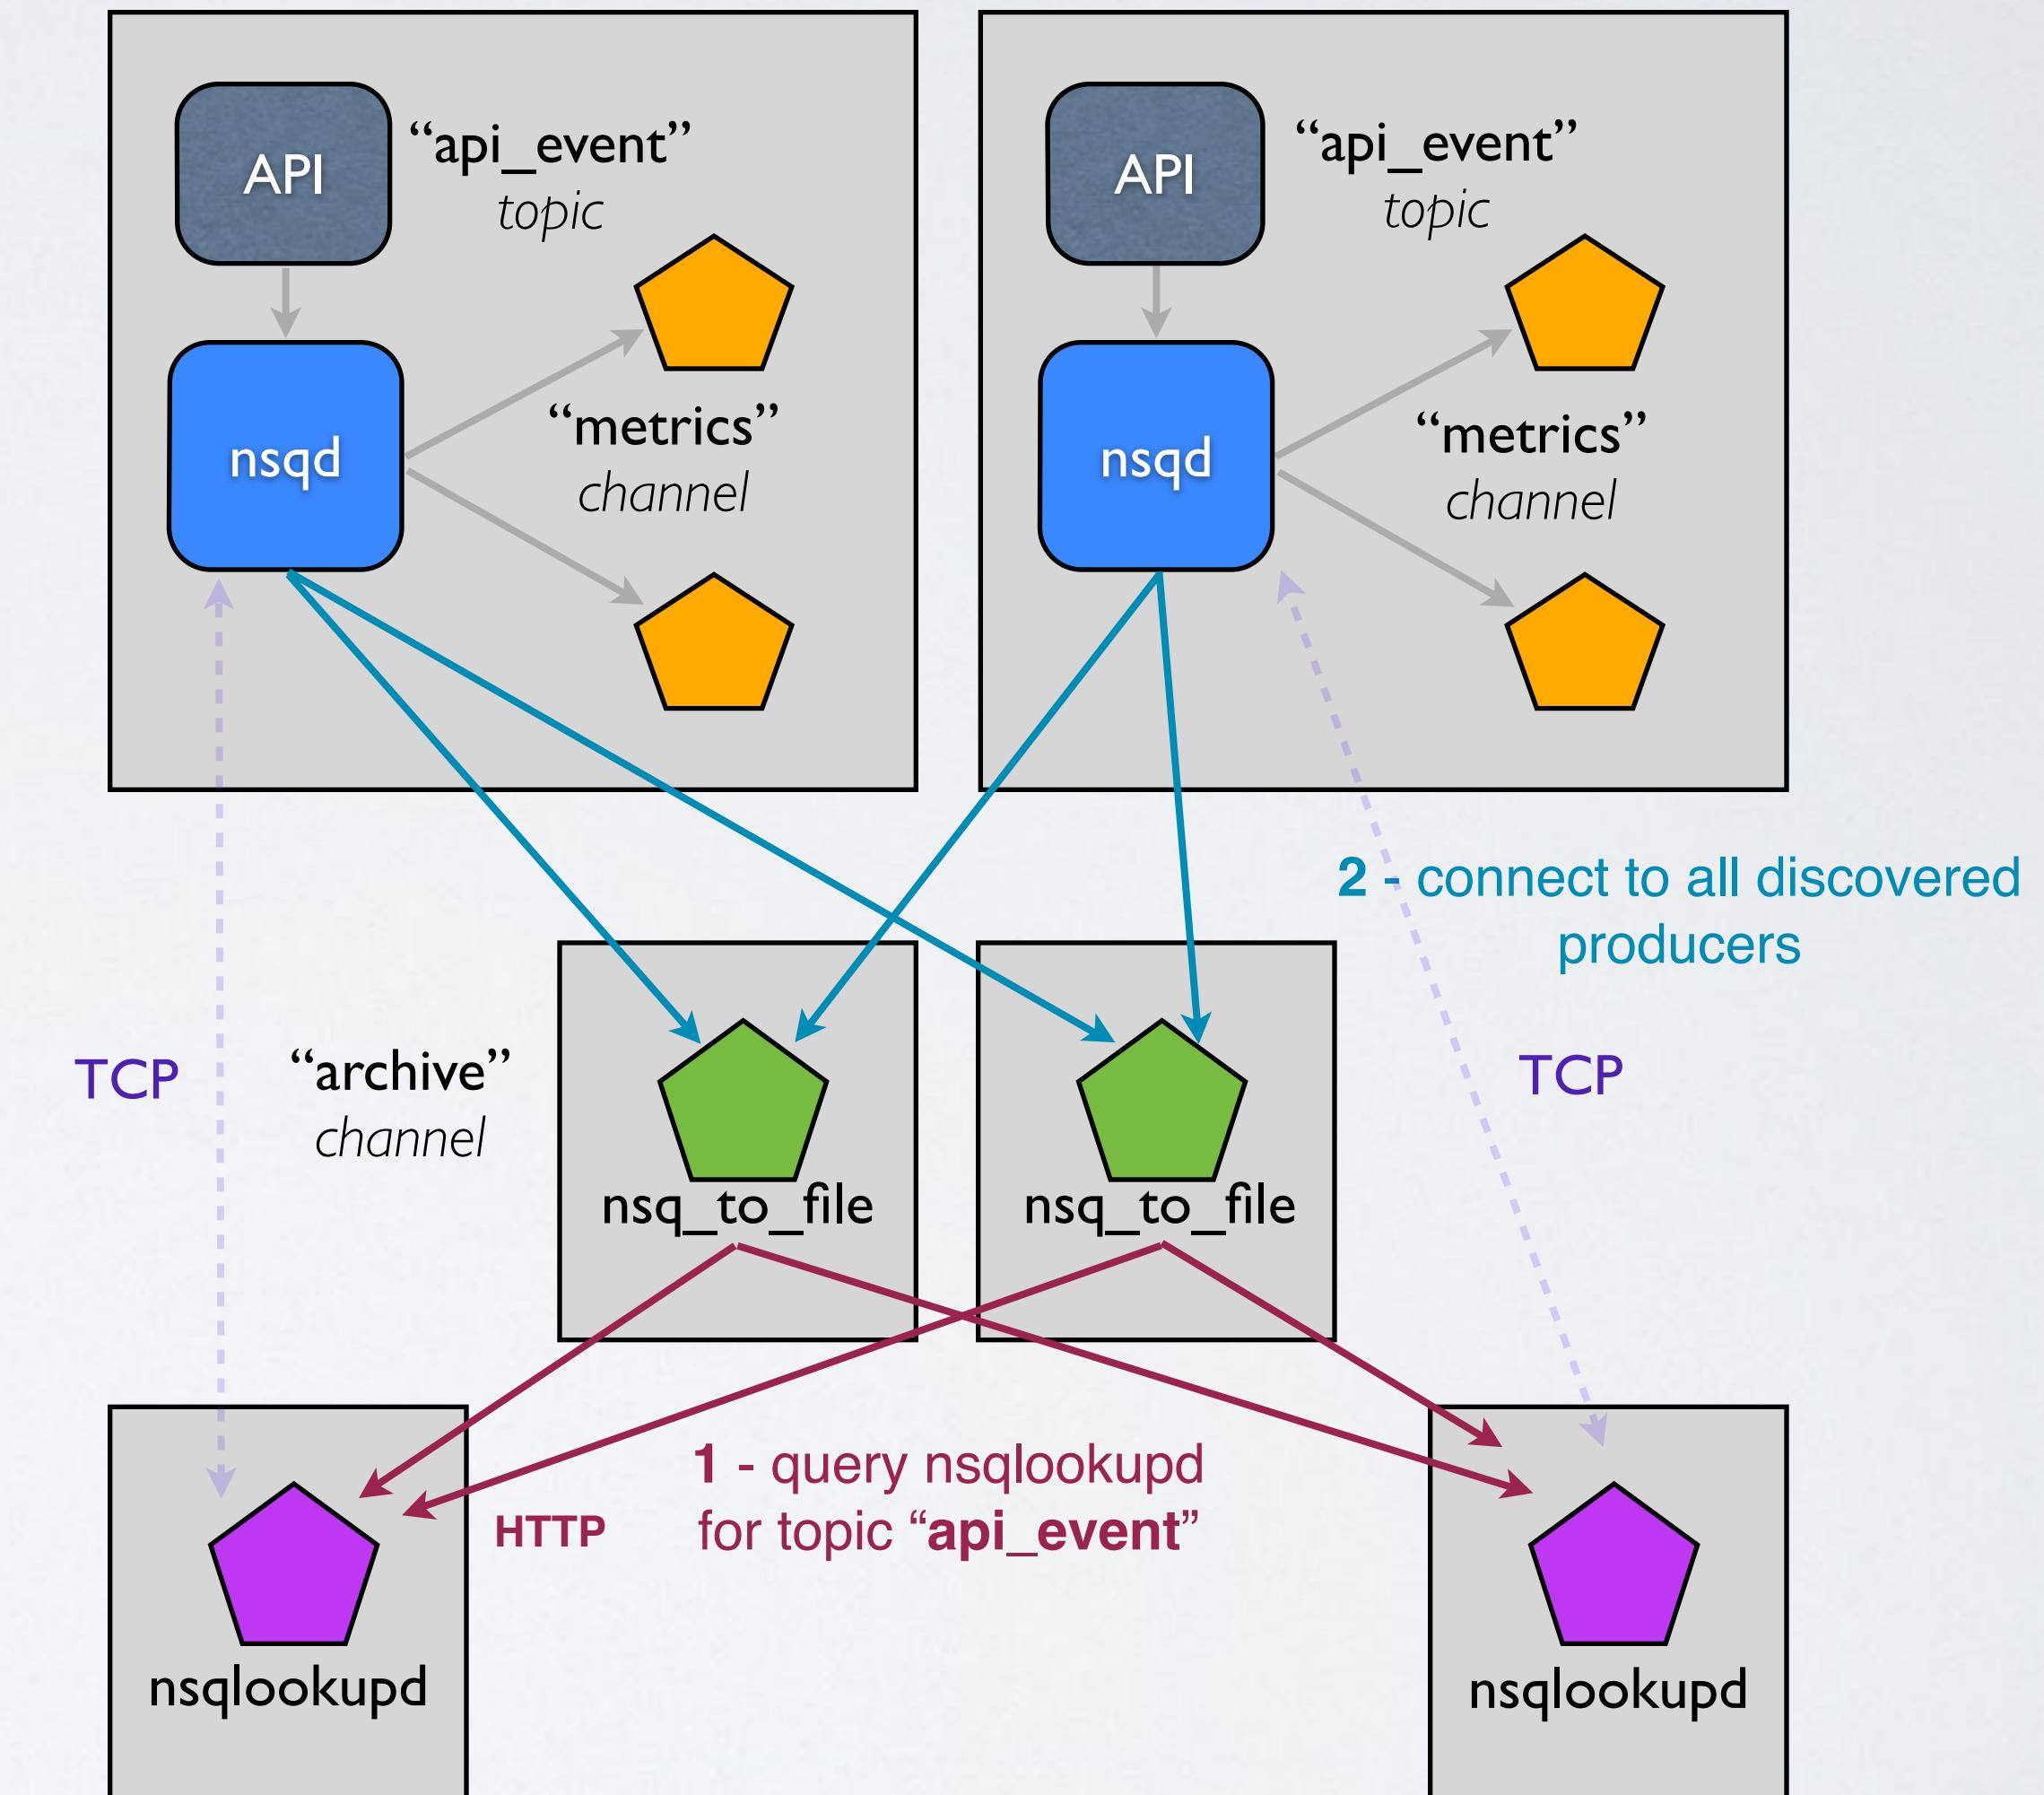
DISCOVERABILITY

- introduce nsqlookupd
- producers and consumers come and go
- other services can discover and subscribe to this topic



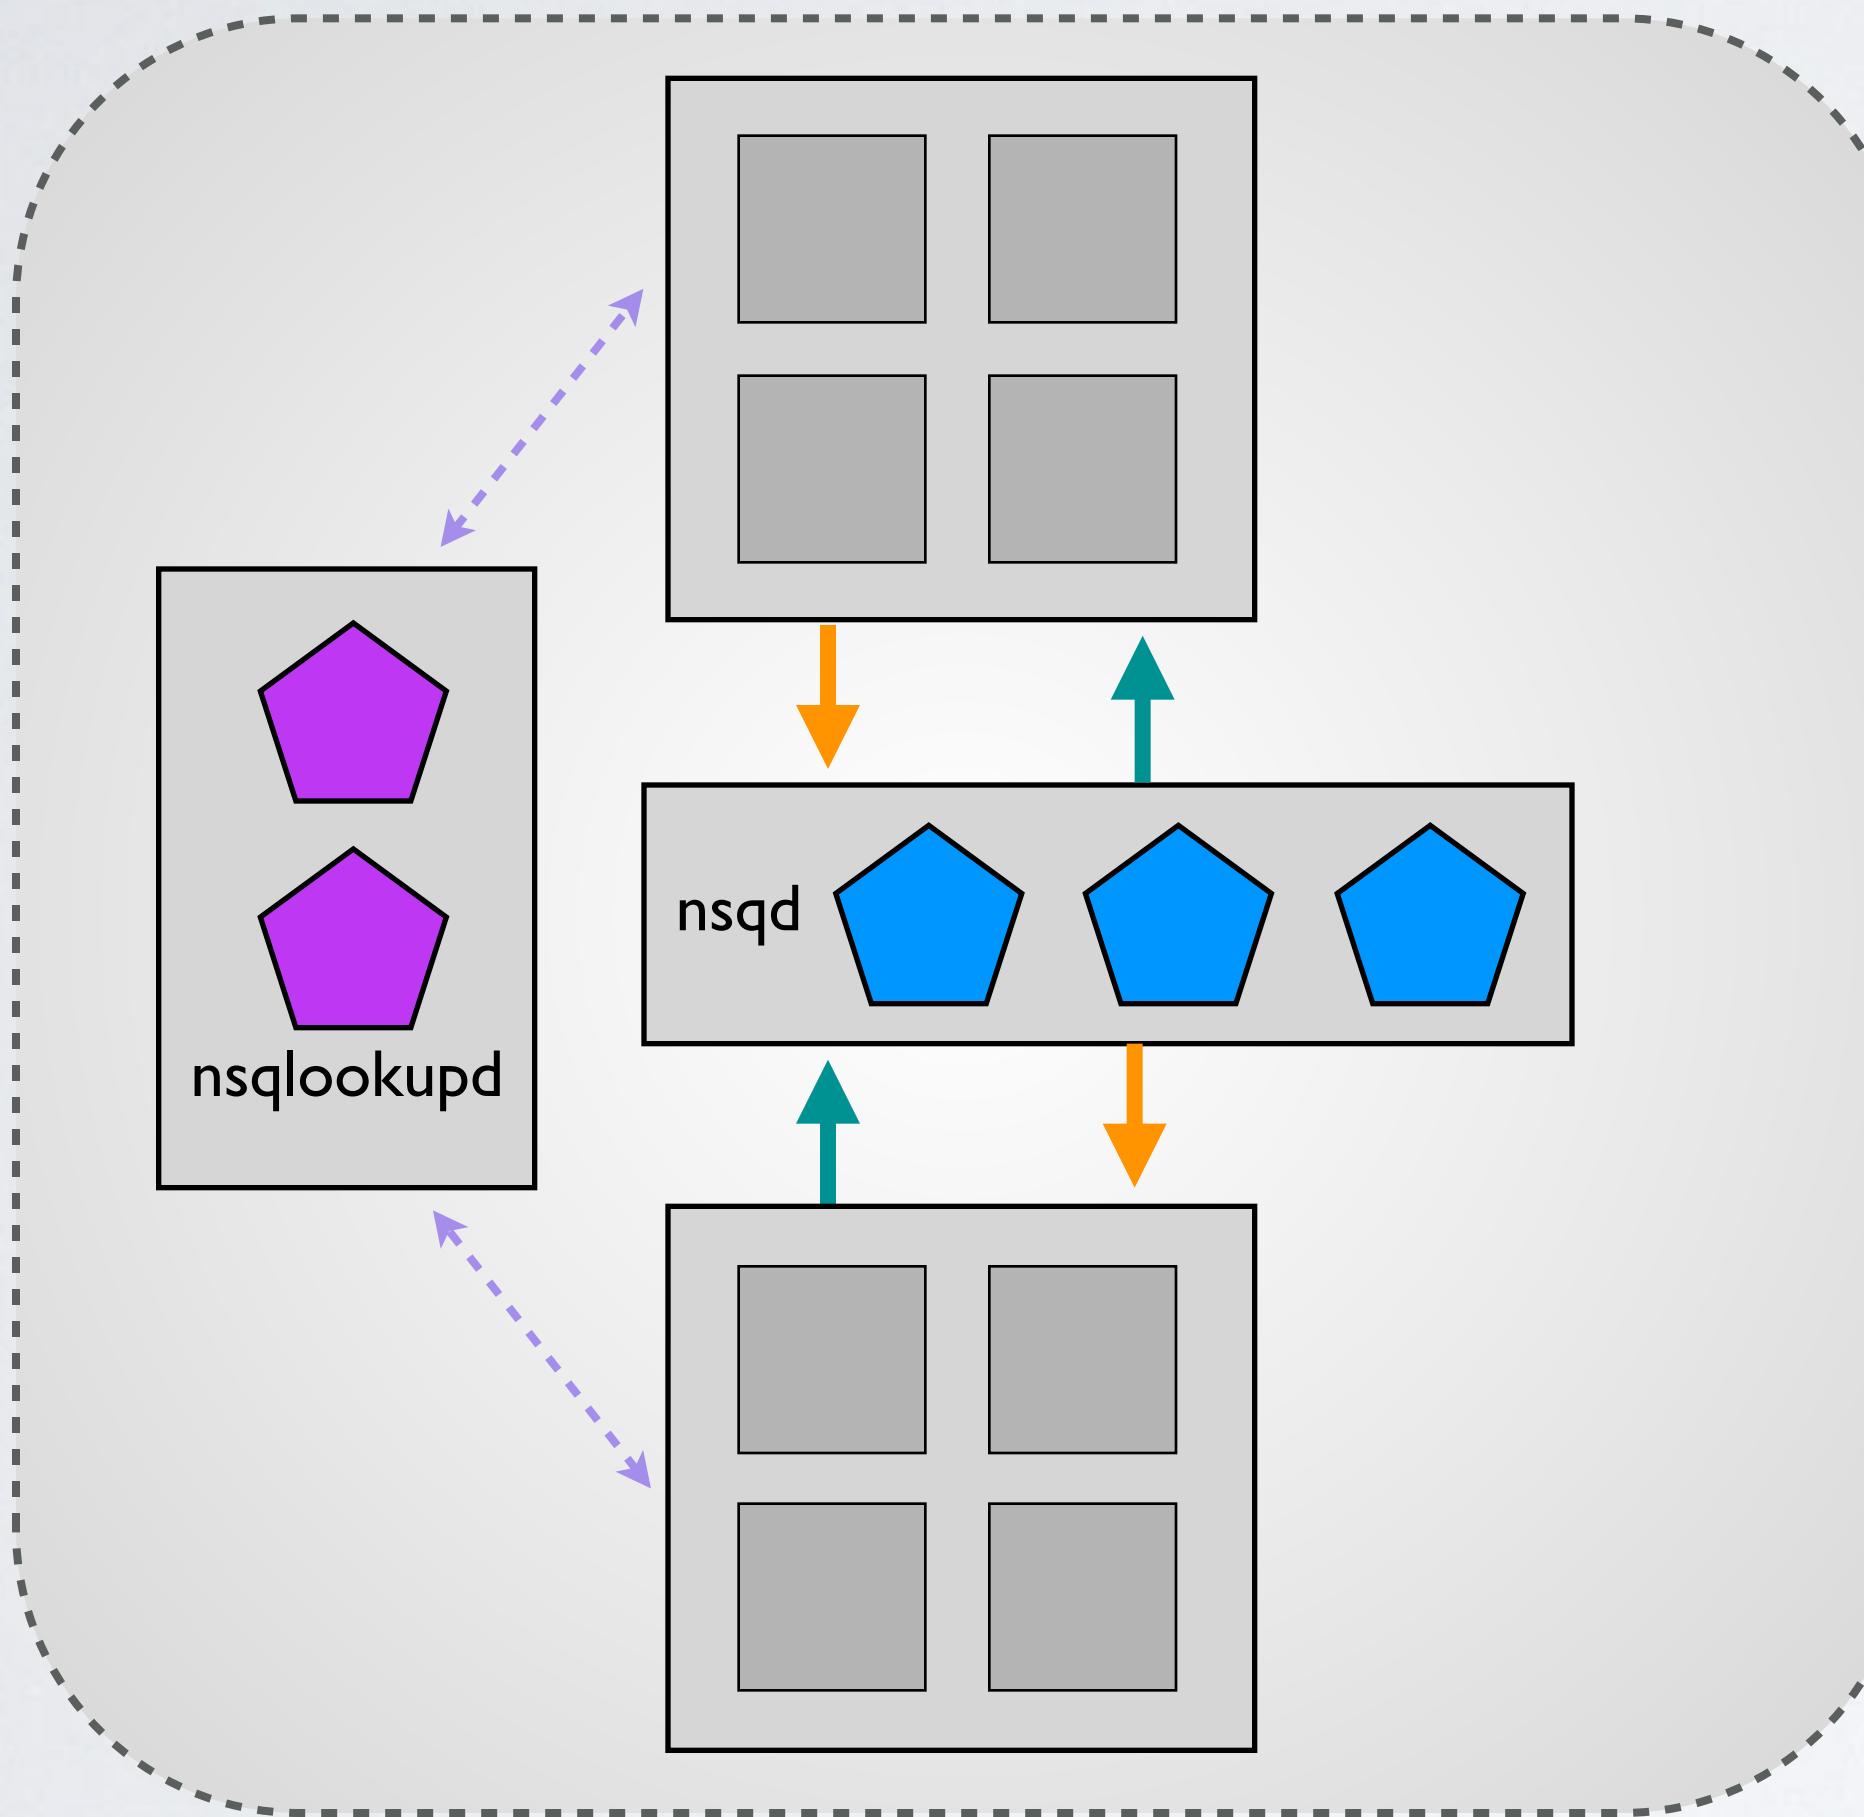
DISCOVERABILITY

- introduce nsqlookupd
- producers and consumers come and go
- other services can discover and subscribe to this topic



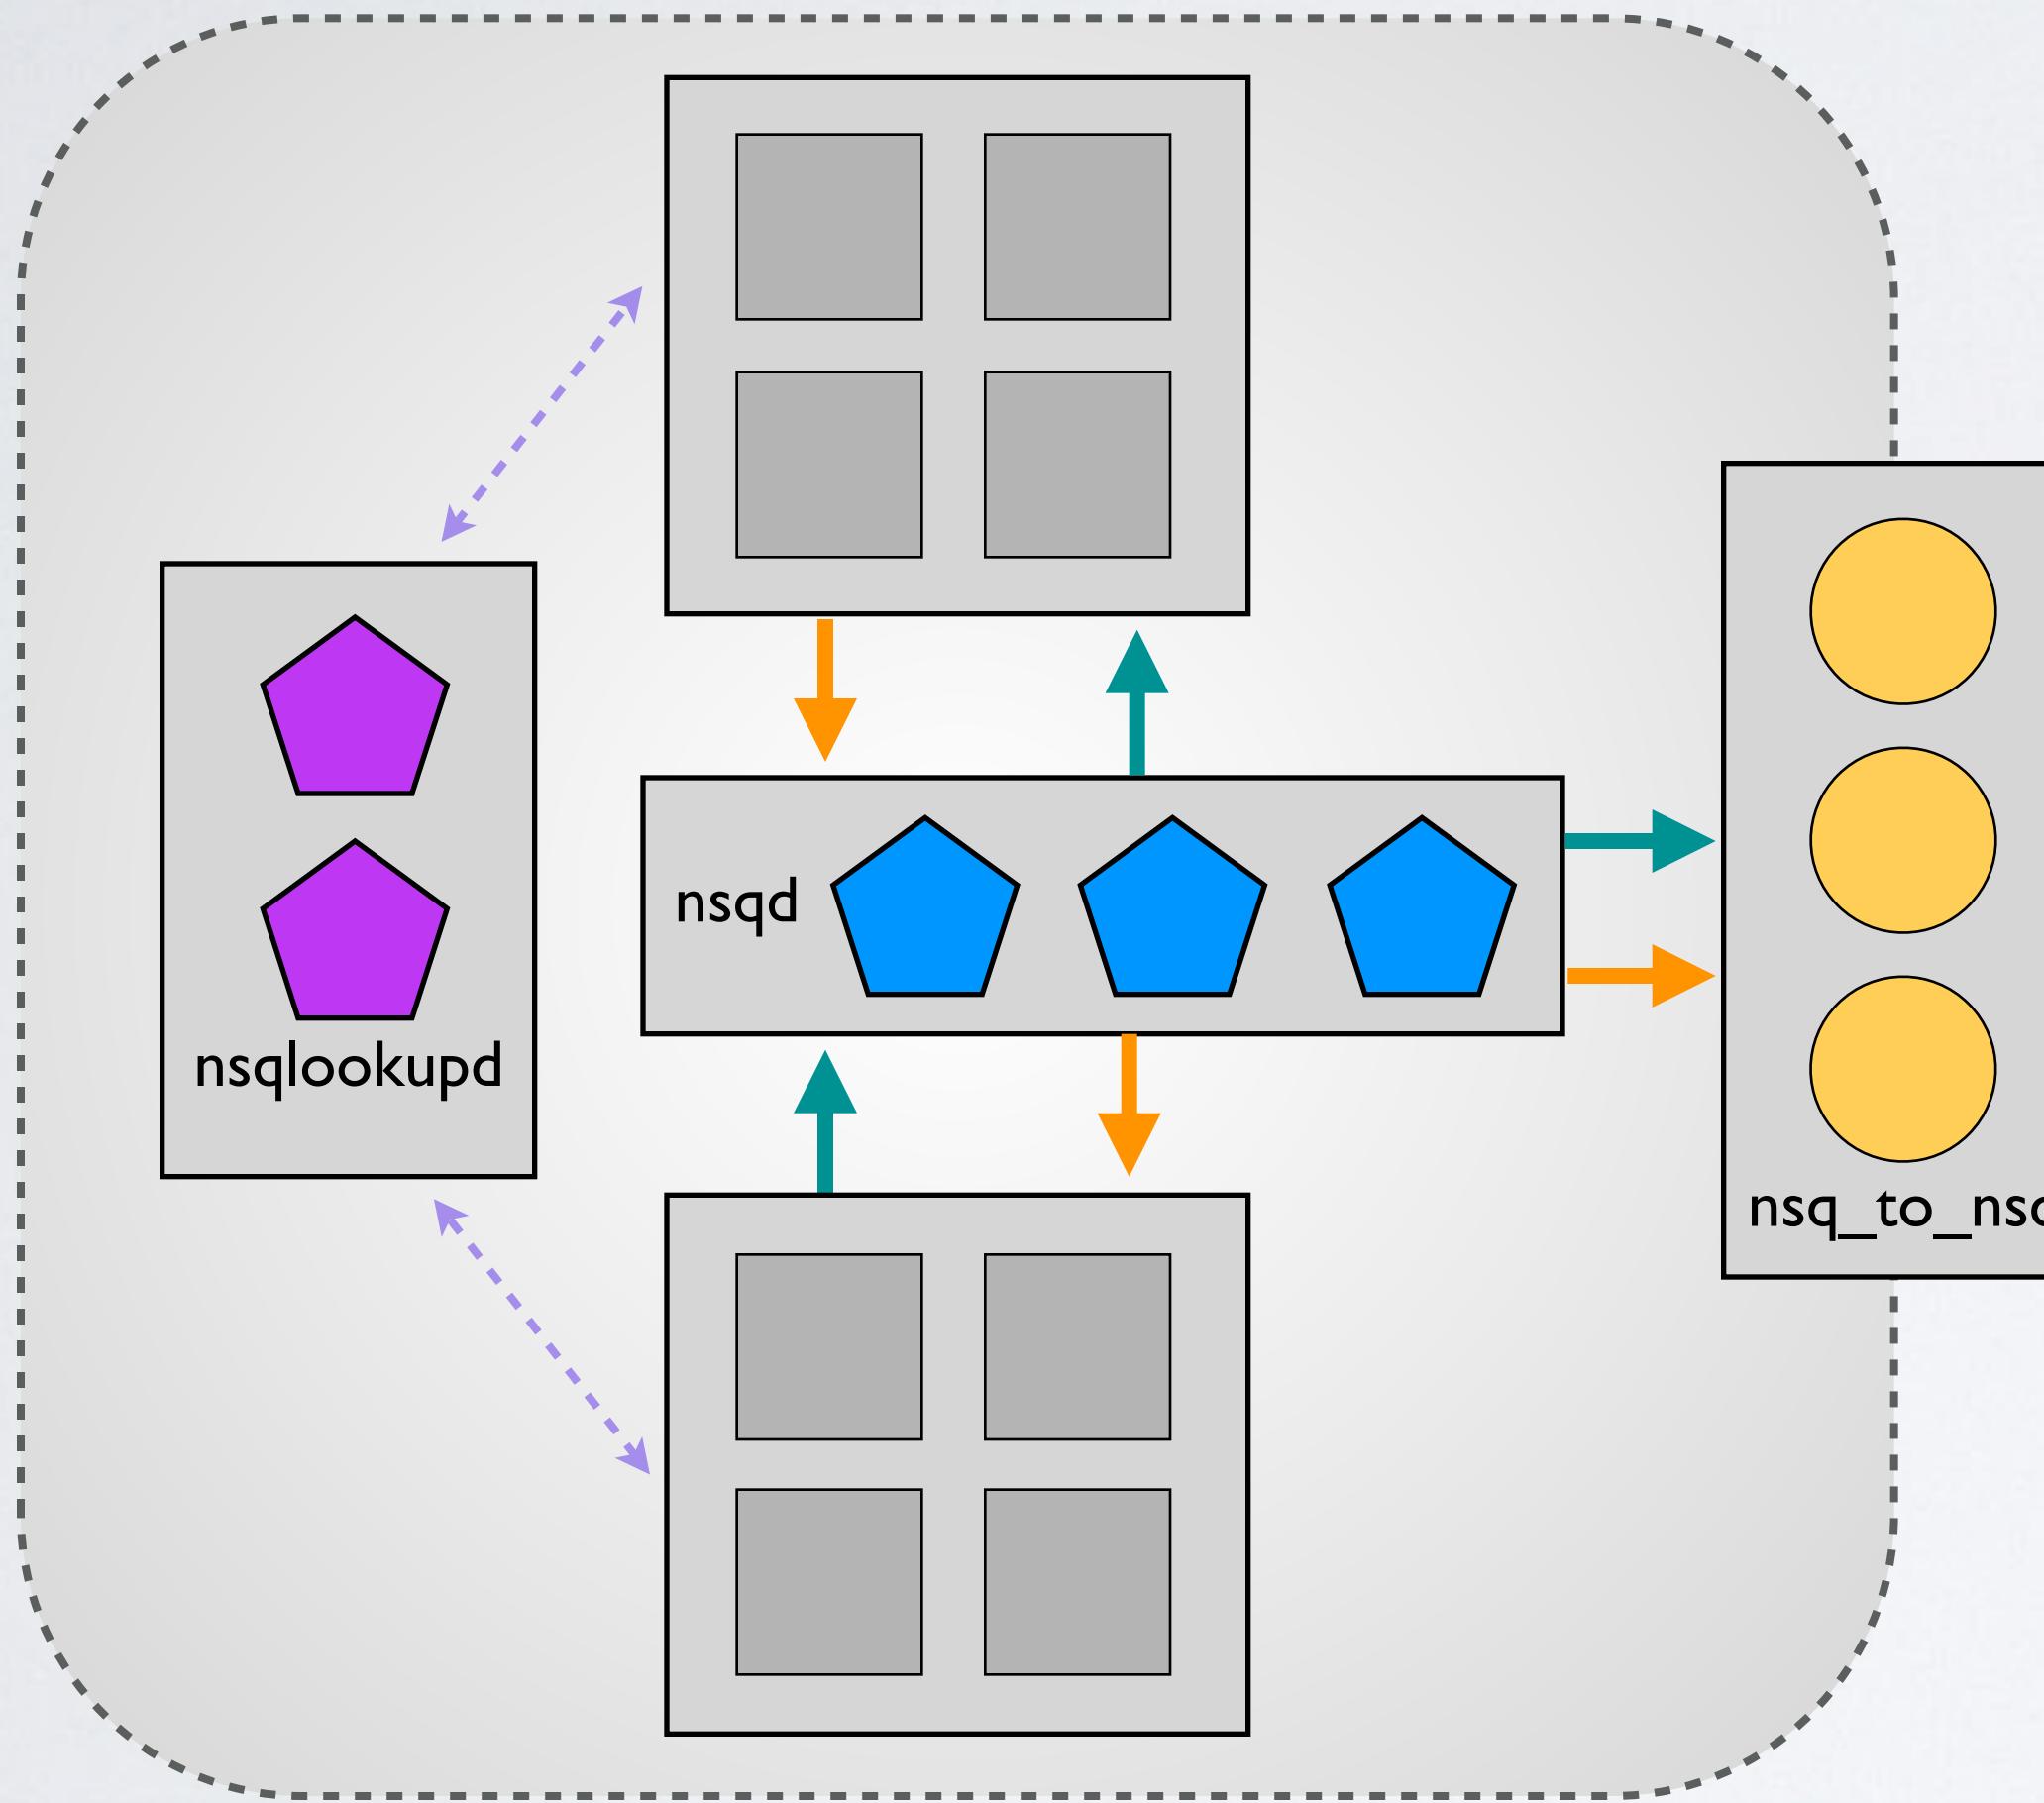
TO INFINITY AND BEYOND

Datacenter A



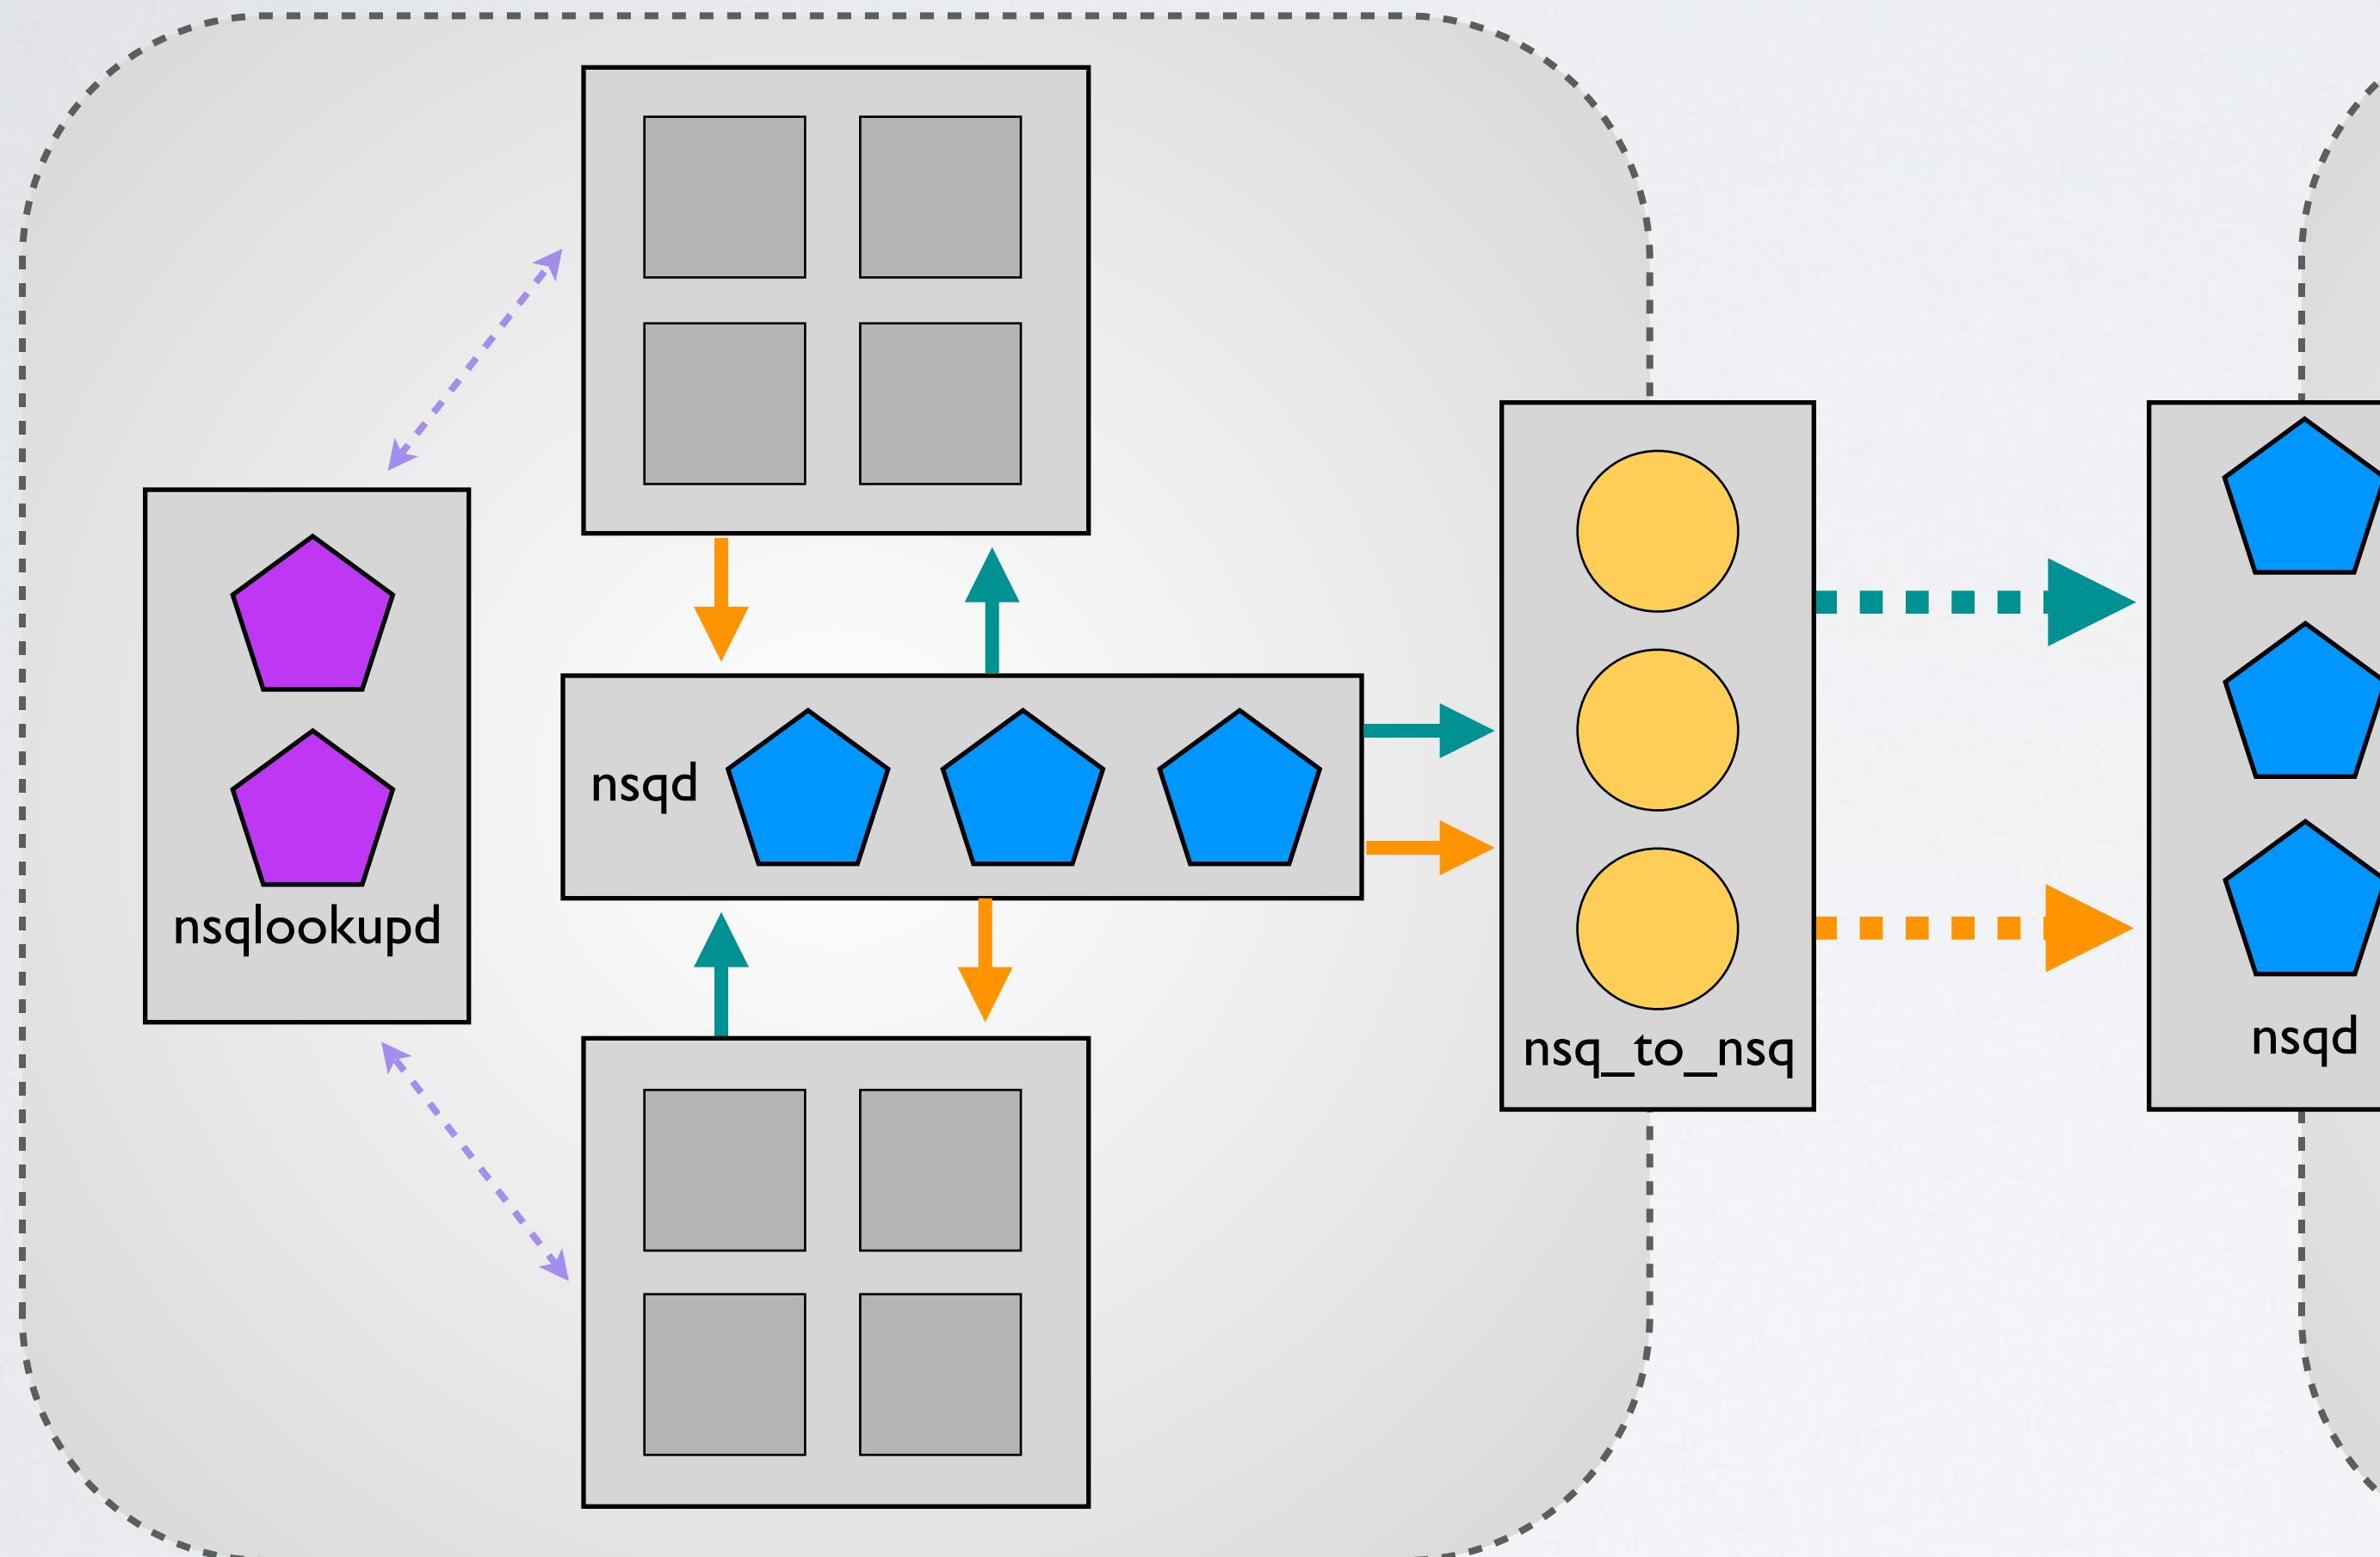
TO INFINITY AND BEYOND

Datacenter A

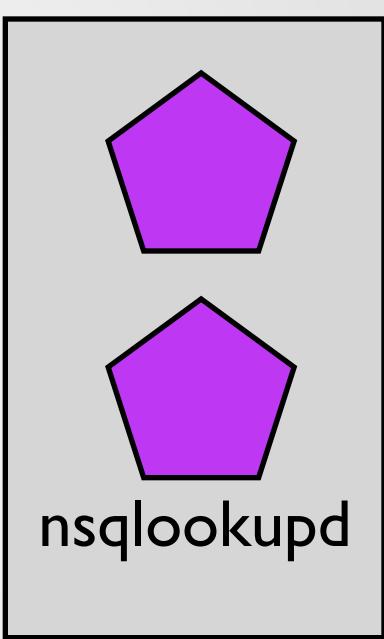


TO INFINITY AND BEYOND

Datacenter A

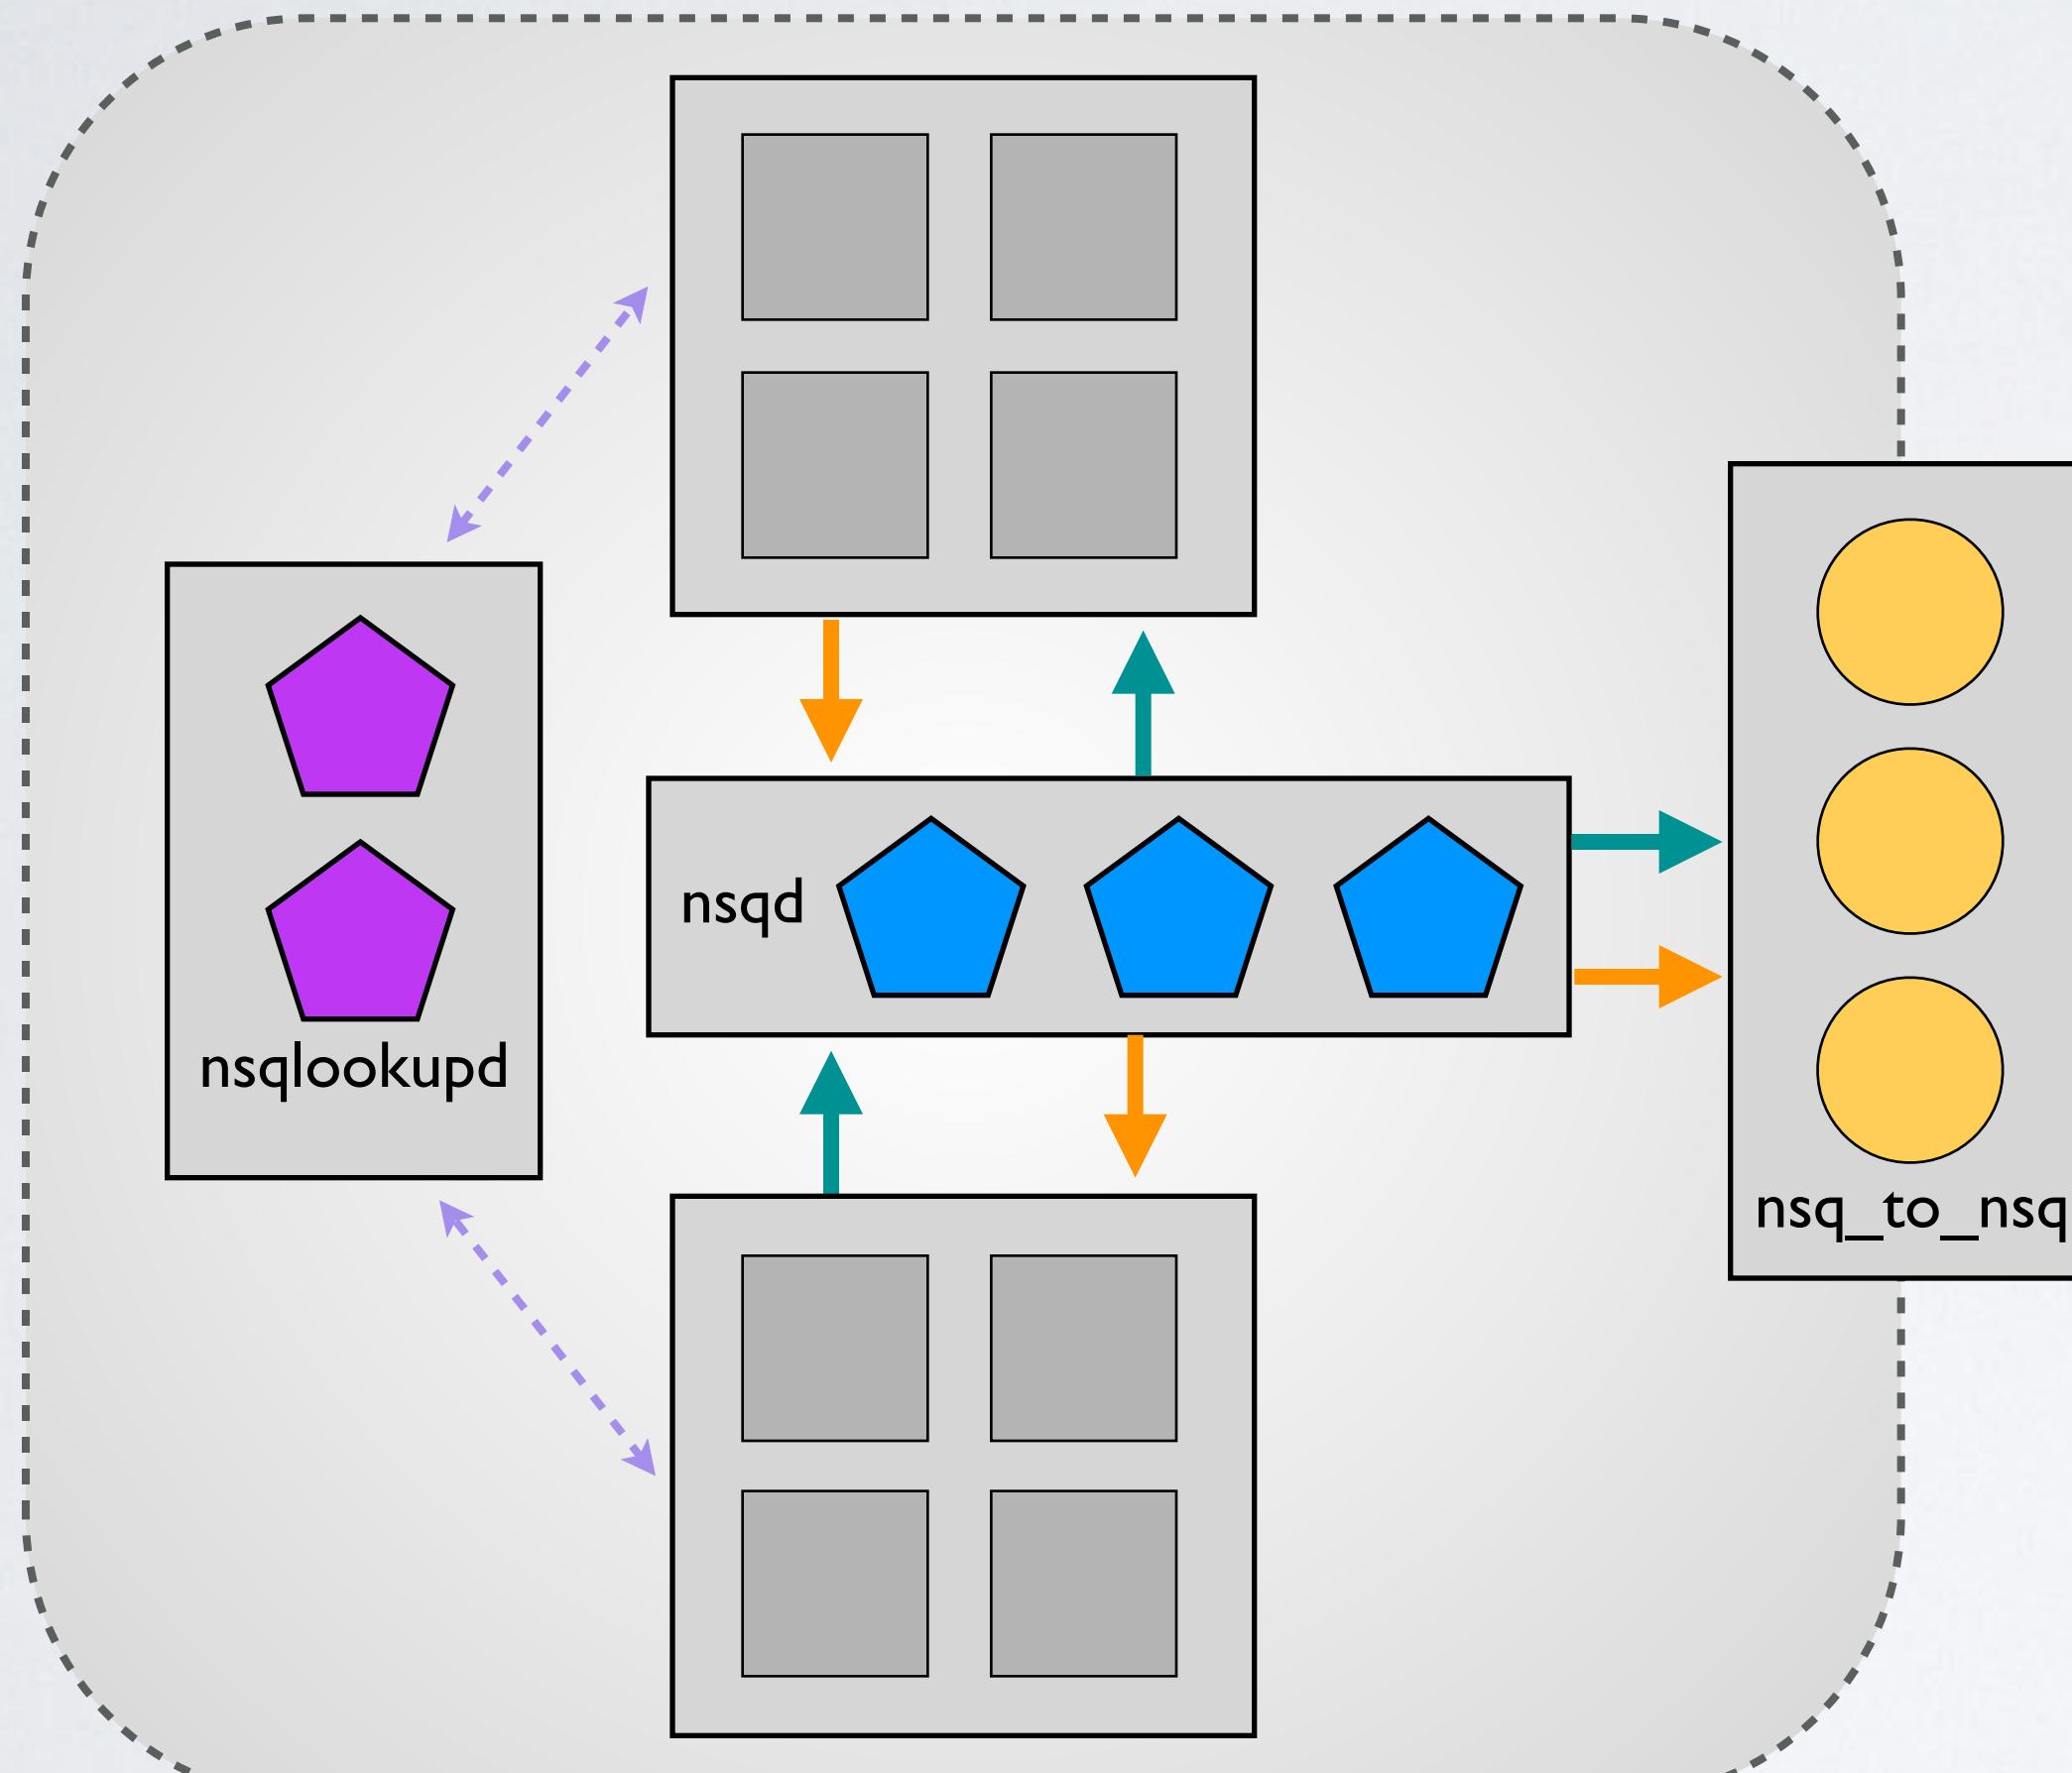


Datacenter B

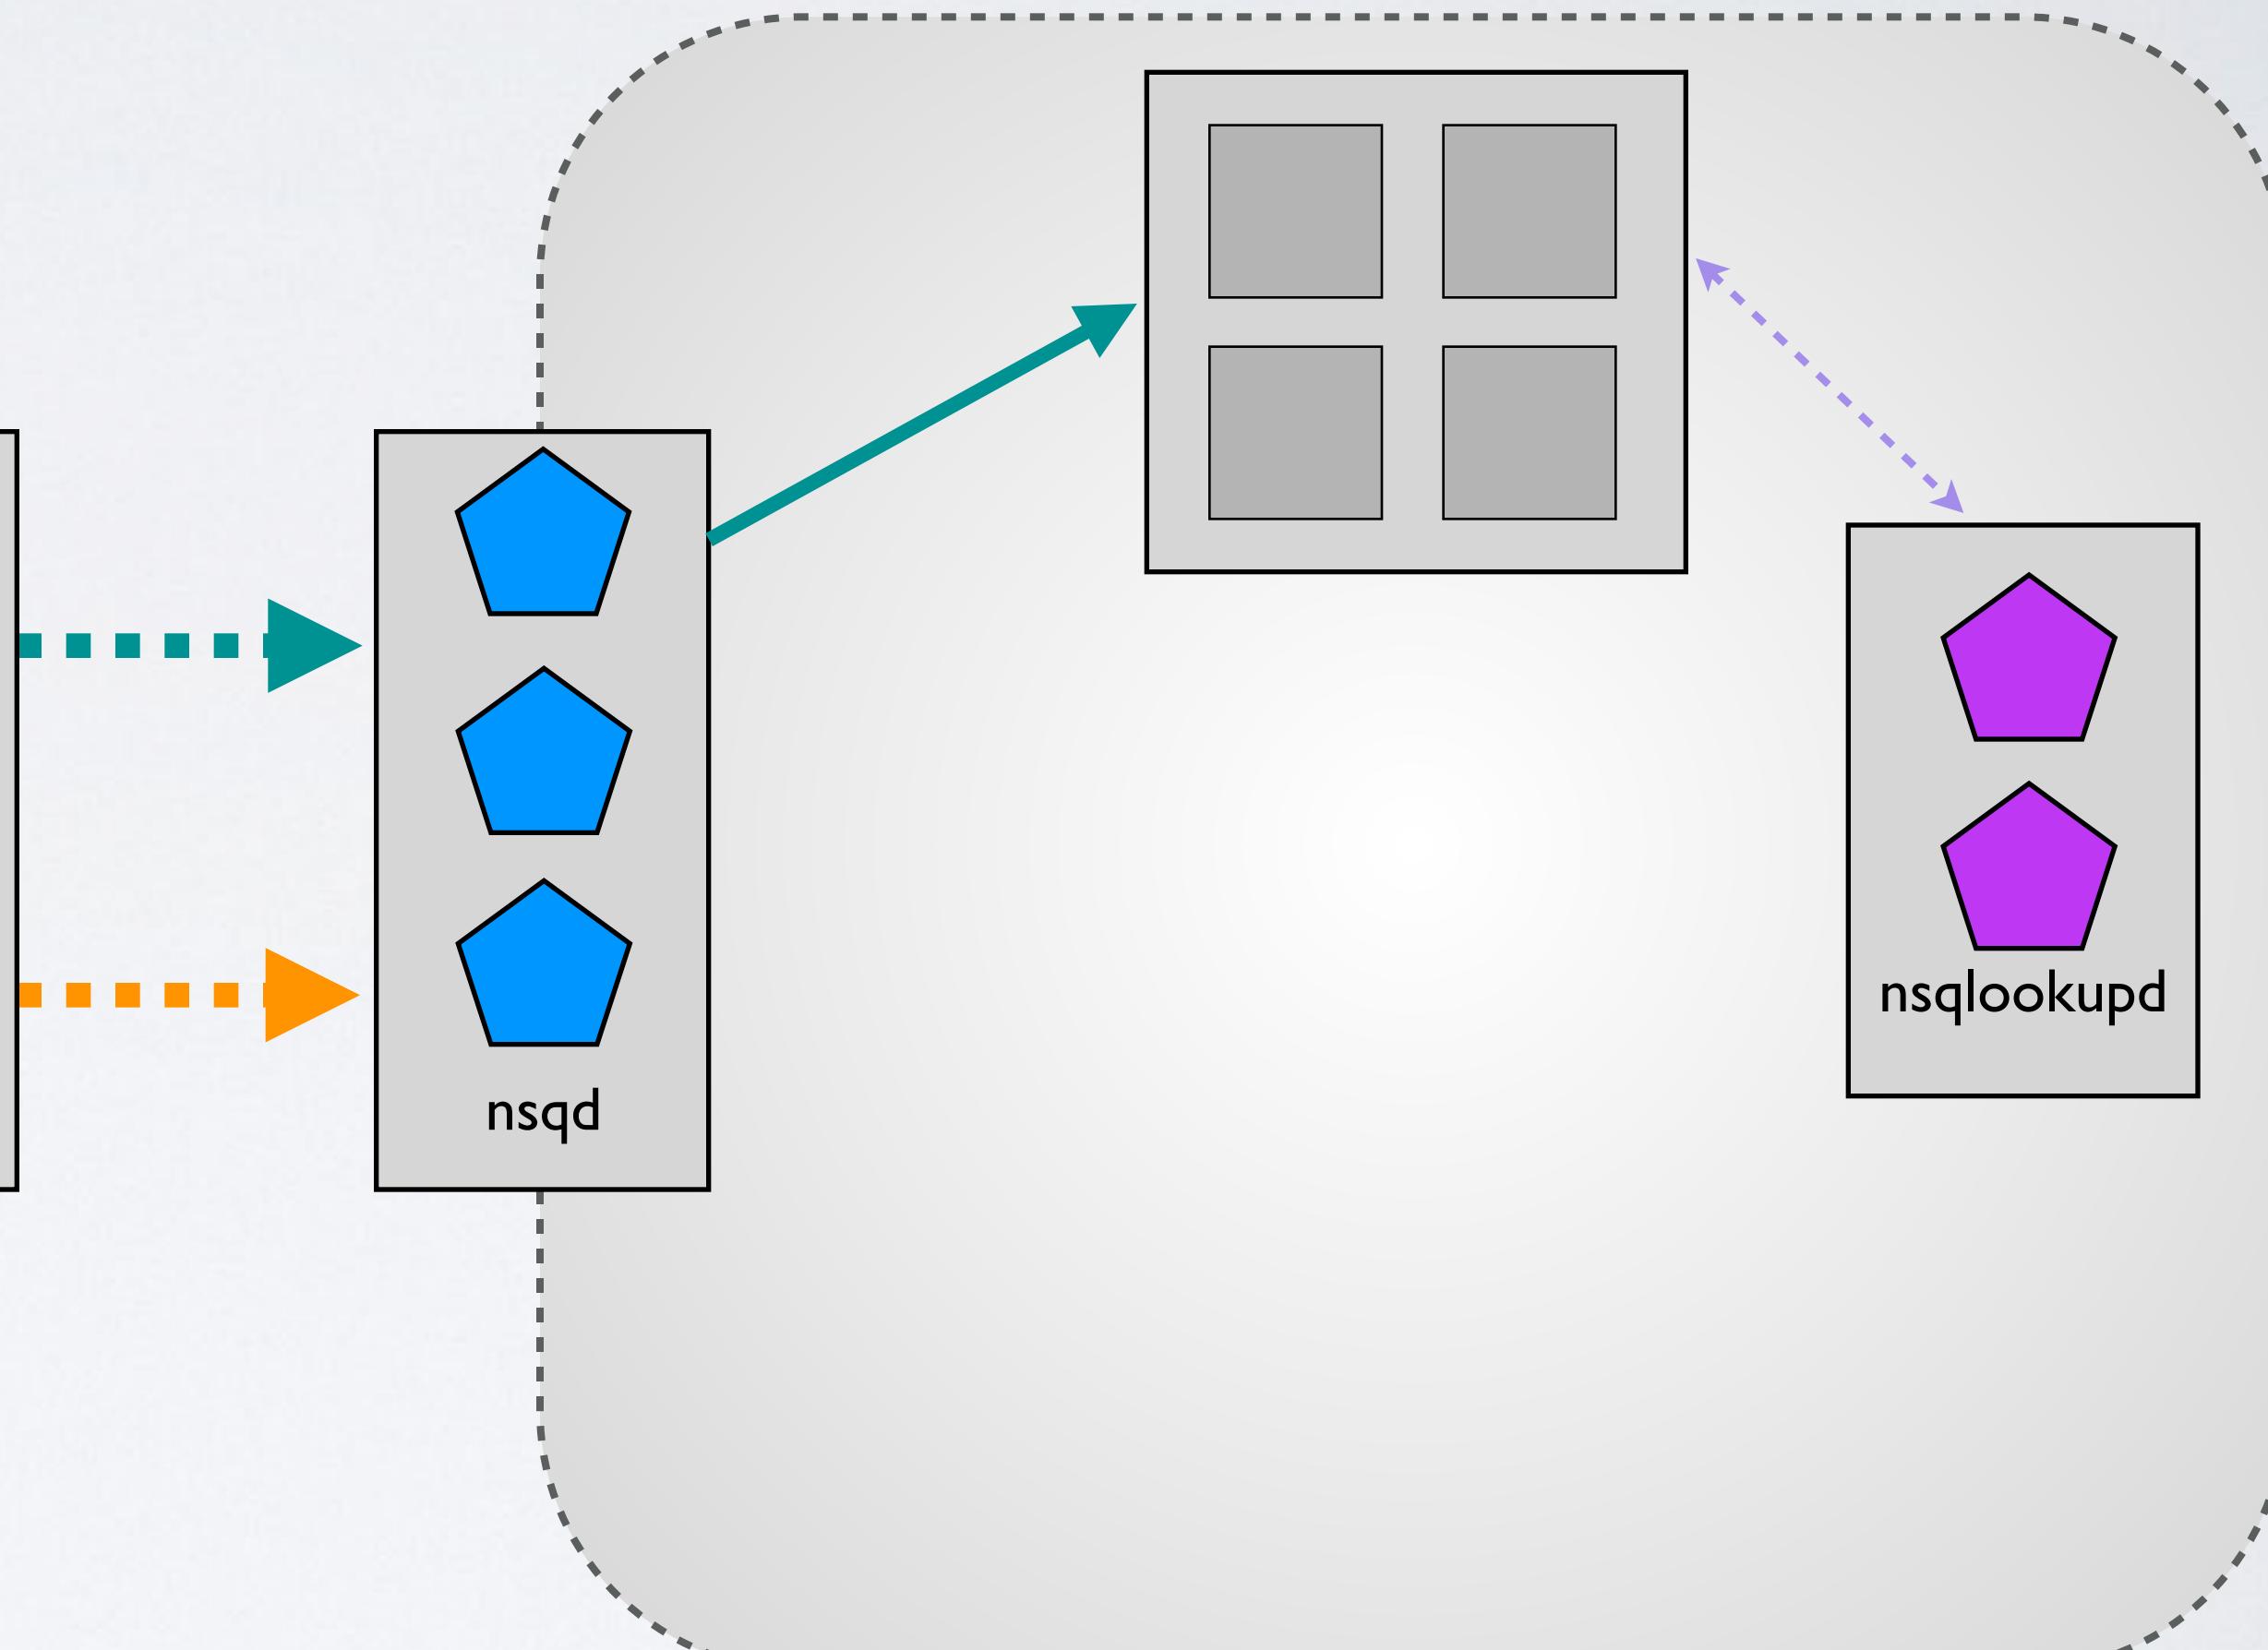


TO INFINITY AND BEYOND

Datacenter A

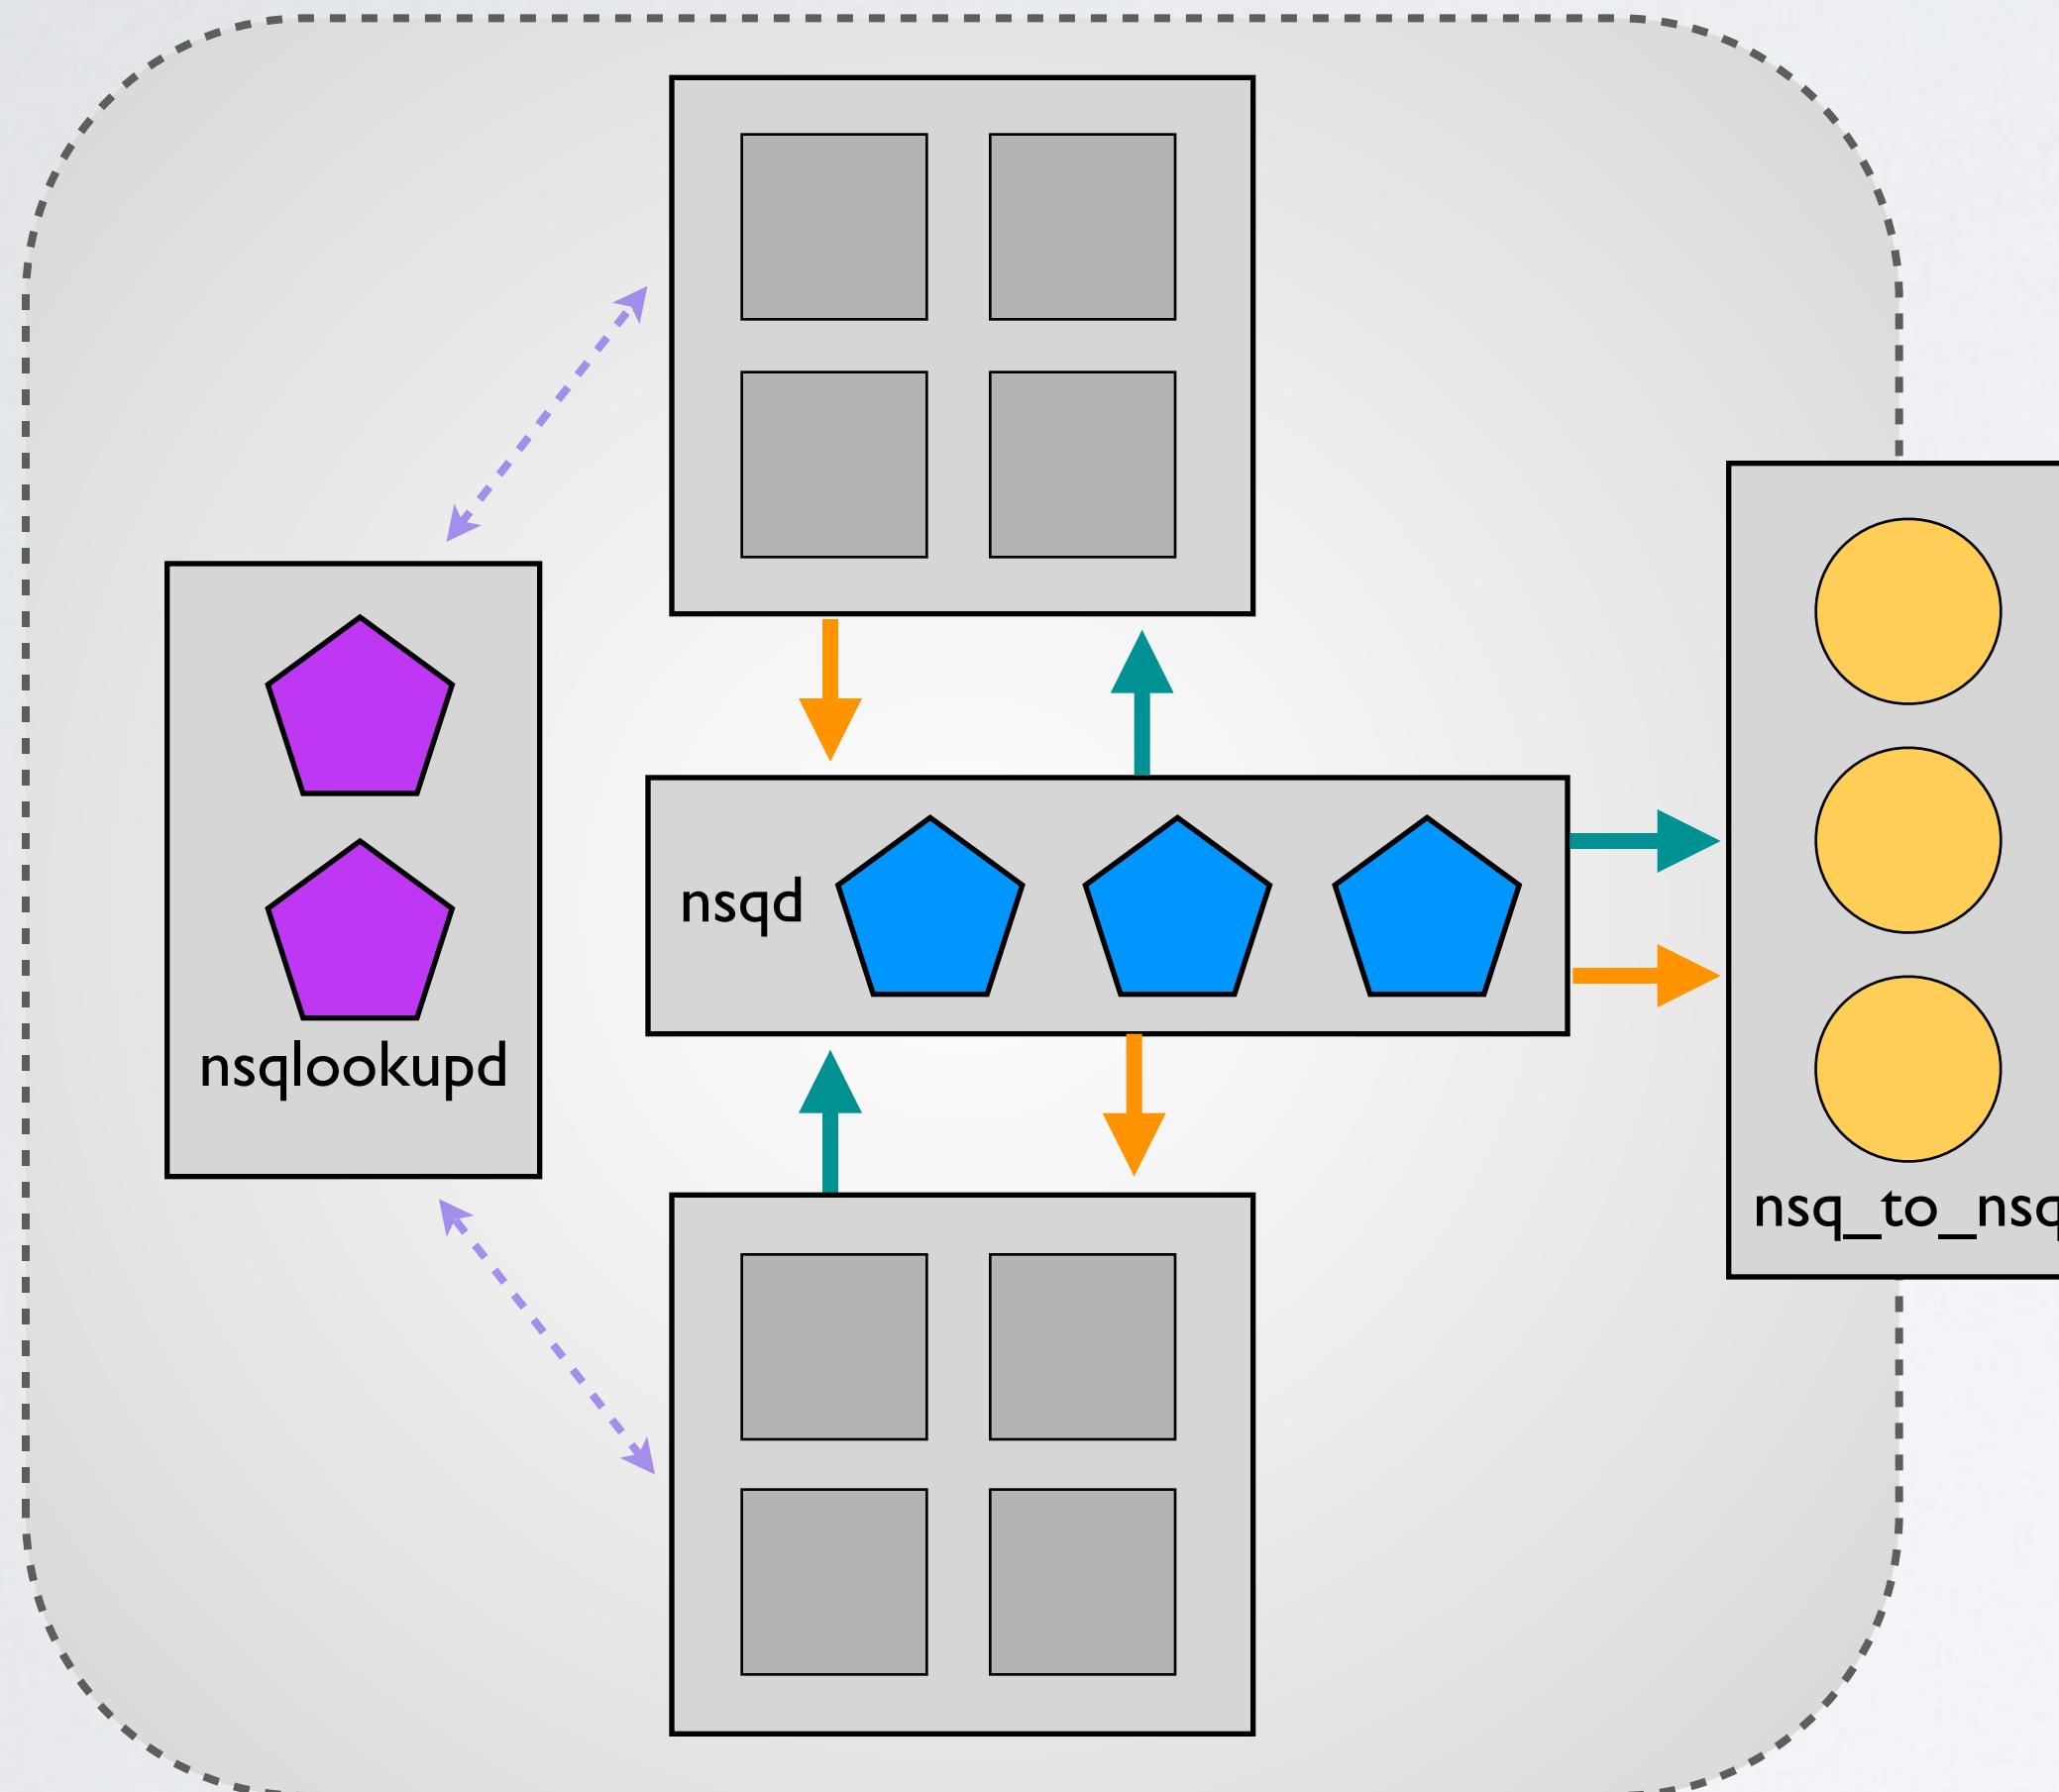


Datacenter B

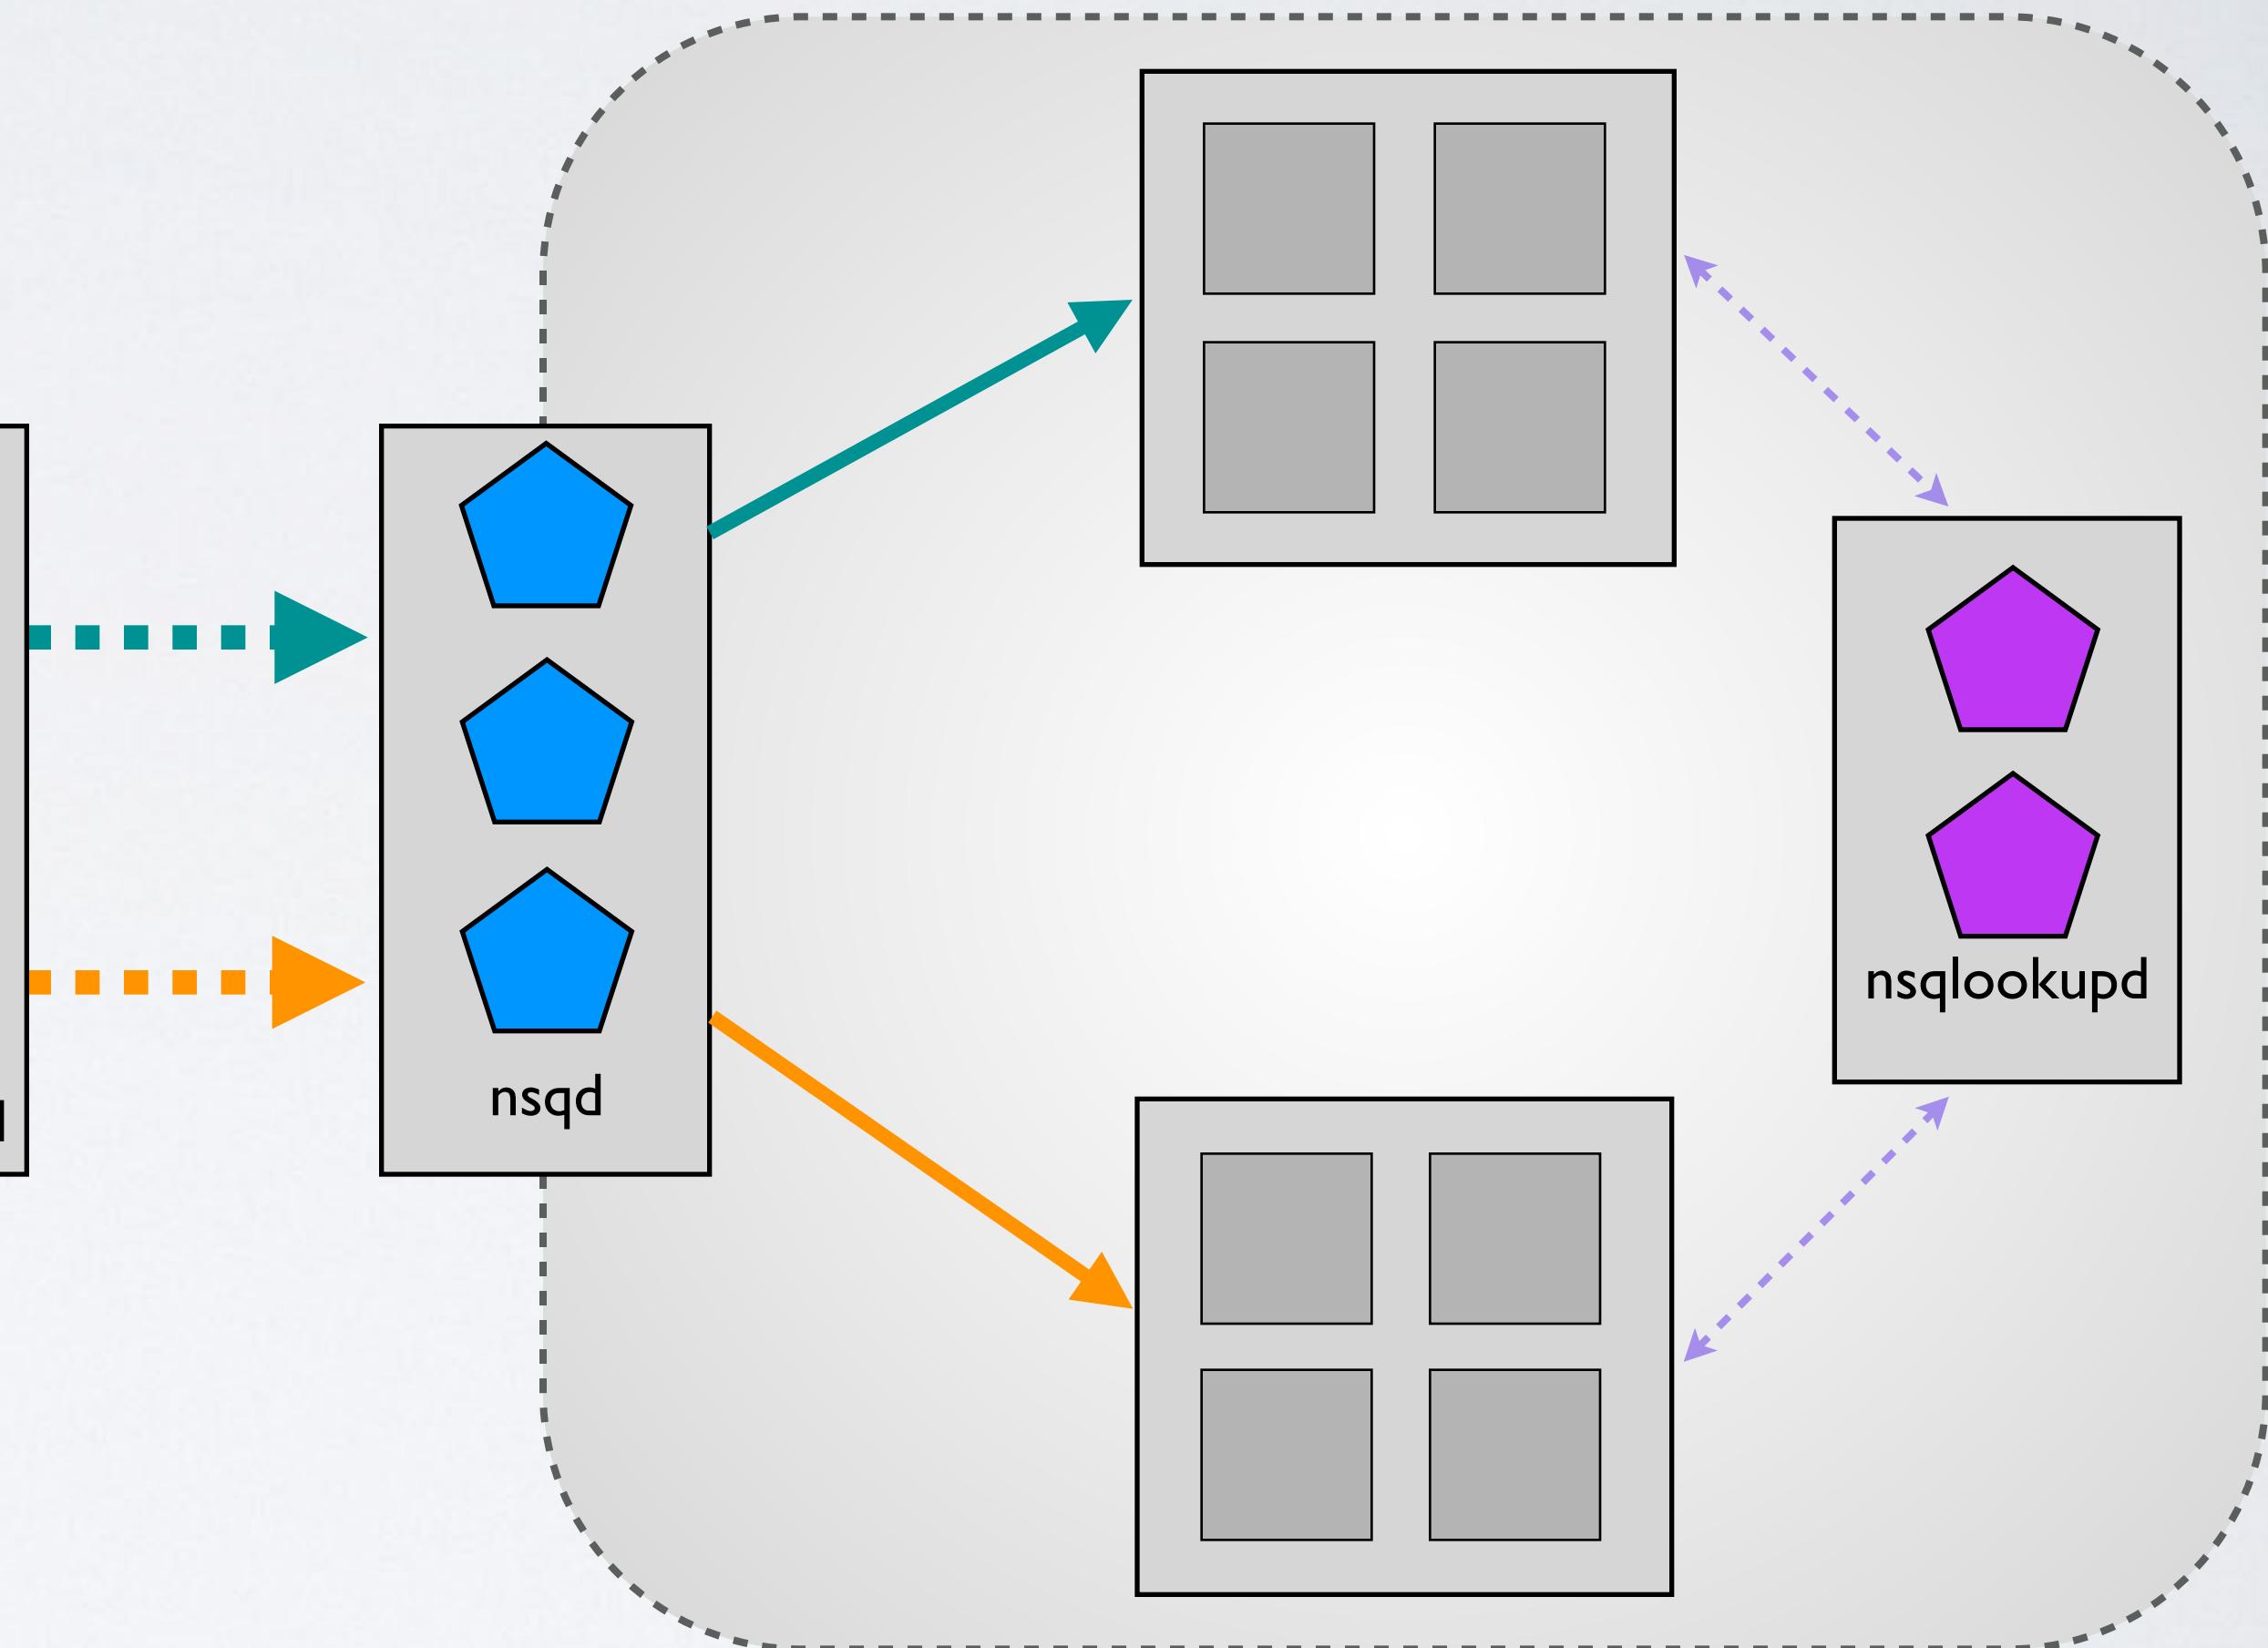


TO INFINITY AND BEYOND

Datacenter A



Datacenter B



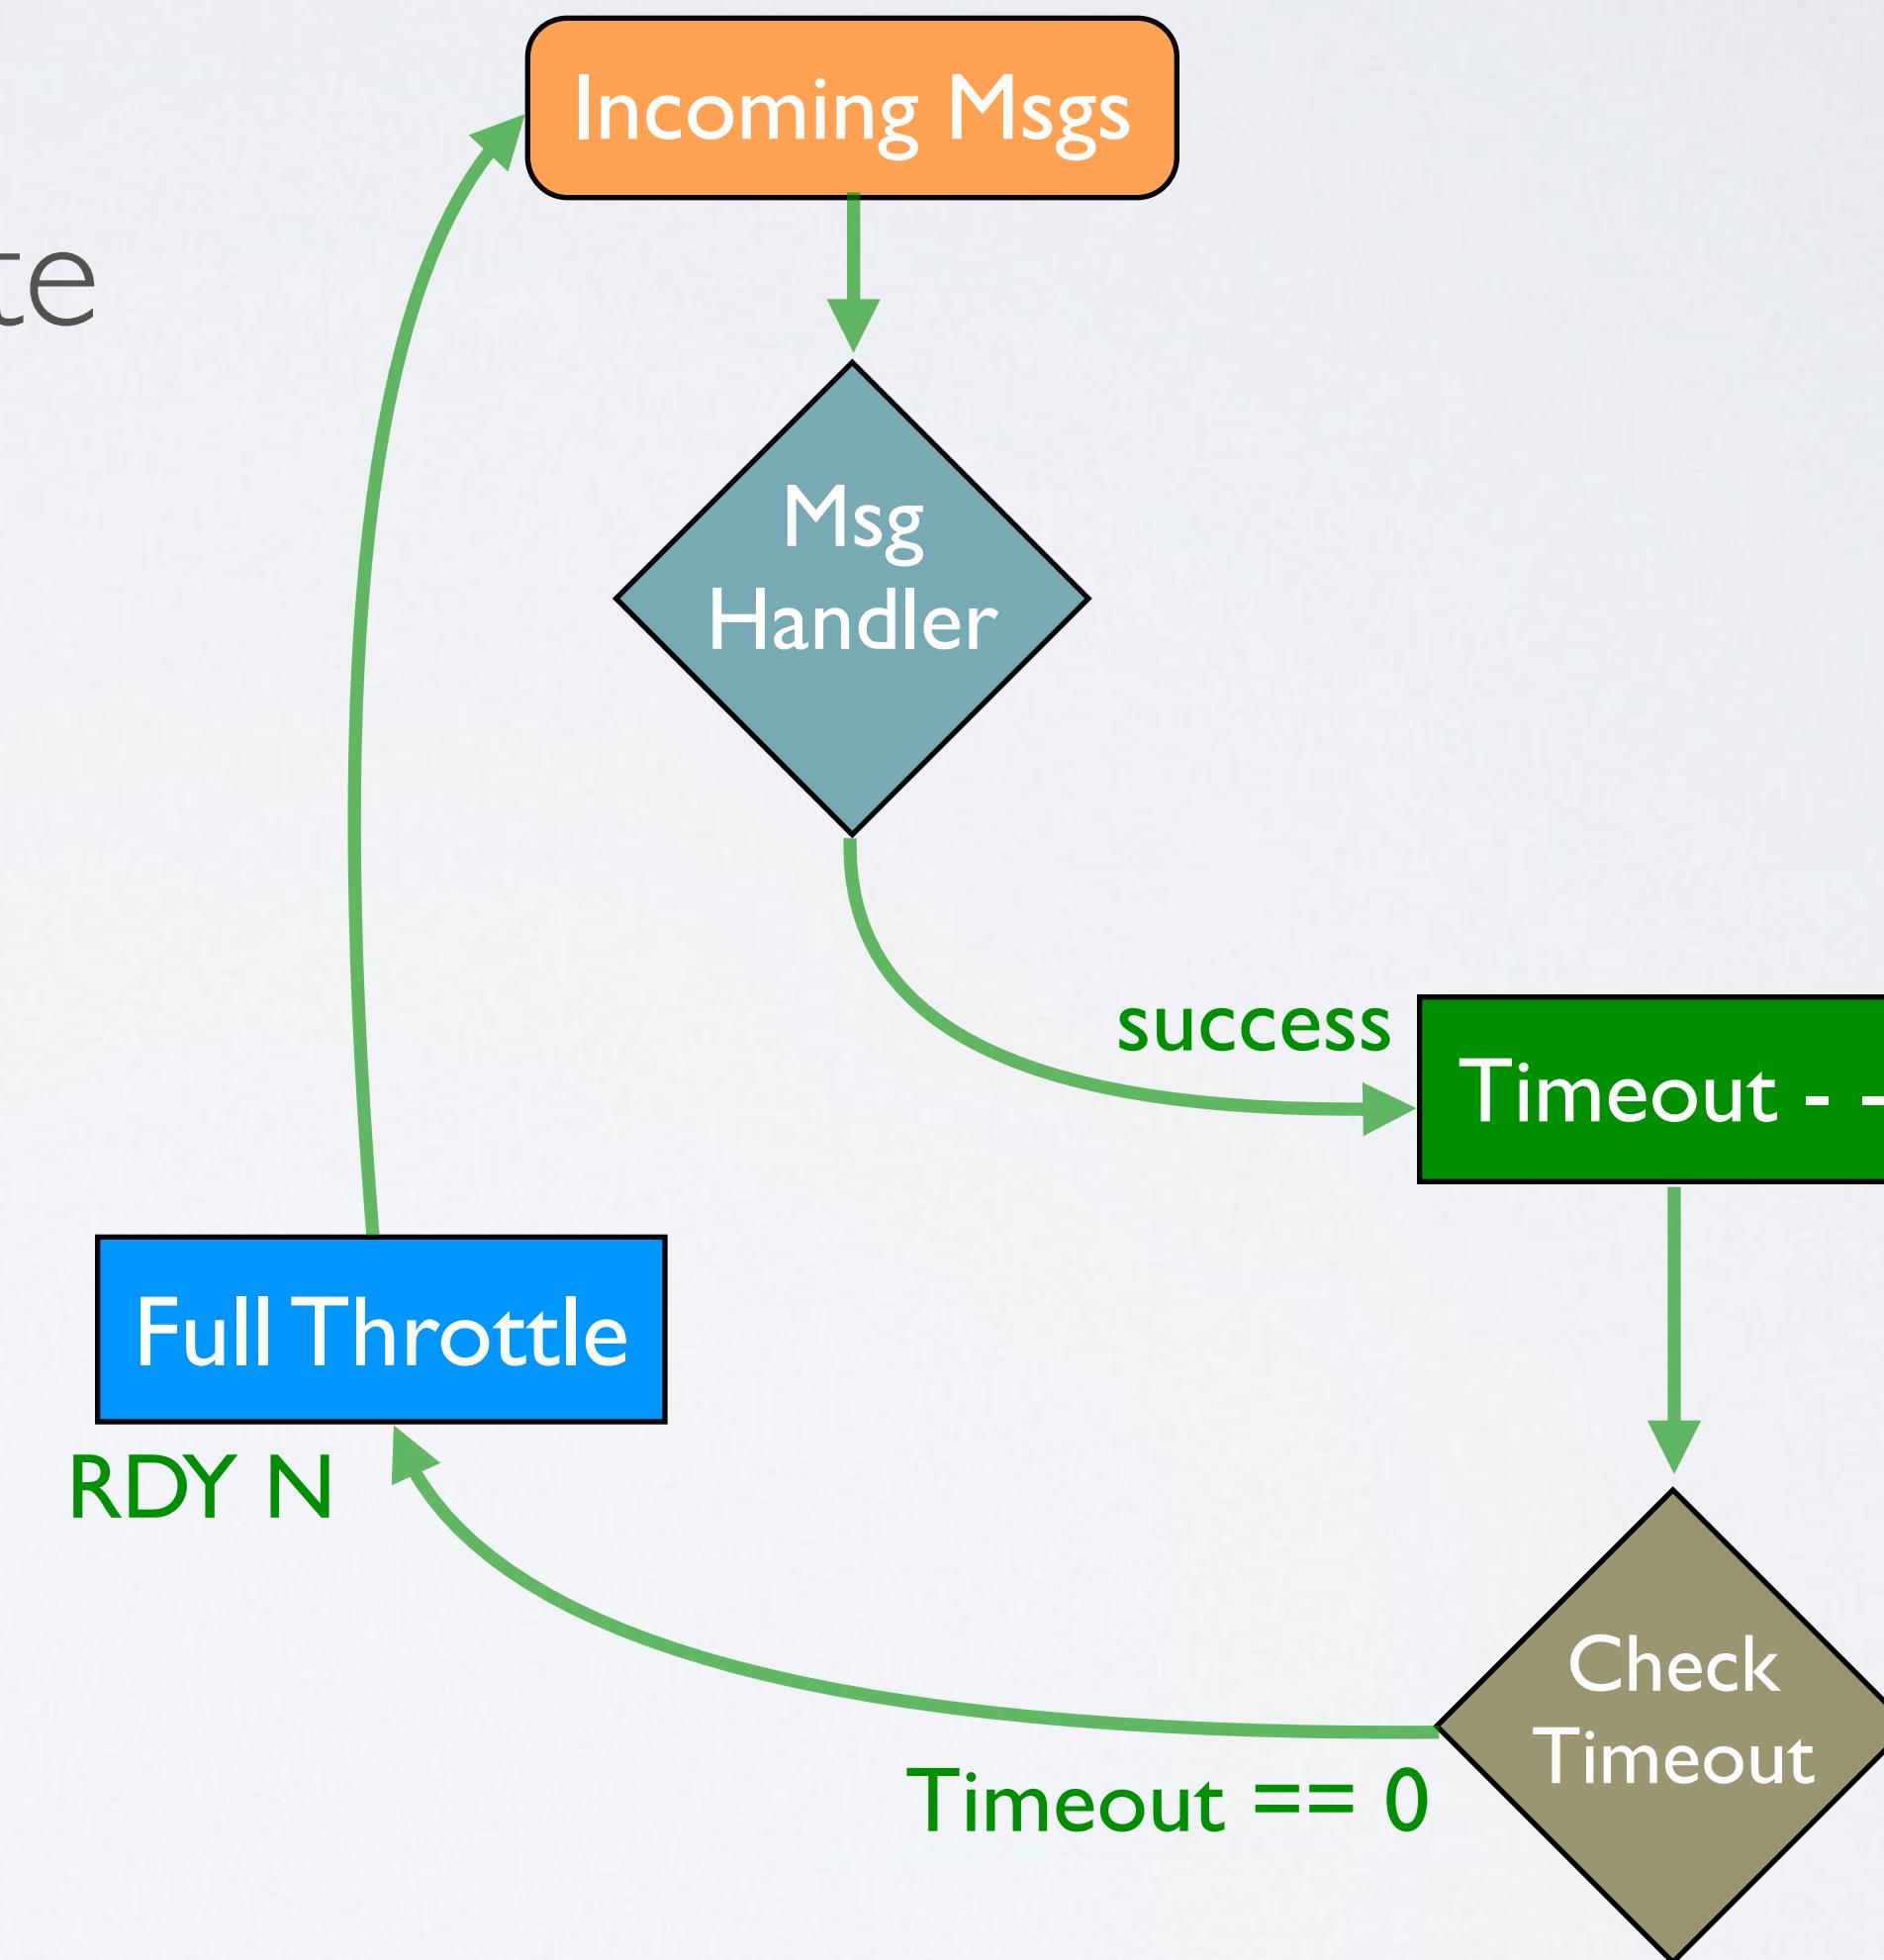
OPS

GUARANTEES

- messages are delivered *at least once*
- messages are *not durable* (by default)
- messages received are *un-ordered*
- consumers eventually find all topic producers

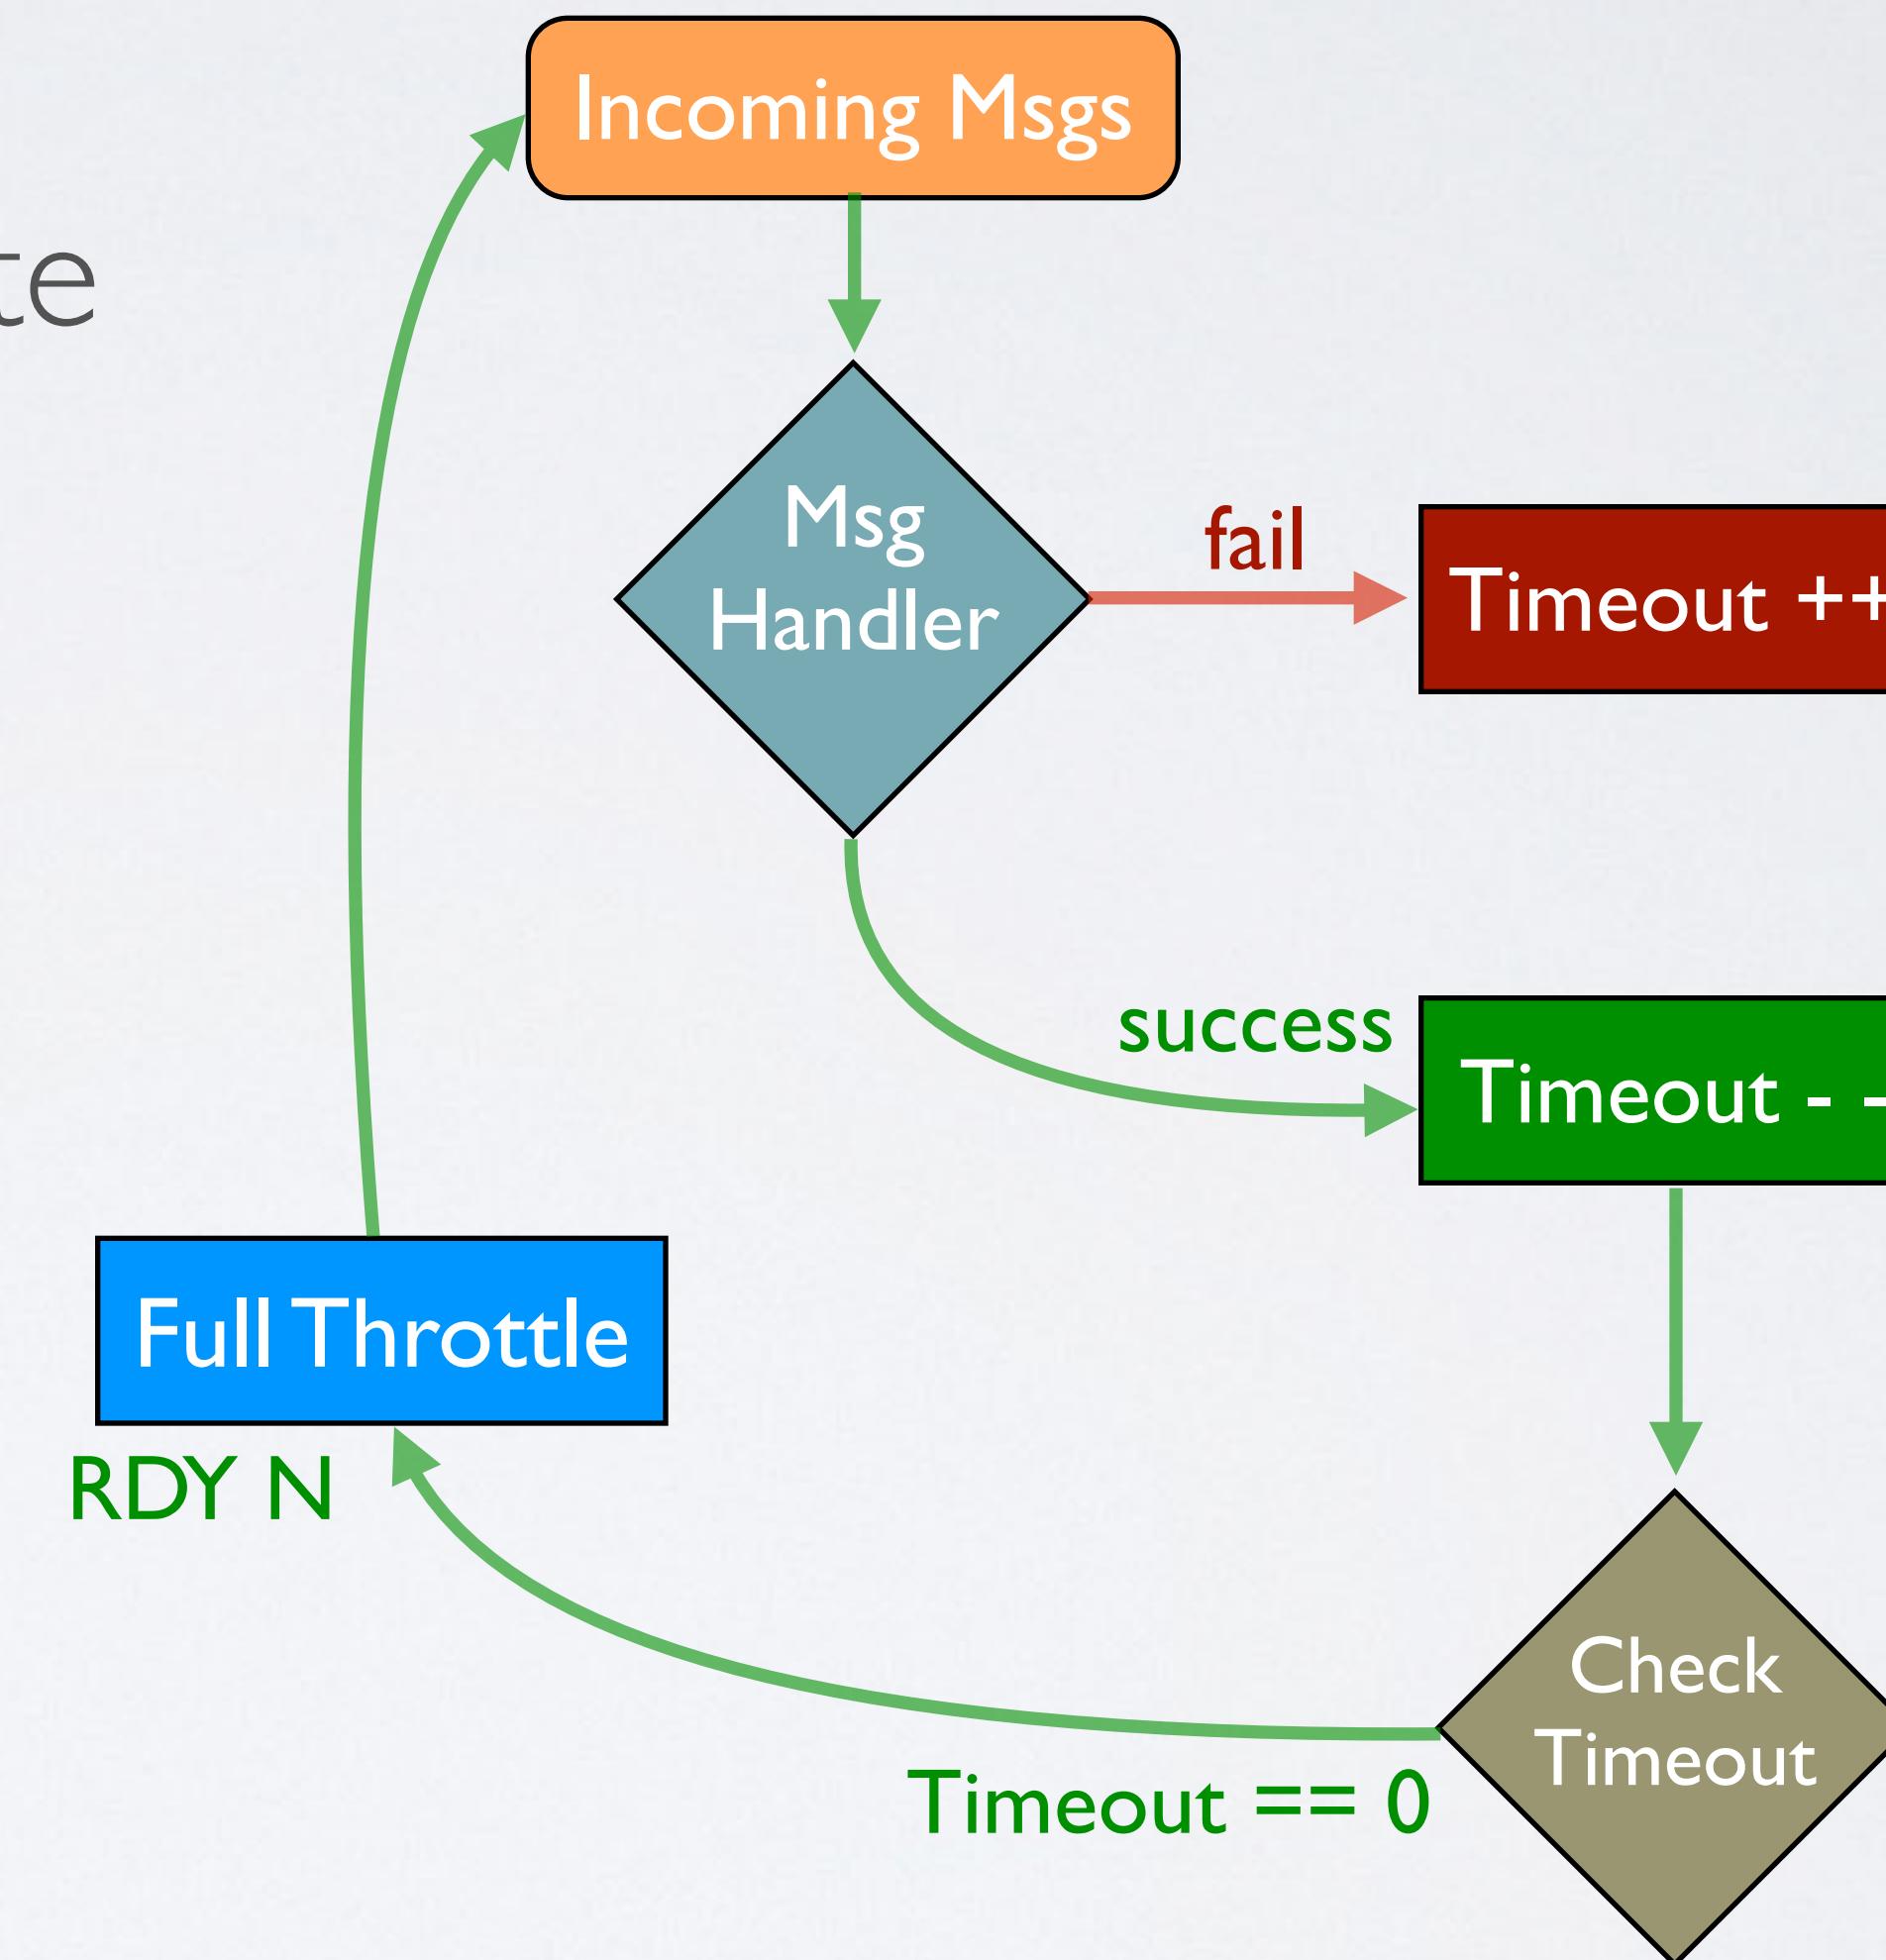
BACK PRESSURE

- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing



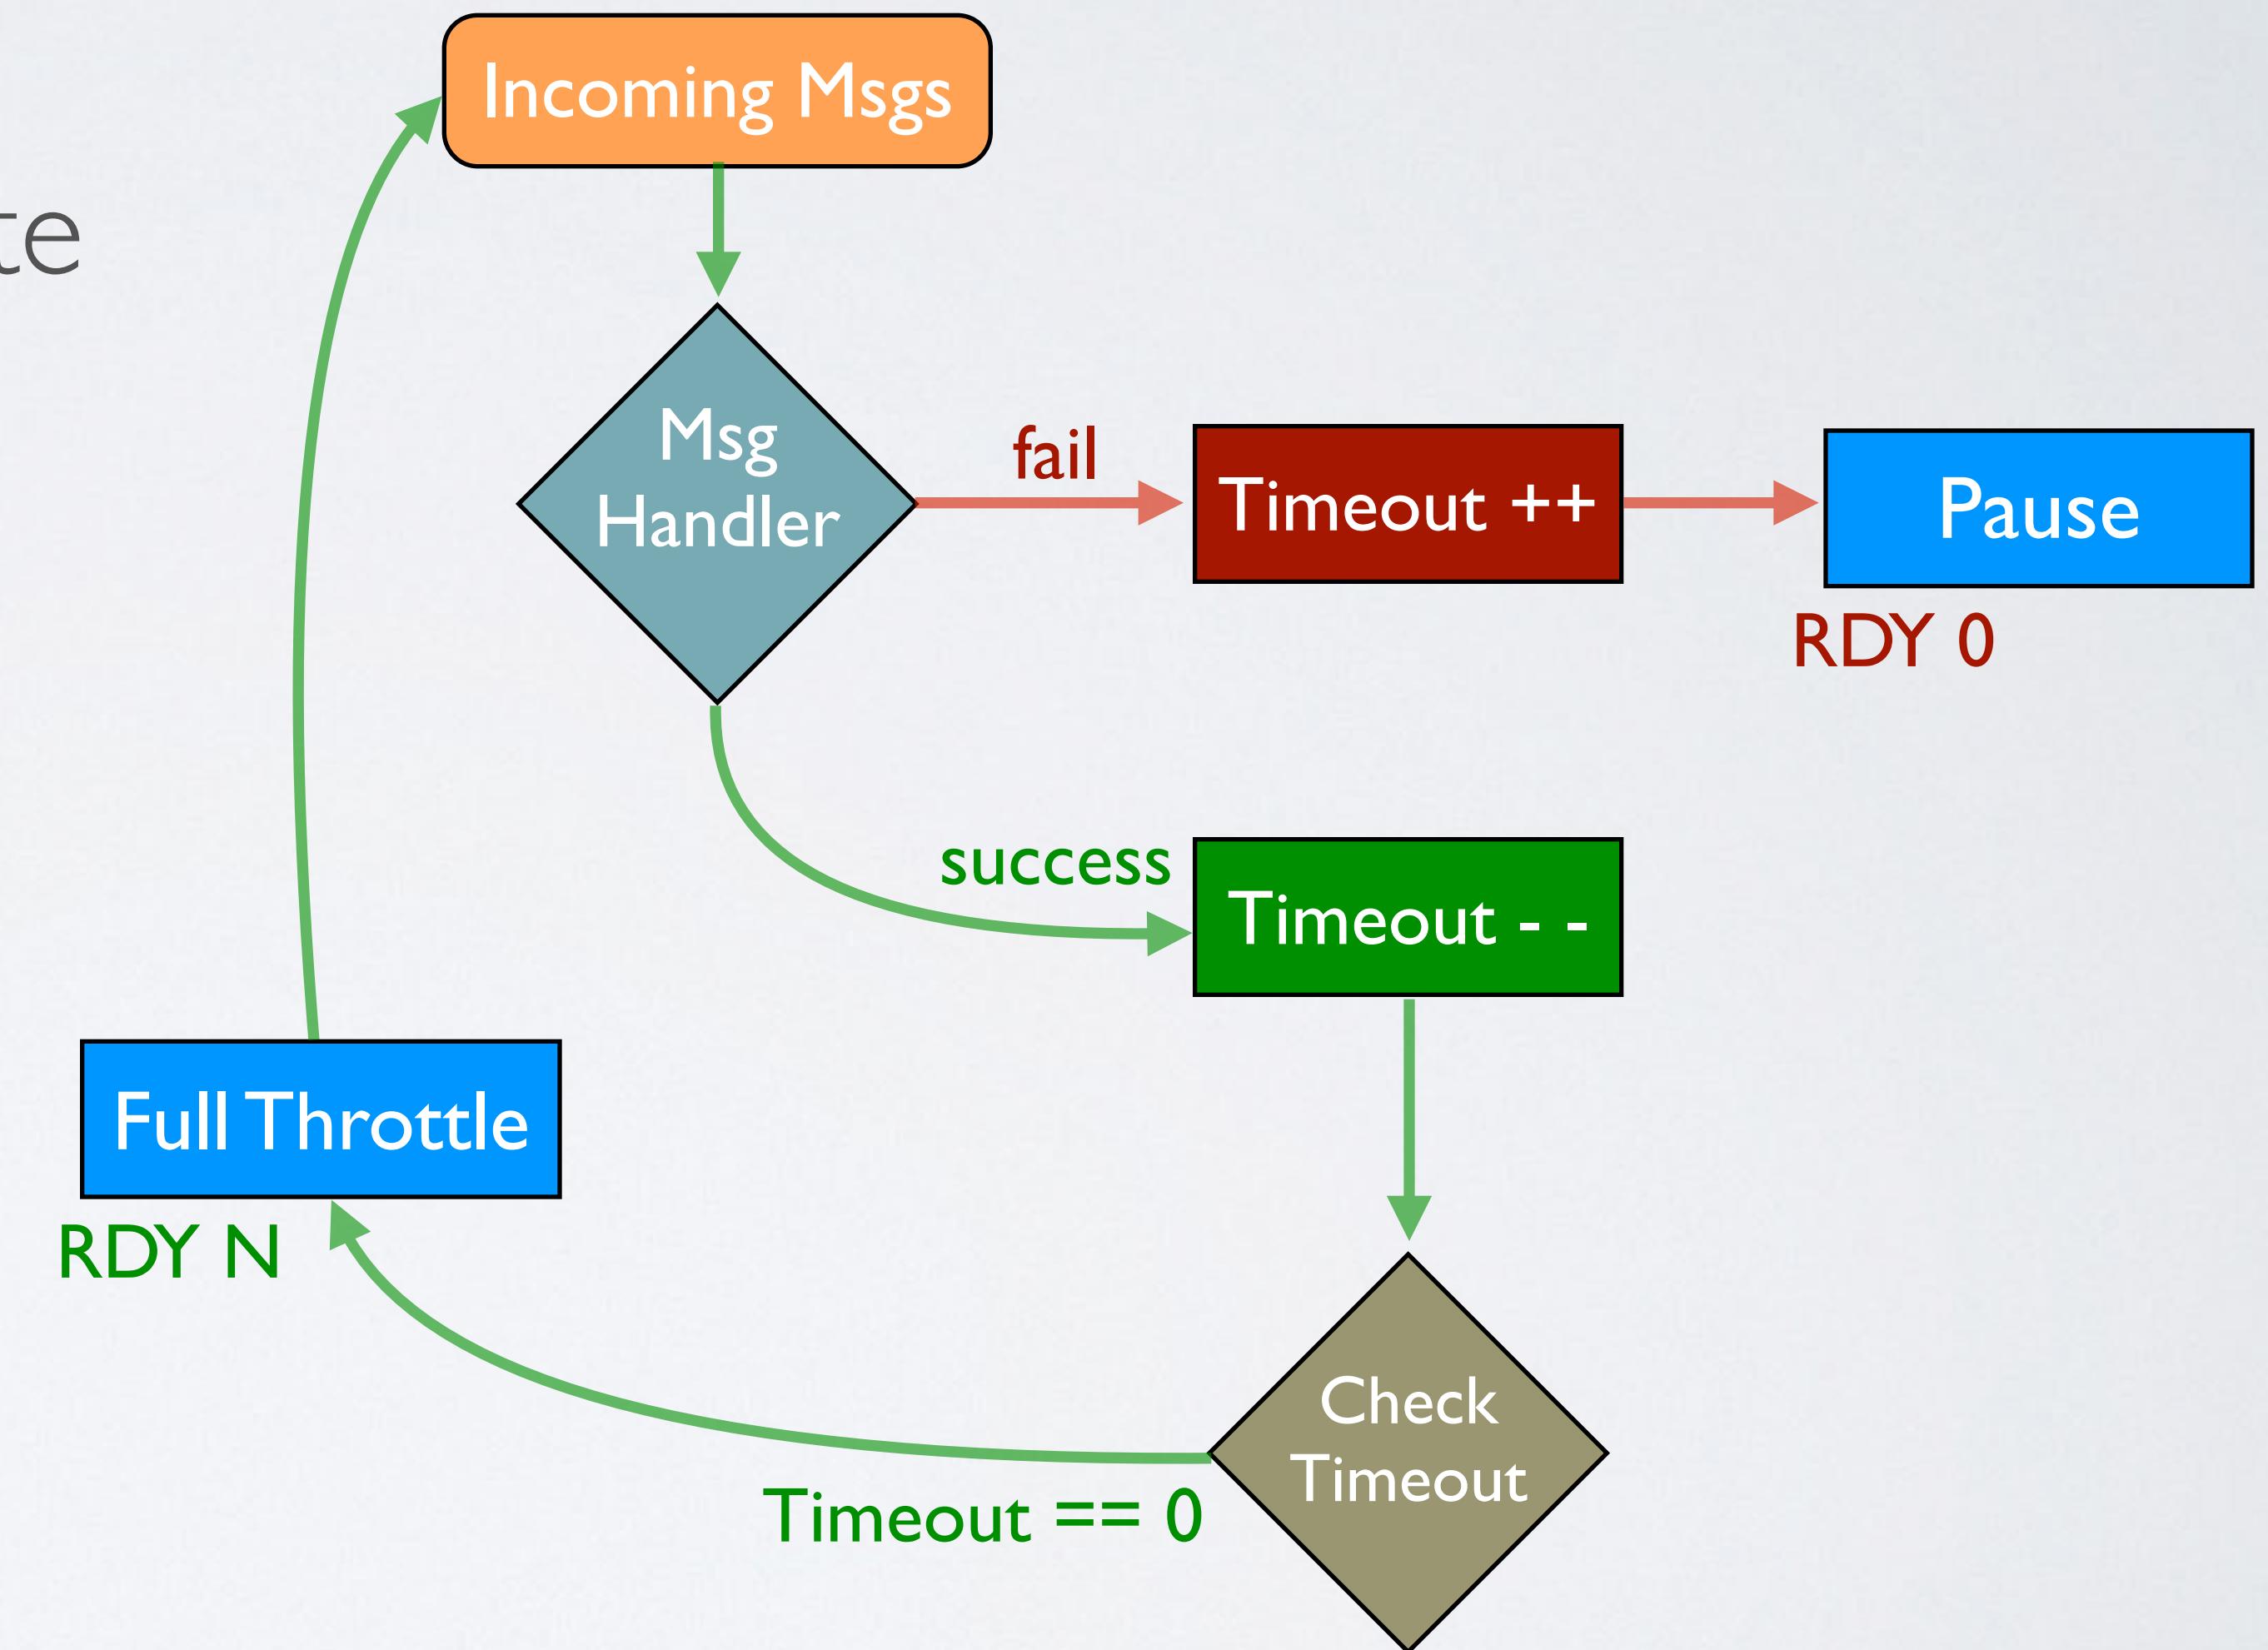
BACK PRESSURE

- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing



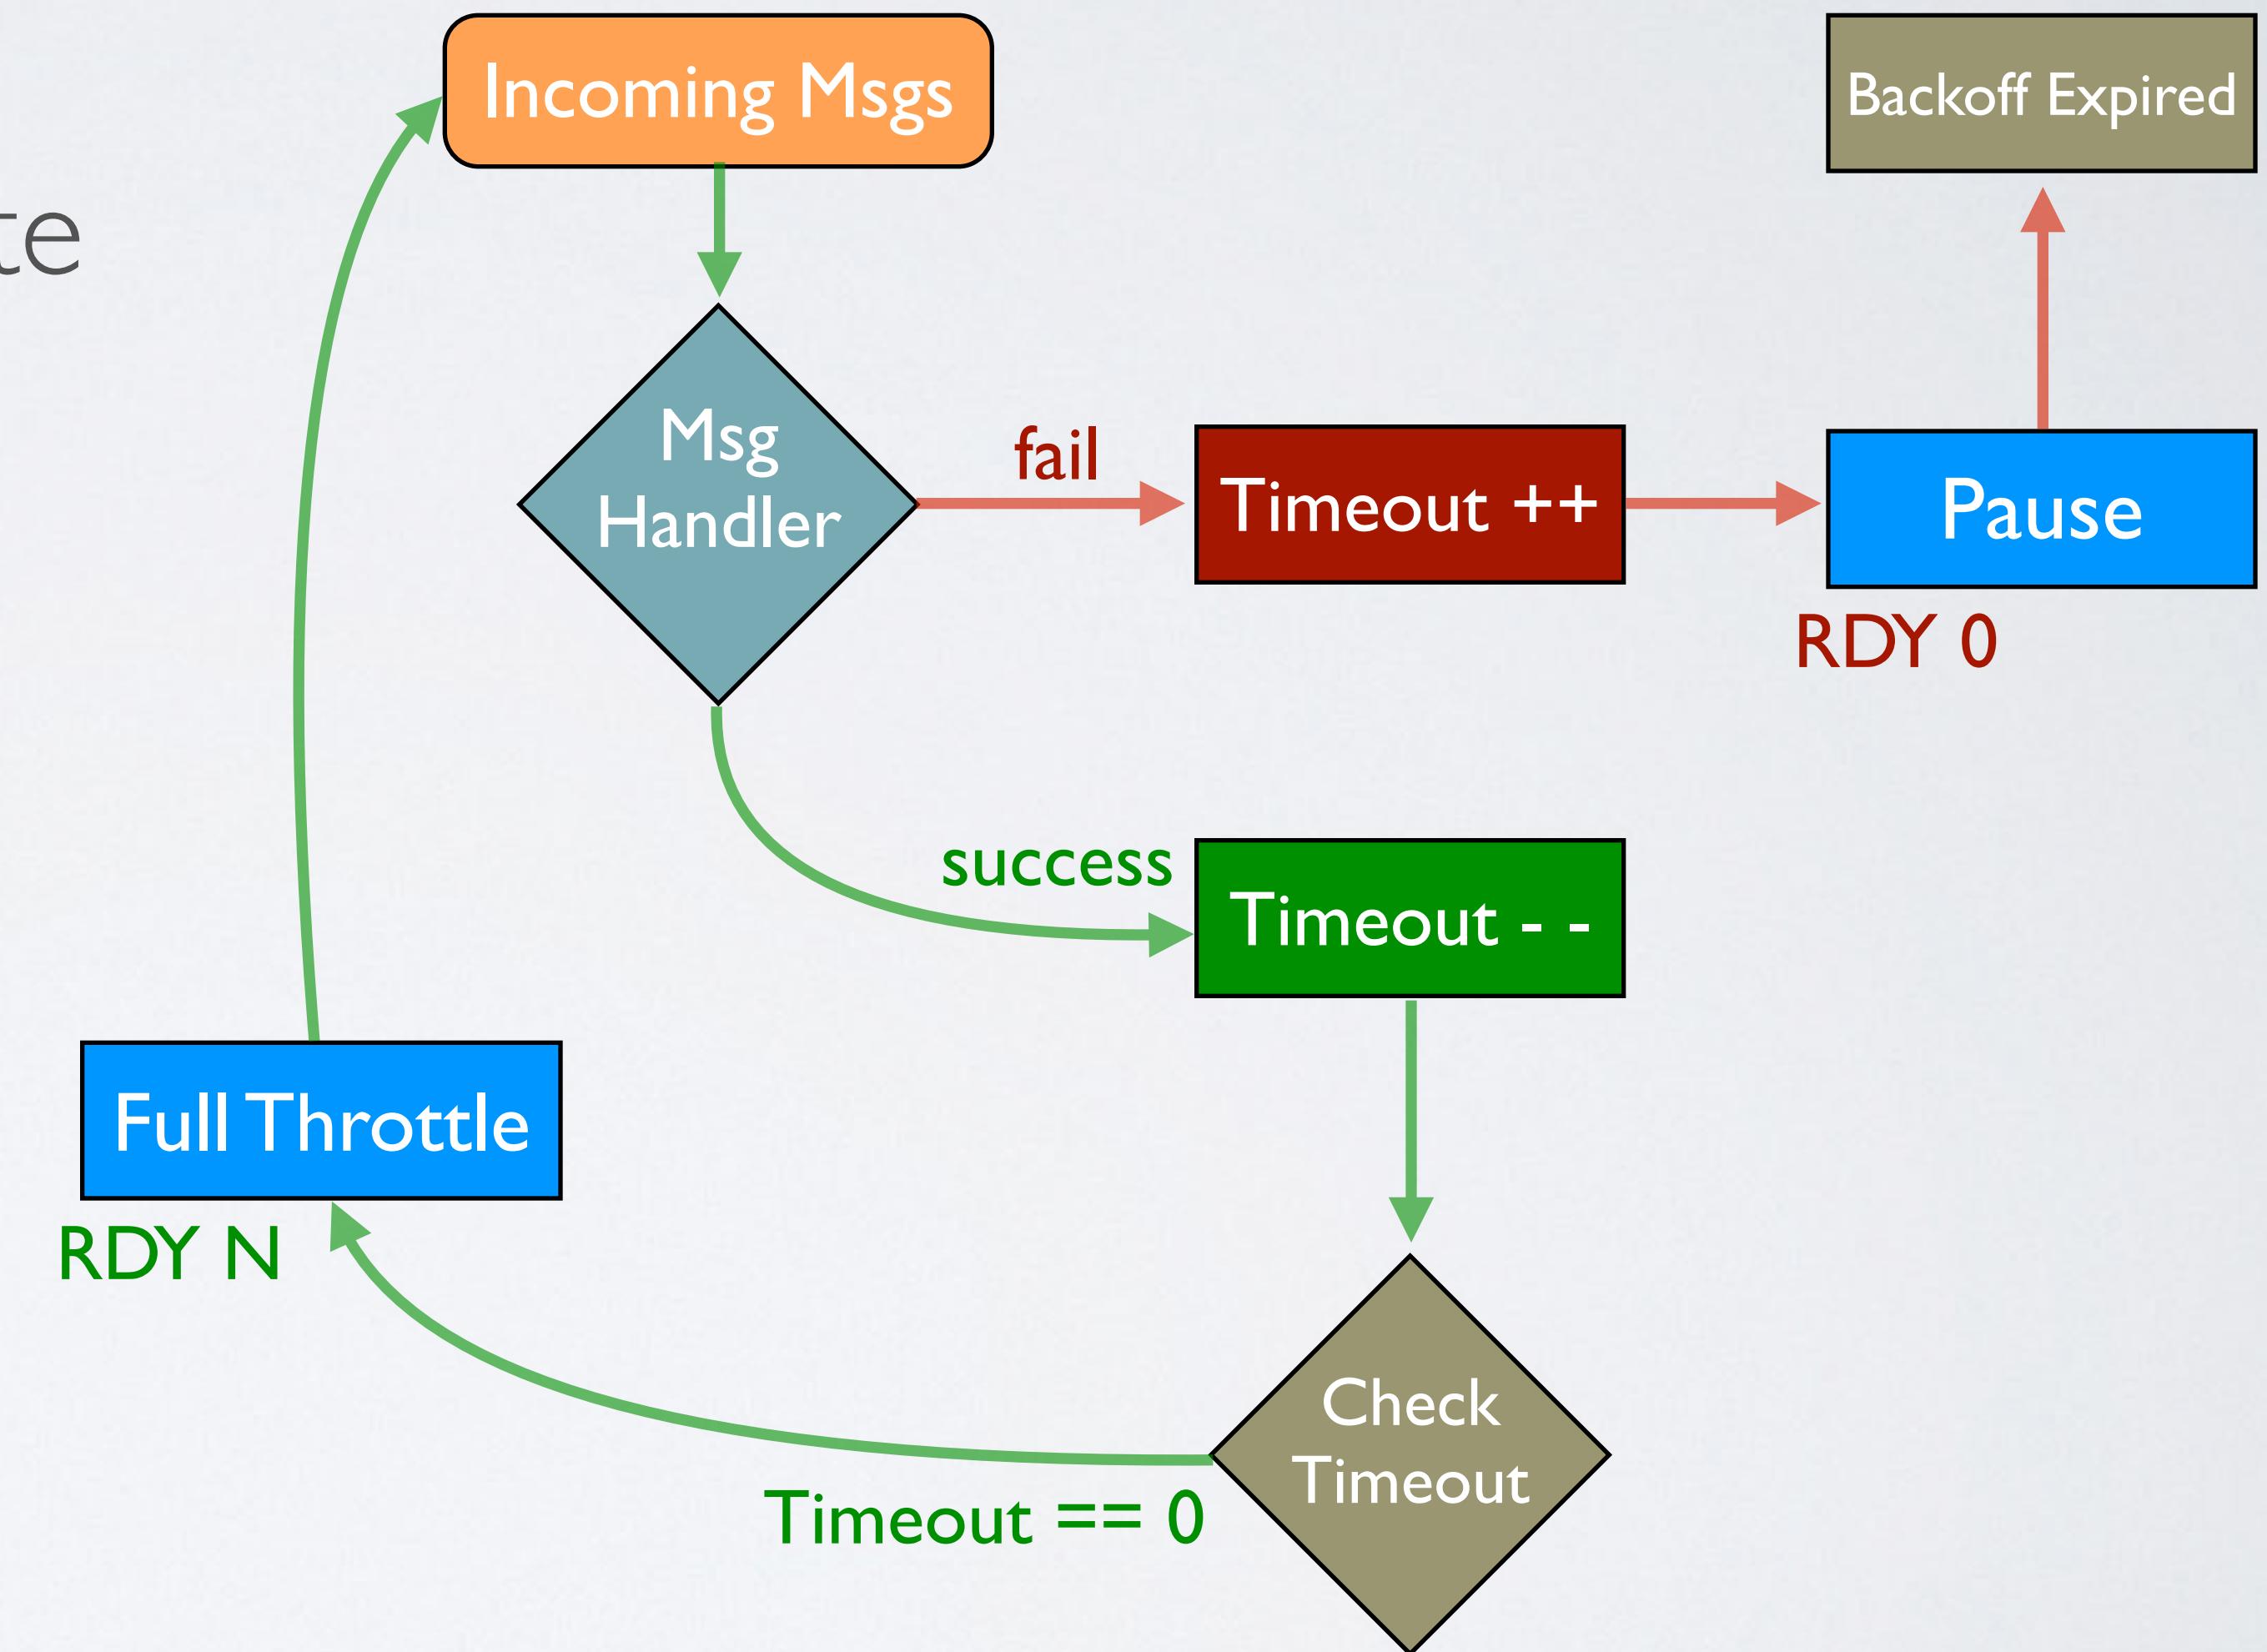
BACK PRESSURE

- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing



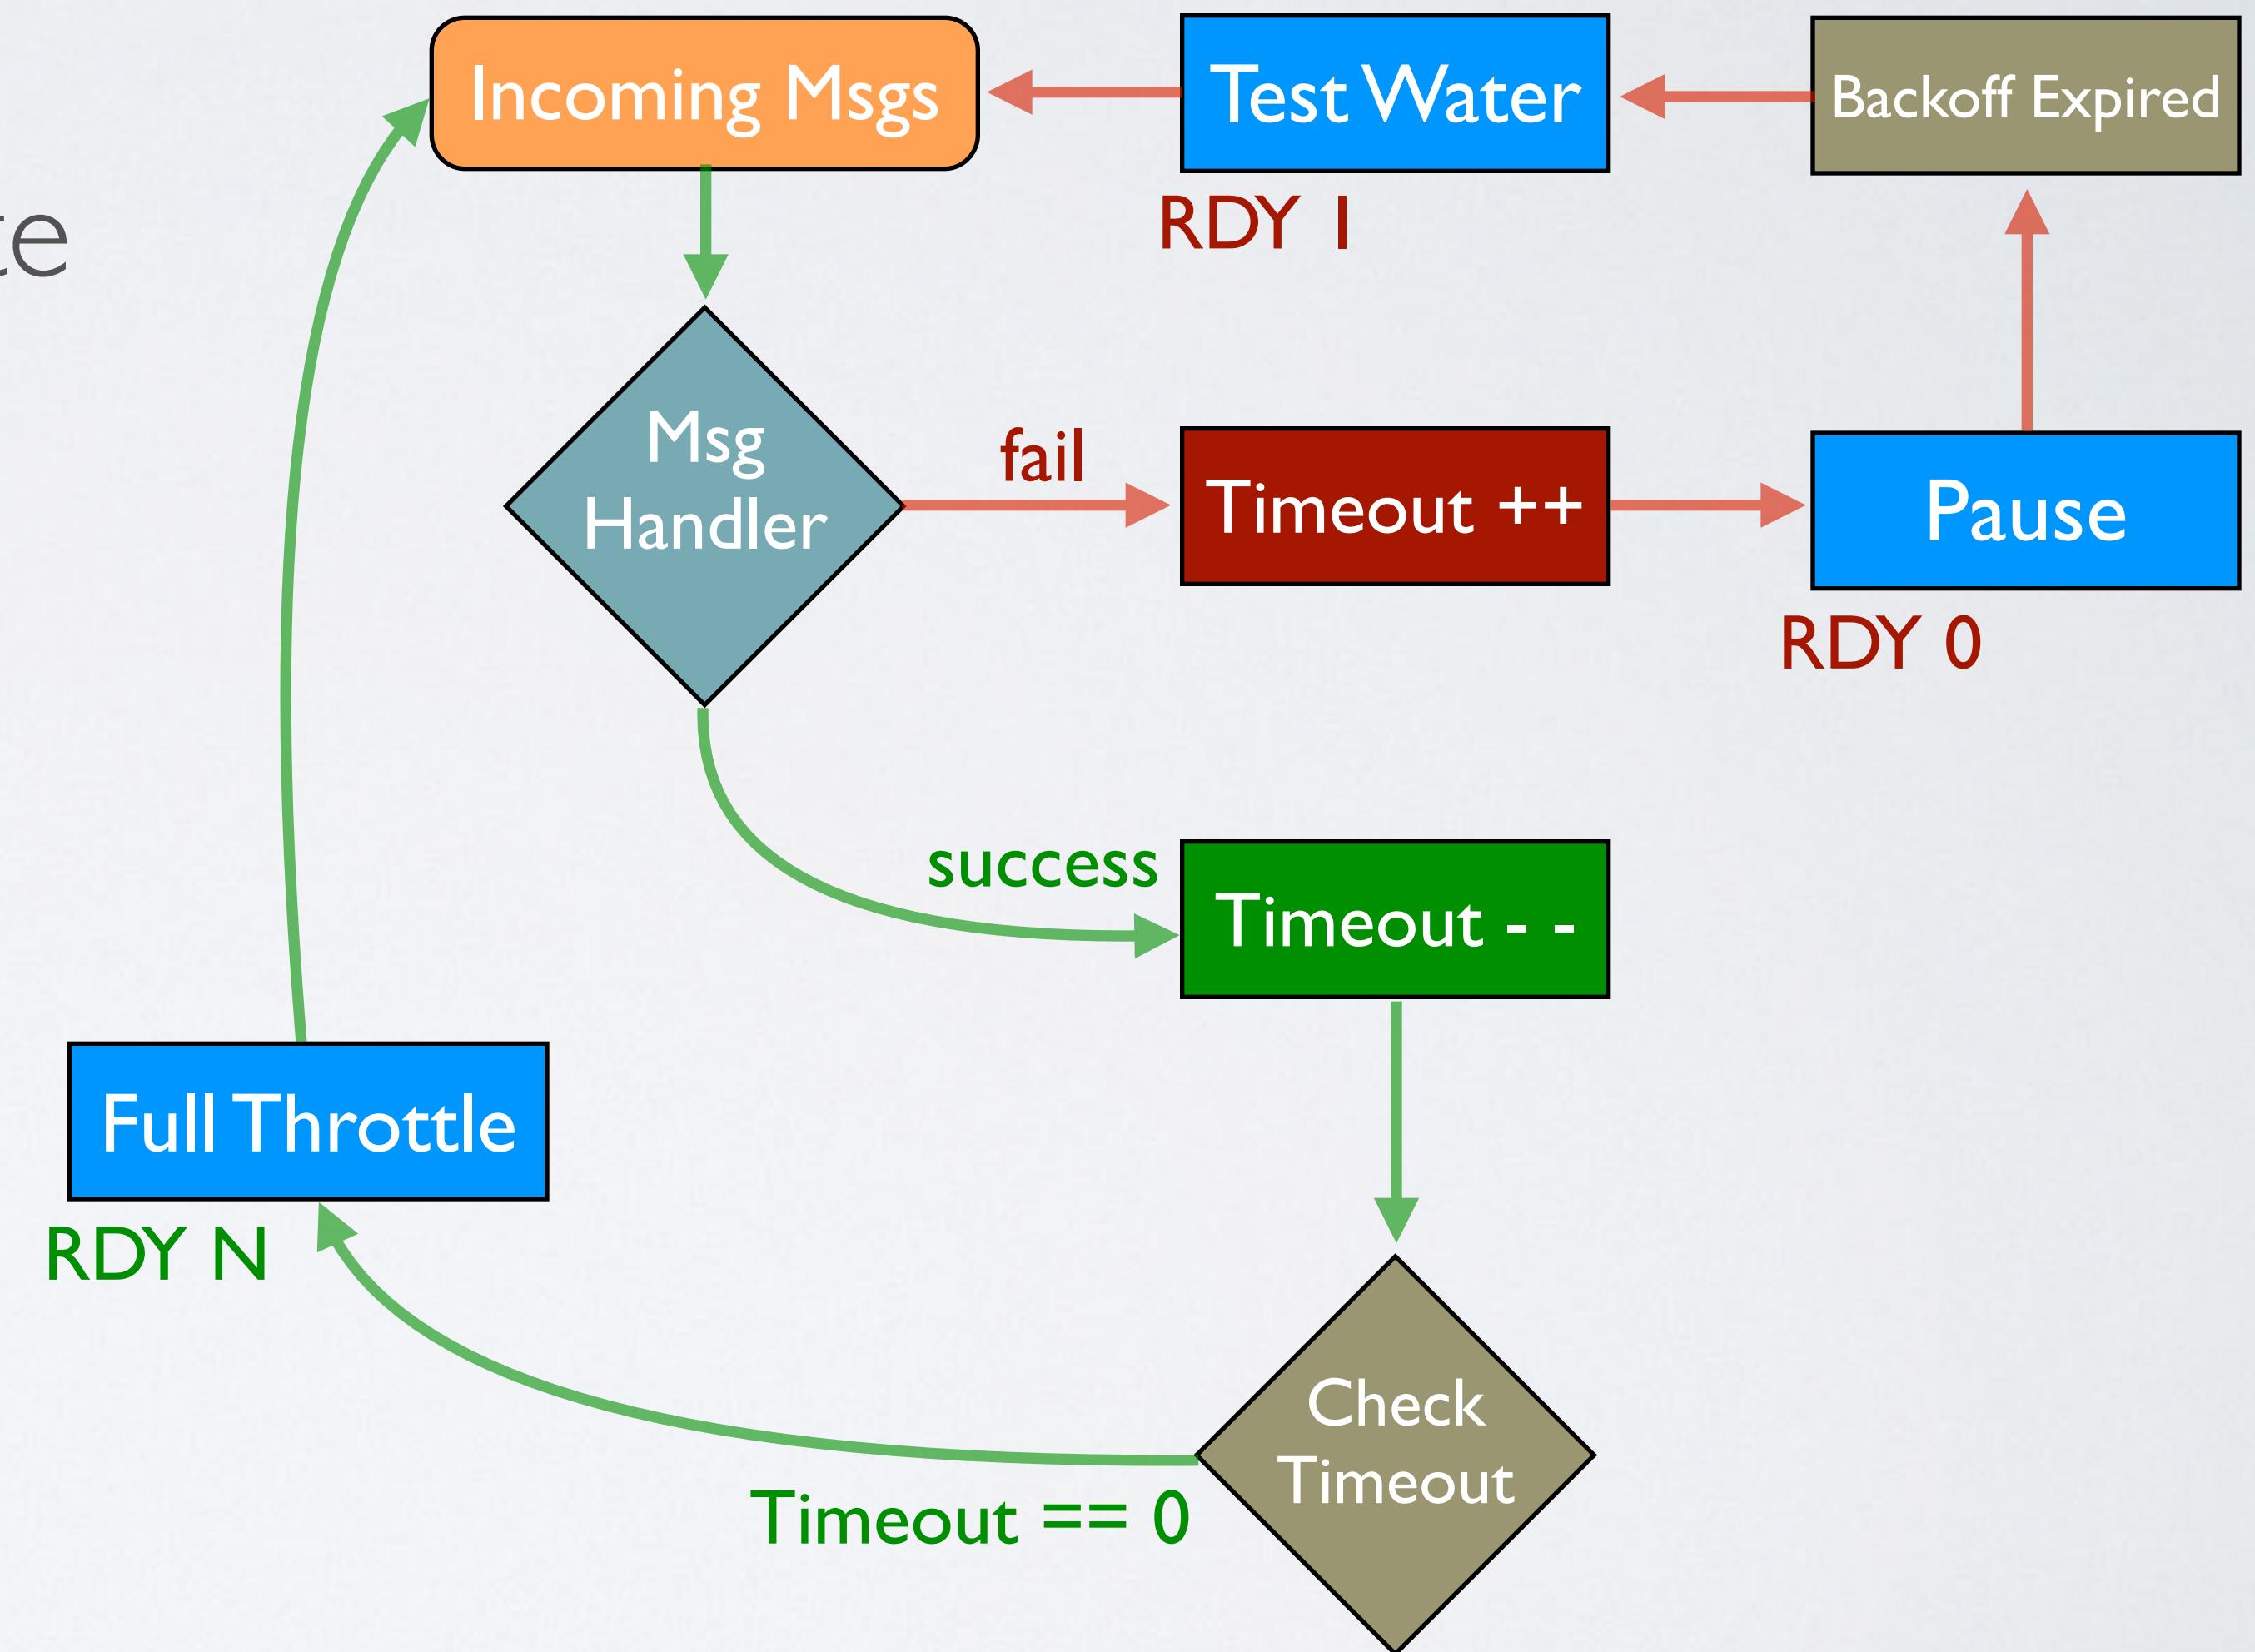
BACK PRESSURE

- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing



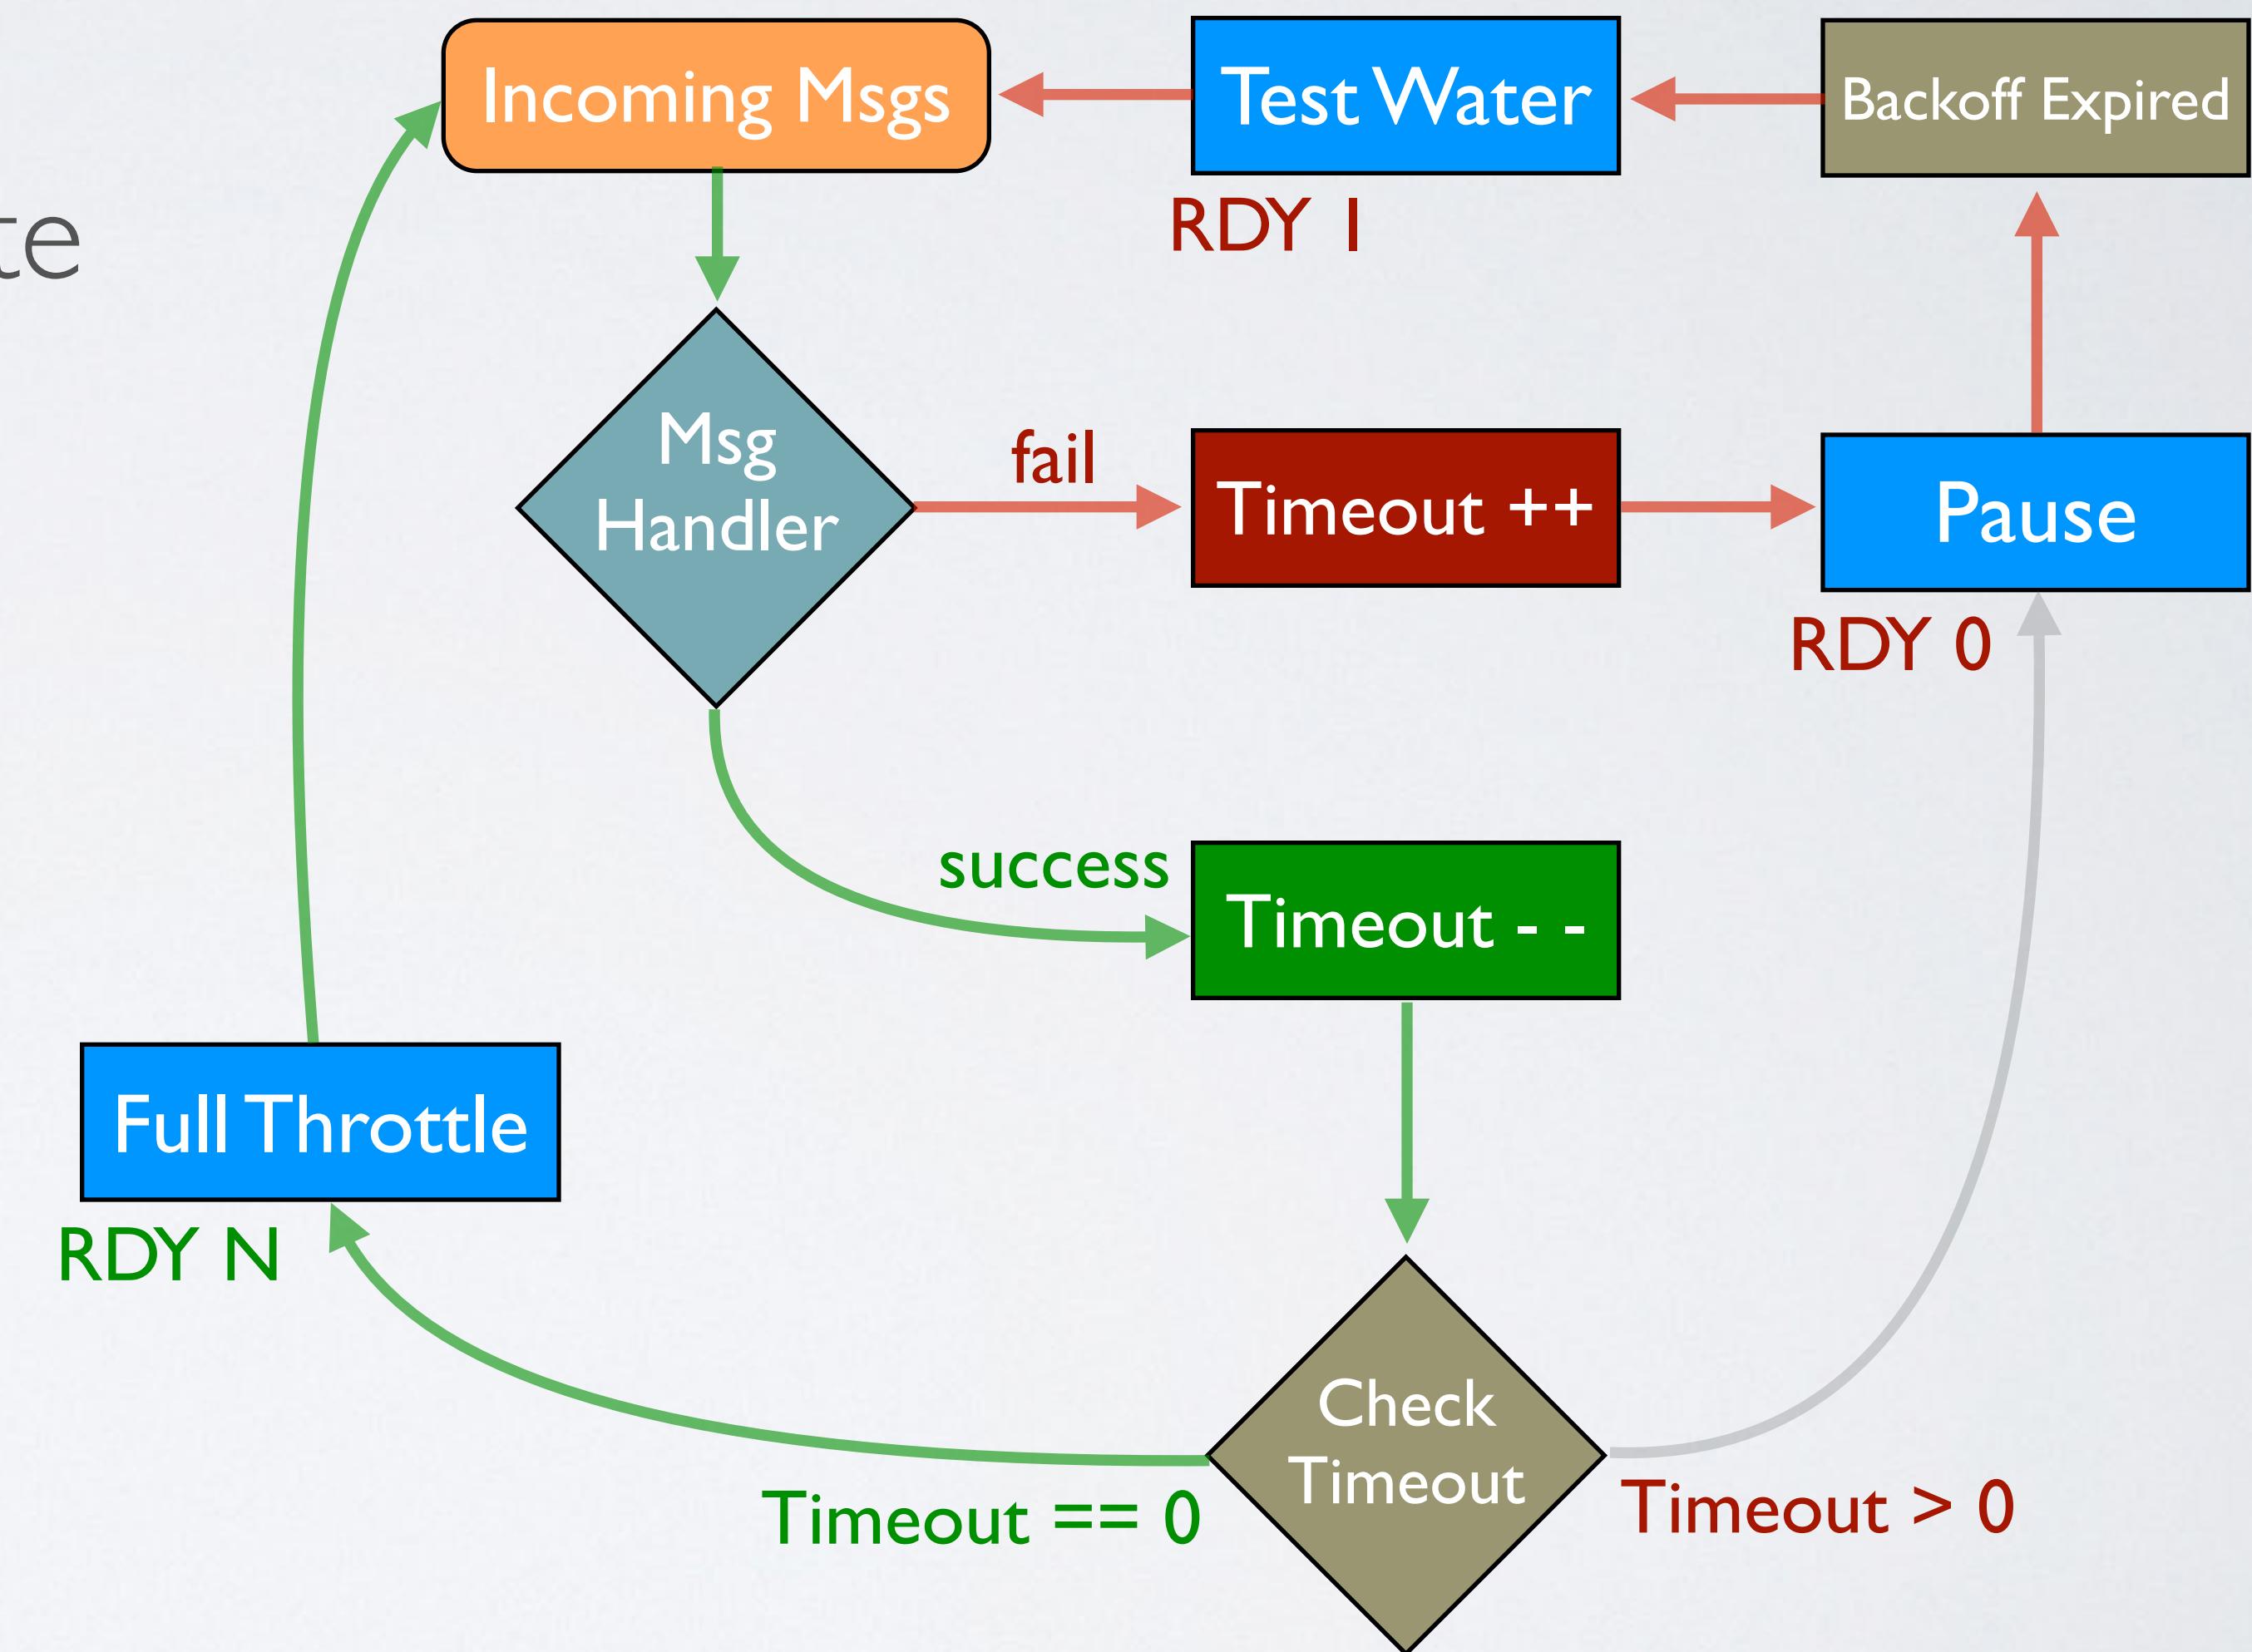
BACK PRESSURE

- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing

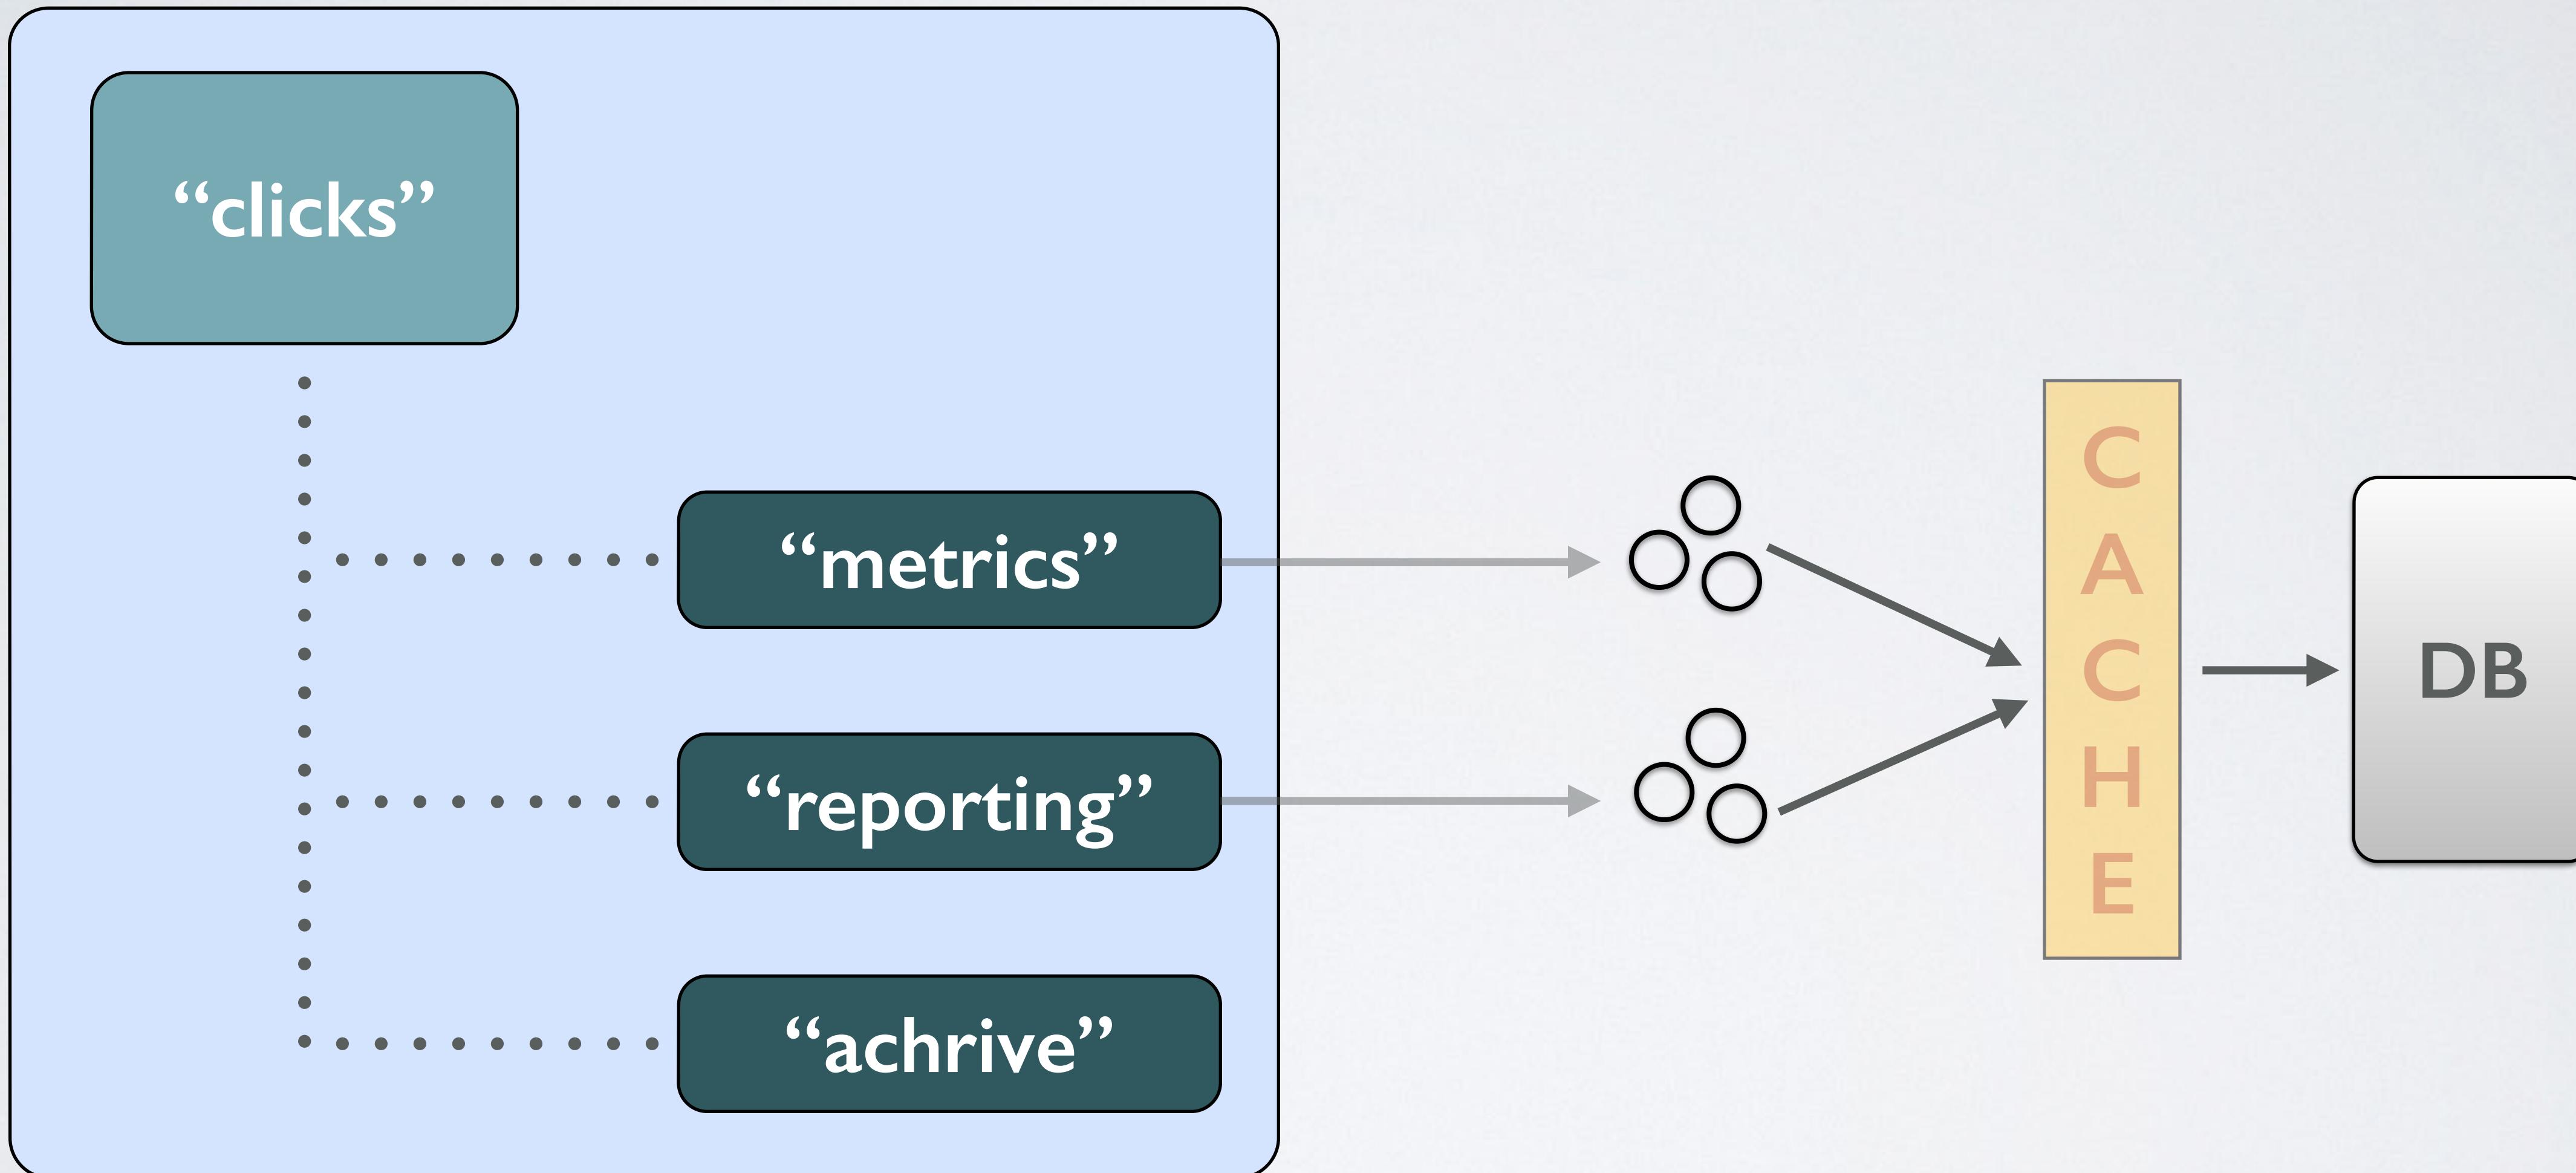


BACK PRESSURE

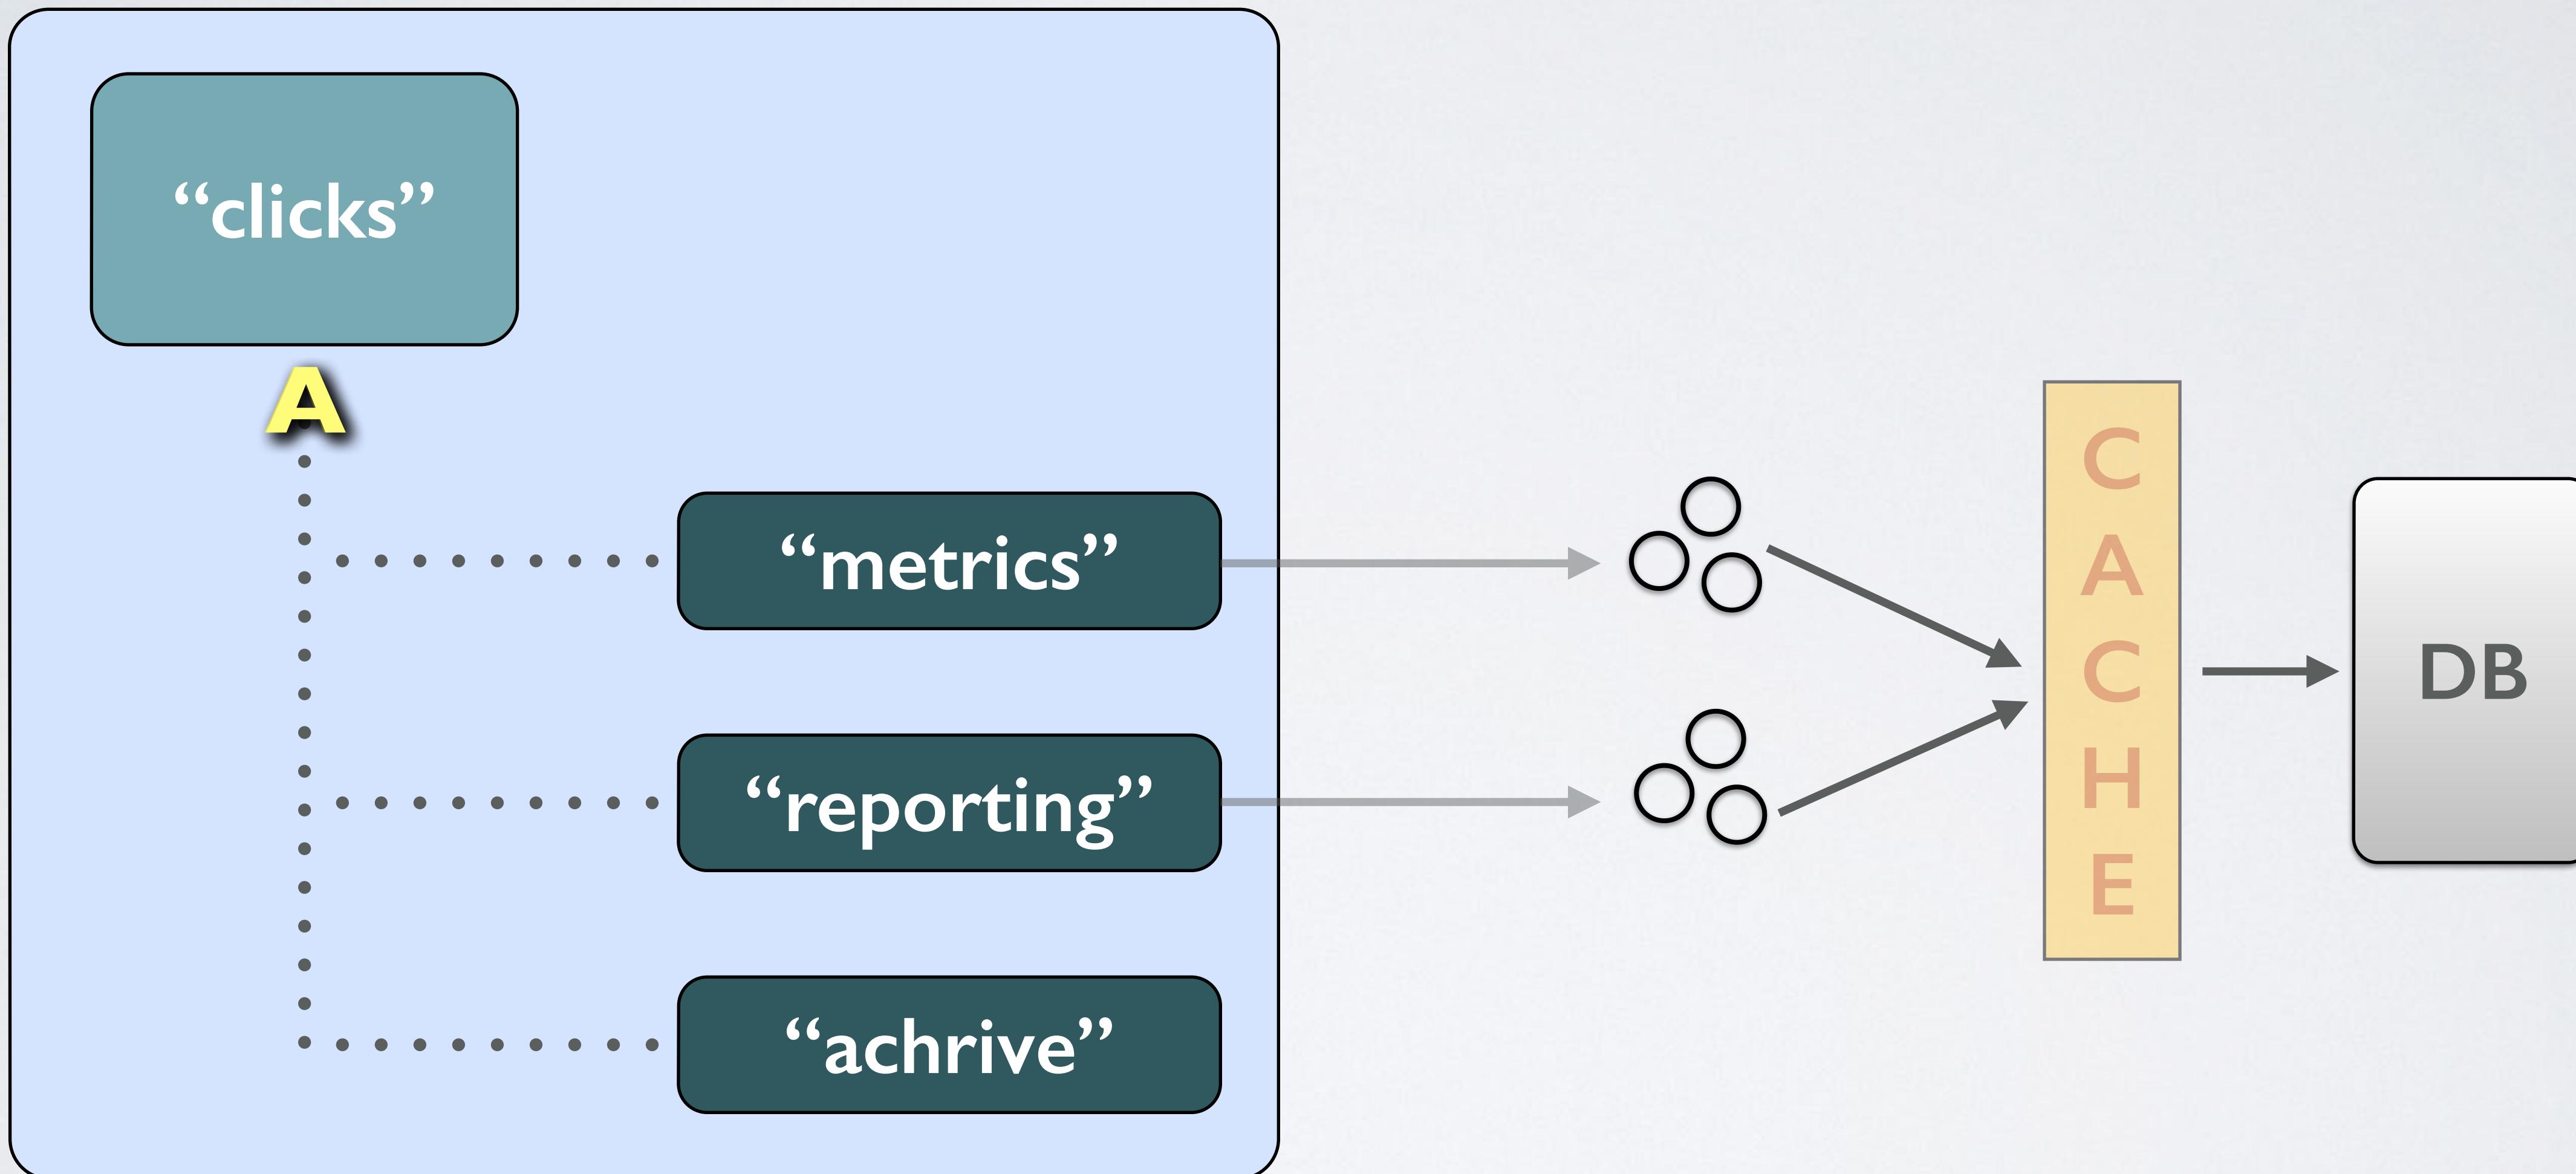
- clients control flow - **RDY** state
 - concurrency knob
- processing fails?
 - slow down rate
- bad message?
 - limit attempts
 - delay re-processing



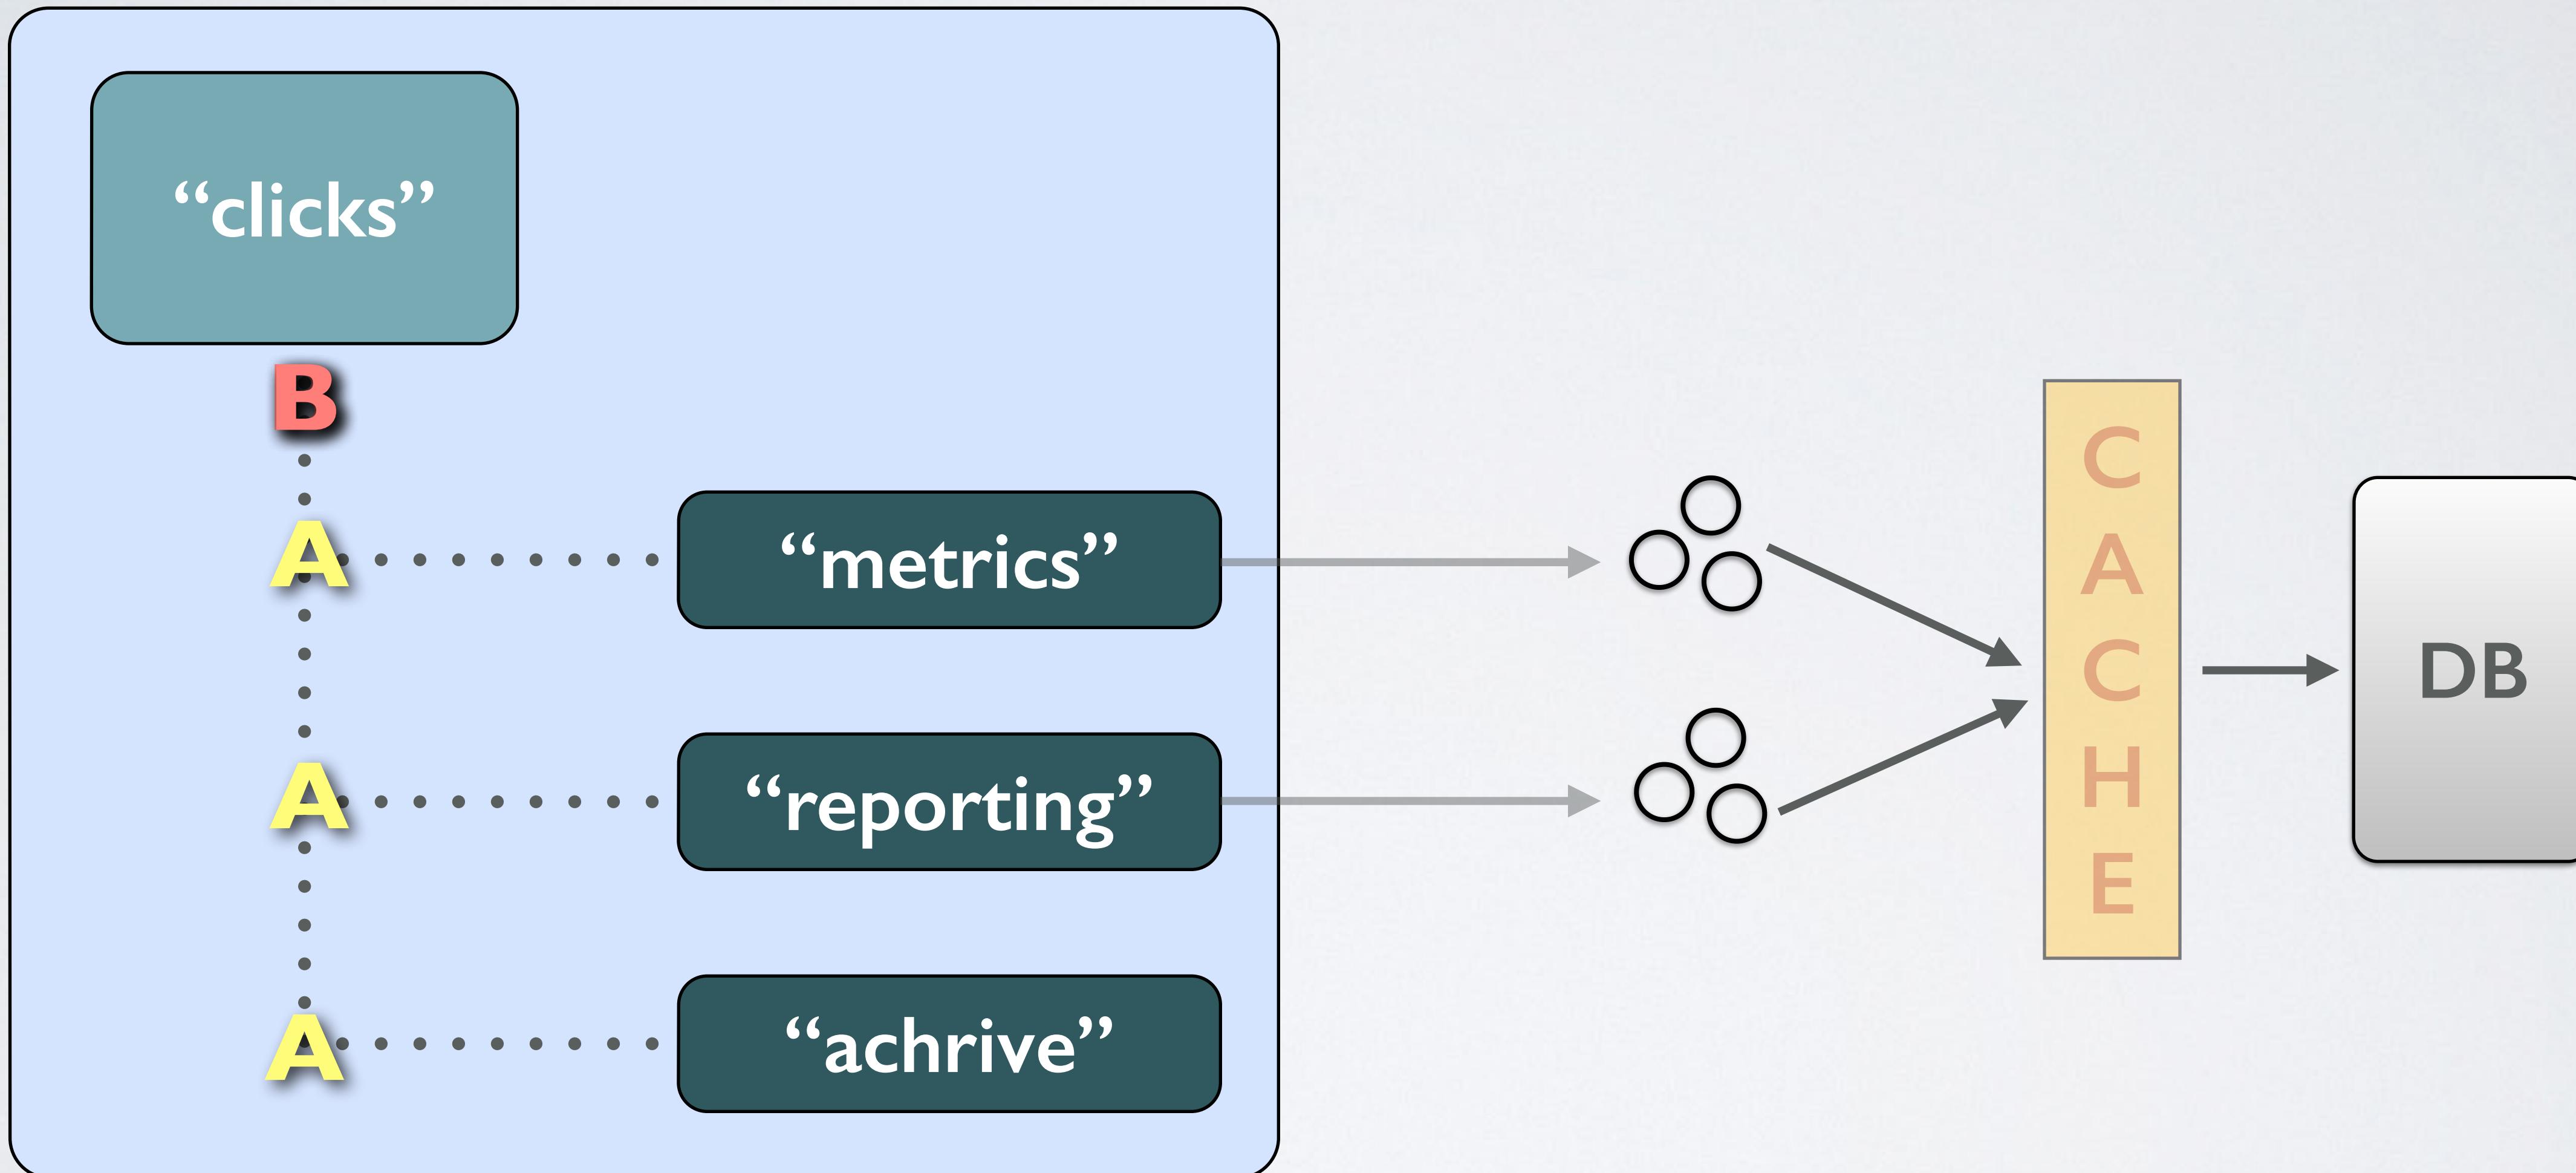
TOPIC & CHANNEL PAUSING



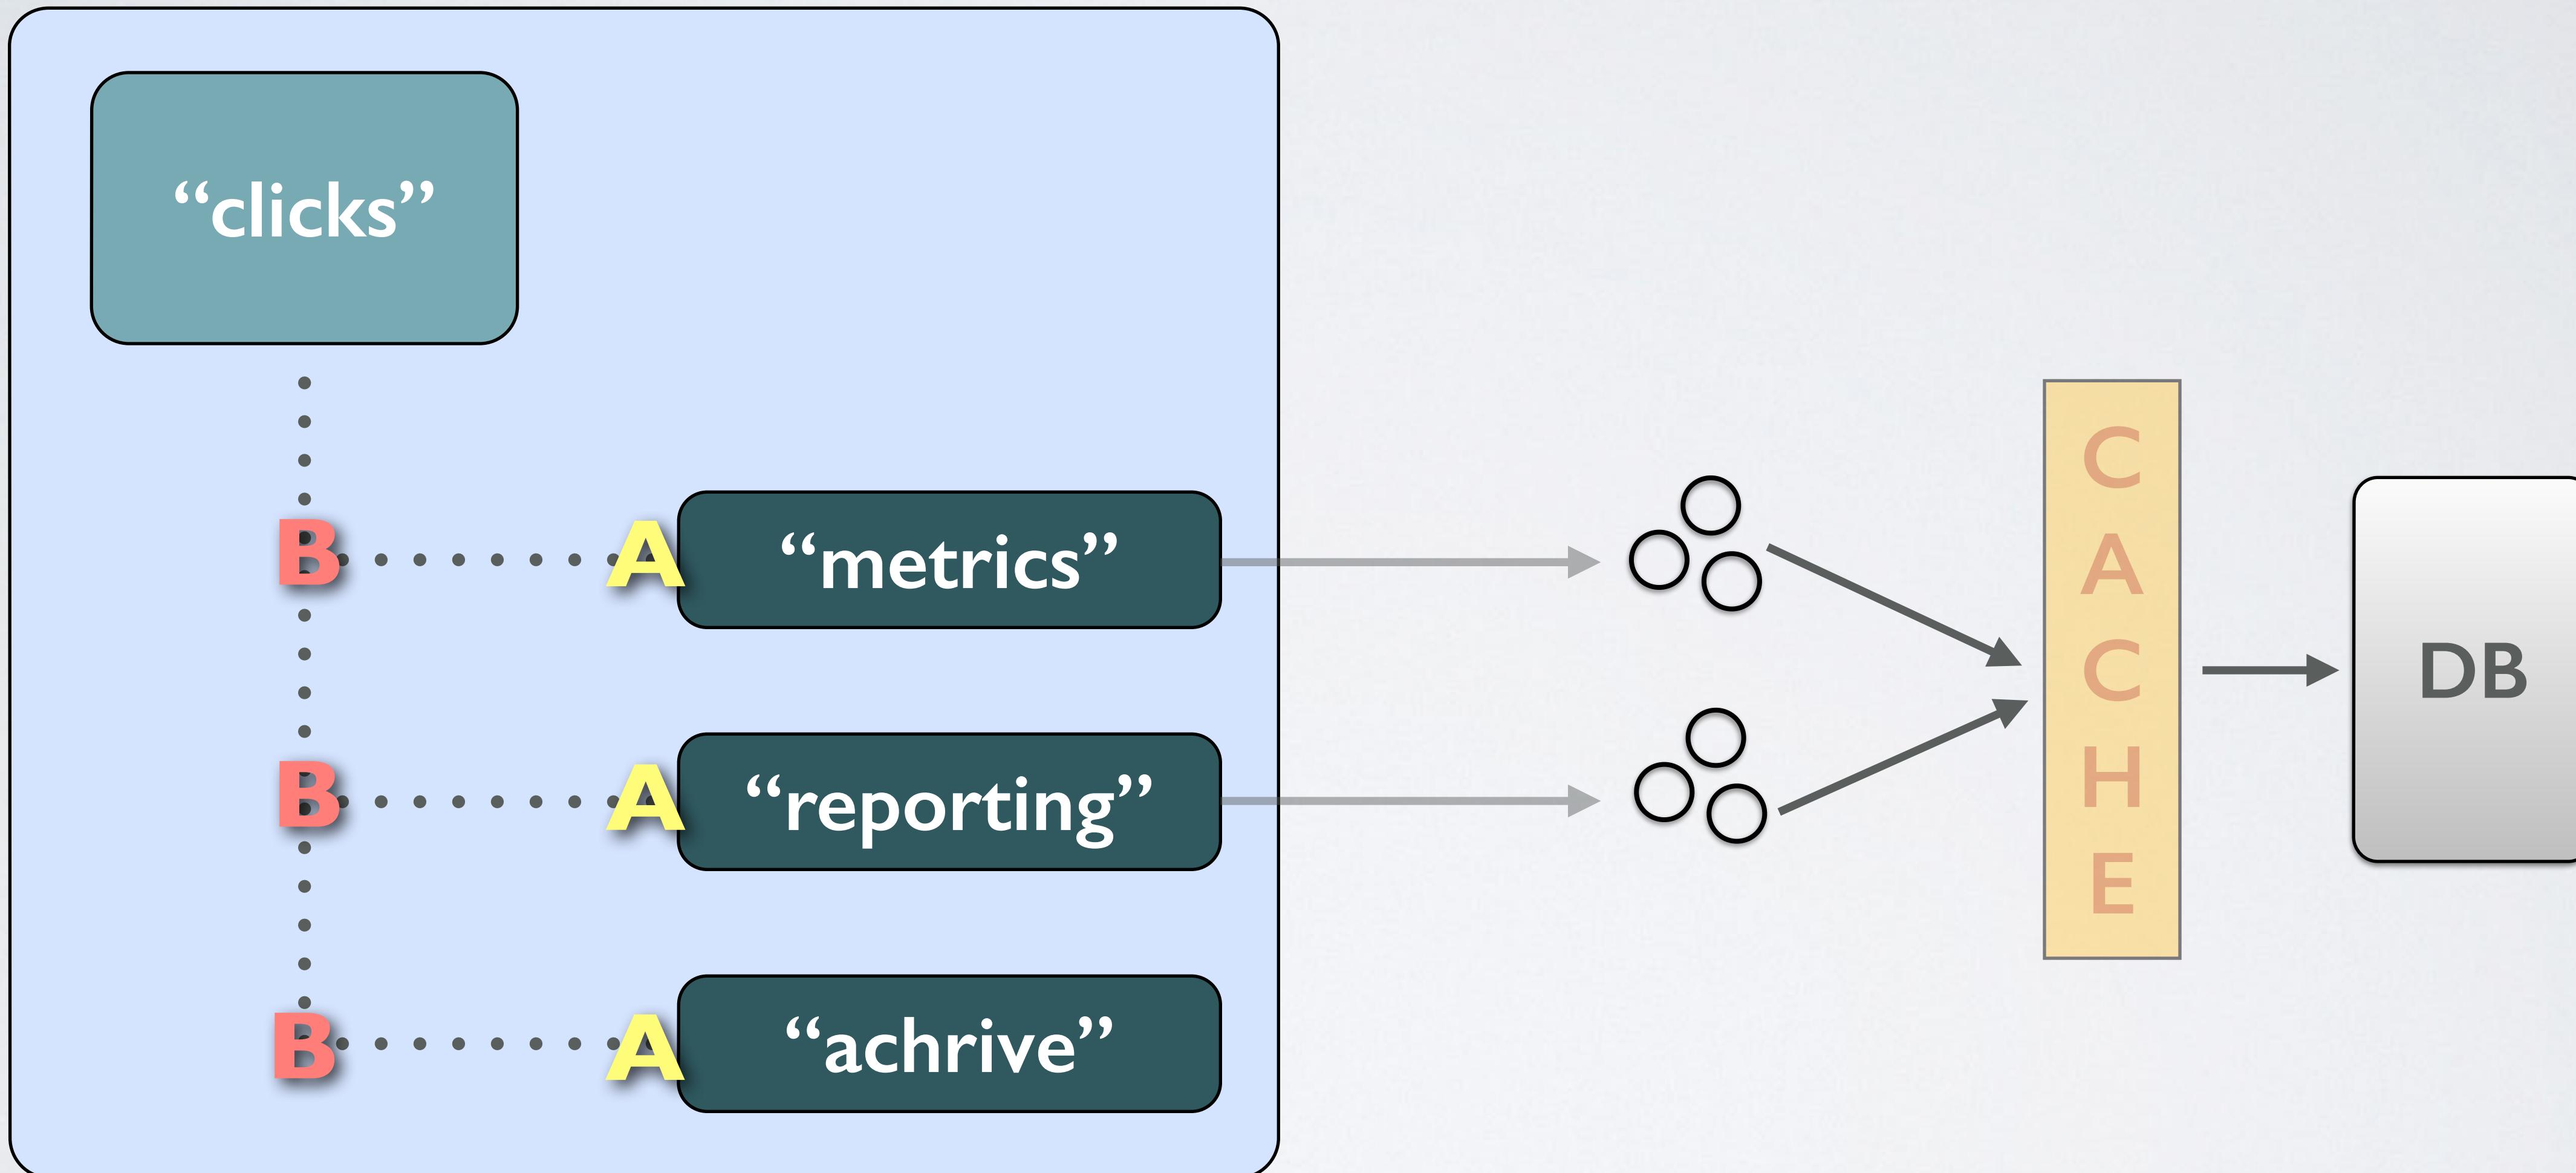
TOPIC & CHANNEL PAUSING



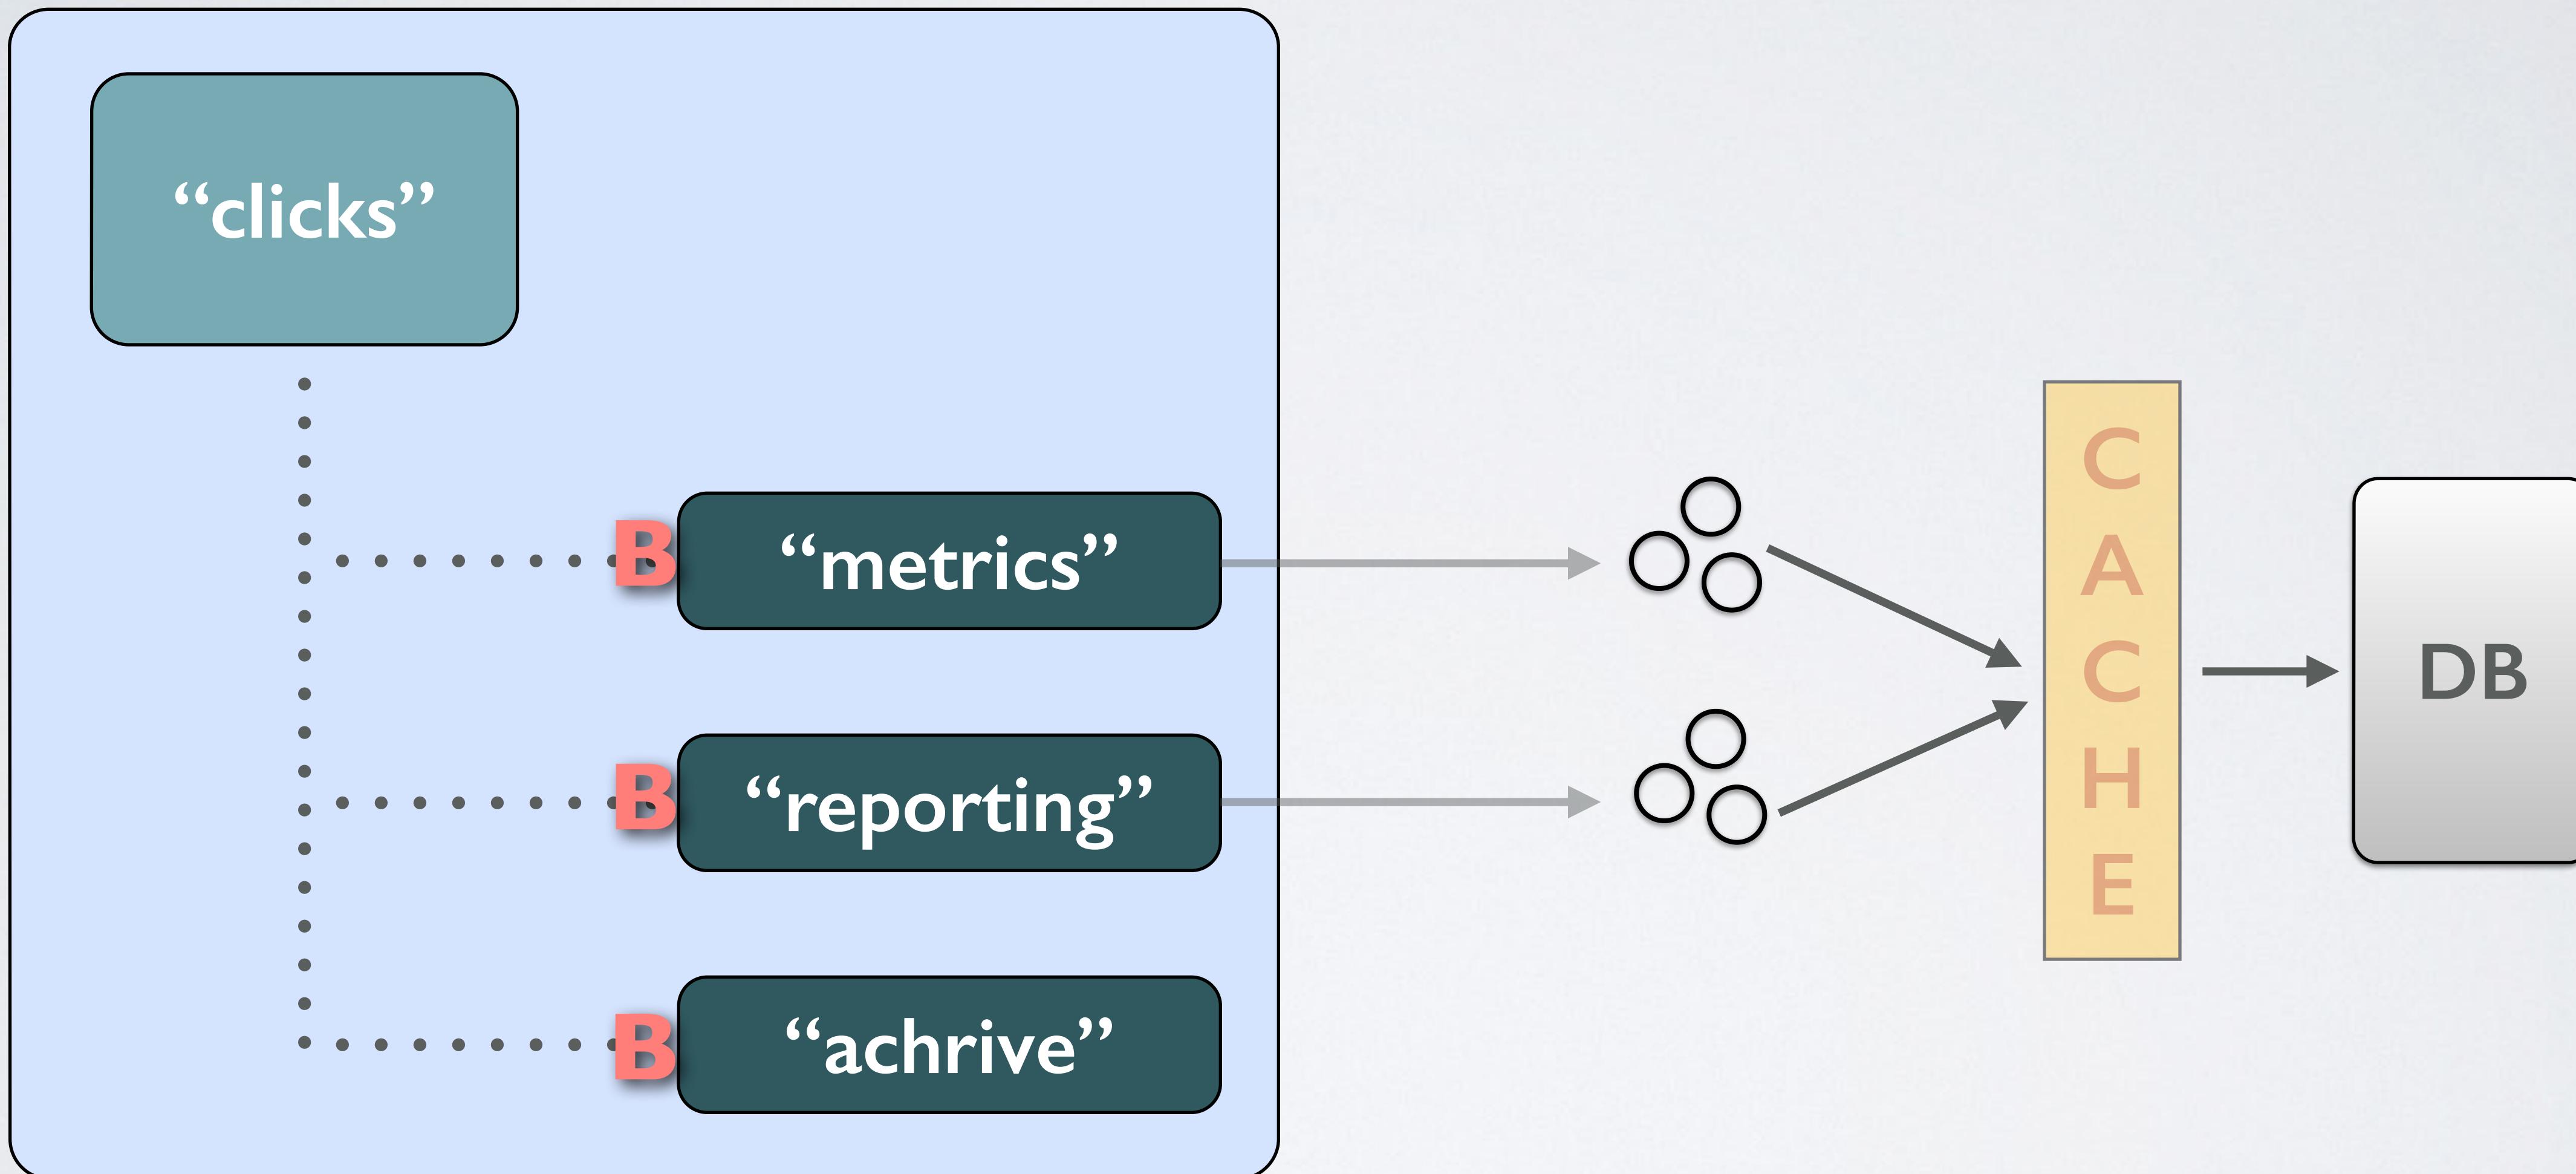
TOPIC & CHANNEL PAUSING



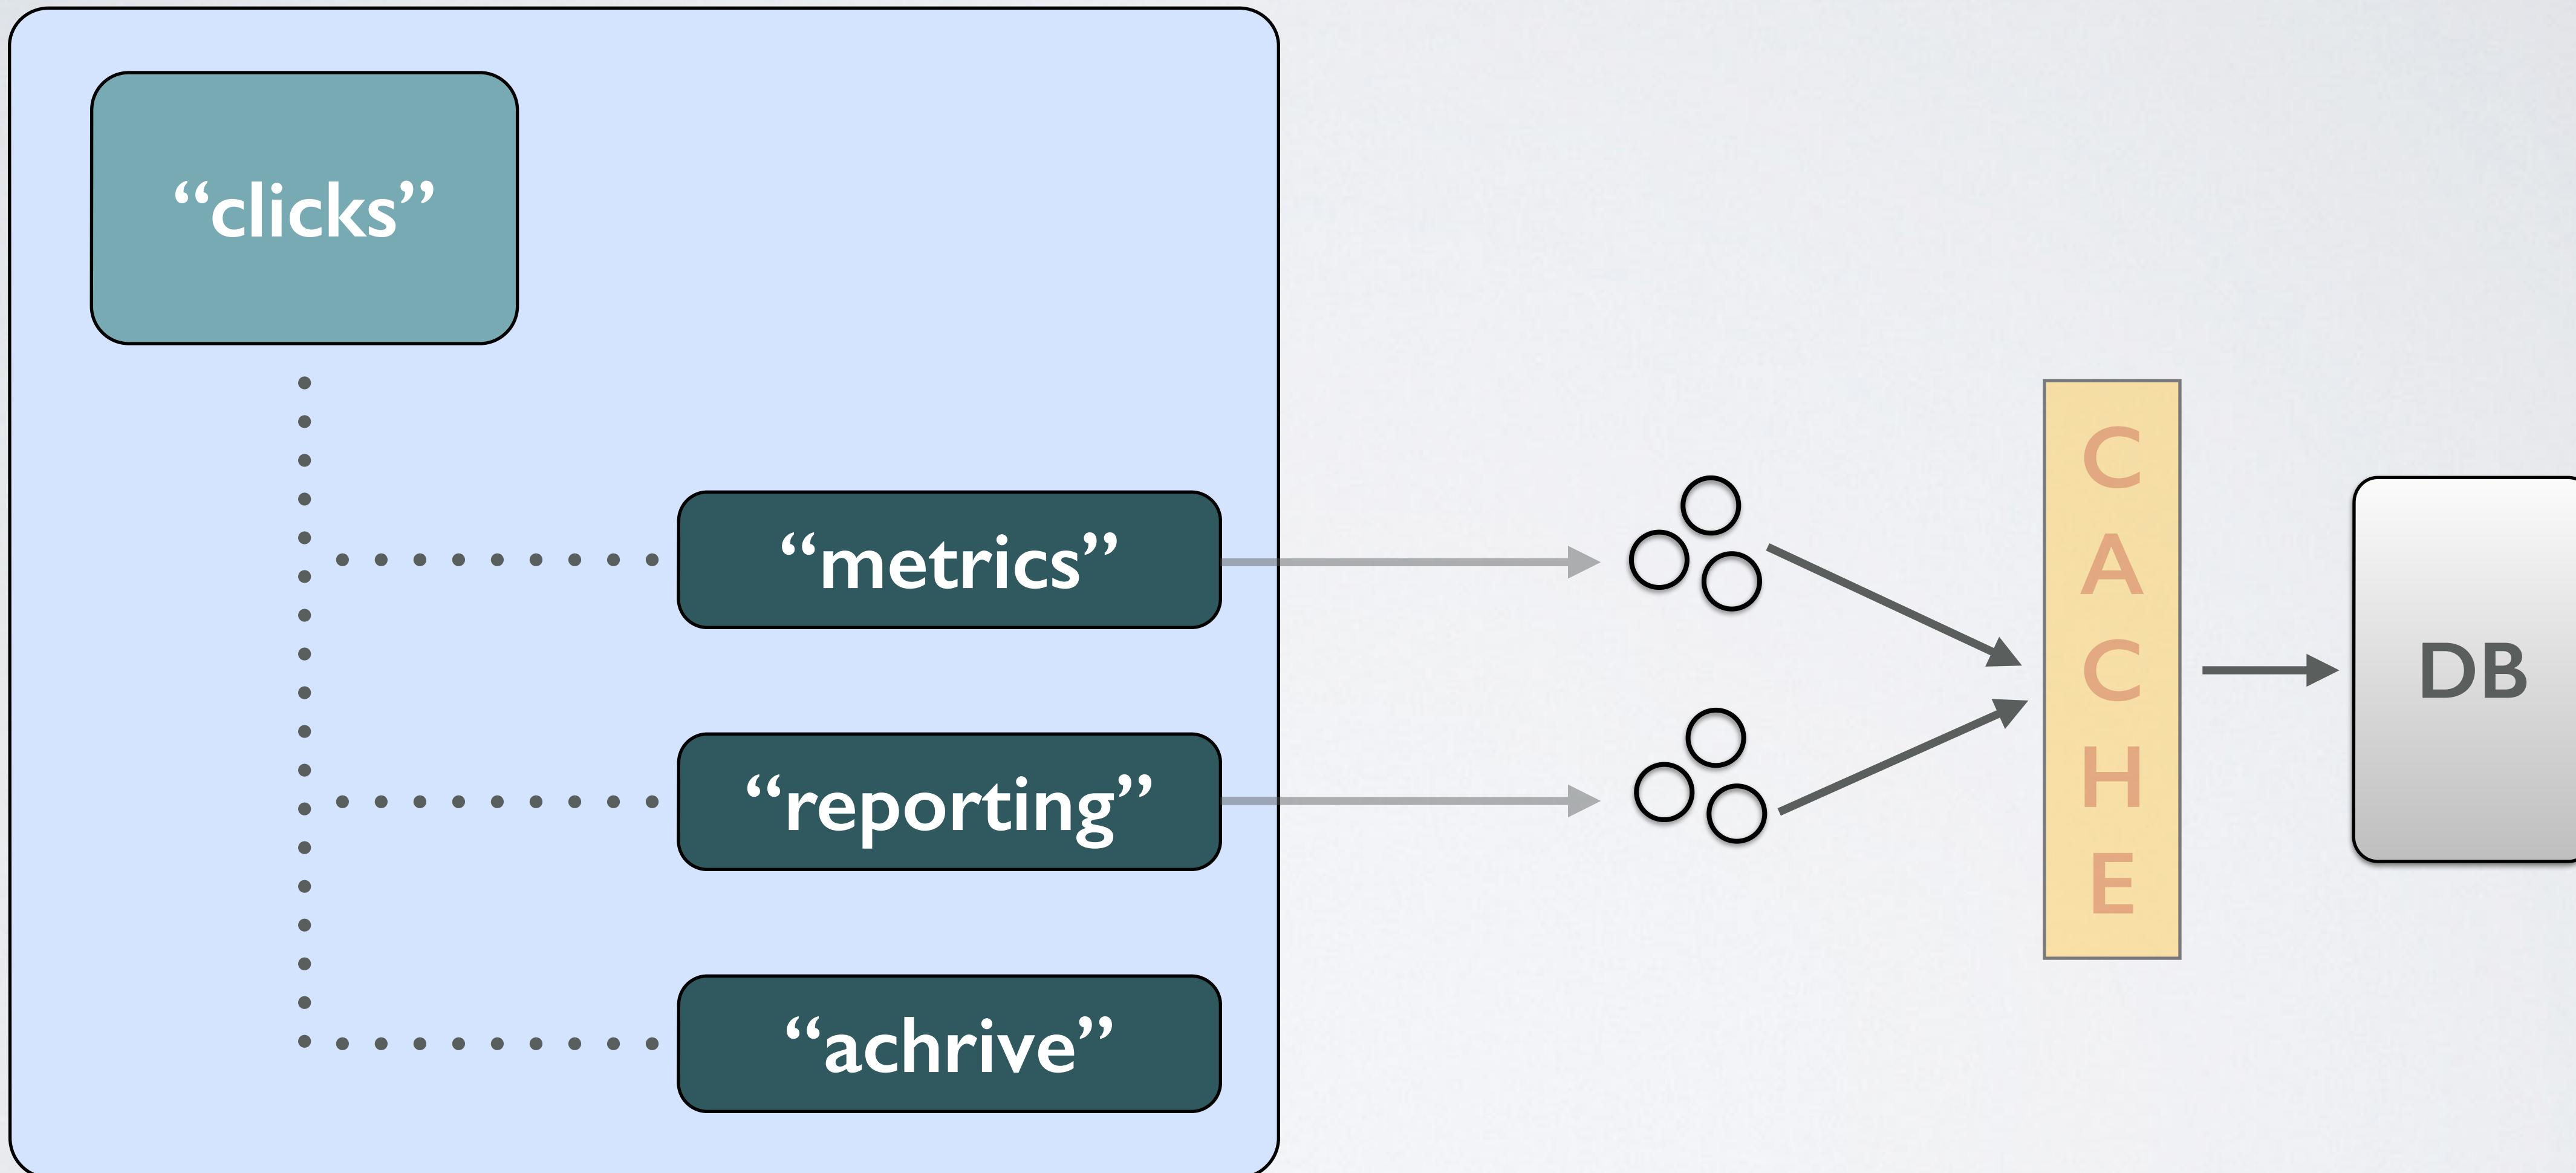
TOPIC & CHANNEL PAUSING



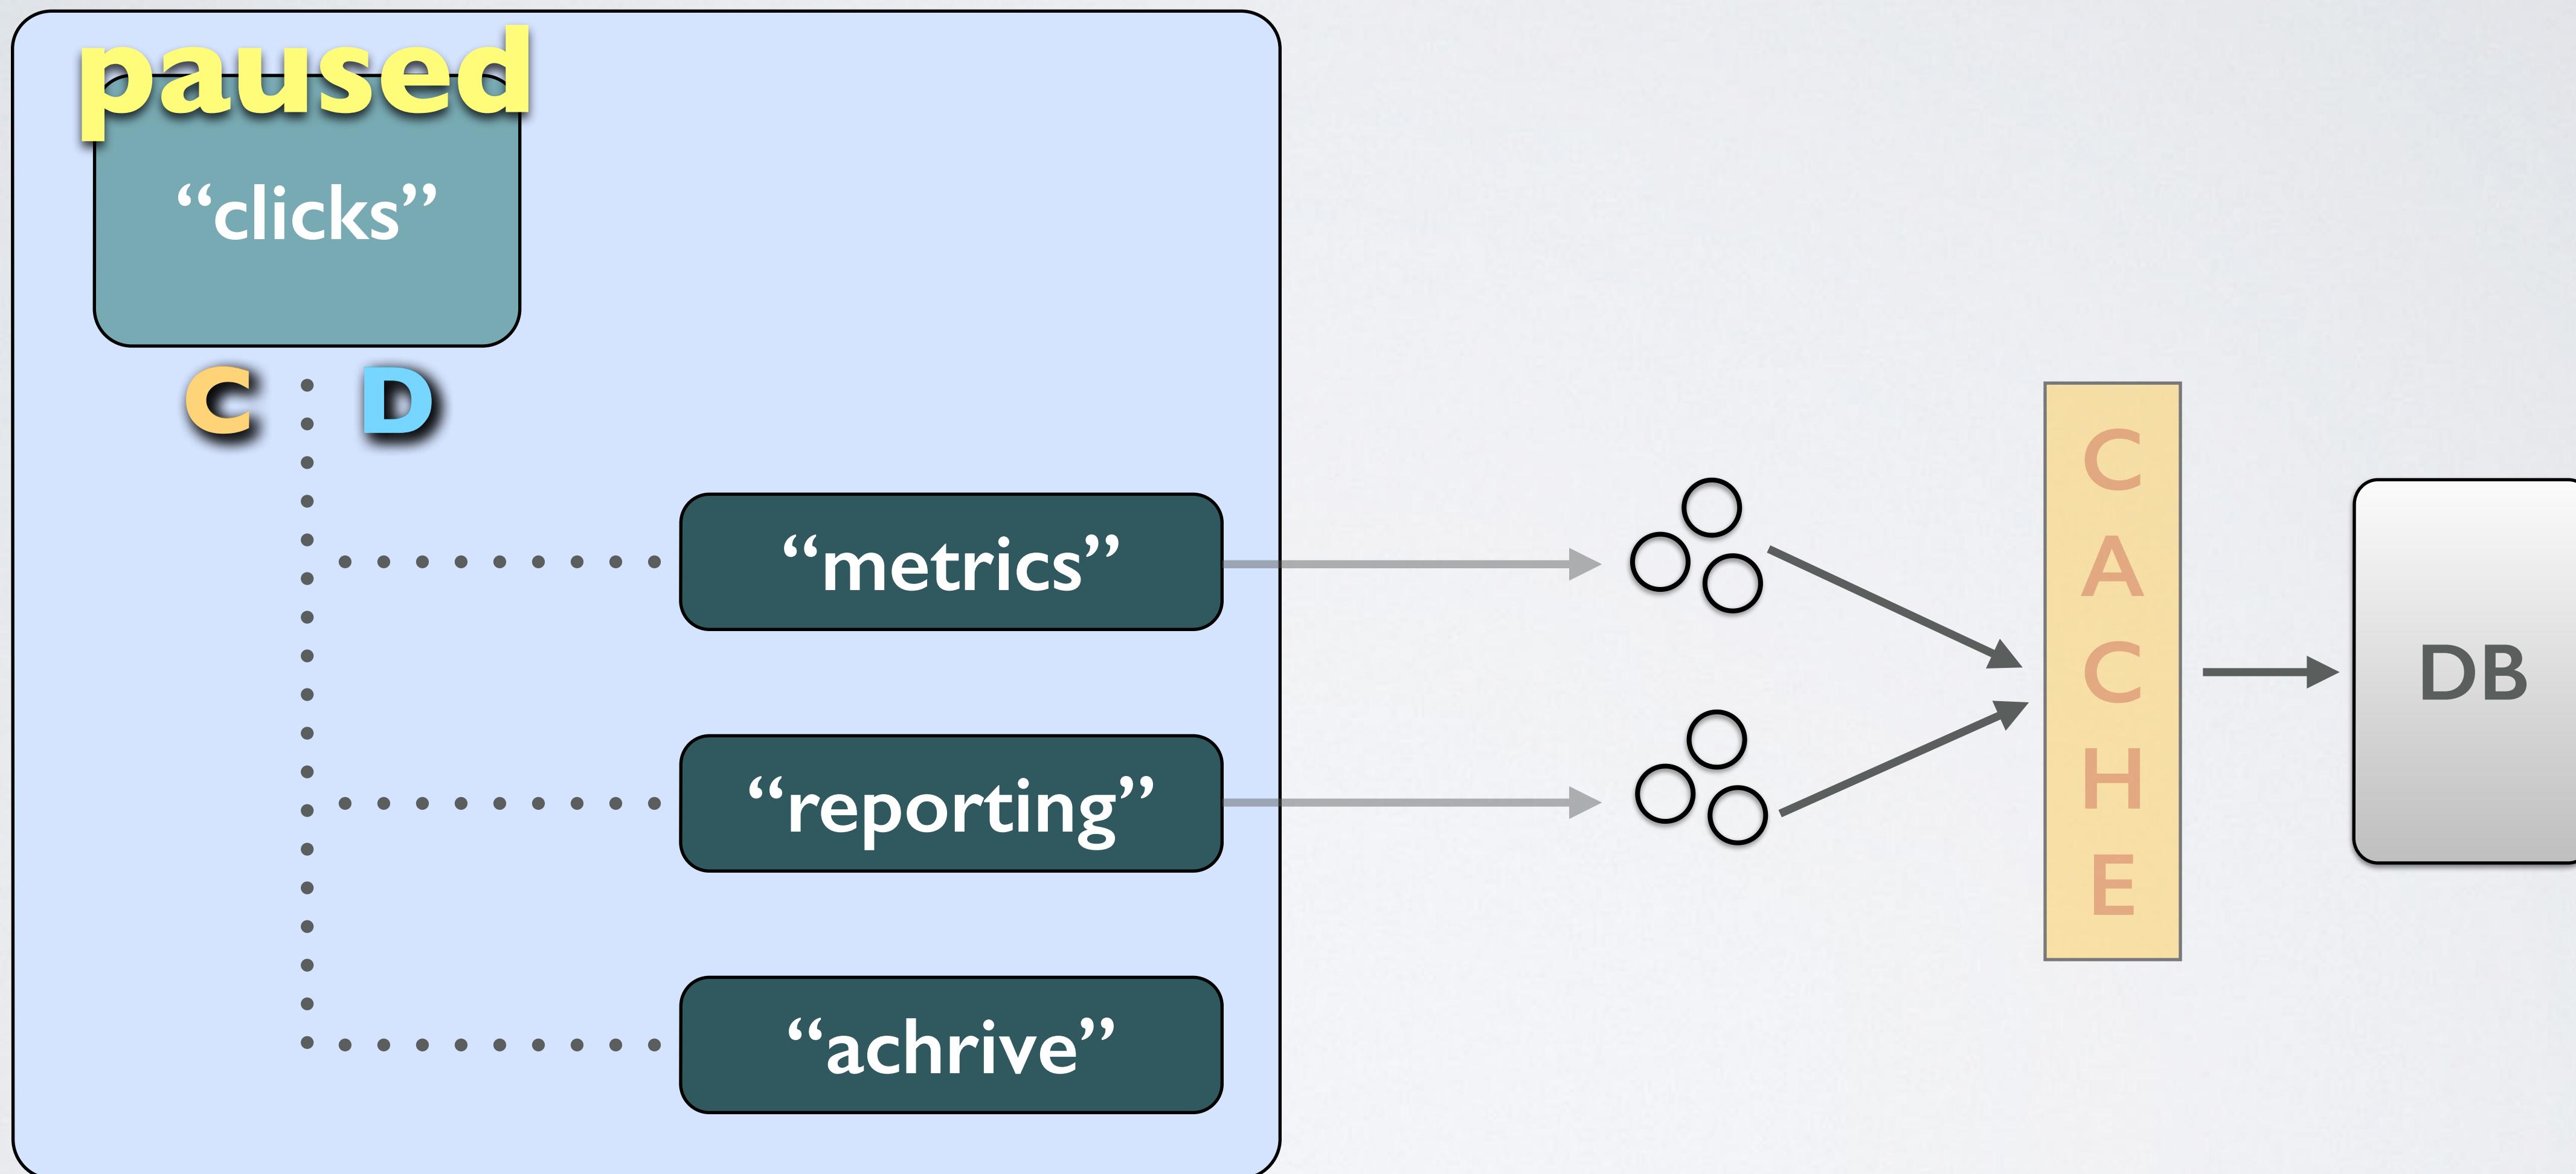
TOPIC & CHANNEL PAUSING



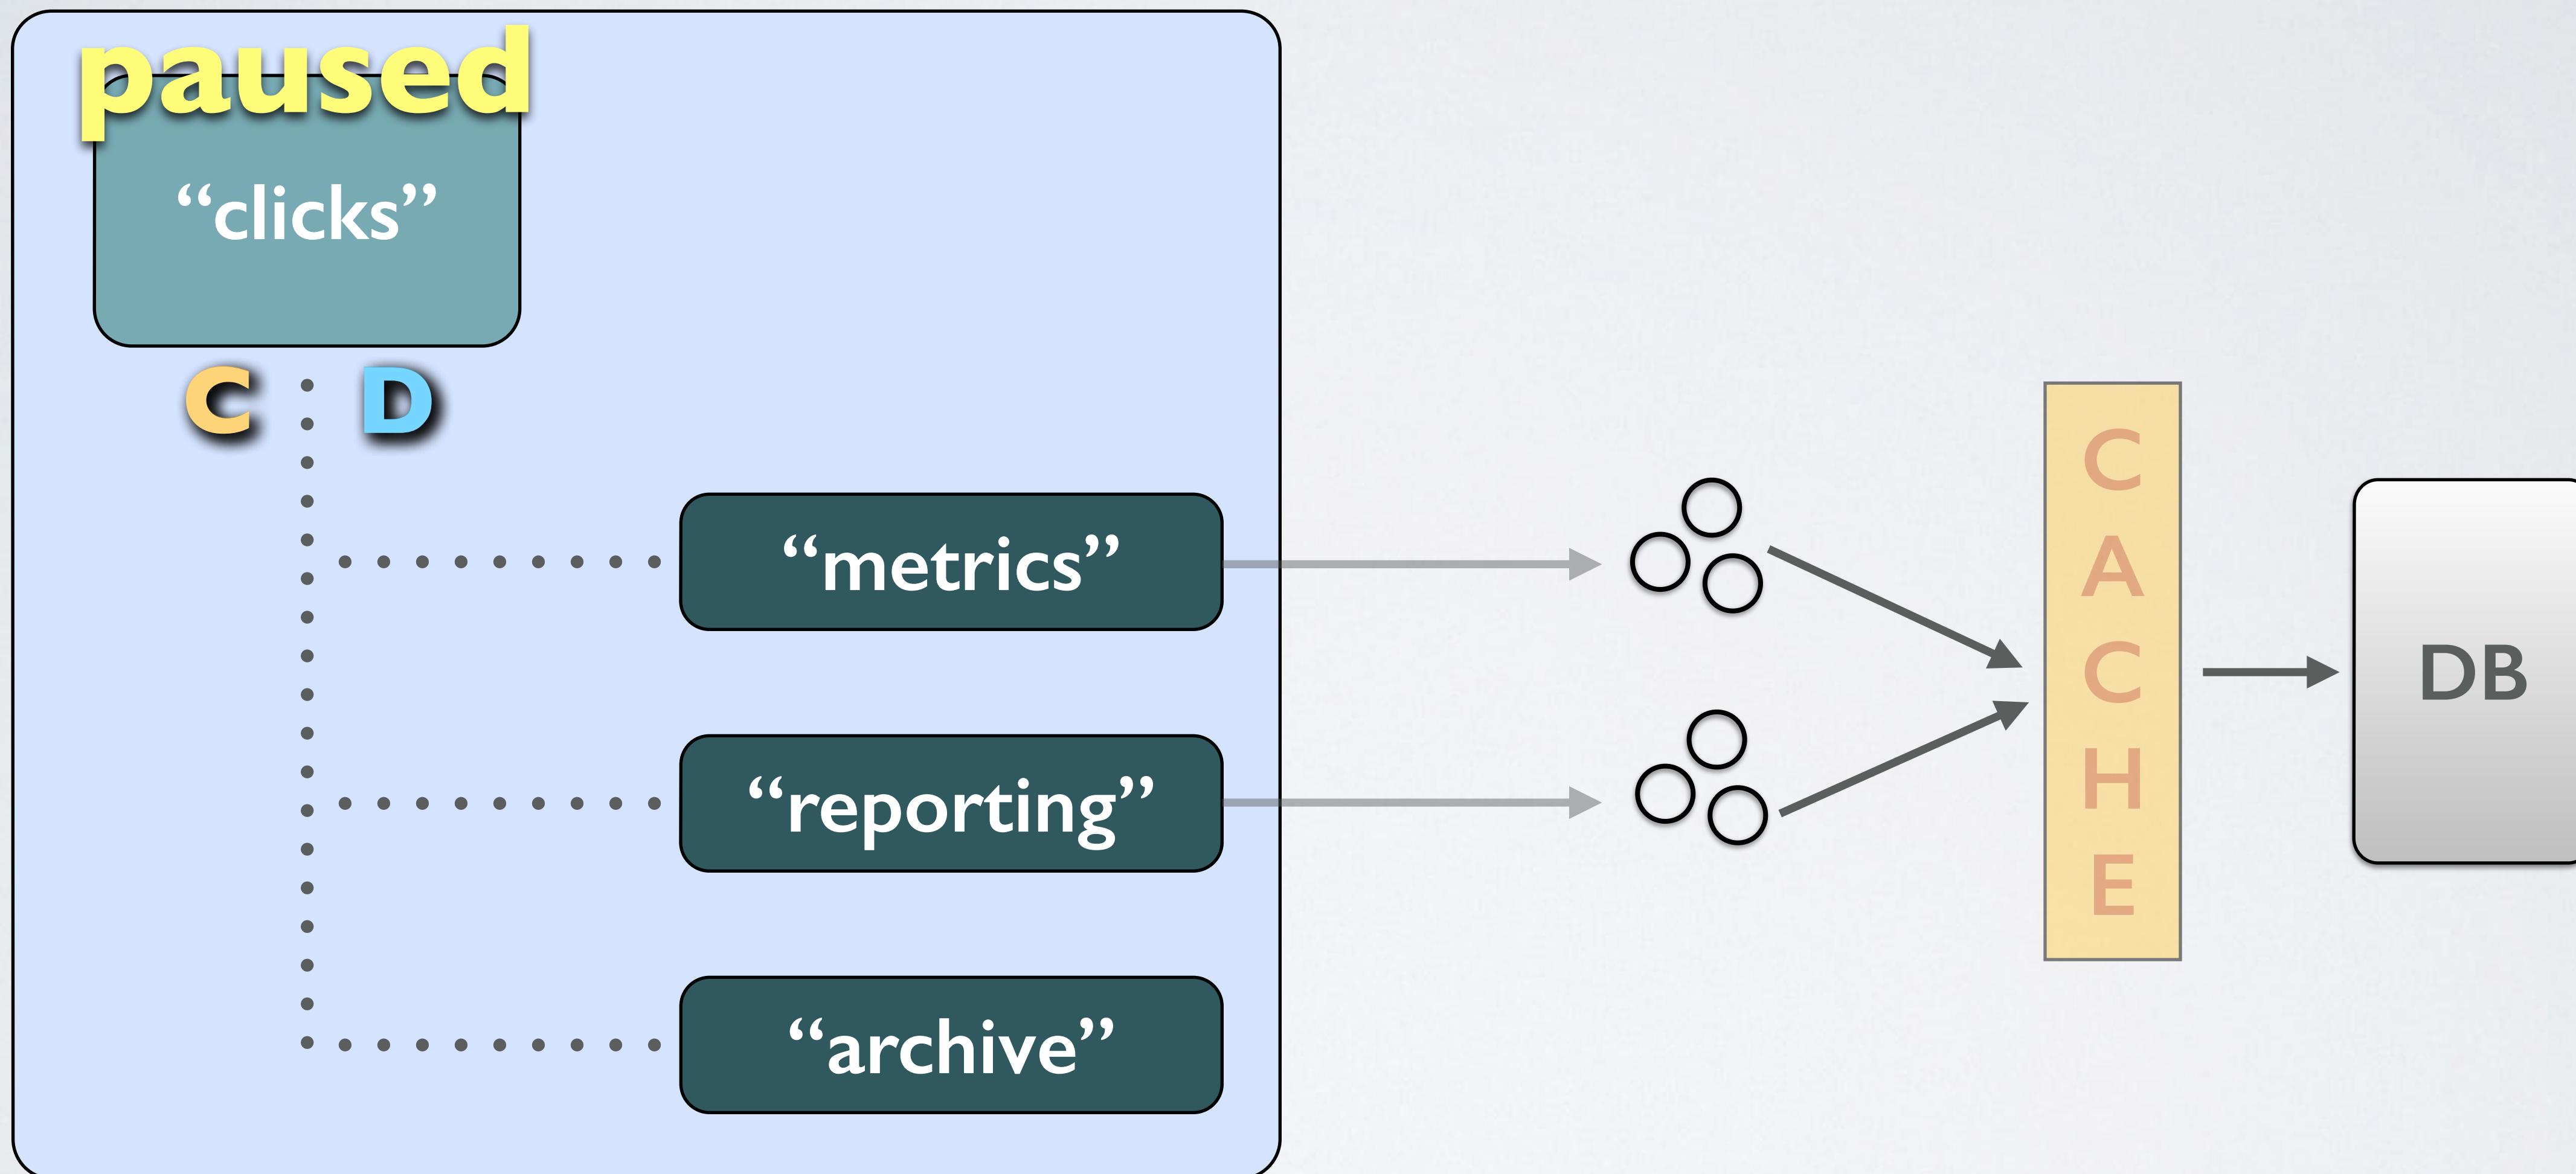
TOPIC & CHANNEL PAUSING



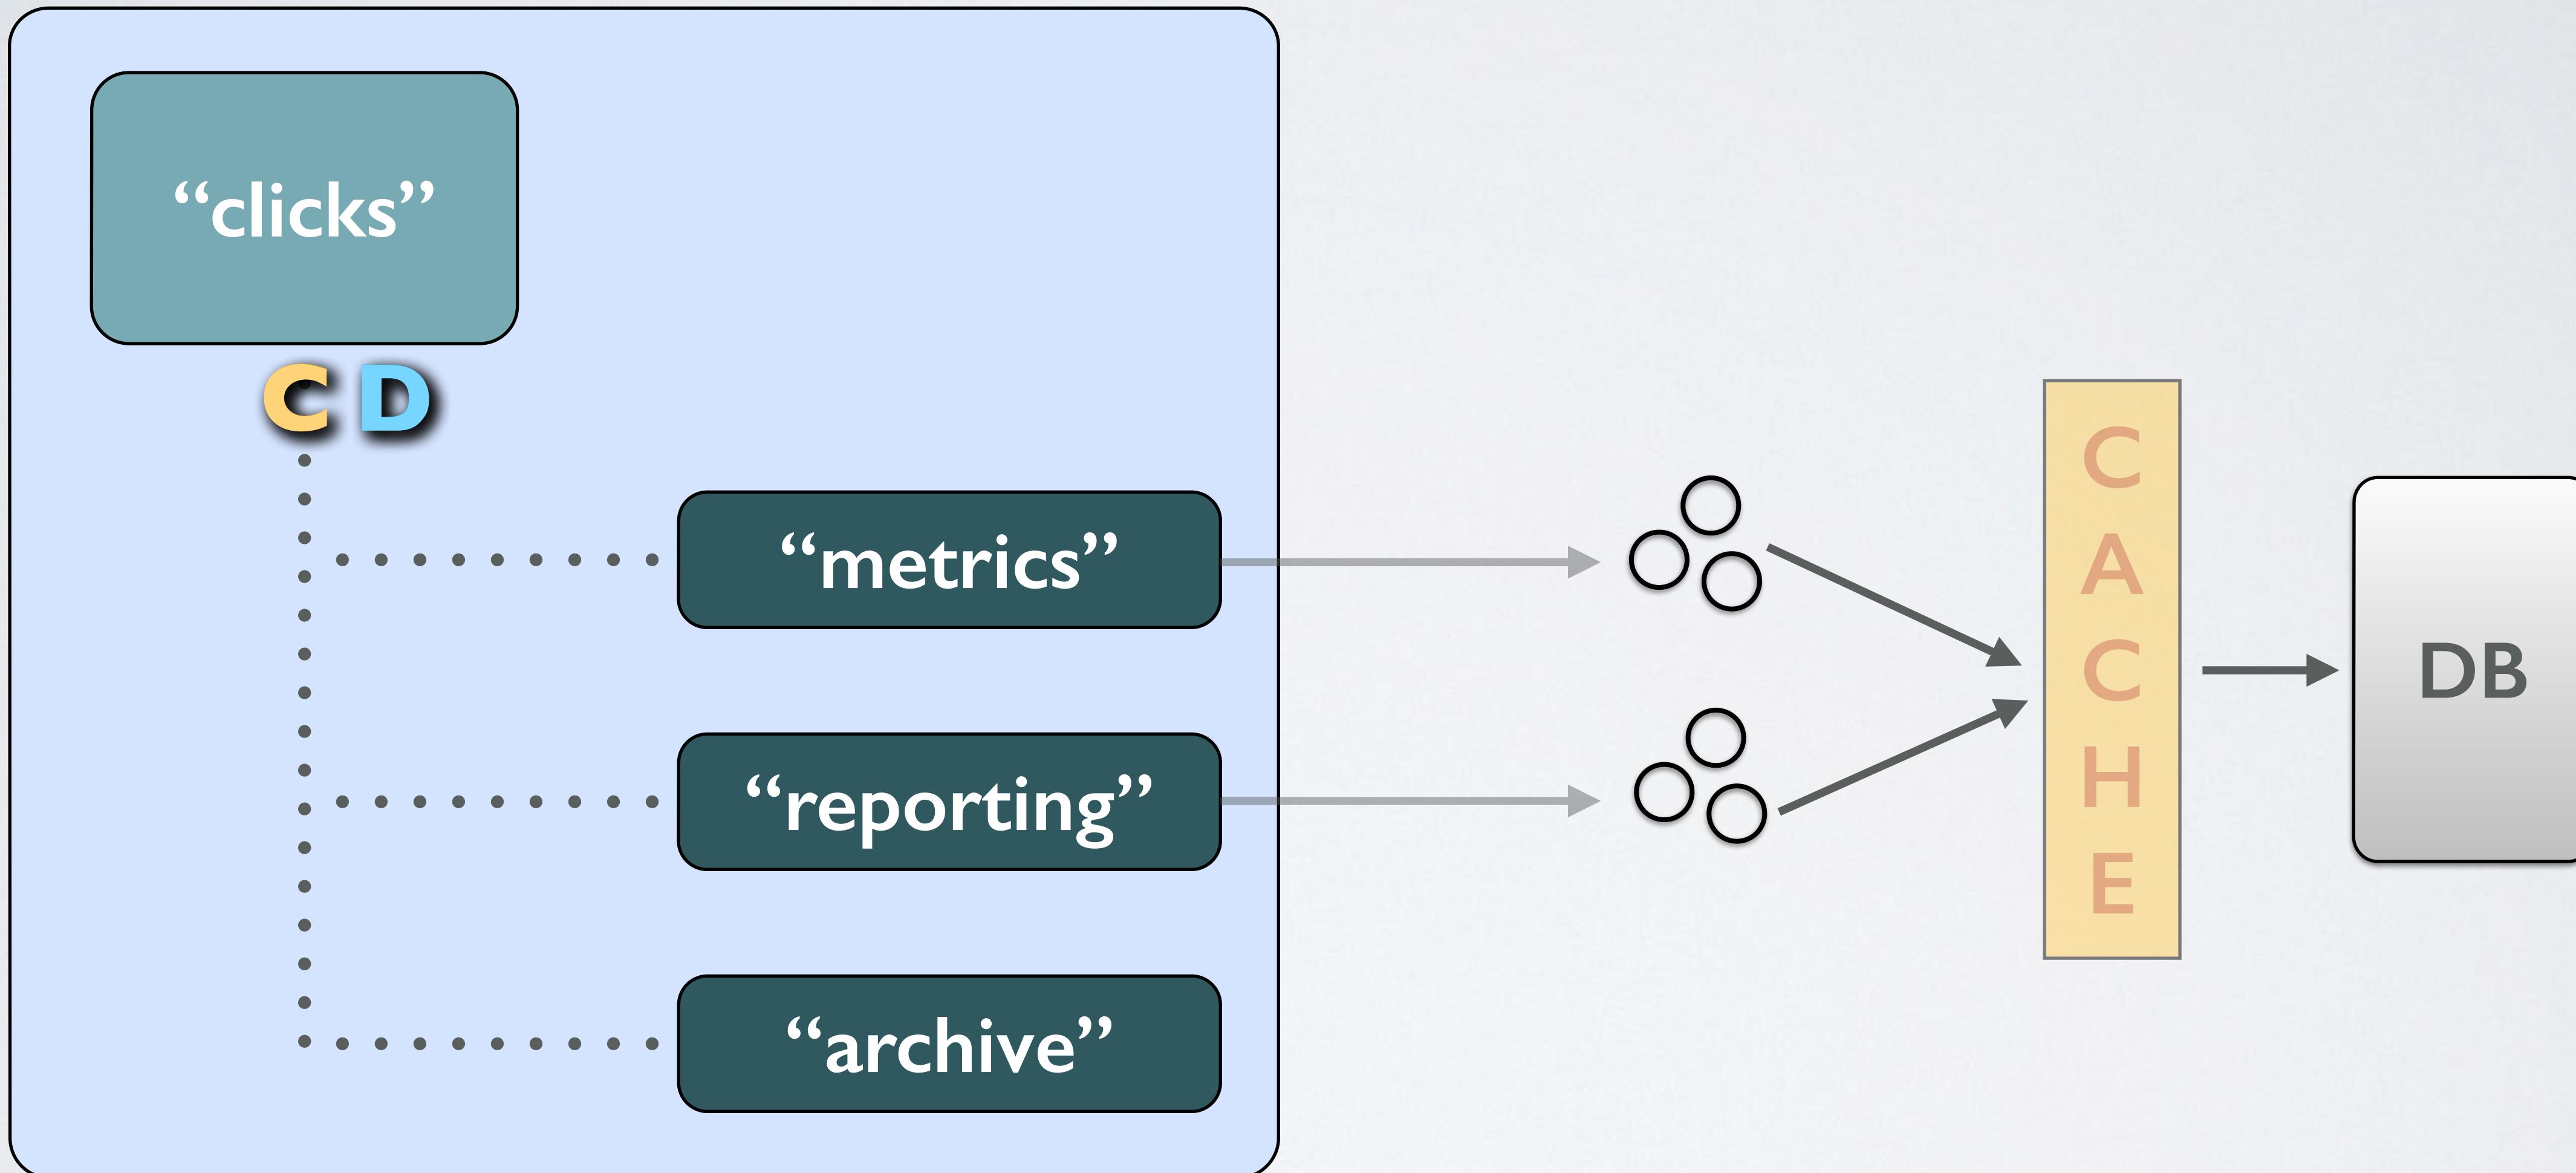
TOPIC & CHANNEL PAUSING



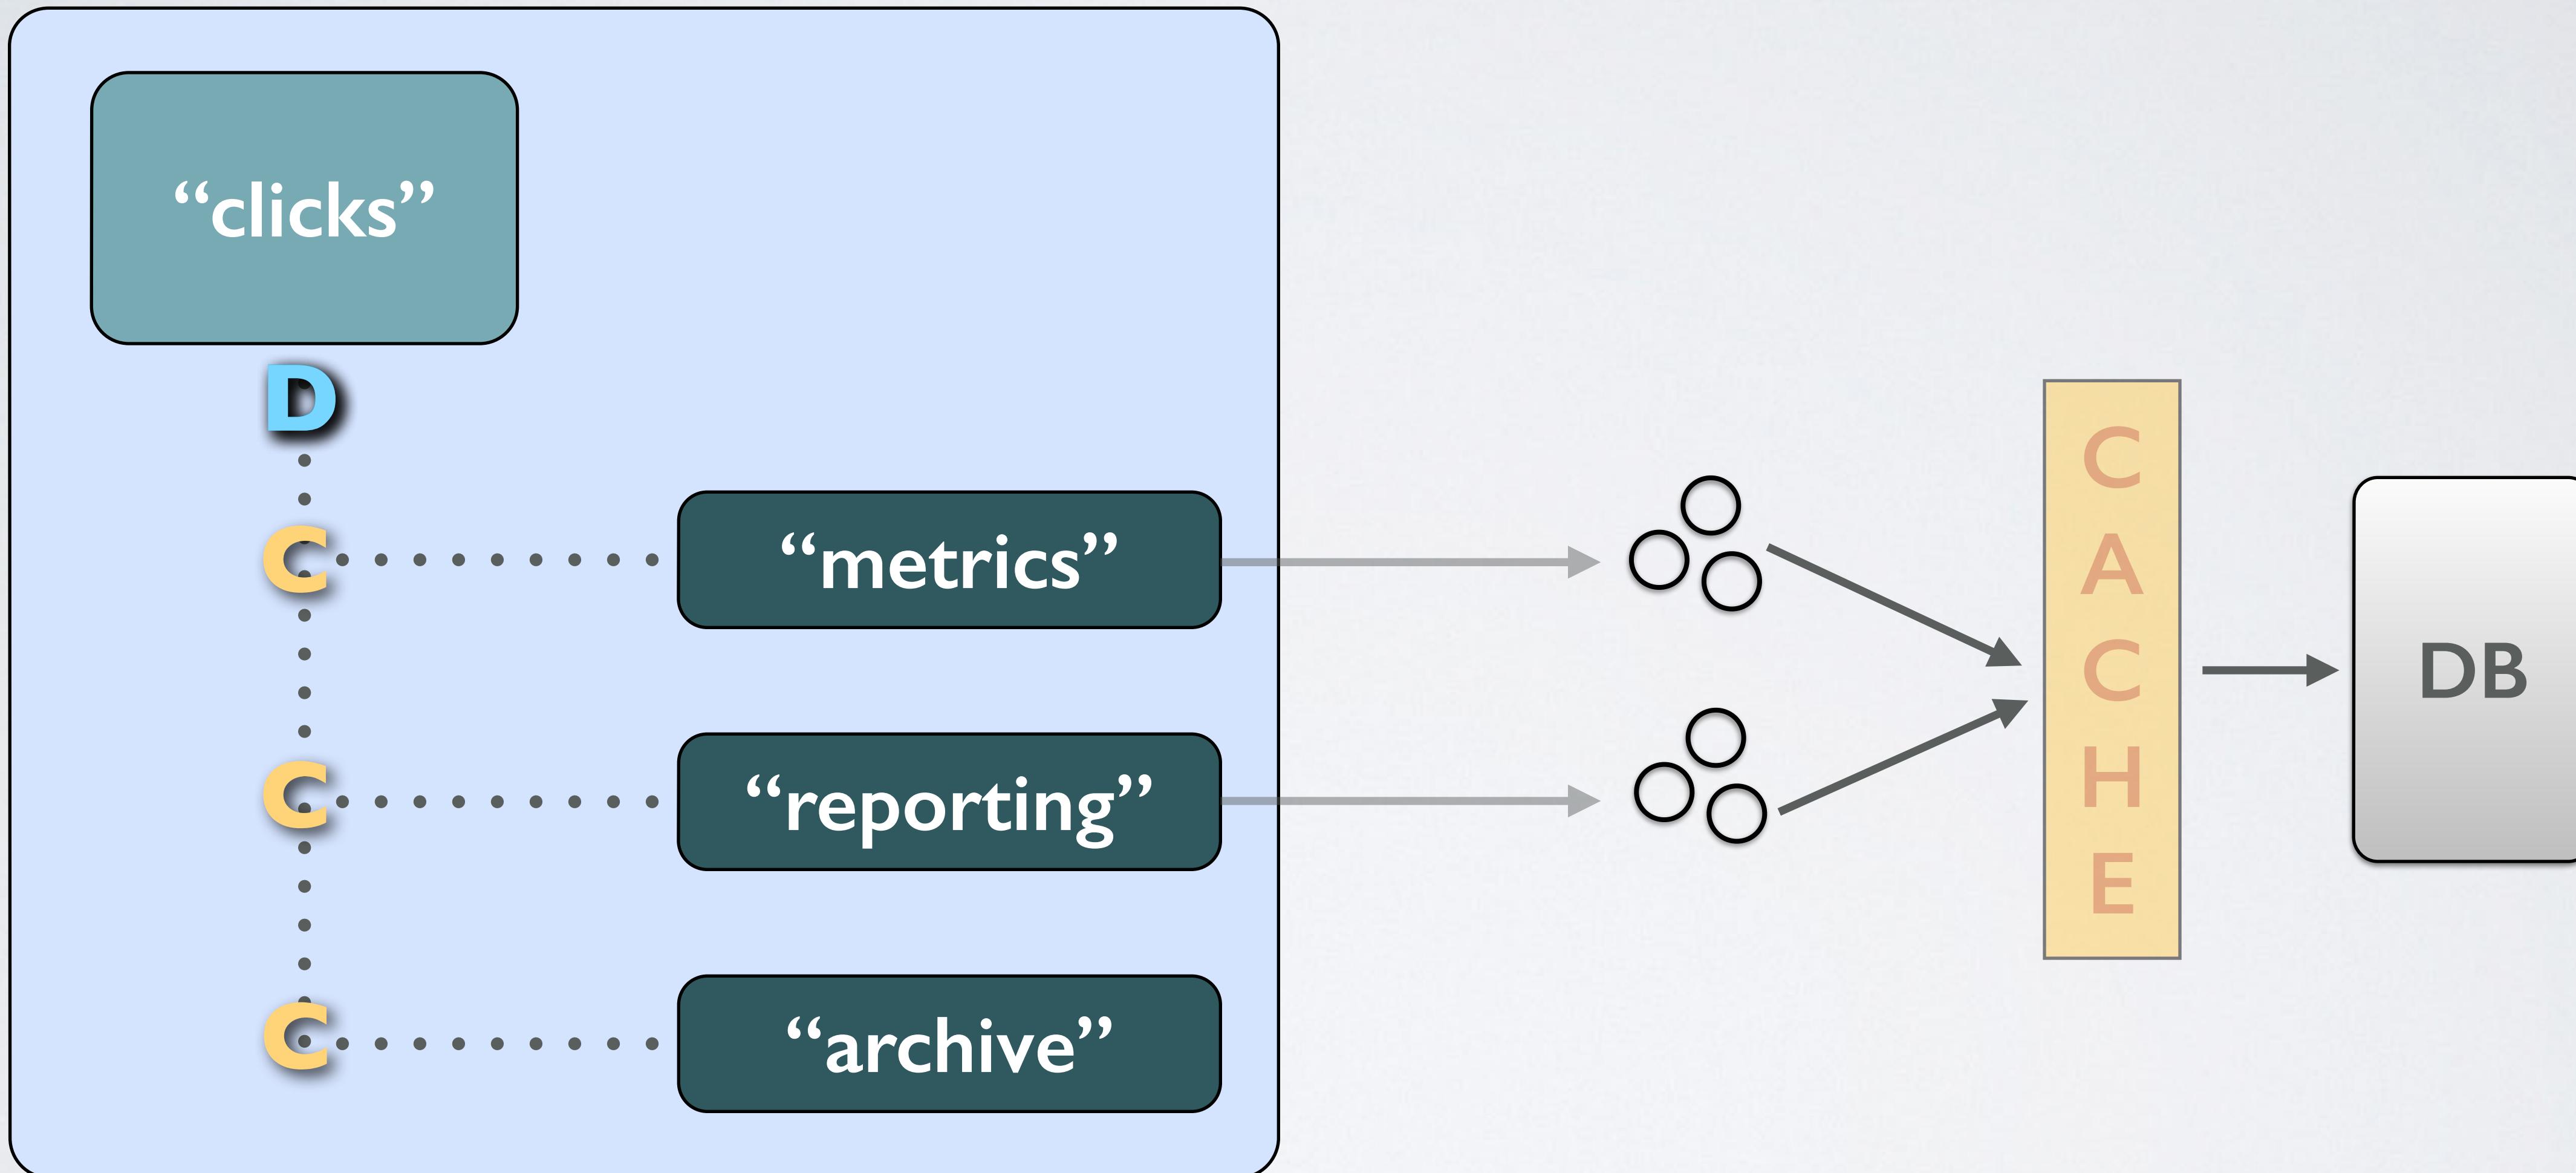
TOPIC & CHANNEL PAUSING



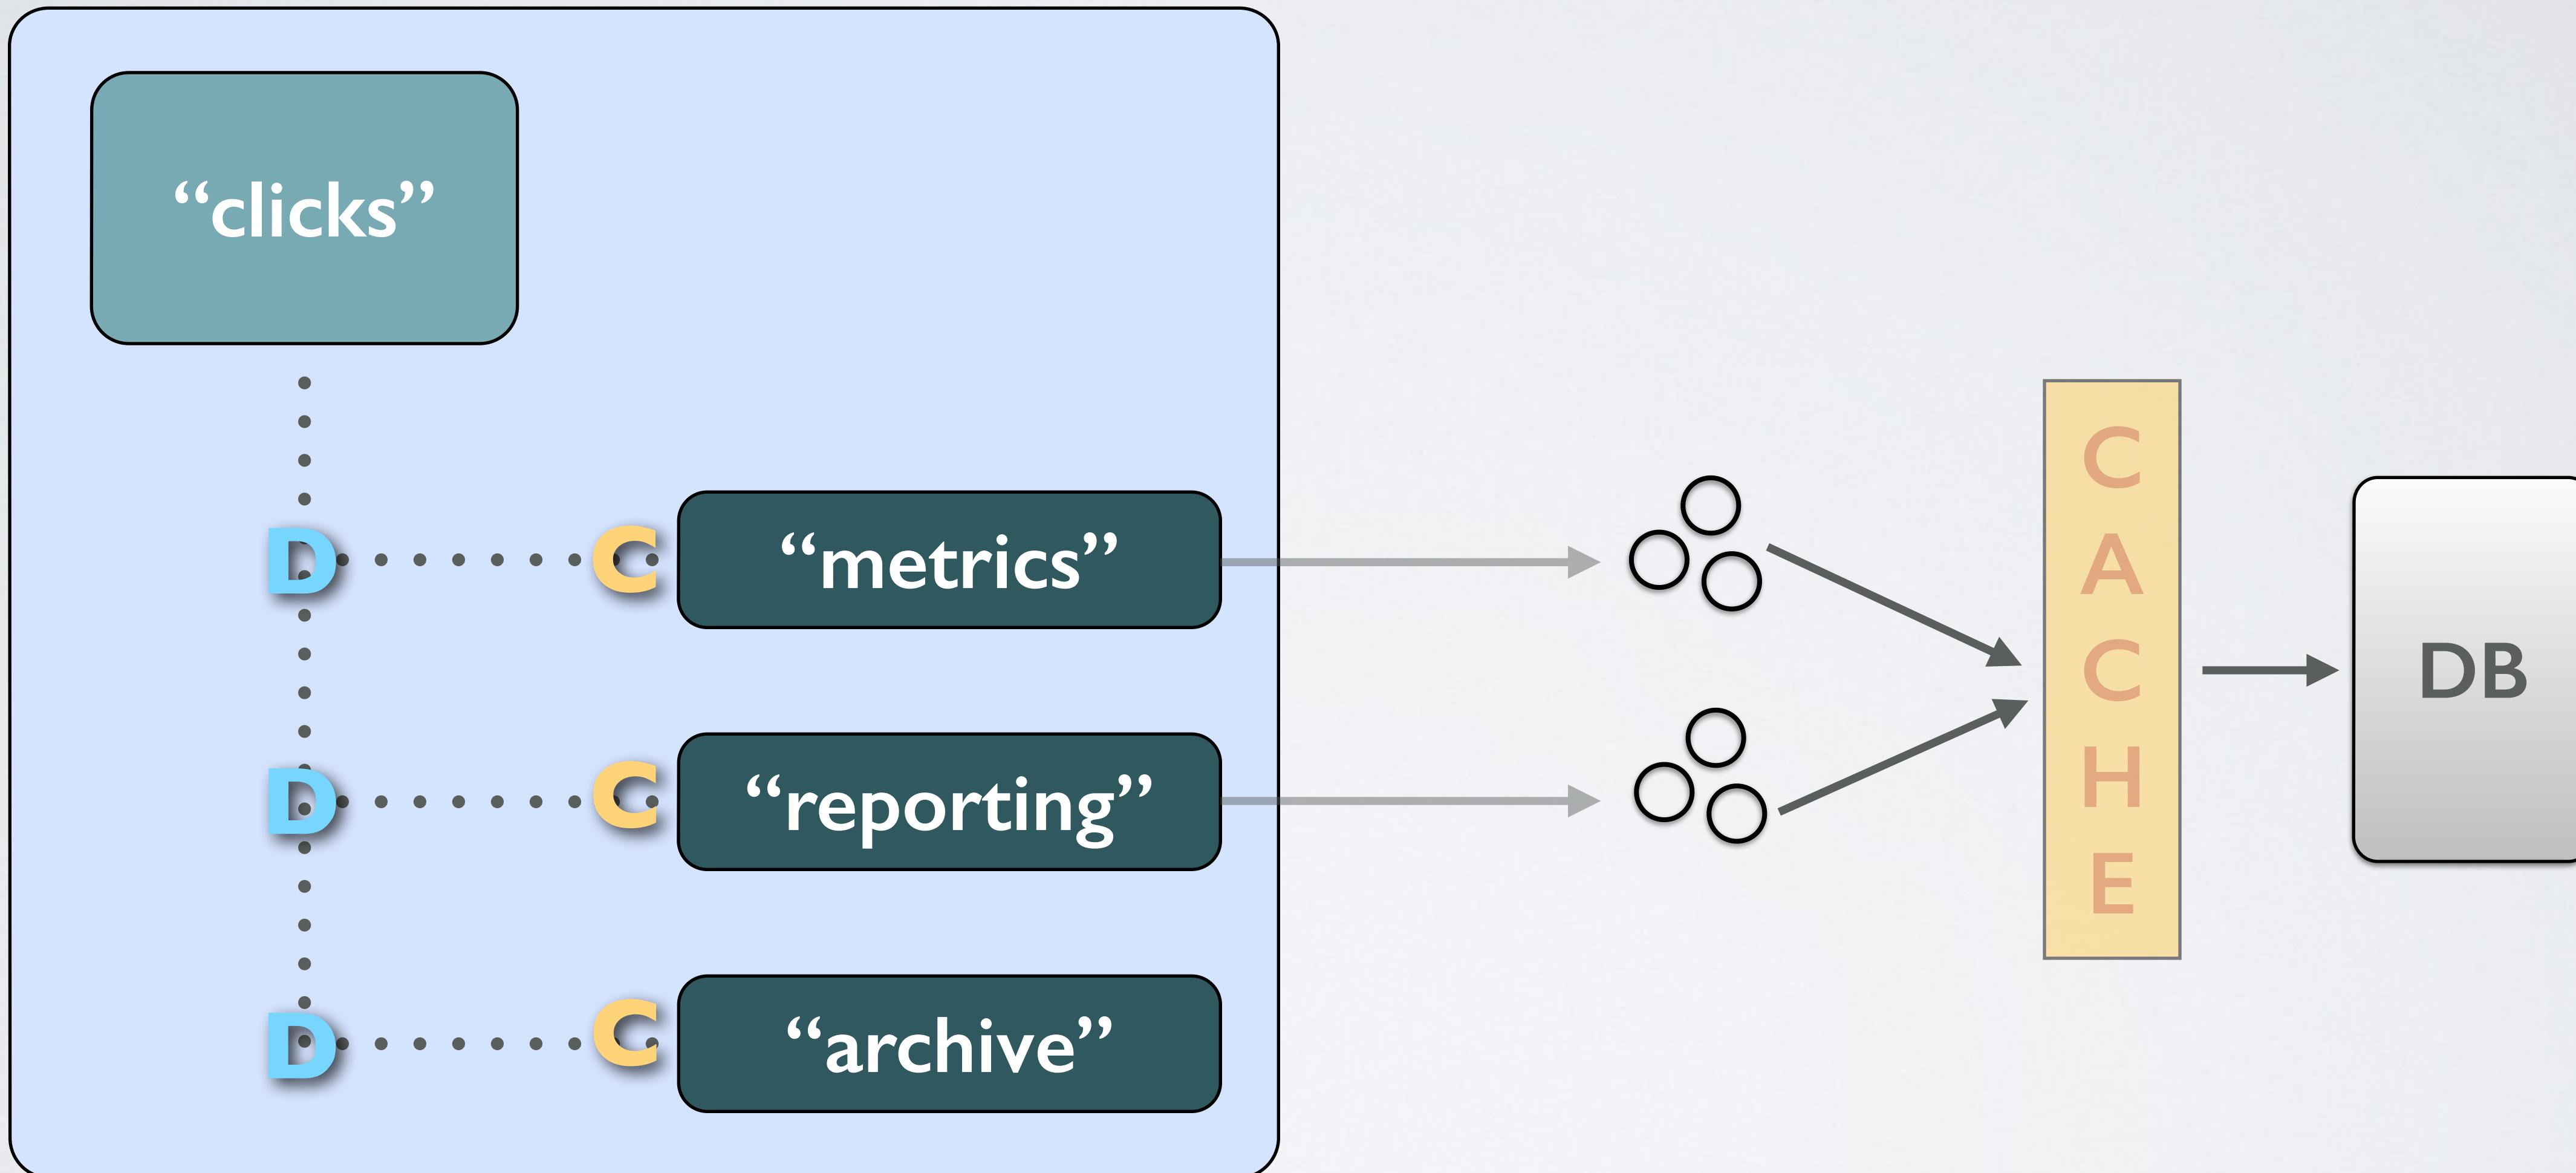
TOPIC & CHANNEL PAUSING



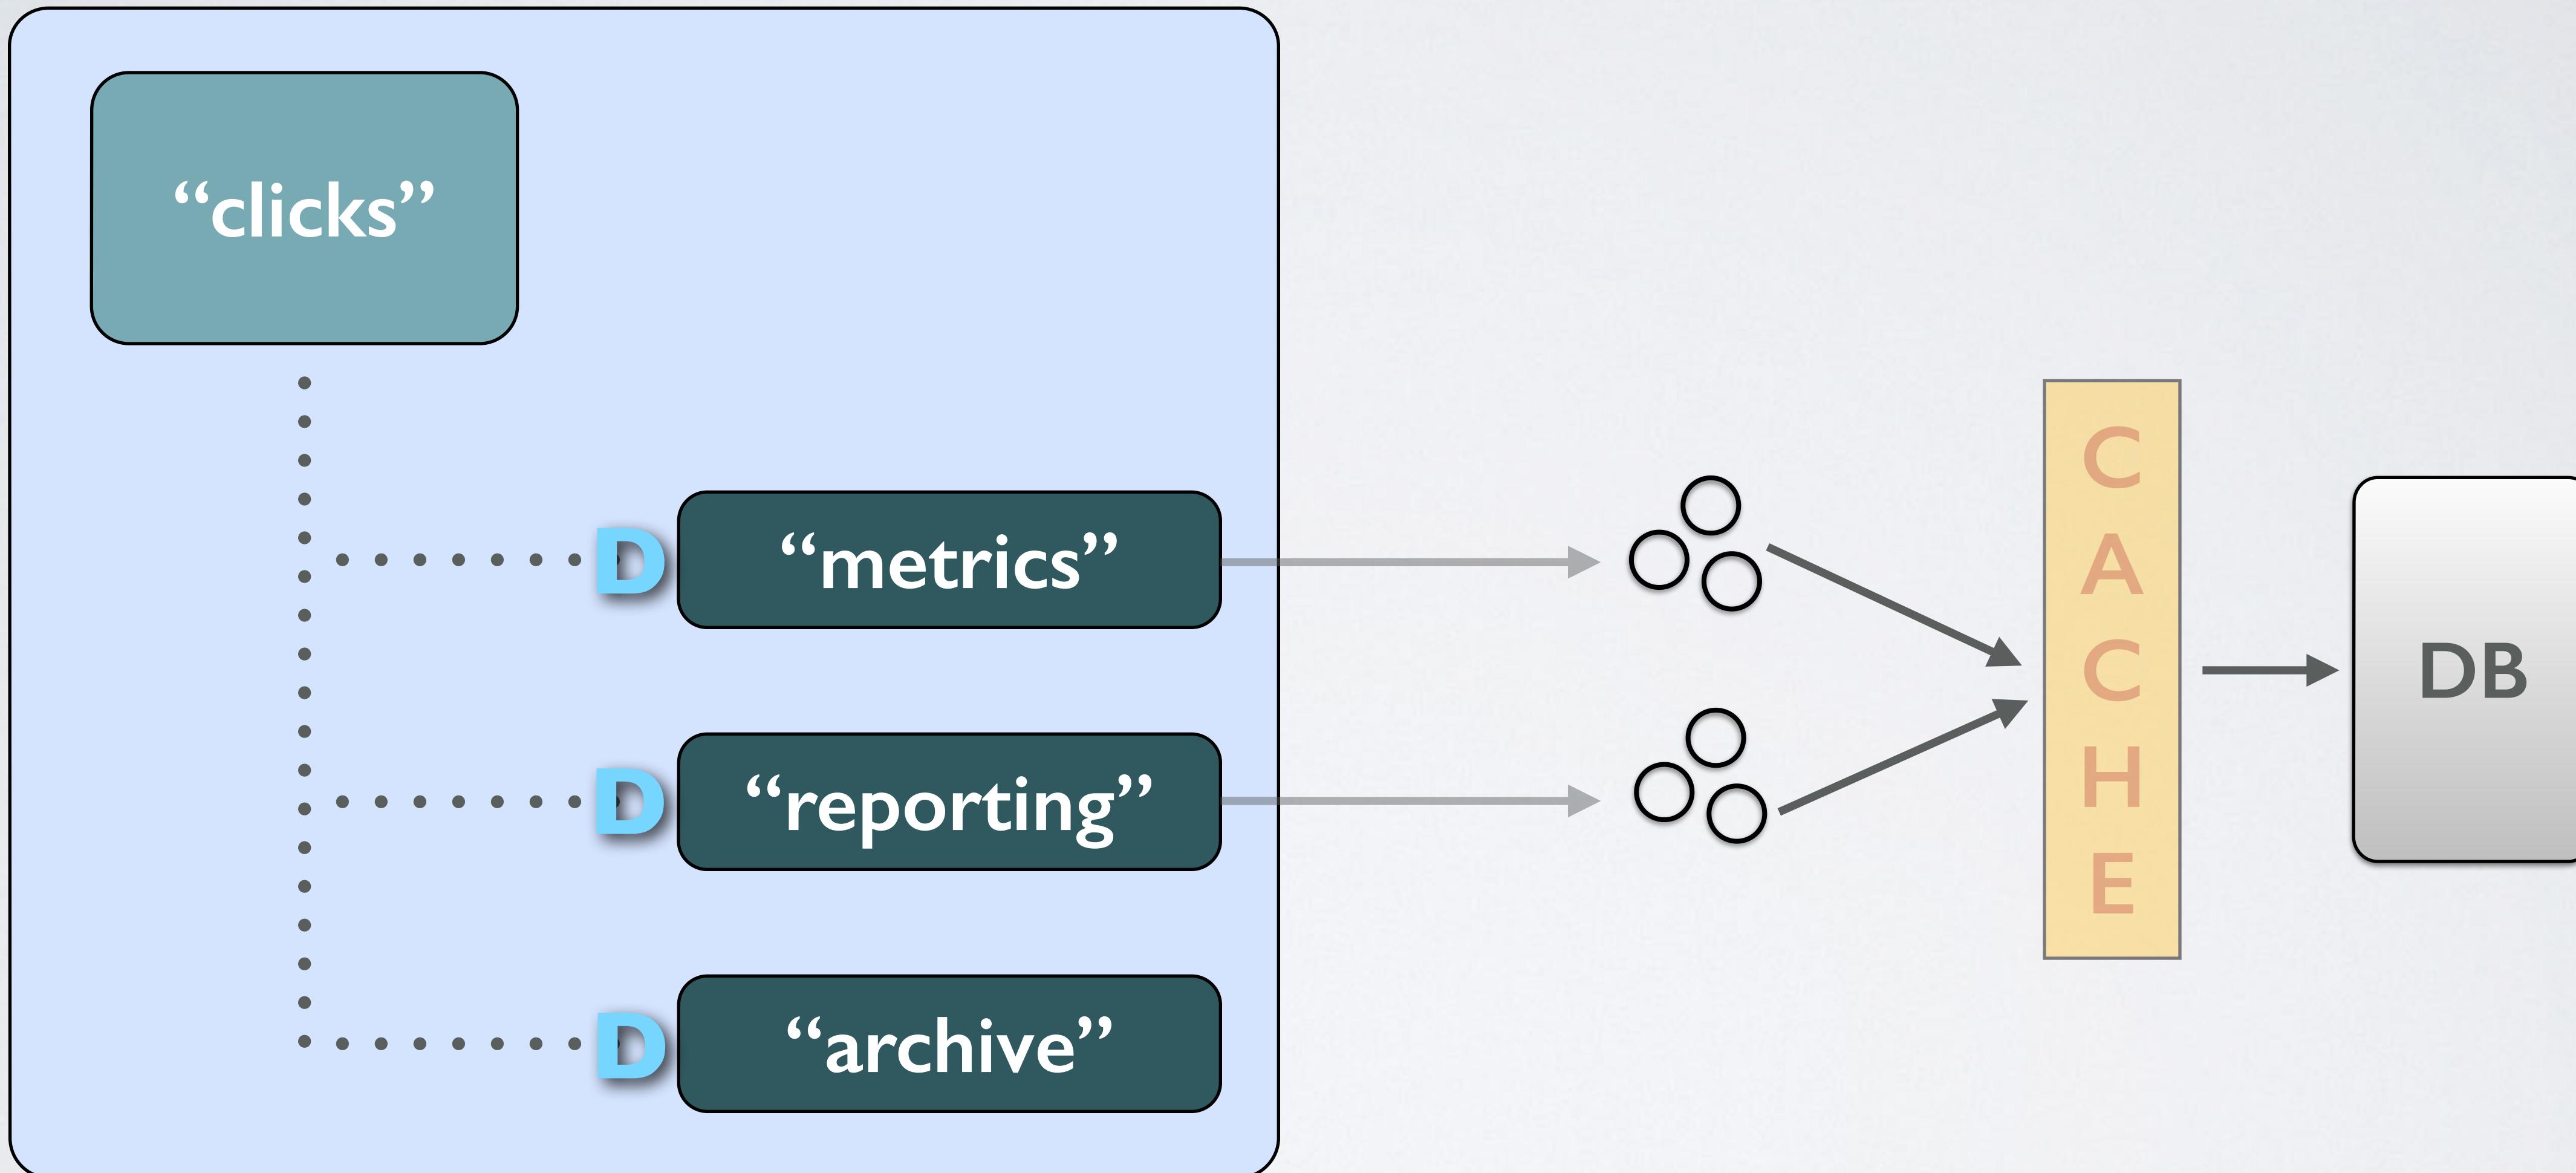
TOPIC & CHANNEL PAUSING



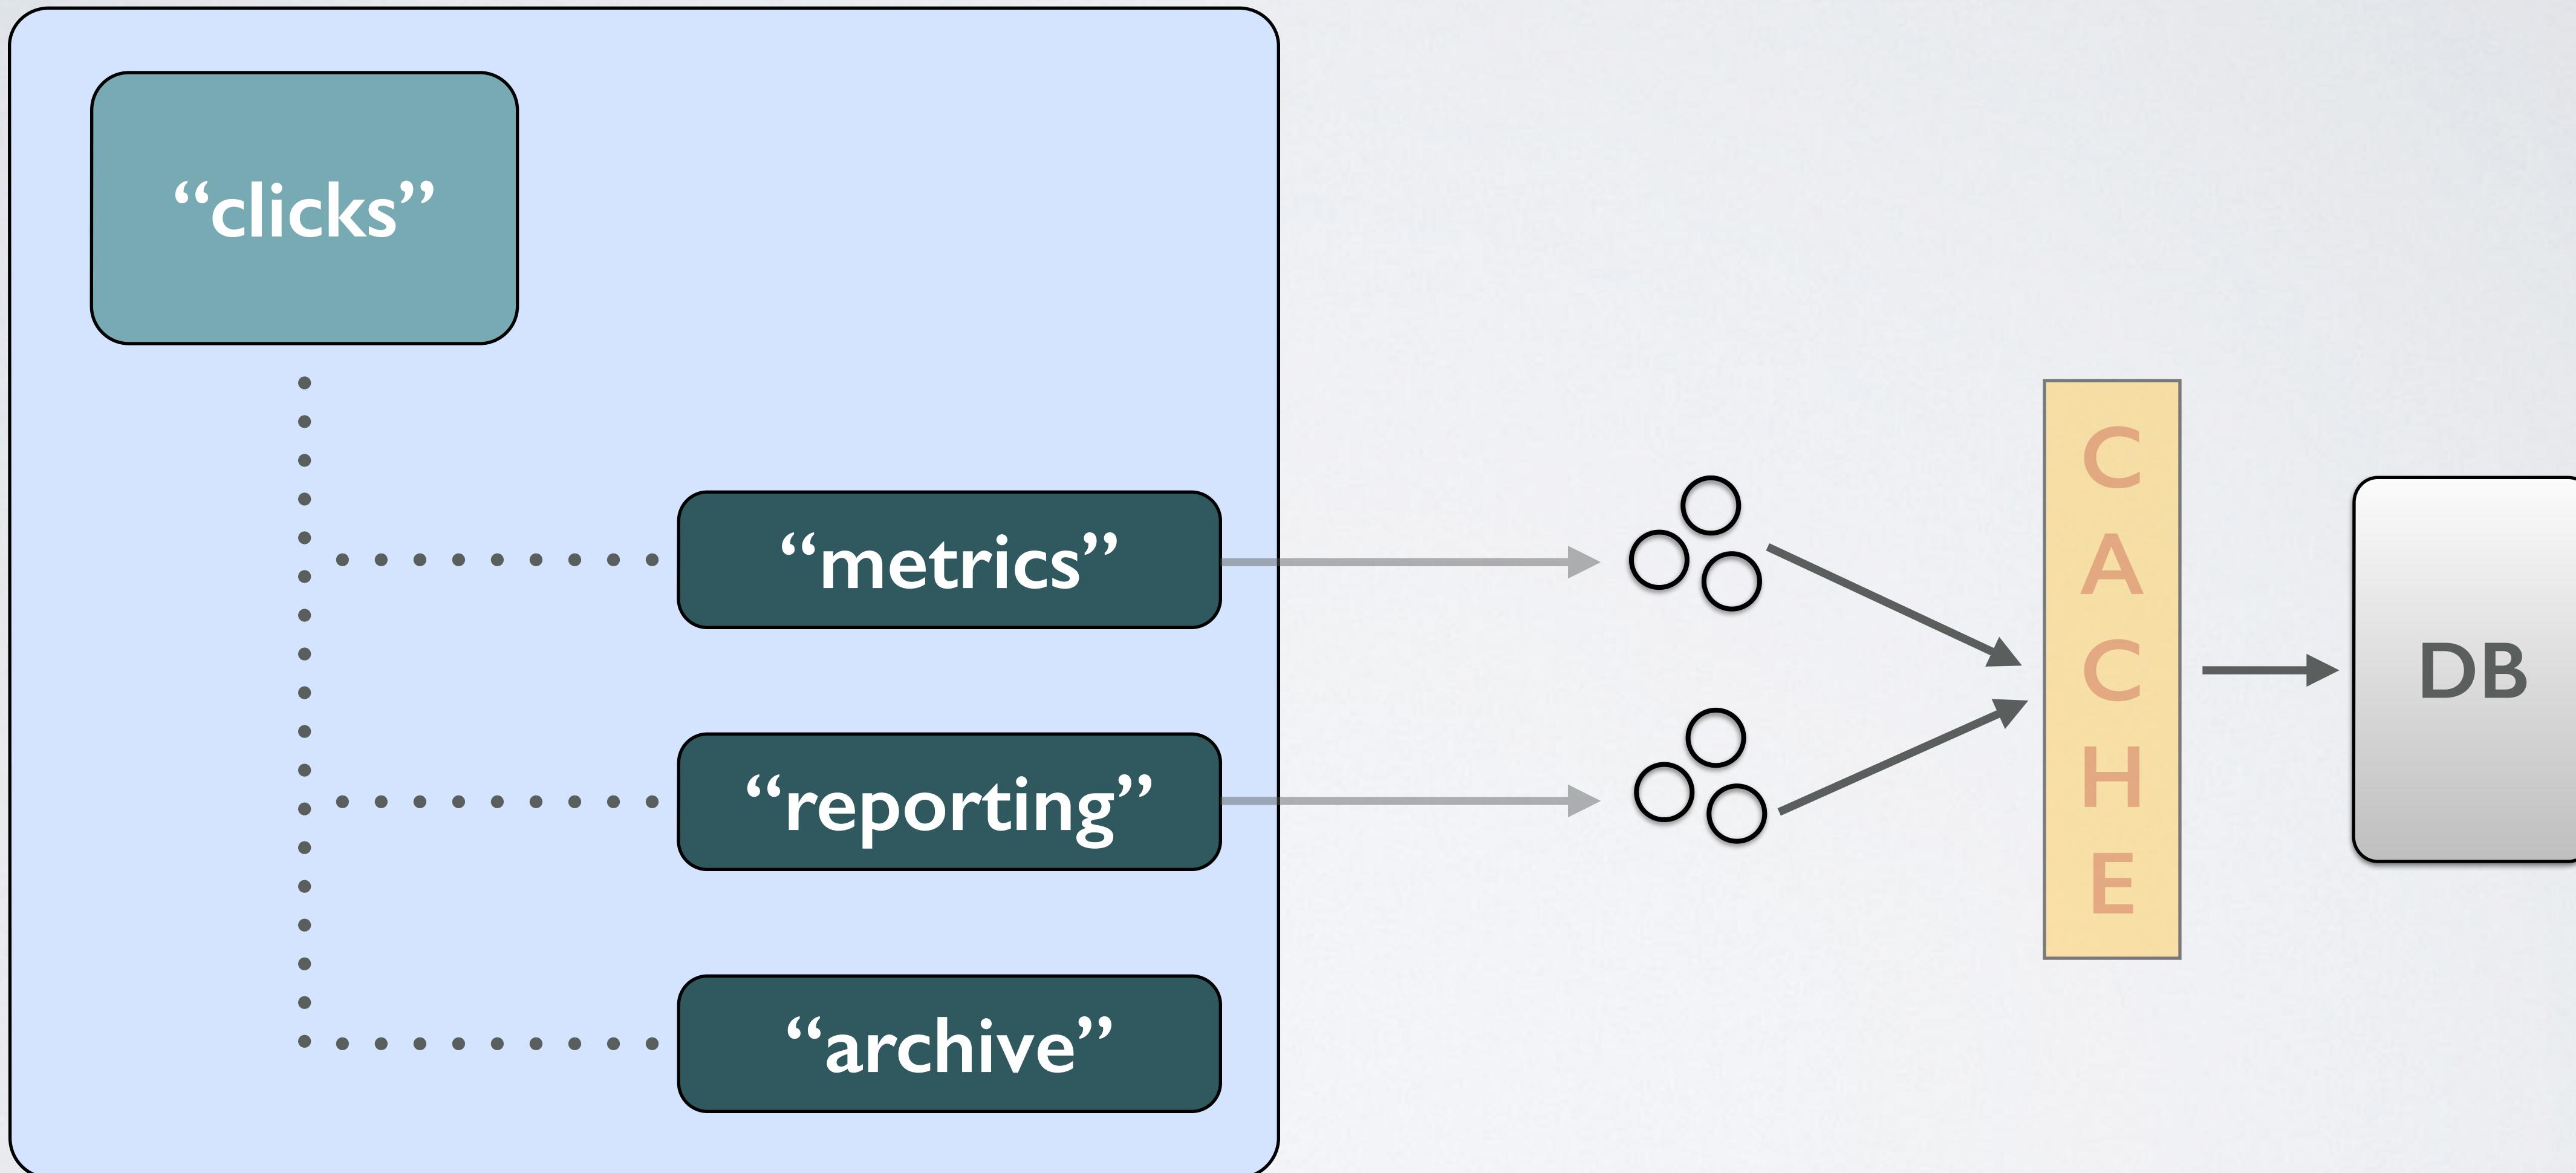
TOPIC & CHANNEL PAUSING



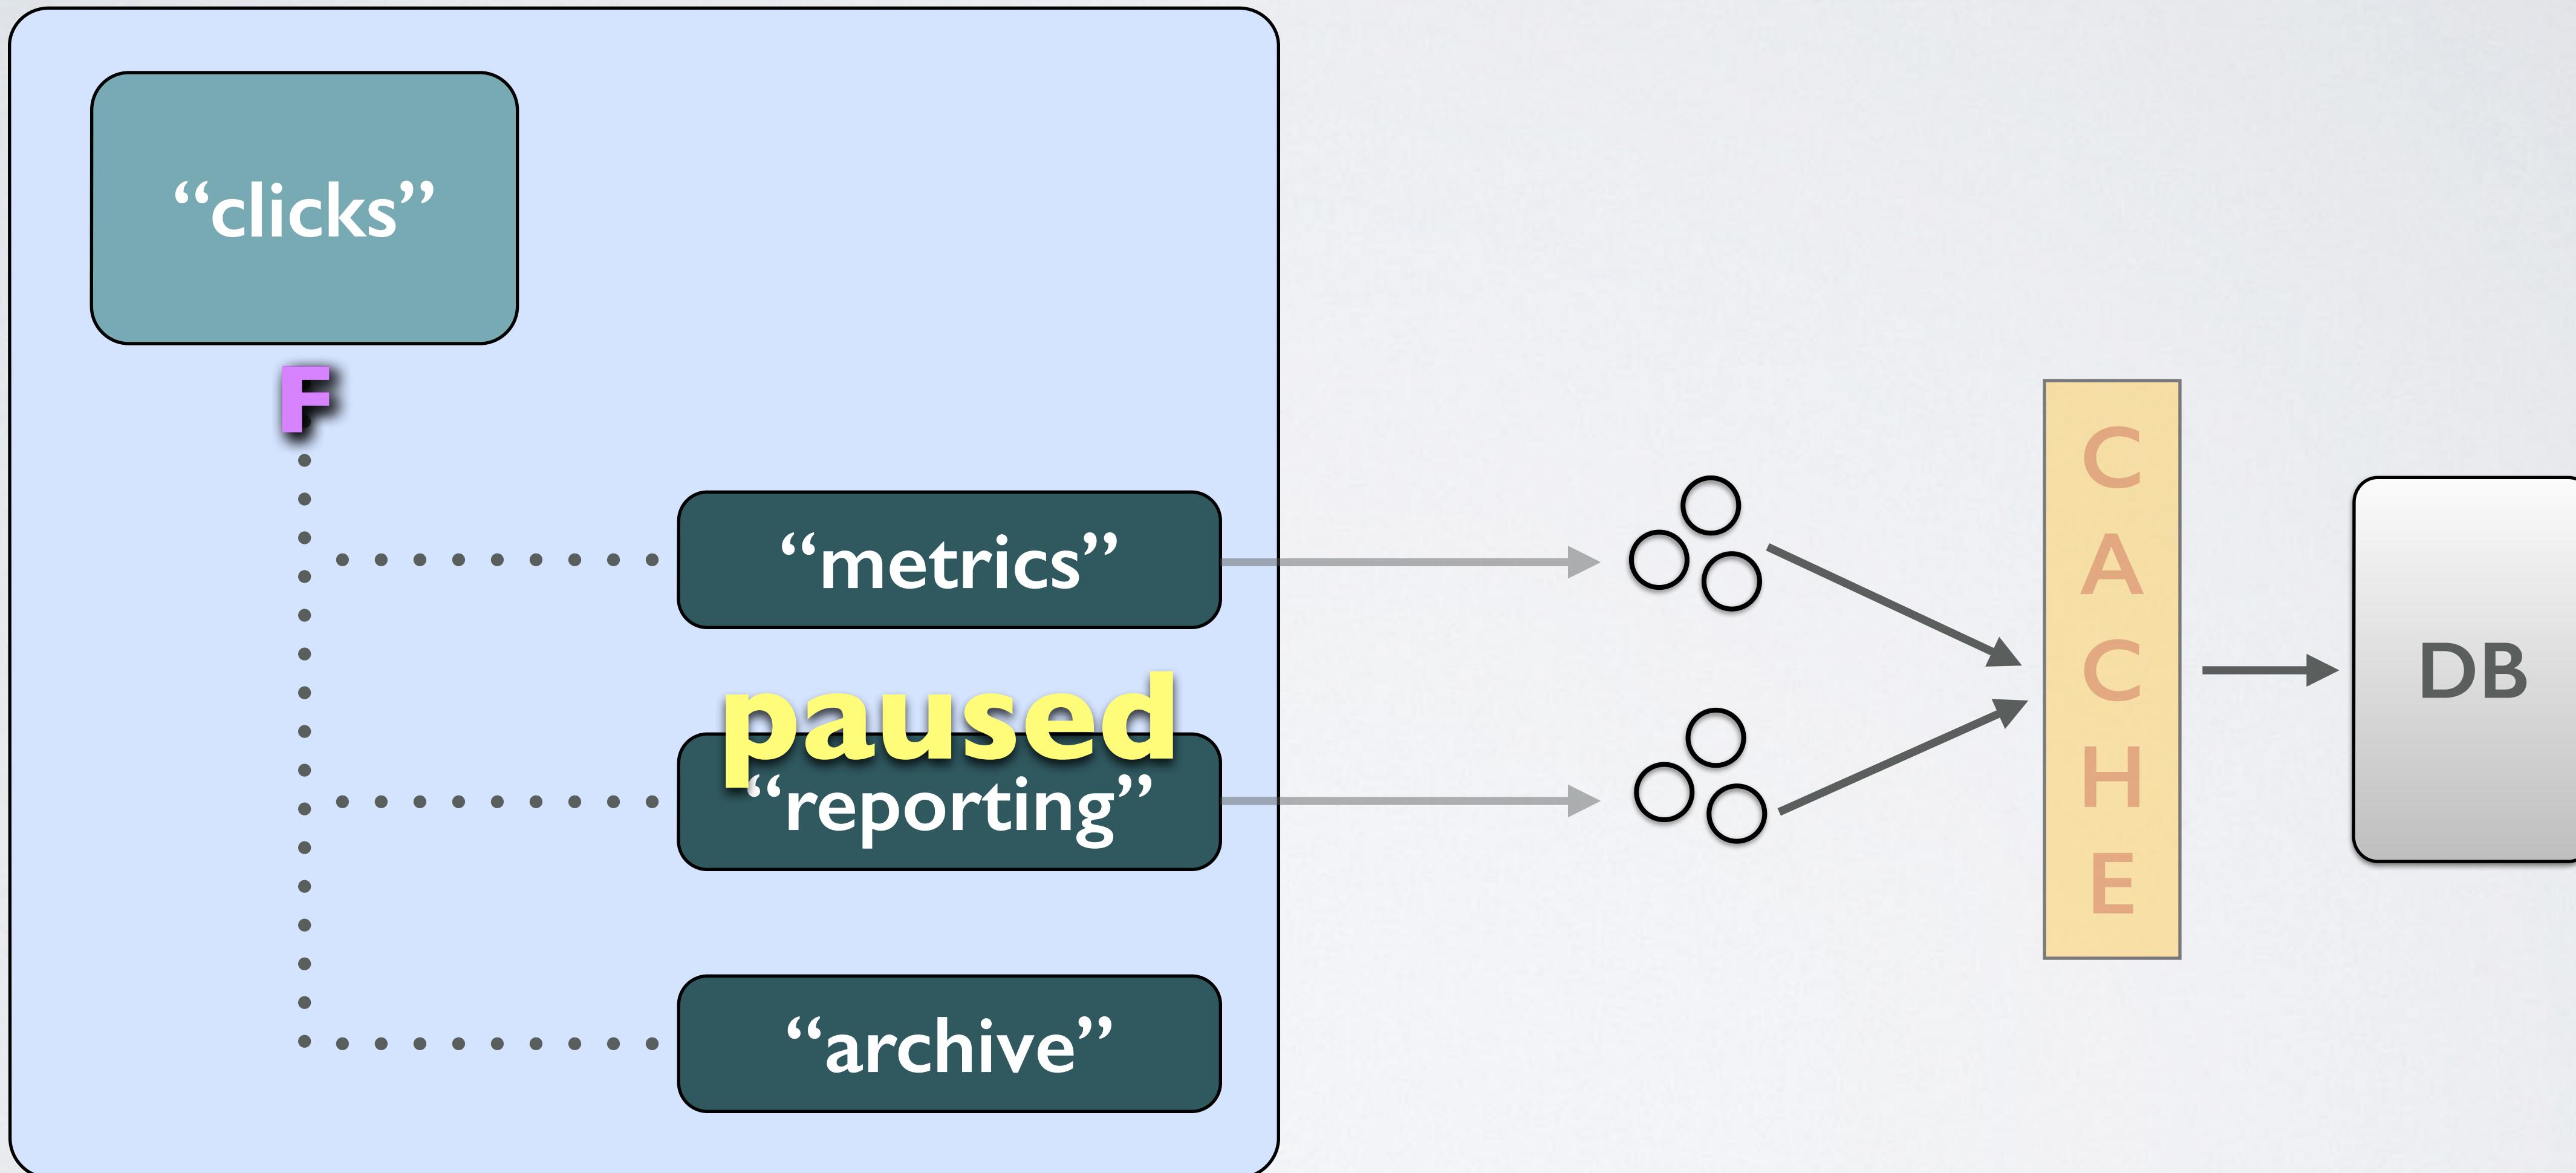
TOPIC & CHANNEL PAUSING



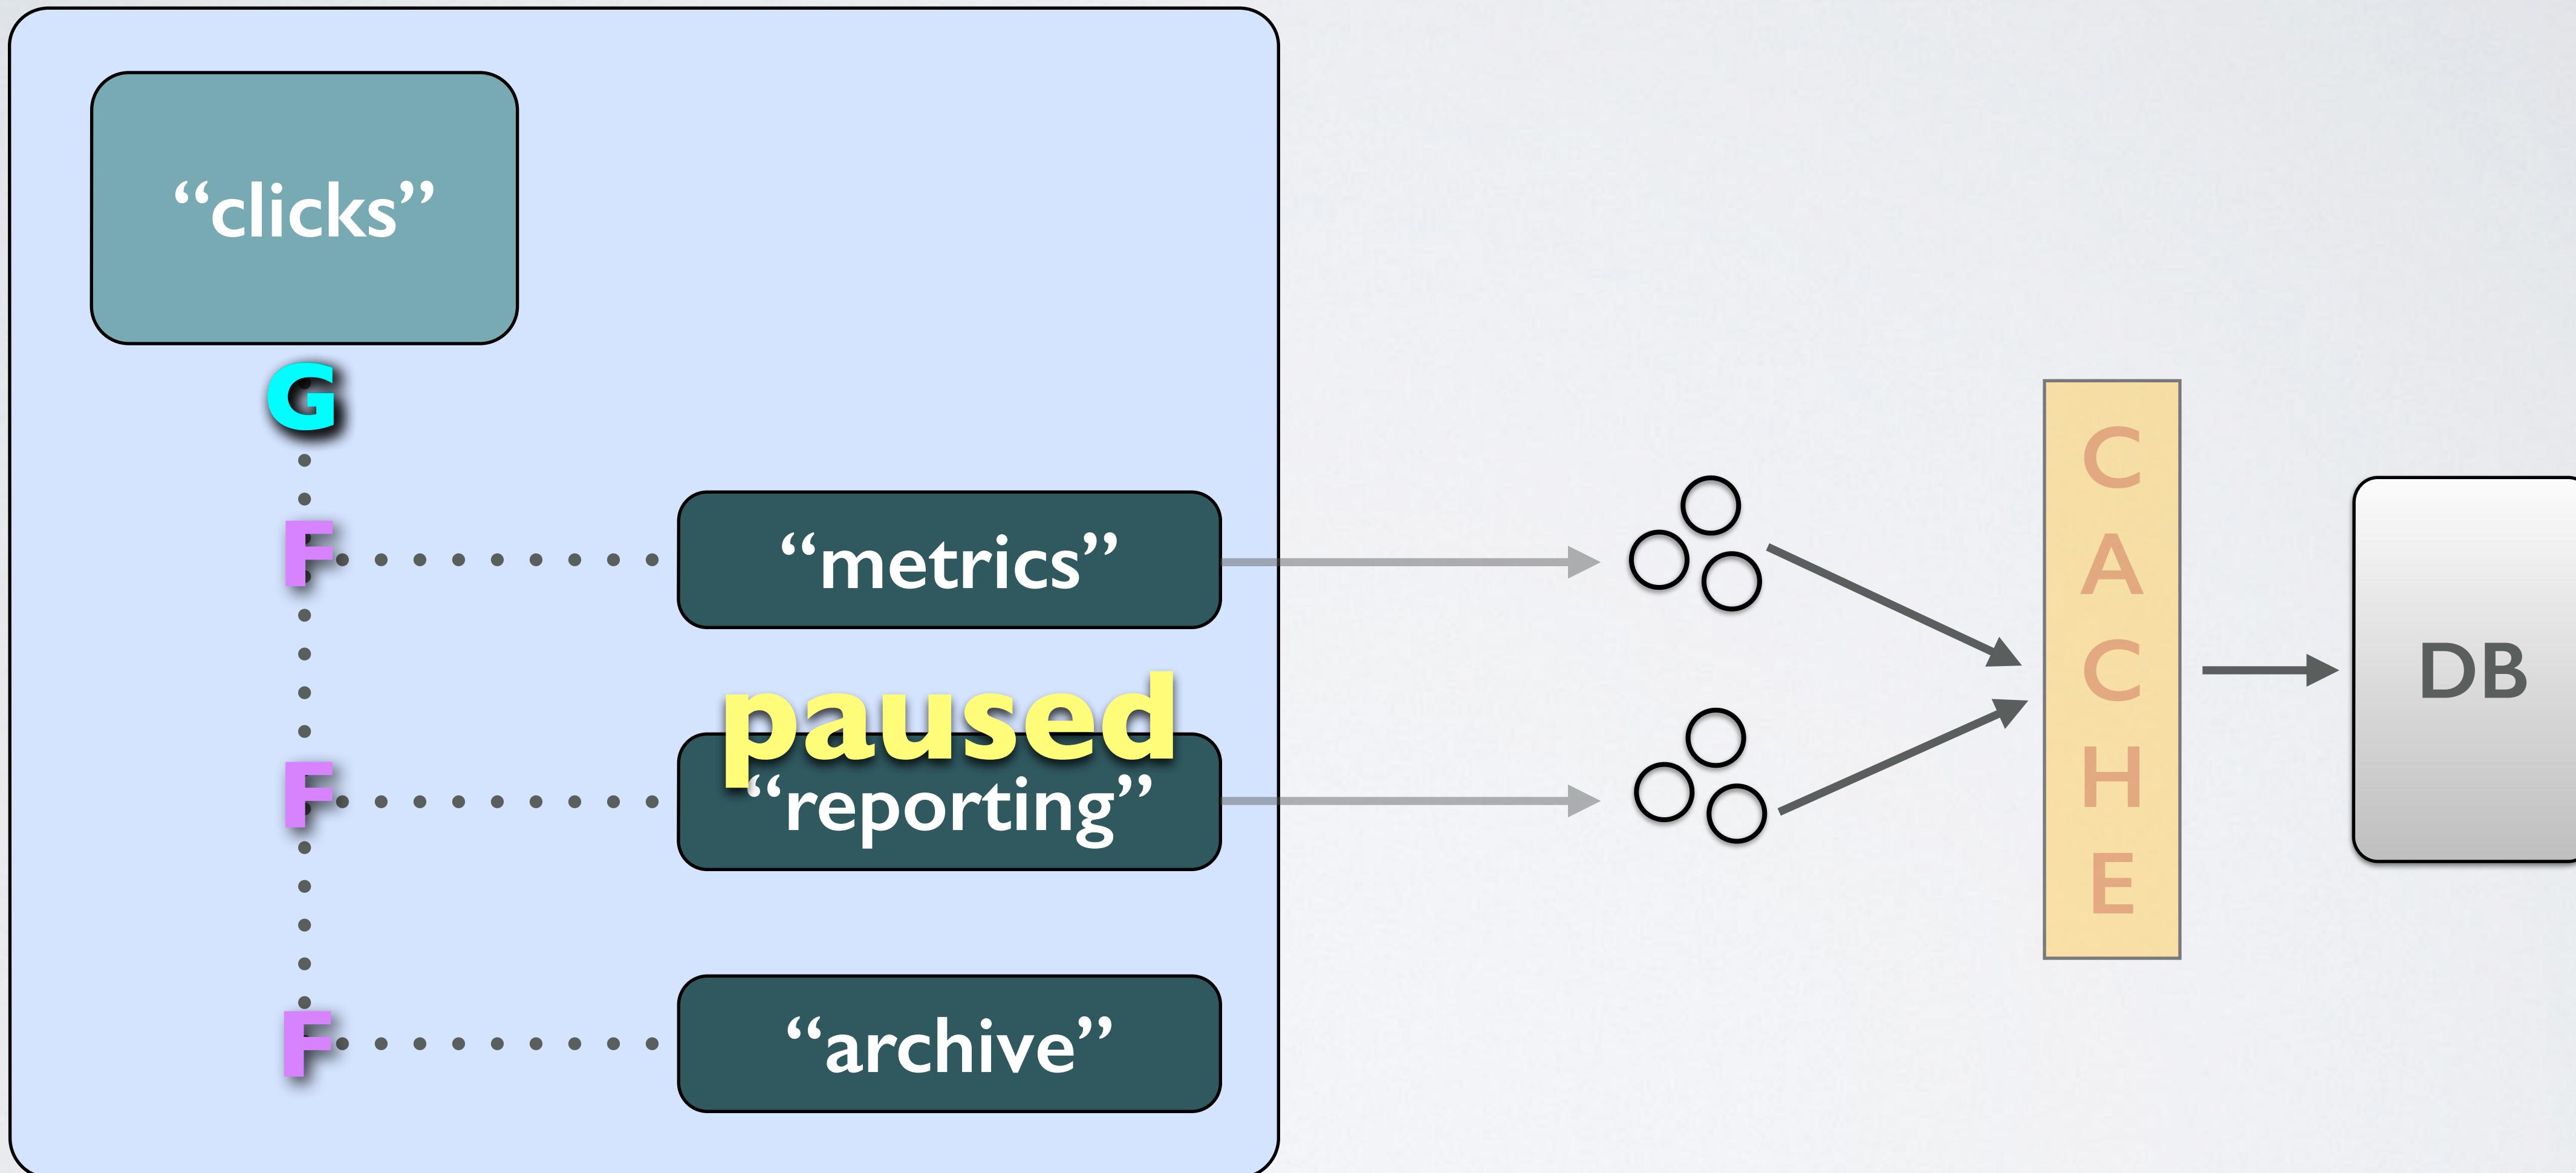
TOPIC & CHANNEL PAUSING



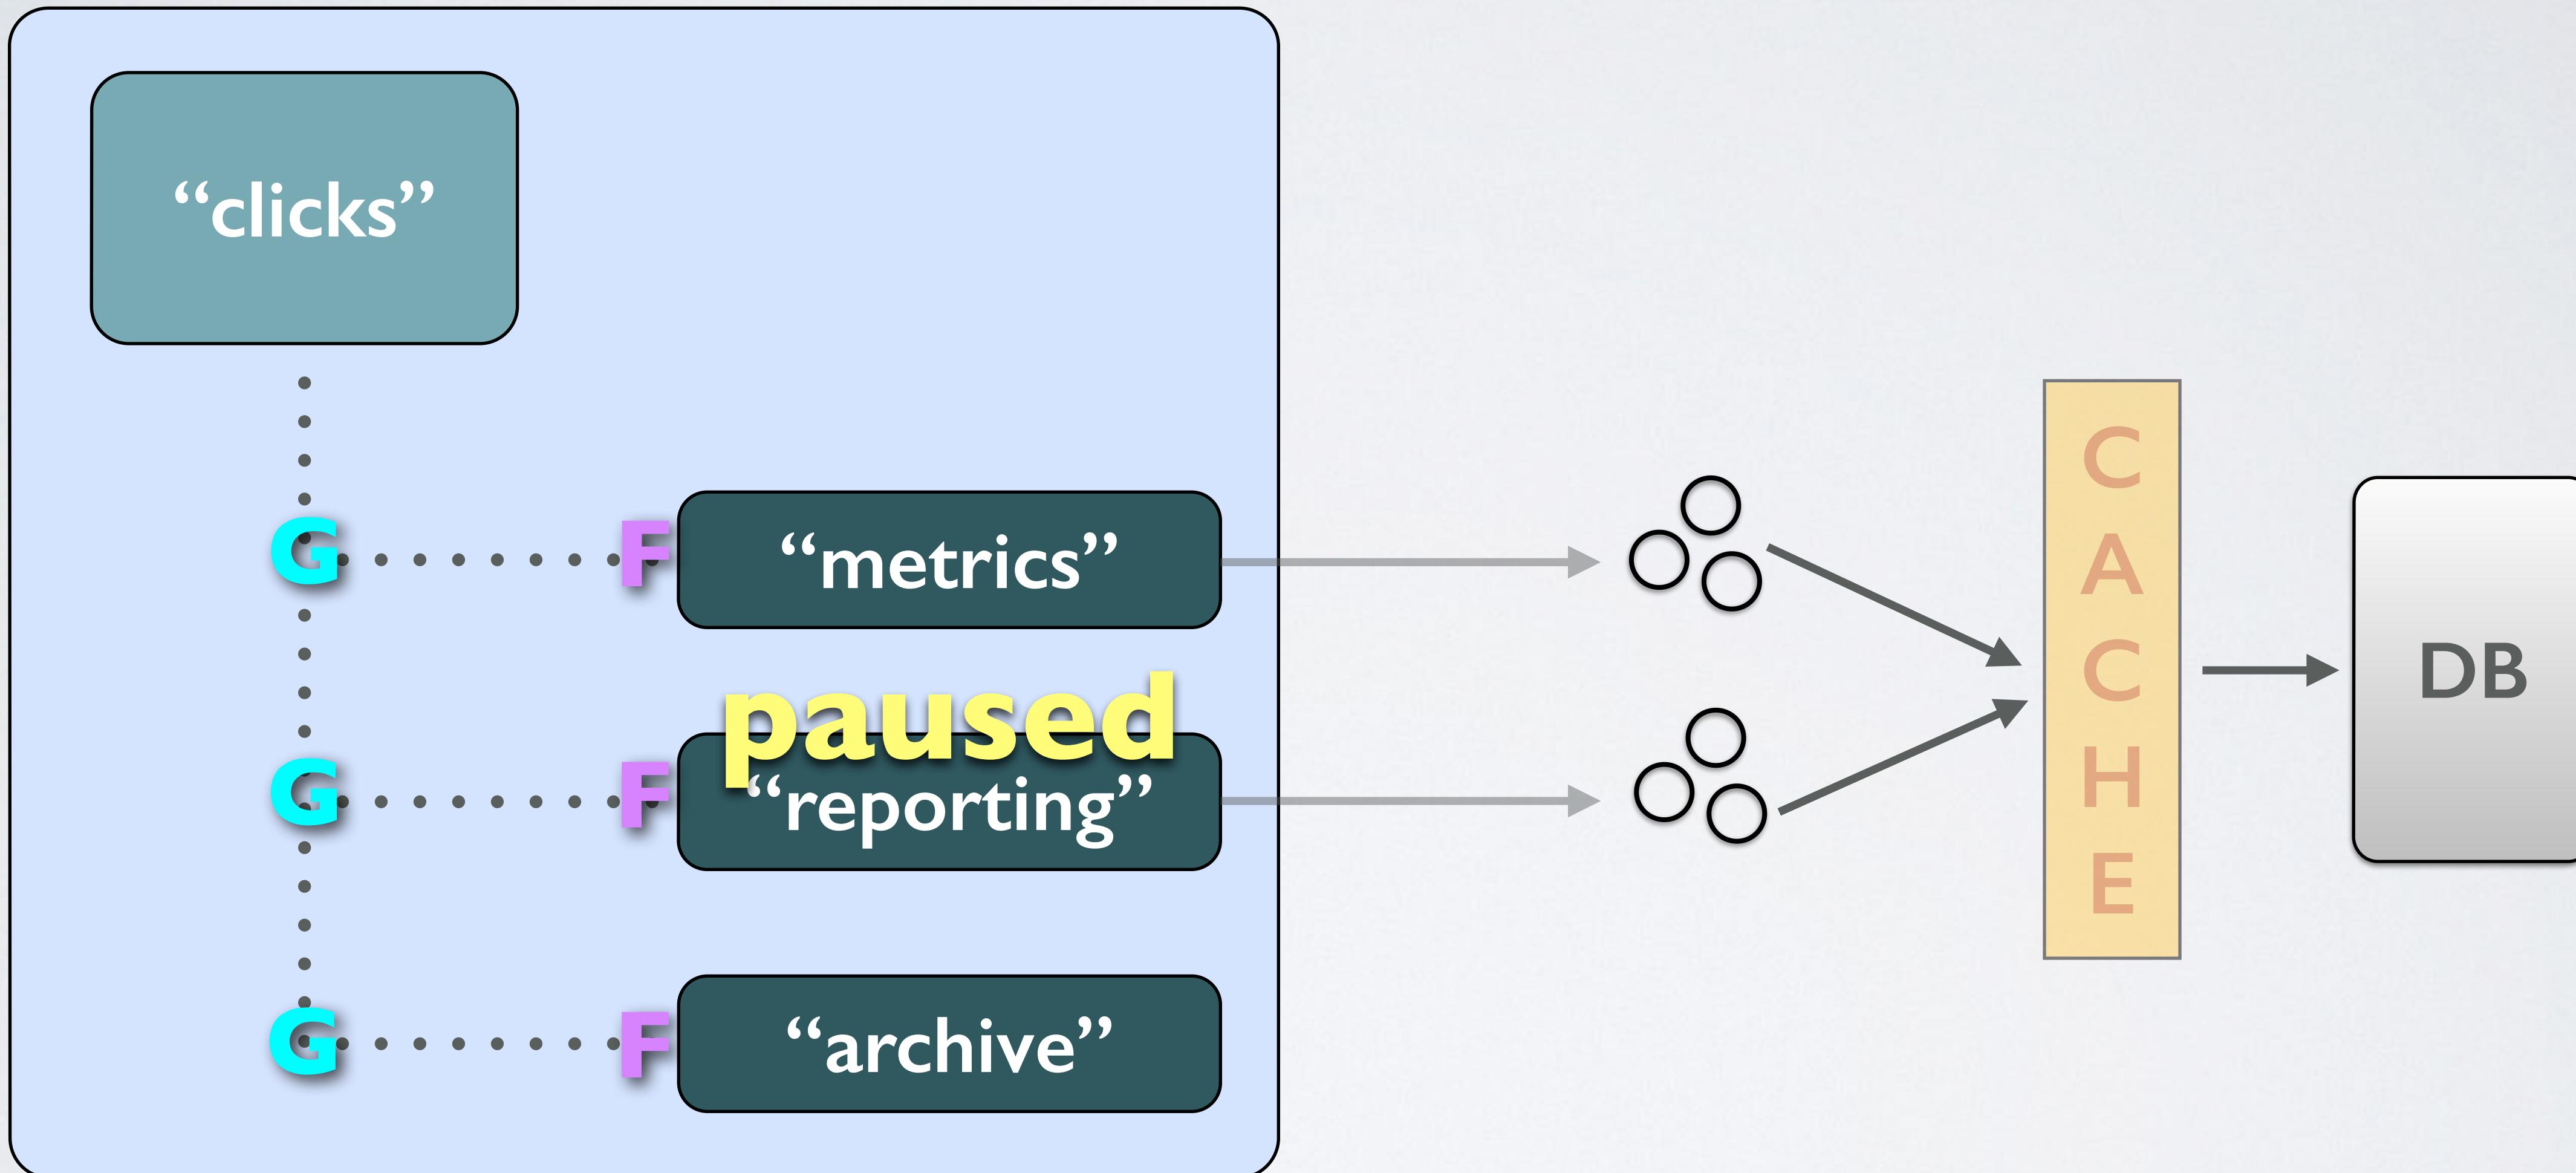
TOPIC & CHANNEL PAUSING



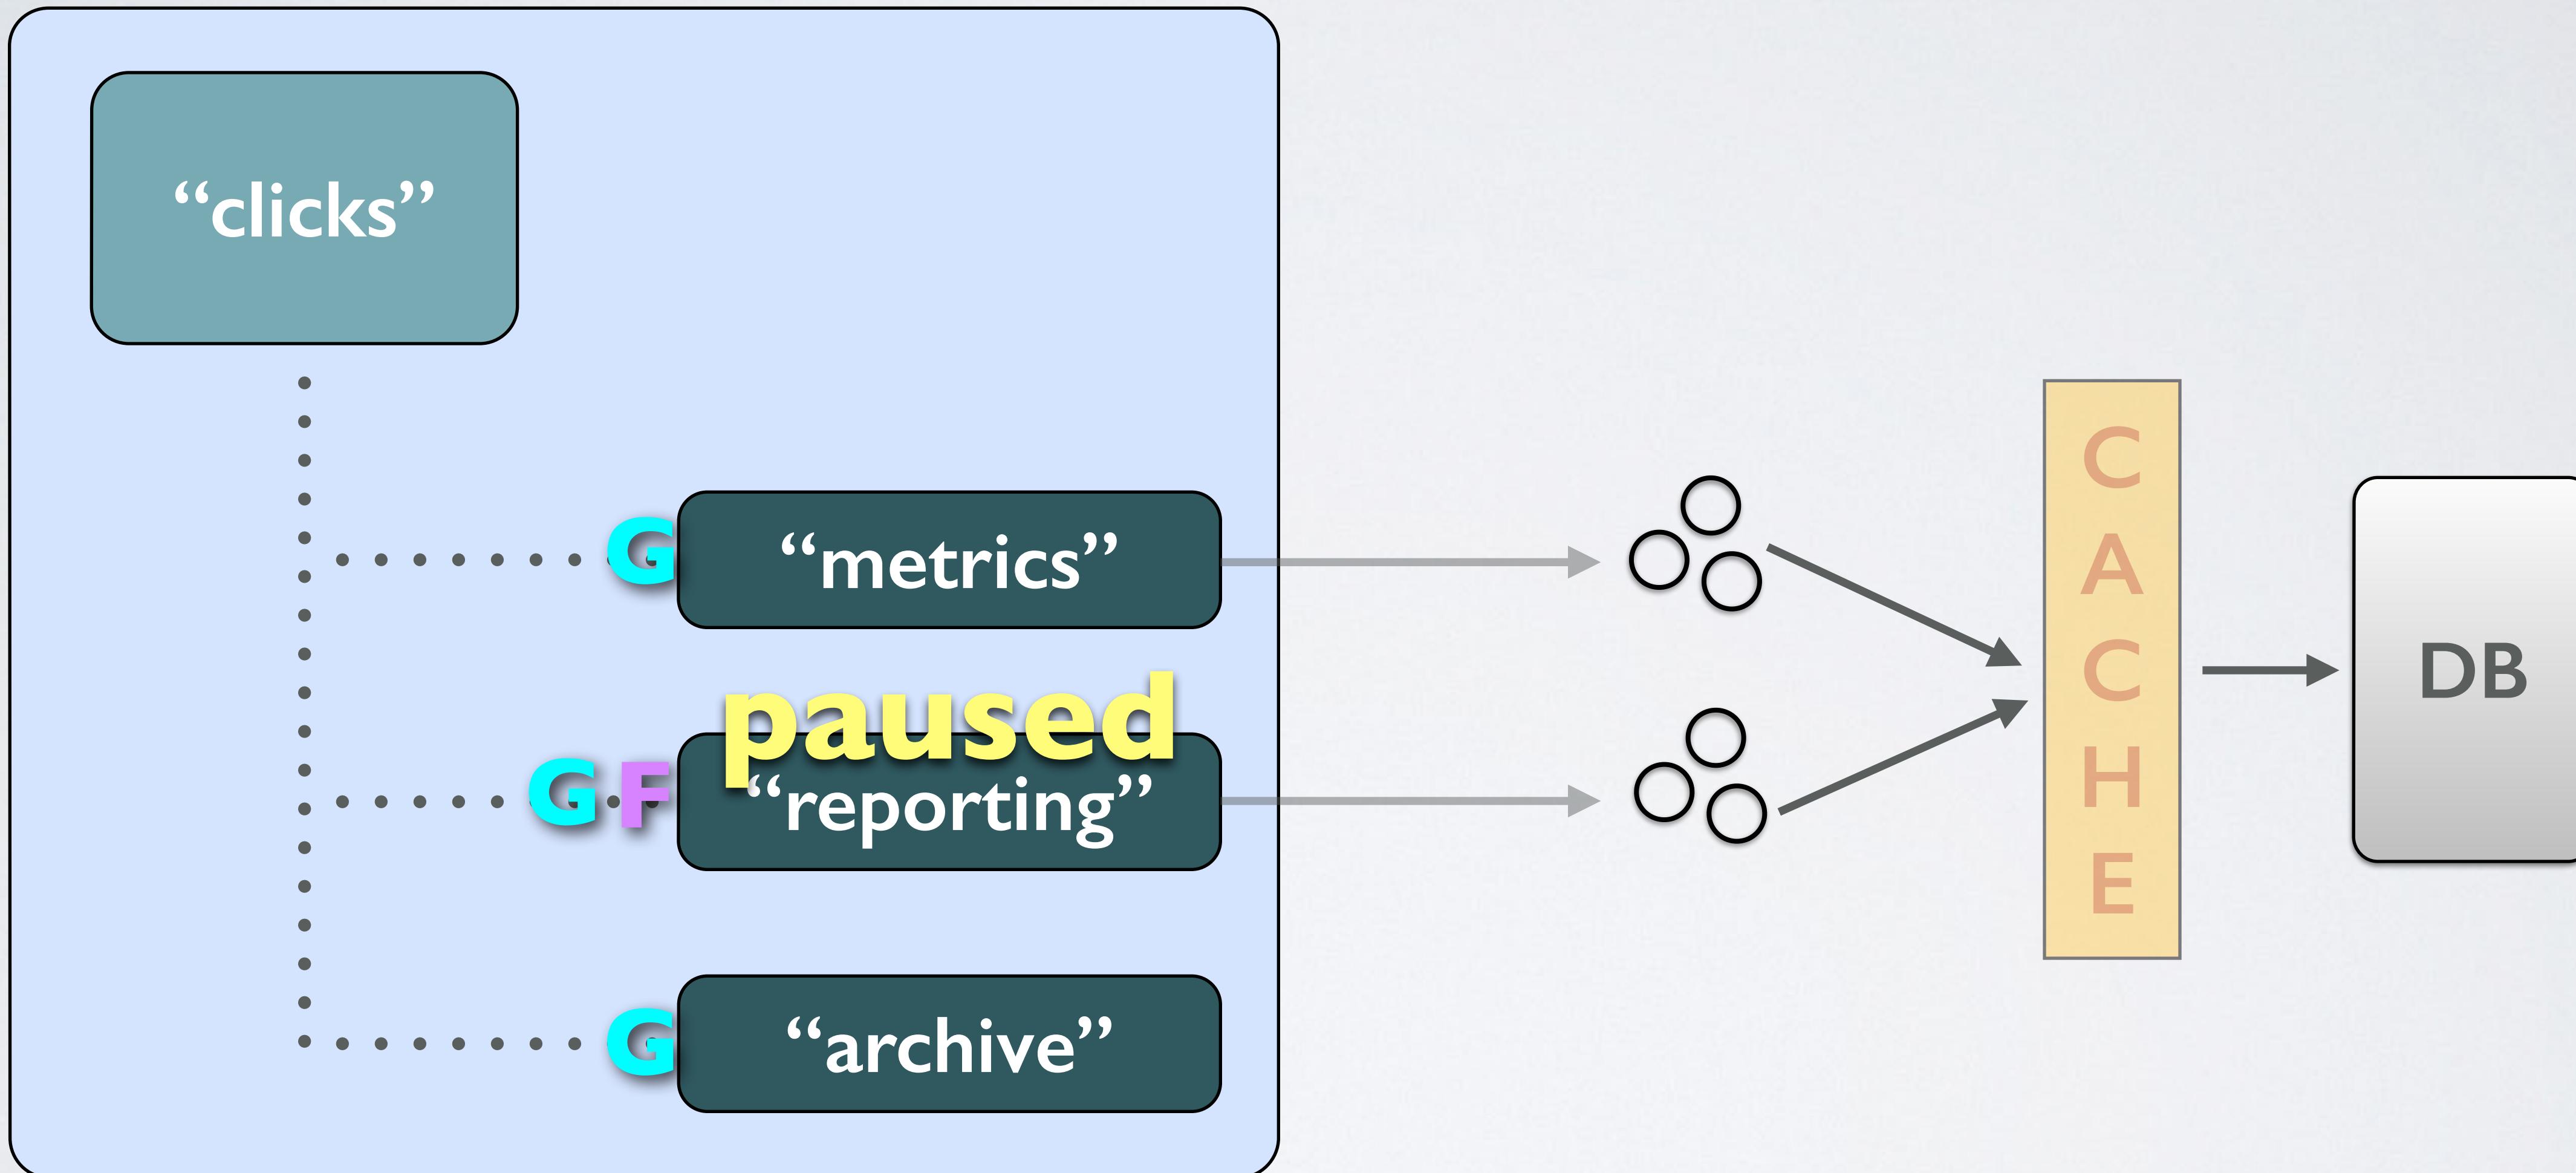
TOPIC & CHANNEL PAUSING



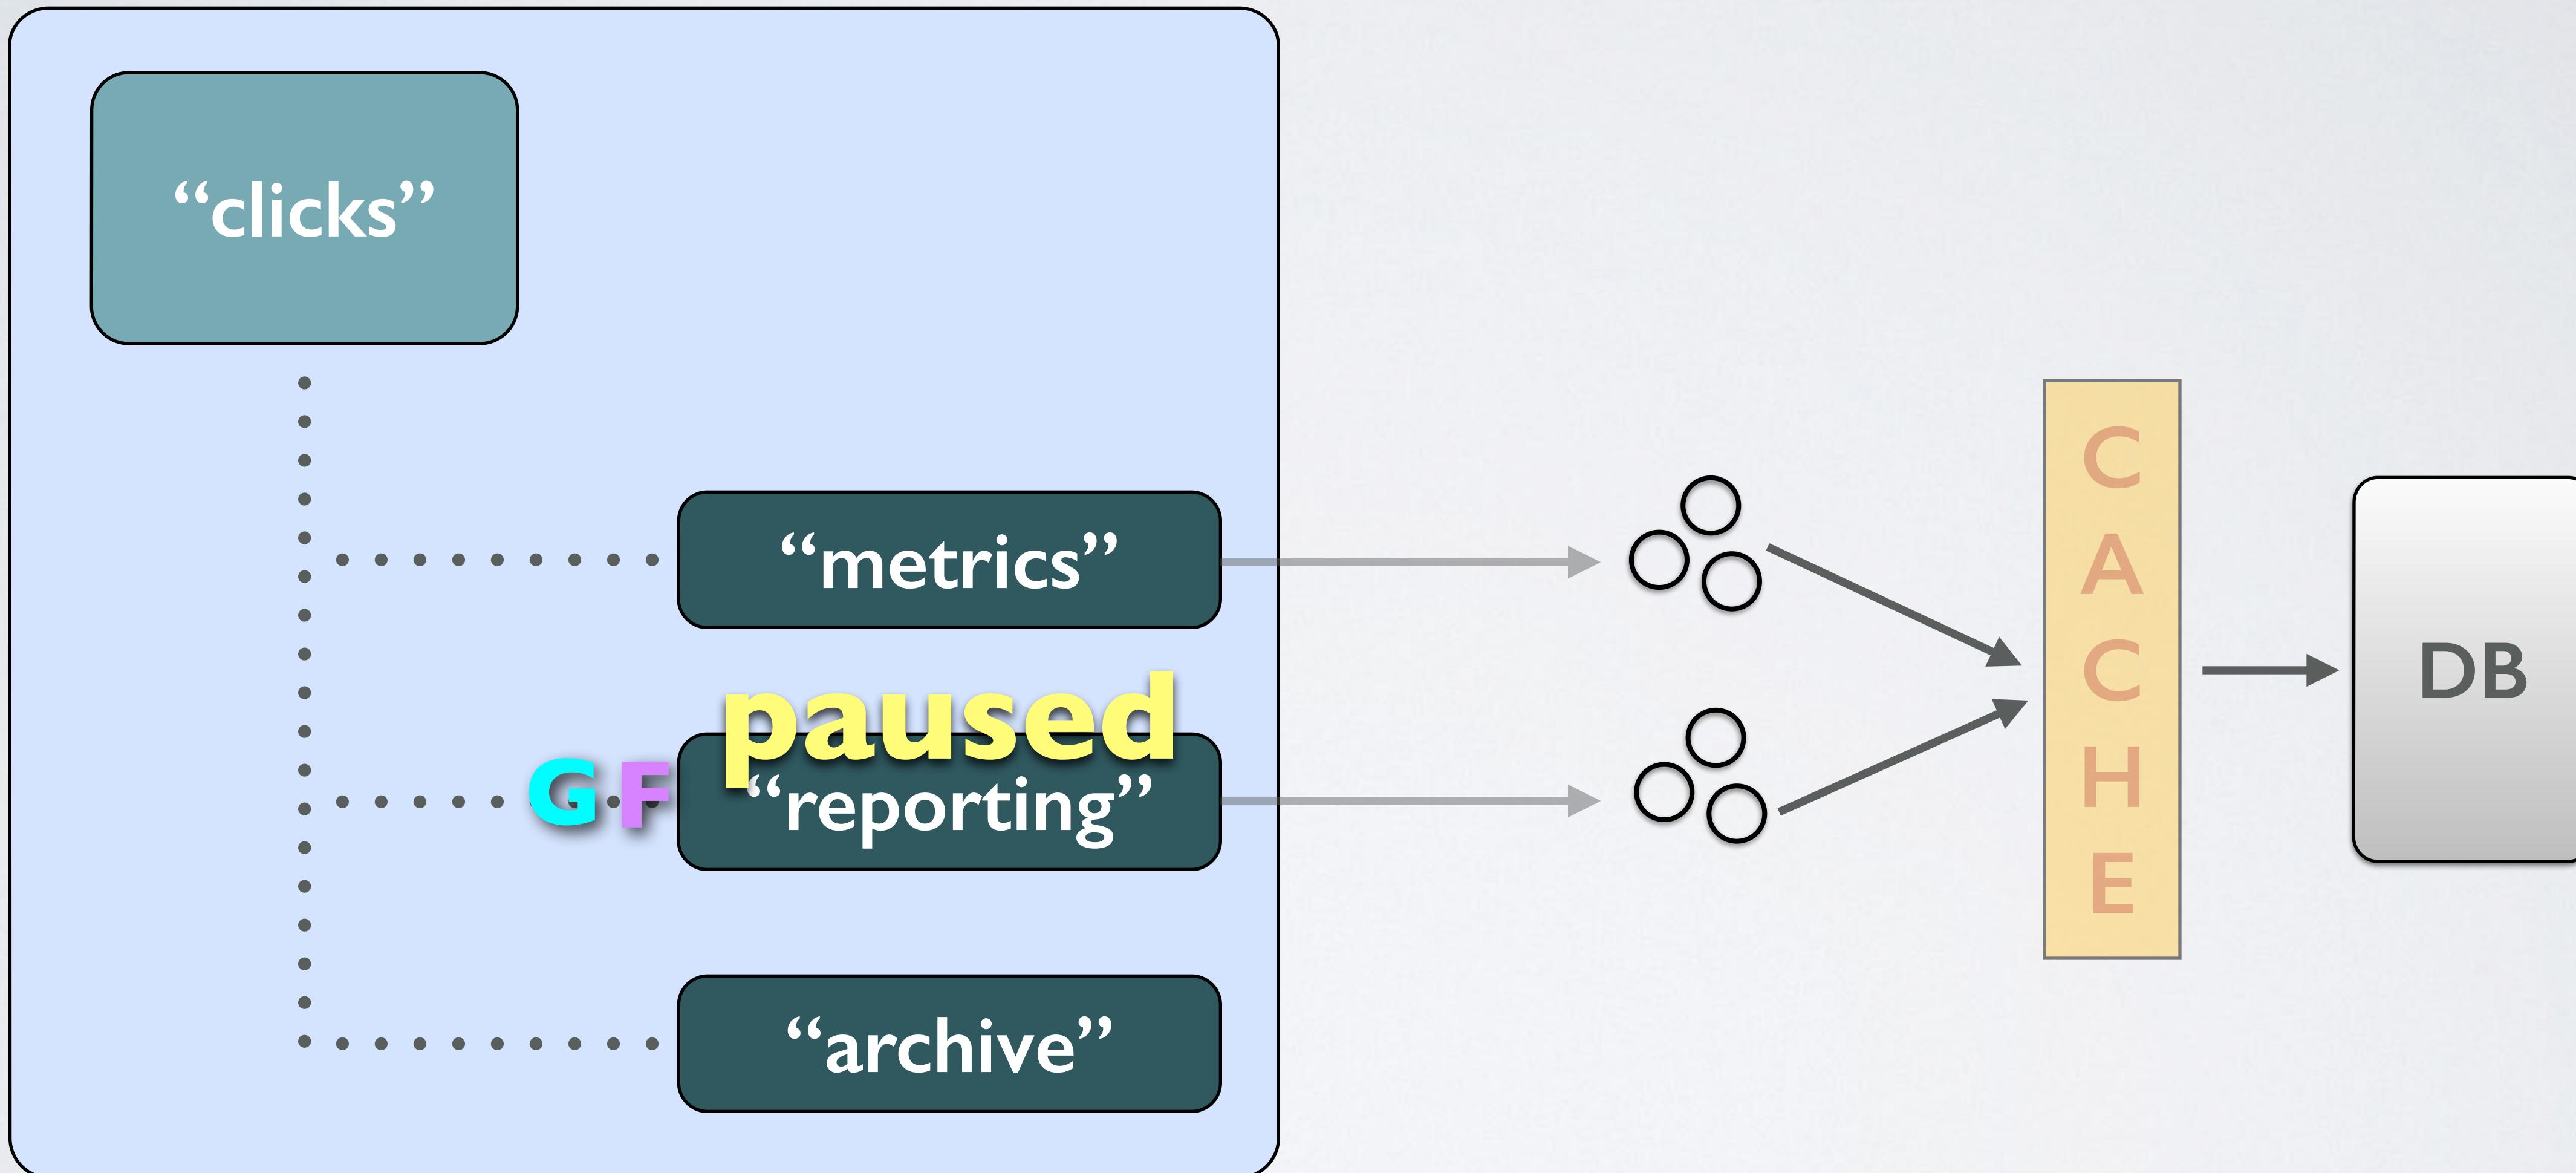
TOPIC & CHANNEL PAUSING



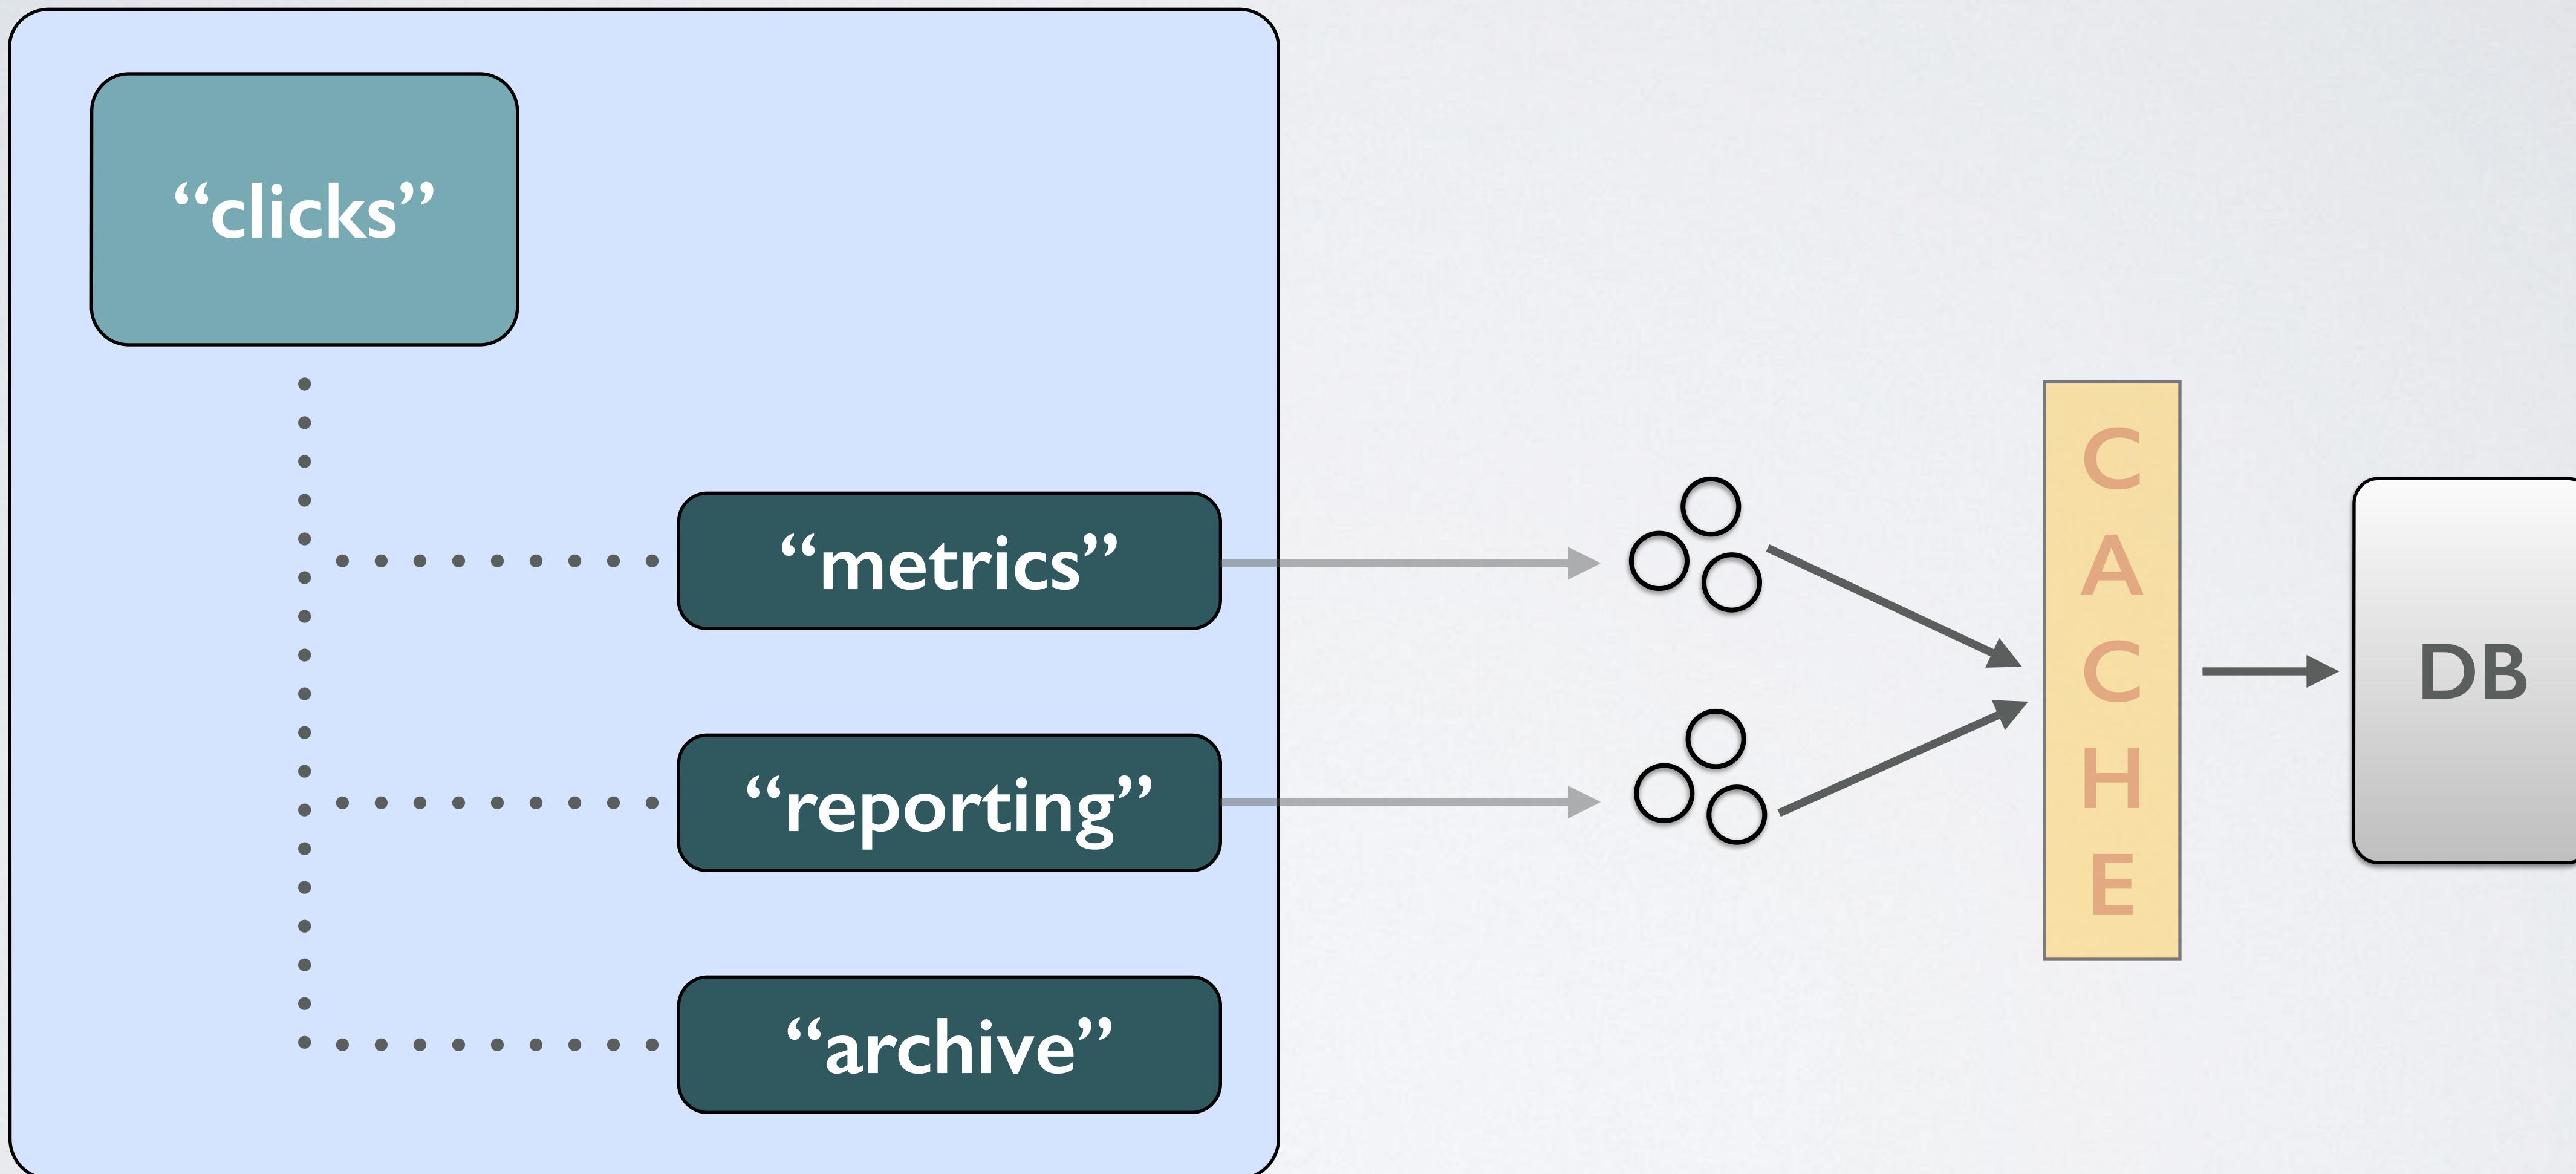
TOPIC & CHANNEL PAUSING



TOPIC & CHANNEL PAUSING

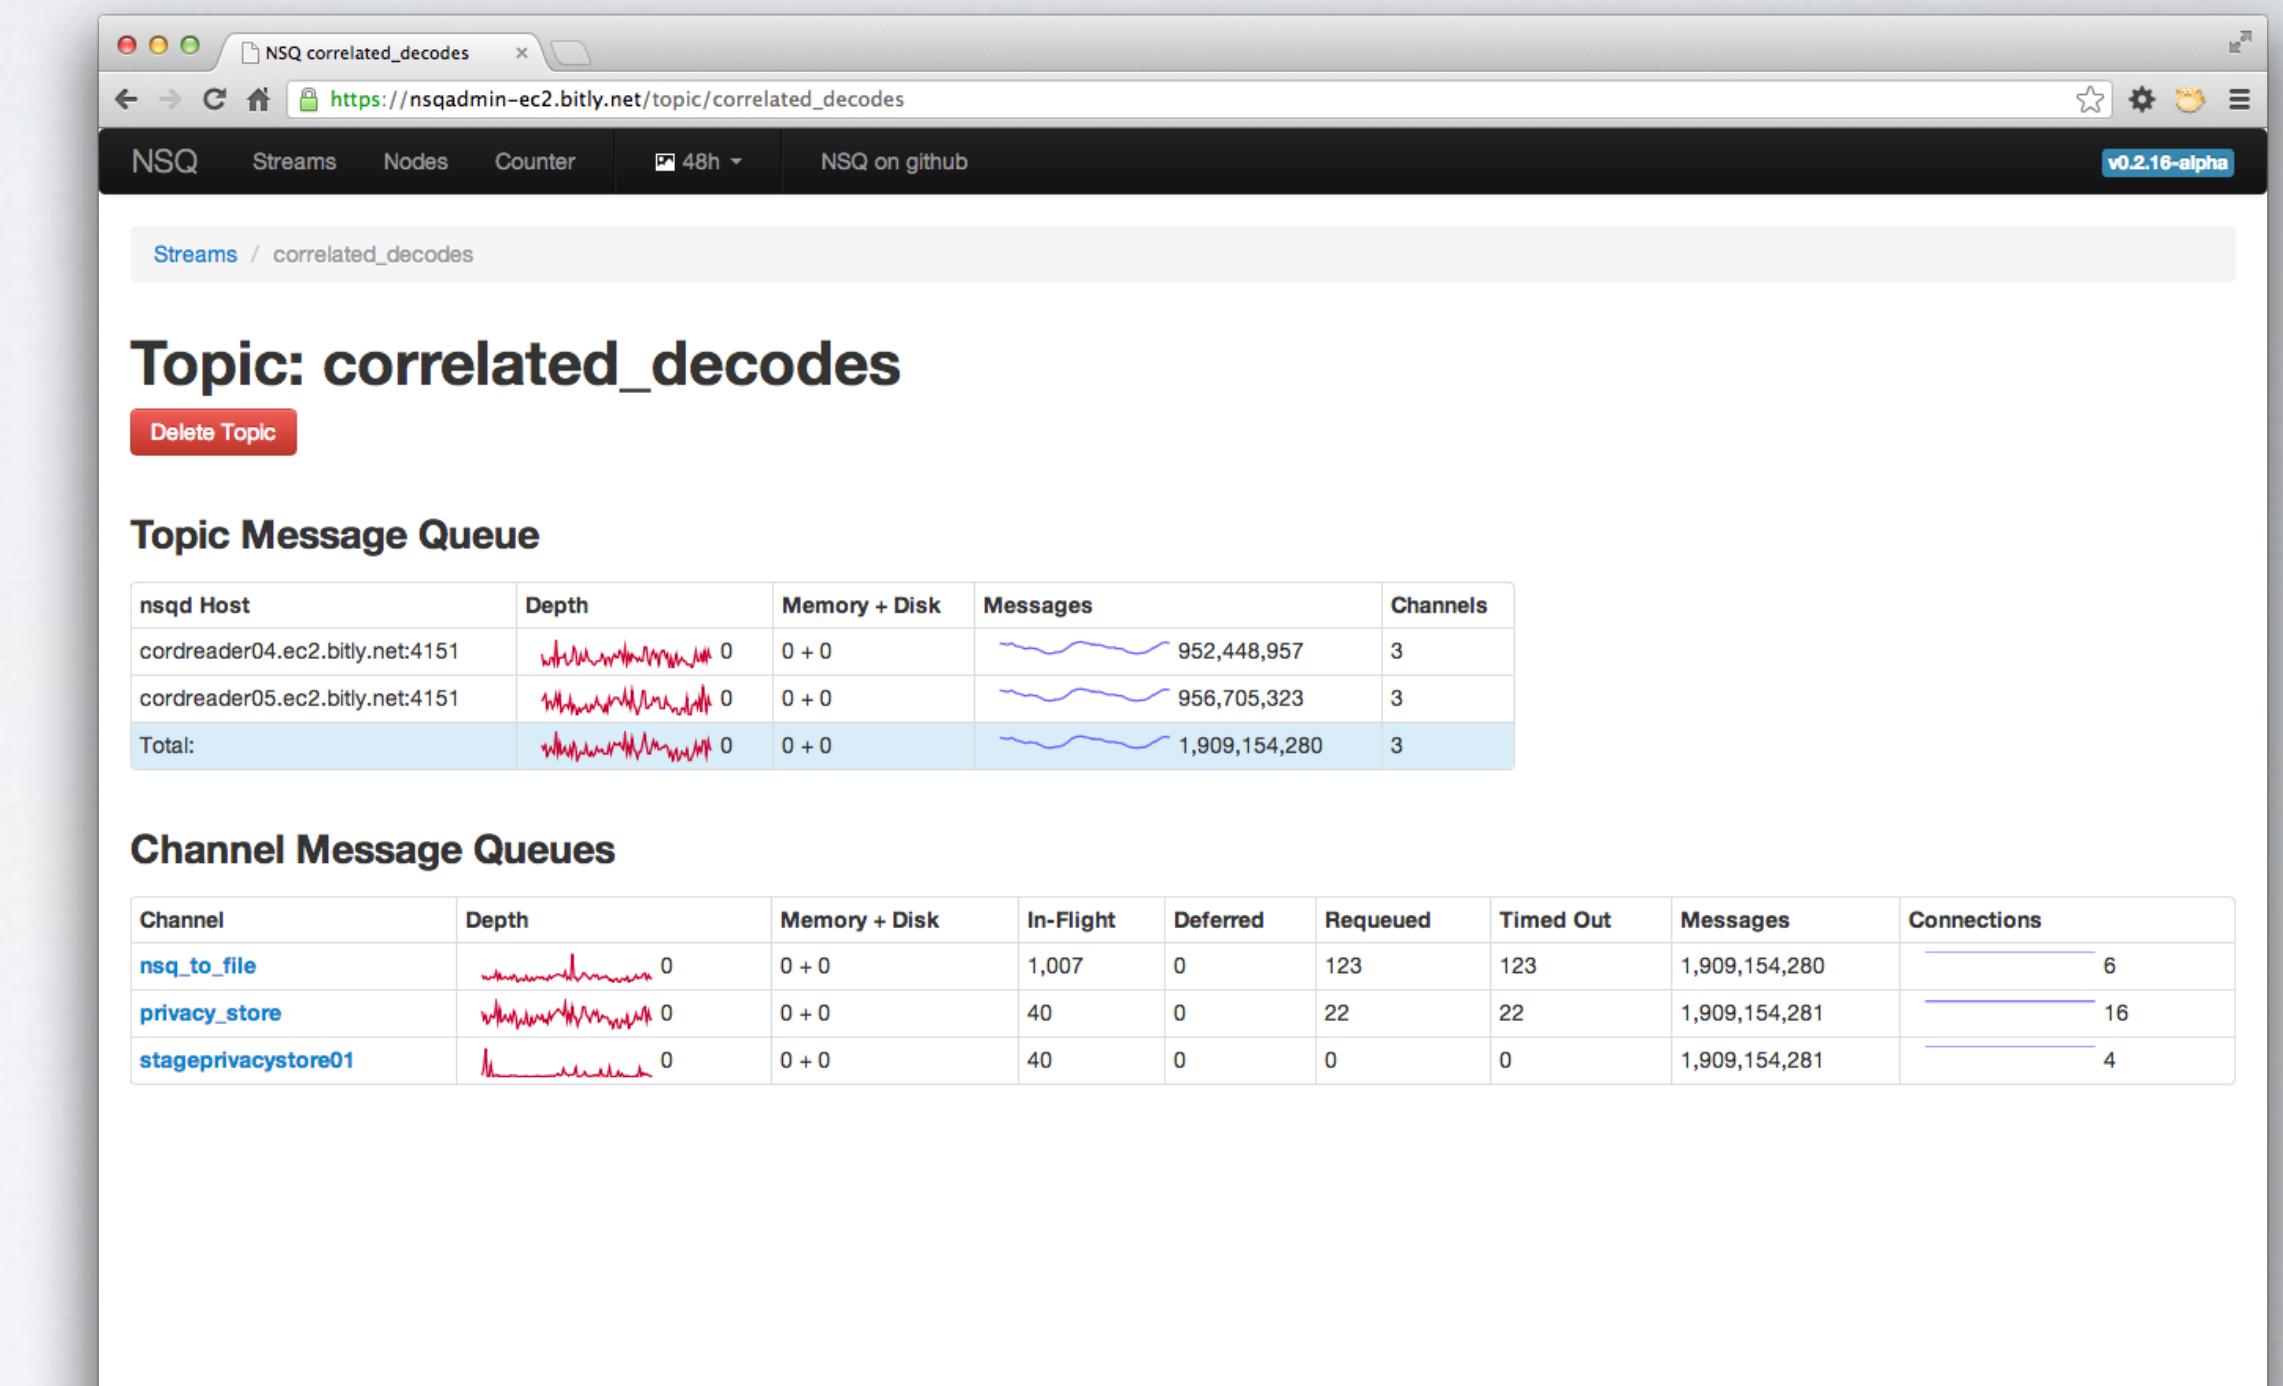


TOPIC & CHANNEL PAUSING



ONE MORE THING

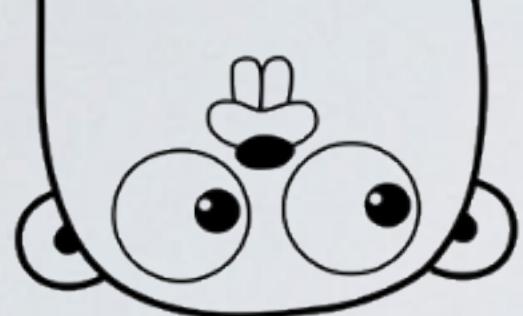
- #ephemeral channels
- channel sampling
- TLS / Snappy
- telemetry over statsd / HTTP



SUMMARY

- github.com/bitly/nsq
- 10,300 lines of Go
- 19 client libraries, 11 languages
- over 2 years in production





SUMMARY

- github.com/bitly/nsq
- 10,300 lines of Go
- 19 client libraries, 11 languages
- over 2 years in production



MOZ



 patterns

bitly

simple**reach**

HAIL O
TM

Path

 Life360

IN PRODUCTION

DRAMA FEVER

 EnergyHub™

Trendrr



heavy water



REONOMY

 **Rakuten**
MEDIAFORGE™



Segment.io

eventful

 **Lytics**

 **FIFO**

Thanks!

@imsnakes

<https://github.com/bitly/nsq>

*shoutout to **@jehiah** (co-author of NSQ)*